
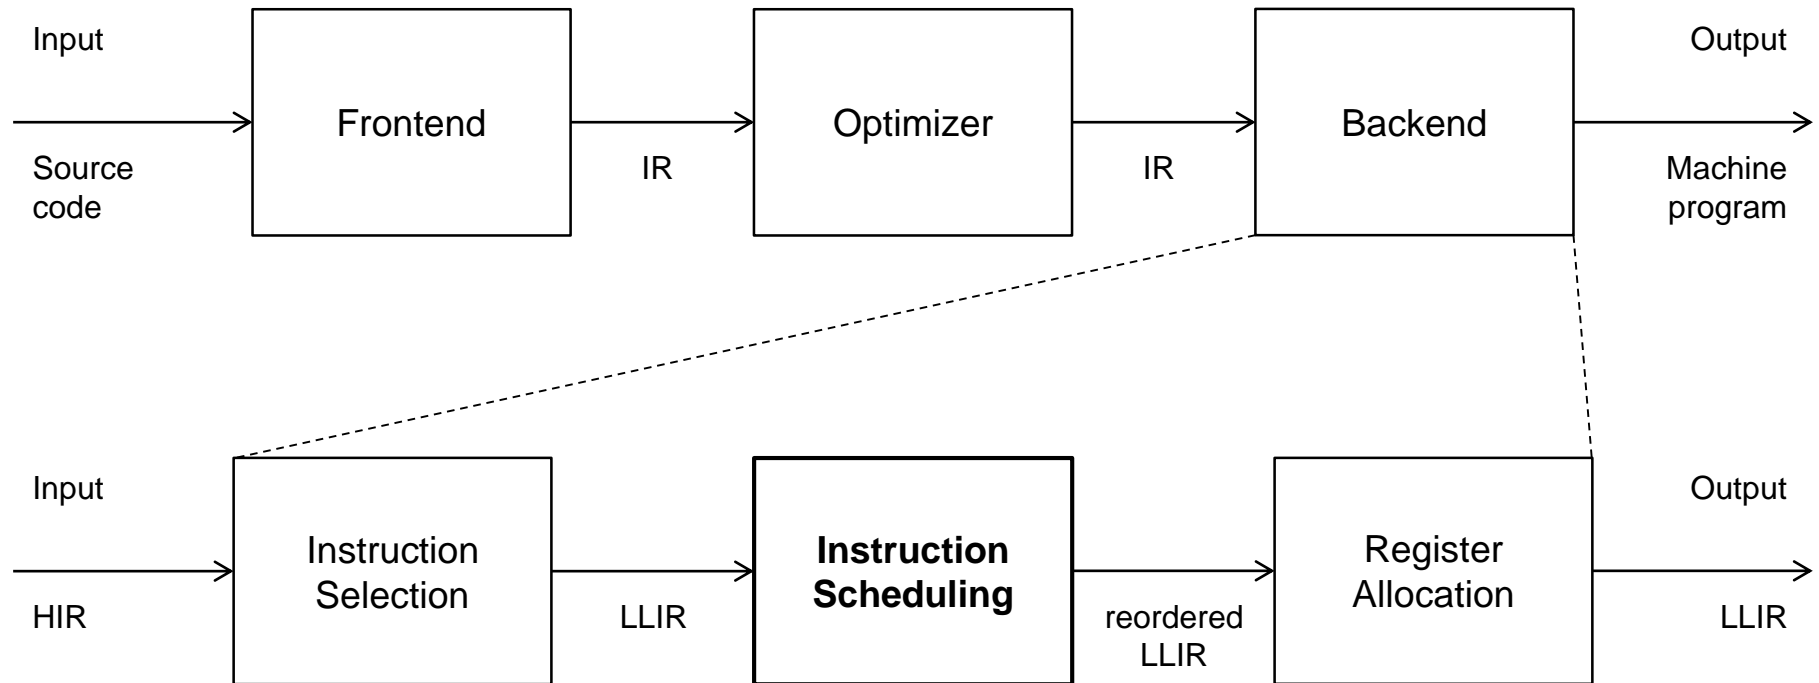


Instruction Scheduling



Instruction Scheduling




Motivation

- Or: why don't we just issue the code as-is?

```
load  r3 ← MEM[r0]
add   r4 ← r3, #1
load  r3 ← MEM[r1]
add   r3 ← r3, #1
add   r0 ← r0, #4
add   r1 ← r1, #4
mul   r3 ← r3, r4
store MEM[r2], r3
add   r2 ← r2, #4
```

LLIR




```
load  r3 ← MEM[r0]
add   r4 ← r3, #1
load  r3 ← MEM[r1]
add   r3 ← r3, #1
add   r0 ← r0, #4
add   r1 ← r1, #4
mul   r3 ← r3, r4
store MEM[r2], r3
add   r2 ← r2, #4
```

reordered
LLIR

Motivation

- Assume a target machine with the following properties
 - pipelined with forwarding, single issue, in-order
 - operation latencies: add, sub: 1 cycle; mul, load: 3 cycles; store: 1 cycle
- Executing the code from before


```
load  r3 ← MEM[r0]
add   r4 ← r3, #1
load  r3 ← MEM[r1]
add   r3 ← r3, #1
add   r0 ← r0, #4
add   r1 ← r1, #4
mul   r3 ← r3, r4
store MEM[r2], r3
add   r2 ← r2, #4
```



Motivation

- Assume a target machine with the following properties
 - pipelined with forwarding, single issue, in-order
 - operation latencies: add, sub: 1 cycle; mul, load: 3 cycles; store: 1 cycle
- Executing the code from before

```
load  r3 ← MEM[r0]
add   r4 ← r3, #1
load  r3 ← MEM[r1]
add   r3 ← r3, #1
add   r0 ← r0, #4
add   r1 ← r1, #4
mul   r3 ← r3, r4
store MEM[r2], r3
add   r2 ← r2, #4
```




cycle	operation
1	load r3 ← MEM[r0]
2	
3	
4	add r4 ← r3, #1
5	load r3 ← MEM[r1]
6	
7	
8	add r3 ← r3, #1
9	add r0 ← r0, #4
10	add r1 ← r1, #4
11	mul r3 ← r3, r4
12	
13	
14	store MEM[r2], r3
15	add r2 ← r2, #4

Motivation

■ Can we do better?

```
load  r3 ← MEM[r0]
add   r4 ← r3, #1
load  r3 ← MEM[r1]
add   r3 ← r3, #1
add   r0 ← r0, #4
add   r1 ← r1, #4
mul   r3 ← r3, r4
store MEM[r2], r3
add   r2 ← r2, #4
```



cycle	operation
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	
15	


target machine properties

- pipelined with forwarding, single issue, in-order
- operation latencies: add, sub: 1 cycle; mul, load: 3 cycles; store: 1 cycle

Motivation

■ Can we do better?

```
load  r3 ← MEM[r0]
add   r4 ← r3, #1
load  r3 ← MEM[r1]
add   r3 ← r3, #1
add   r0 ← r0, #4
add   r1 ← r1, #4
mul   r3 ← r3, r4
store MEM[r2], r3
add   r2 ← r2, #4
```



cycle	operation
1	load r4 ← MEM[r0]
2	load r3 ← MEM[r1]
3	add r0 ← r0, #4
4	add r4 ← r4, #1
5	add r3 ← r3, #1
6	mul r3 ← r3, r4
7	add r1 ← r1, #4
8	
9	store MEM[r2], r3
10	add r2 ← r2, #4

target machine properties

- pipelined with forwarding, single issue, in-order
- operation latencies: add, sub: 1 cycle; mul, load: 3 cycles; store: 1 cycle

Motivation

■ Comparison

cycle	operation
1	load r3 \leftarrow MEM[r0]
2	
3	
4	add r4 \leftarrow r3, #1
5	load r3 \leftarrow MEM[r1]
6	
7	
8	add r3 \leftarrow r3, #1
9	add r0 \leftarrow r0, #4
10	add r1 \leftarrow r1, #4
11	mul r3 \leftarrow r3, r4
12	
13	
14	store MEM[r2], r3
15	add r2 \leftarrow r2, #4

versus

cycle	operation
1	load r4 \leftarrow MEM[r0]
2	load r3 \leftarrow MEM[r1]
3	add r0 \leftarrow r0, #4
4	add r4 \leftarrow r4, #1
5	add r3 \leftarrow r3, #1
6	mul r3 \leftarrow r3, r4
7	add r1 \leftarrow r1, #4
8	
9	store MEM[r2], r3
10	add r2 \leftarrow r2, #4

33% improvement

- Can we do even better? What is the optimal schedule?
- Should we do better? After all, we got out-of-order processors by now


Instruction Scheduling

- Solves the problem: how do we get a good schedule?
i.e., defines an ordering of the operations in a partially ordered list of operations
 - main constraint
 - ▶ preserve meaning of the code
(control flow, data flow)
 - “good” schedule?
 - ▶ typically “shortest” in terms of execution time
 - ▶ under additional constraints (i.e., register pressure)
 - under consideration of H/W properties
 - ▶ operation latencies
 - ▶ processor pipeline
 - ▶ # of functional units (FU) available
 - ▶ memory hierarchy
 - ▶ ...

Data Dependence Graph

■ $DDG = (V, E)$


- nodes V representing each operation and augmented with
 - ▶ operation type
 - ▶ operation latency (delay)
- edges E representing data dependences between operations
 - ▶ forward (def-use)
 - ▶ anti (use-def)
 - ▶ output (def-def)
- root nodes = no successors
leave nodes = no predecessors
- latencies on nodes or edges?
- latency of anti/output dependences?



Renaming

■ Dealing with anti/output dependences

- anti/output dependences are artificial dependences that constrain the scheduler
- can be eliminated by renaming
 - ▶ effect on register pressure?
 - ▶ can we eliminate all anti/output dependences?



Instruction Schedule

■ $S(n): n \in V \rightarrow t \in \mathbb{N}^+$

a mapping from an operation n to a non-negative integer t denoting the cycle in which the operation should be issued.

Constraints:

- $S(n) \geq 1$ (and at least one operation o with $S(o) = 1$)
- if $(n_1, n_2) \in E$ then $S(n_1) + \text{delay}(n_1) \leq S(n_2)$
- for each t , there are no more operations with $S(n) = t$ than the H/W can support

■ Length of the schedule

$$L(S) = \max_{n \in V} (S(n) + \text{delay}(n))$$


S

cycle t	operation n
1	load $r4 \leftarrow \text{MEM}[r0]$
2	load $r3 \leftarrow \text{MEM}[r1]$
3	add $r0 \leftarrow r0, \#4$
4	add $r4 \leftarrow r4, \#1$
5	add $r3 \leftarrow r3, \#1$
6	mul $r3 \leftarrow r3, r4$
7	add $r1 \leftarrow r1, \#4$
8	
9	store $\text{MEM}[r2], r3$
10	add $r2 \leftarrow r2, \#4$

$$L(S) = 11$$

Instruction Schedule

- Path length
starting at the root, annotate each node with its accumulated delay
- the **critical path** is the longest path over all paths in the data dependence graph
- the minimal schedule cannot be shorter than the critical path





Instruction Scheduling Techniques

- Local instruction scheduling is an NP-complete problem
(scheduling → job shop scheduling → TSP)
- Classification of techniques

scheduling	cycle	
	operation	
search	greedy	
	backtrack	
flow analysis	linear	
	graph	
scheduling unit	acyclic	basic block
		trace
		DAG
	cyclic	

Cycle vs. Operation Scheduling



- operation scheduling more powerful than cycle-based scheduling in the presence of long-latency operations
 - however, much more complicated to implement

Linear vs. Graph-Based Techniques

■ Linear techniques

- runtime $O(n)$
- produces the schedule by one or more passes over the input LLIR
- most common technique: critical-path scheduling
 - ▶ three passes: ASAP, ALAP, non-critical operations
- limitation: unable to consider global properties of operations

■ Graph-based techniques

- runtime: $O(n^2)$ for DAG creation plus scheduling
- prevalent technique: list scheduling ($O(n \log n)$)
 - ▶ greedy: select one operation and schedules it

List Scheduling

- The prevalent scheduling heuristics are based on **list scheduling**
- Method
 - rename (optional)
 - build data dependence graph
 - assign priorities to operations
 - iteratively select and schedule an operation

List Scheduling

■ The list scheduling algorithm

```
t := 1
ready := { leaves of DDG }
active := {}
while (ready  $\cup$  active  $\neq$  {}) do
  for each operation o in active do
    if (S(o) + delay(o) < t) then
      active := active \ {o}
      for each successor s of o in DDG do
        if (s is ready) then
          ready := ready  $\cup$  {s}

  if (ready  $\neq$  {}) then
    o := pick the operation from ready
      with the highest priority
    if (o can be scheduled on the H/W units) then
      ready := ready \ {o}
      active := active  $\cup$  {o}
      S(op) := t

  t := t + 1
end
```

< or \leq ?

List Scheduling

- Picking an operation from `ready`
 - if `ready` never contains more than one operation, the generated schedule is optimal
 - if more than one operations are ready, the choice of the next-to-be-scheduled operation is critical to the performance of the algorithm
 - ▶ pick the operation with the highest priority
 - ▶ most algorithms use several priorities to break ties
- How do we compute these priorities?

Priority Functions

■ Priority function in list scheduling

```
procedure ListScheduling (...);  
begin  
  ...  
  o := pick one operation from ready  
       using some priority function;  
  ...  
end;
```

- common priority functions:
 - ▶ height: distance from exit node
gives priority to amount of work left
 - ▶ slackness: inverse of slack
gives priority to operations on the critical path
 - ▶ register use: number of source operands
reduces the number of live registers
 - ▶ uncover: fanout (number of children)
free up nodes quickly
 - ▶ original instruction order

List Scheduling

■ Priorities based on the DDG

- *Estart*: earliest start time (ASAP – as soon as possible)

$$Estart(op) = \begin{cases} 0 & \text{if } op \text{ has no predecessors} \\ \max_{p \in pred(op)} Estart(p) + latency(p) & \text{otherwise} \end{cases}$$

- *Lstart*: latest start time (ALAP – as late as possible)

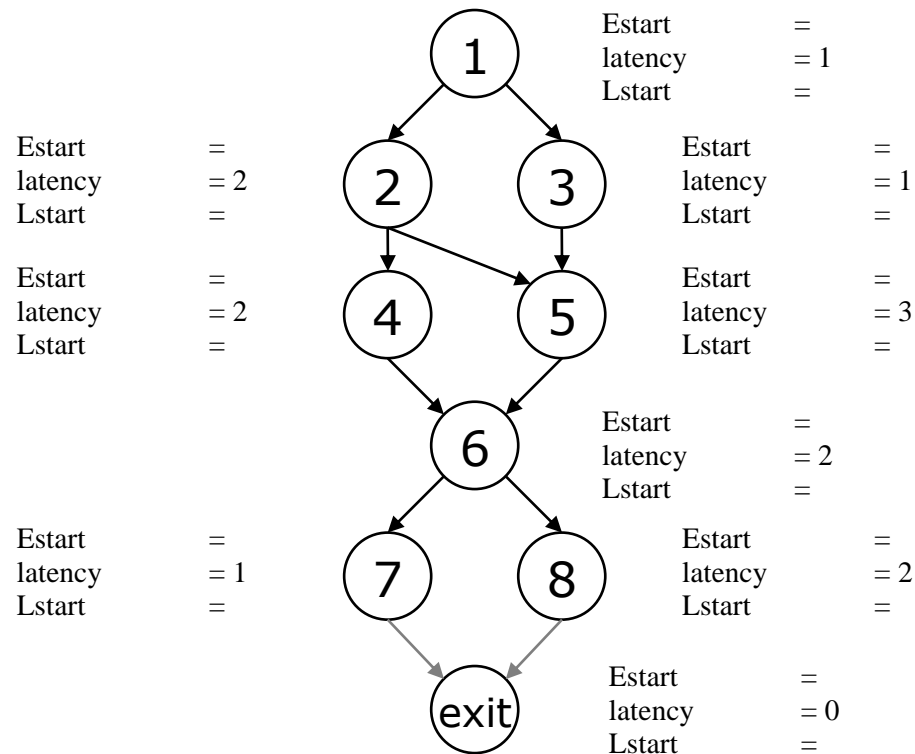
$$Lstart(op) = \begin{cases} Estart(op) & \text{if } op \text{ has no successors} \\ \min_{s \in succ(op)} Lstart(s) - latency(op) & \text{otherwise} \end{cases}$$

- *slack*: scheduling freedom

$$slack(op) = Lstart(op) - Estart(op)$$

List Scheduling

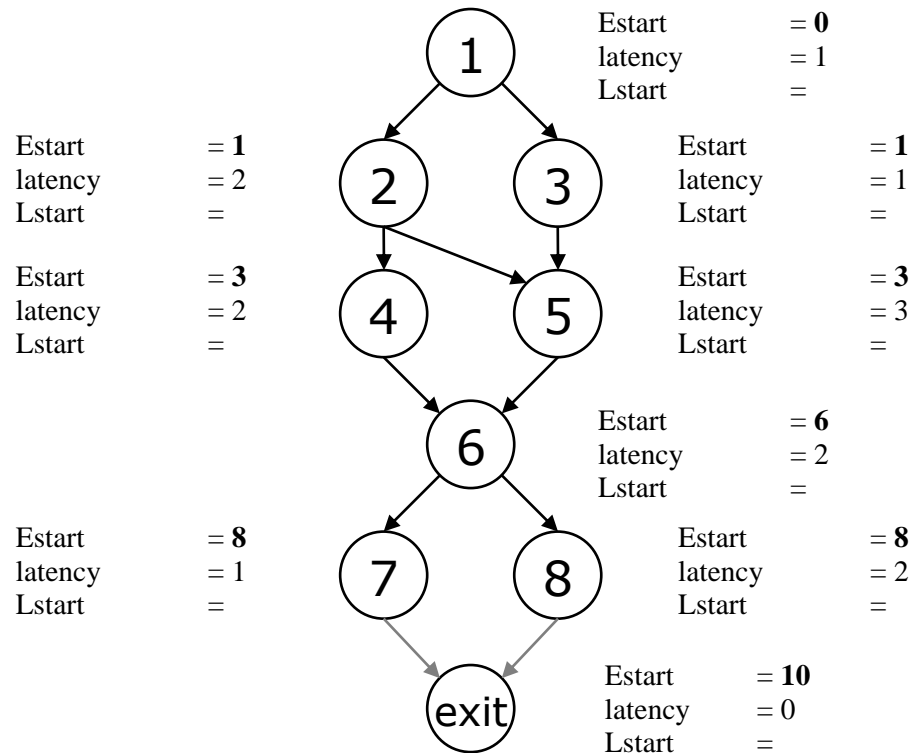
■ Computing *Estart*, *Lstart*, *slack*



List Scheduling

■ Computing *Estart*, *Lstart*, *slack*

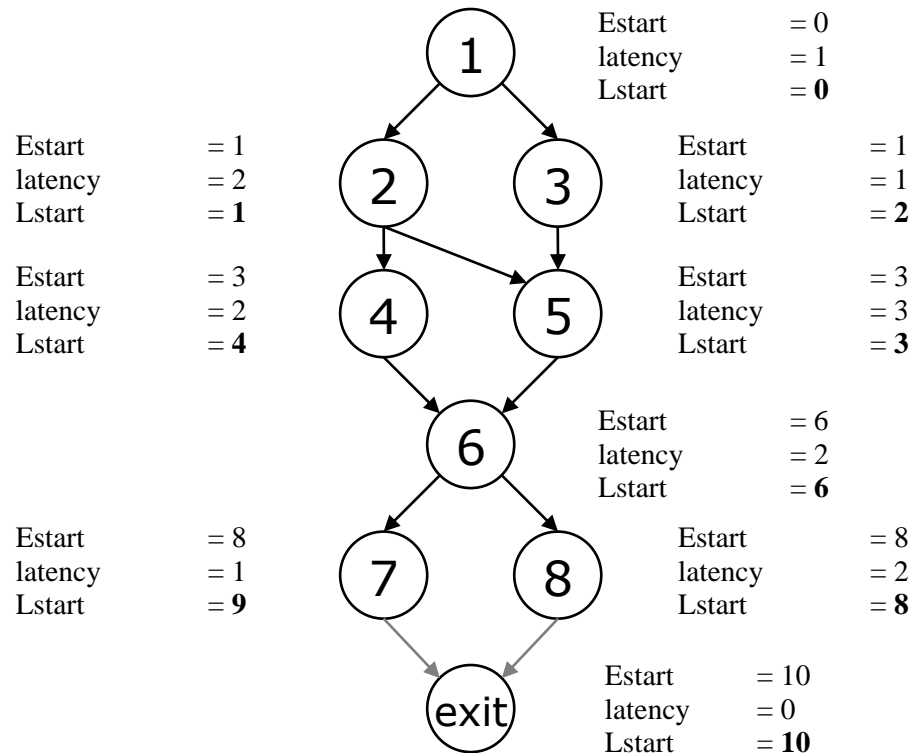
$$Estart(op) = \begin{cases} 0 & \text{if } op \text{ has no predecessors} \\ \max_{p \in pred(op)} Estart(p) + latency(p) & \text{otherwise} \end{cases}$$



List Scheduling

■ Computing $Estart$, $Lstart$, $slack$

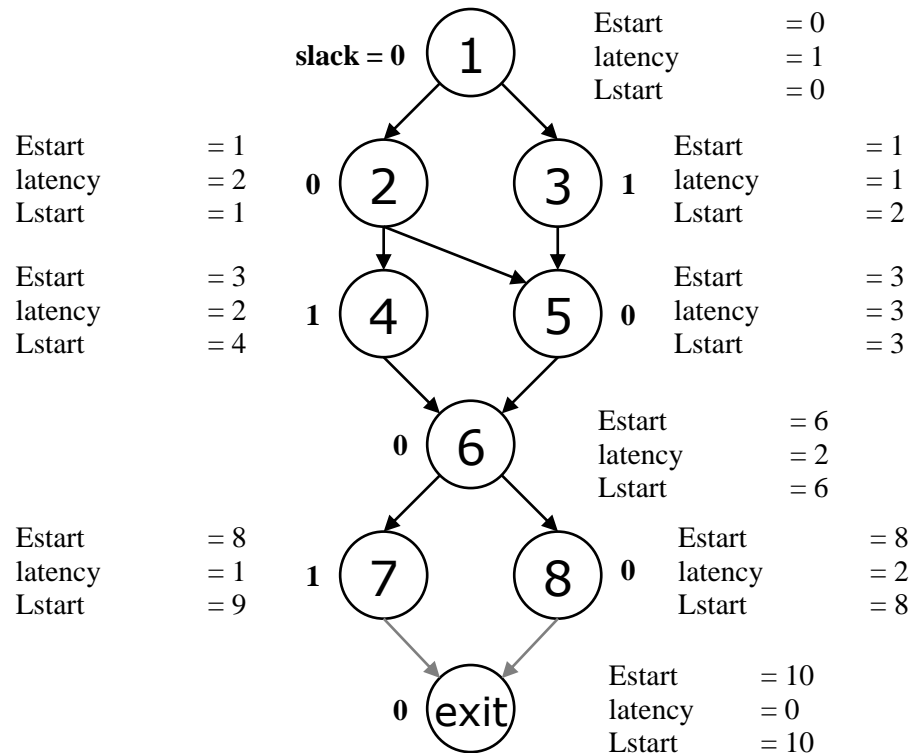
$$Lstart(op) = \begin{cases} Estart(op) & \text{if } op \text{ has no successors} \\ \min_{s \in succ(op)} Lstart(s) - latency(op) & \text{otherwise} \end{cases}$$



List Scheduling


■ Computing $Estart$, $Lstart$, $slack$

$$slack(op) = Lstart(op) - Estart(op)$$



List Scheduling

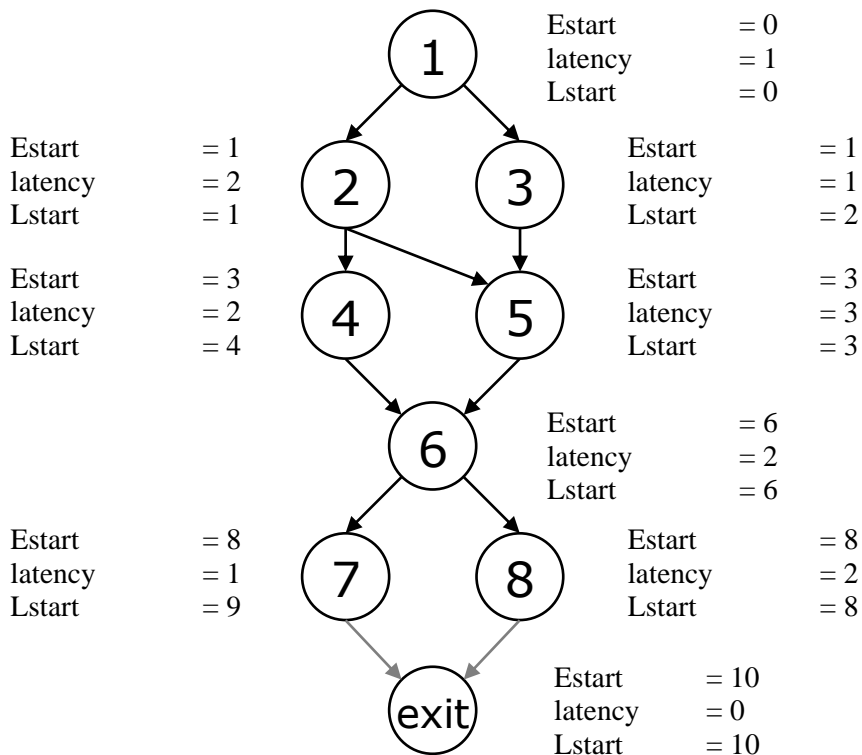
- Another way to look at the critical path
 - sequence of critical operations
 - critical operation: $slack(op) = 0$



Priority Function: Height-Based

■ Height-based priority function

- gives priority to amount of work left
- $priority(op) = Lstart(exit) - Lstart(op) + 1$

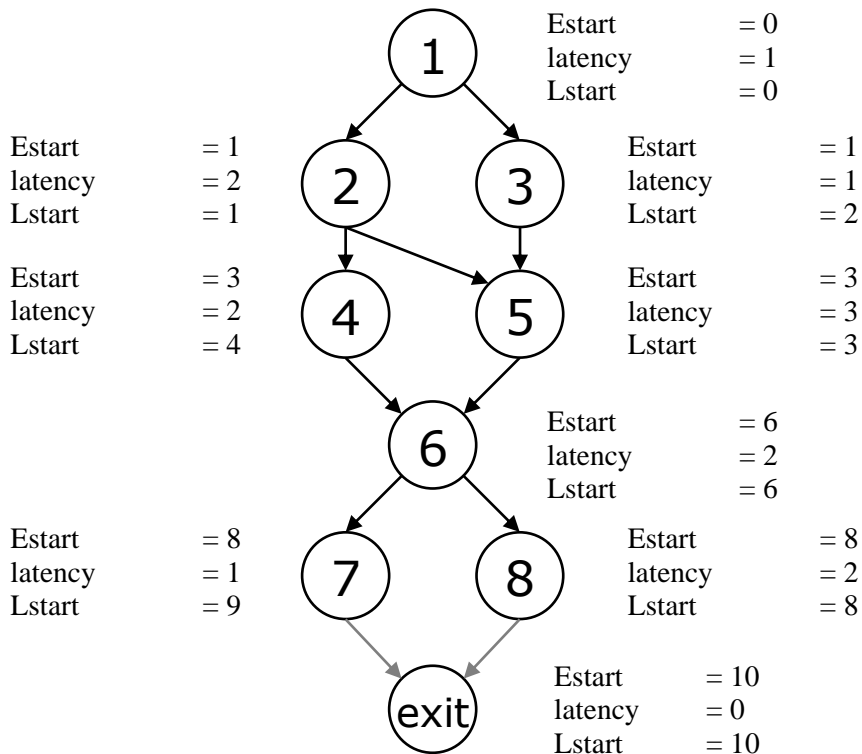


op	priority
1	
2	
3	
4	
5	
6	
7	
8	
exit	

Priority Function: Height-Based

■ Height-based priority function

- gives priority to amount of work left
- $priority(op) = Lstart(exit) - Lstart(op) + 1$



op	priority
1	11
2	10
3	9
4	7
5	8
6	5
7	2
8	3
exit	1

Example


■ Applying height-based list scheduling to the previous example

```

1 load  r4 ← MEM[r0]
2 add   r4 ← r4, #1
3 load  r3 ← MEM[r1]
4 add   r3 ← r3, #1
5 add   r0 ← r0, #4
6 add   r1 ← r1, #4
7 mul   r3 ← r3, r4
8 store MEM[r2], r3
9 add   r2 ← r2, #4
    
```

target machine properties

- pipelined with forwarding, single issue, in-order
- operation latencies: add, sub: 1 cycle;
mul, load: 3 cycles; store: 1 cycle




Example

■ Applying height-based list scheduling to the previous example

```
1 load  r4 ← MEM[r0]
2 add   r4 ← r4, #1
3 load  r3 ← MEM[r1]
4 add   r3 ← r3, #1
5 add   r0 ← r0, #4
6 add   r1 ← r1, #4
7 mul   r3 ← r3, r4
8 store MEM[r2], r3
9 add   r2 ← r2, #4
```


target machine properties

- pipelined with forwarding, single issue, in-order
- operation latencies: add, sub: 1 cycle;
mul, load: 3 cycles; store: 1 cycle



Example

- Applying height-based list scheduling to the previous example



op	priority
1	9
3	9
2	6
4	6
7	5
5	2
6	2
8	2
9	2
exit	1

Example

■ Applying height-based list scheduling to the previous example

Initialization

```
t := 1
ready := { 1, 3, 5, 6 }
active := {}
```

Iteration:


```
t = 1, ready = {1,3,5,6}, active = {}
active is empty → pass
```

```
ready is not empty
o := 1 (with priority 9)
S(1) = 1
```

```
t = 2, ready = {3,5,6}, active = {1}
active is not empty
o = 1: S(1)+delay(1) ≤ t? no
```

```
ready is not empty
o := 3 (with priority 9)
ready := {5,6}, active = {1,3}
S(3) = 2
```

op	priority
1	9
3	9
2	6
4	6
7	5
5	2
6	2
8	2
9	2
exit	1



S	
t	op
1	1

S	
t	op
1	1
2	3

Example

■ Applying height-based list scheduling to the previous example

Iteration:

t = 3, ready = {5,6}, active = {1,3}

active is not empty

o = 1: $S(1) + \text{delay}(1) \leq 3$? no

o = 3: $S(3) + \text{delay}(3) \leq 3$? No

ready is not empty

o := 5 (with priority 2)

ready := {6}, active = {1,3,5}

$S(5) = 3$

t = 4, ready = {6}, active = {1,3,5}

active is not empty

o = 1: $S(1) + \text{delay}(1) \leq 4$? yes

active := active \ {1}

ready := ready \cup {2}

o = 3: $S(3) + \text{delay}(3) \leq 4$? No

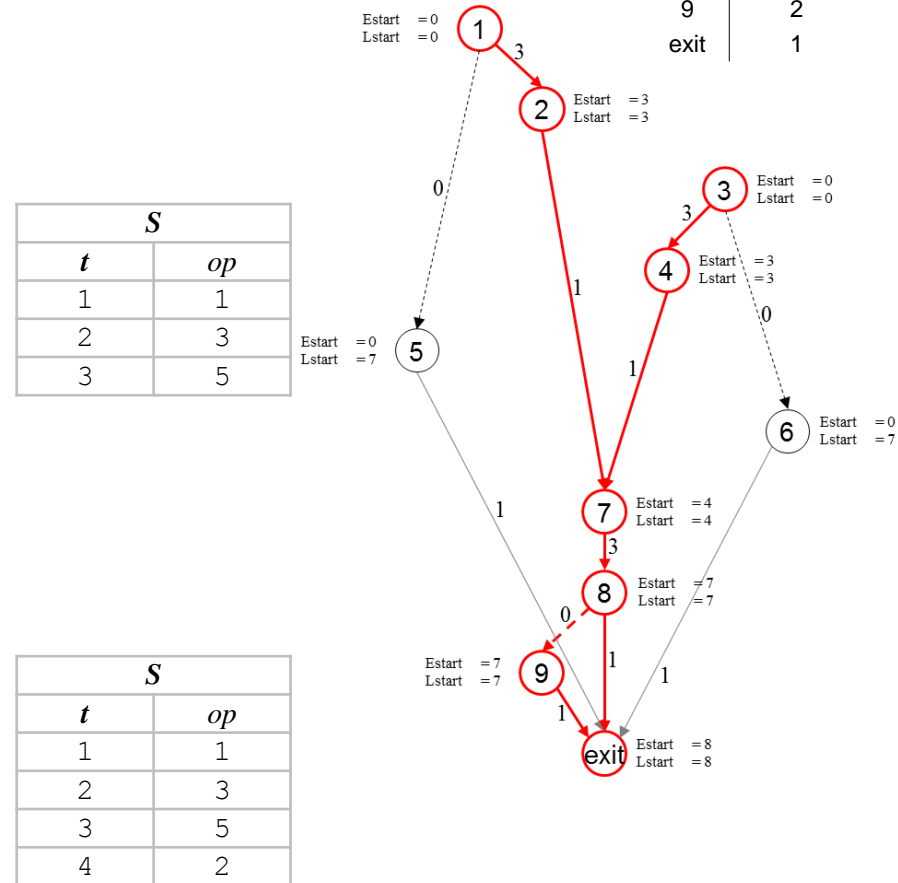
ready is not empty

o := 2 (with priority 6)

ready := {6}, active = {3,5,2}

$S(2) = 4$

op	priority
1	9
3	9
2	6
4	6
7	5
5	2
6	2
8	2
9	2
exit	1



Example

■ Applying height-based list scheduling to the previous example

Iteration:

$t = 5$, **ready** = {6}, **active** = {3,5,2}

active is not empty

o = 3: $S(3) + \text{delay}(3) \leq 5$? yes

active := active \ {3}

ready := ready \cup {4}

o = 5: $S(5) + \text{delay}(5) \leq 4$? yes

active := active \ {5}

no data dependent successors

o = 2: $S(2) + \text{delay}(2) \leq 4$? yes

active := active \ {2}

successor (7) not ready due to 4

ready is not empty


o := 4 (with priority 6)

ready := {6}, active = {4}

$S(4) = 5$

S	
t	op
1	1
2	3
3	5
4	2
5	4

op	priority
1	9
3	9
2	6
4	6
7	5
5	2
6	2
8	2
9	2
exit	1



Example

■ Applying height-based list scheduling to the previous example

Iteration:

t = 6, ready = {6}, active = {4}

active is not empty

o = 4: $S(4) + \text{delay}(4) \leq 6$? ne

active := active \ {4}

ready := ready \cup {7}

ready is not empty

o := 7 (with priority 5)

ready := {6}, active = {7}

$S(7) = 6$

t = 7, ready = {6}, active = {7}

active is not empty

o = 7: $S(7) + \text{delay}(7) \leq 7$? No

ready is not empty


o := 6 (with priority 2)

ready := {}, active = {6, 7}

$S(6) = 7$

<i>S</i>	
<i>t</i>	<i>op</i>
1	1
2	3
3	5
4	2
5	4
6	7

<i>S</i>	
<i>t</i>	<i>op</i>
1	1
2	3
3	5
4	2
5	4
6	7
7	6



Example

■ Applying height-based list scheduling to the previous example

Iteration:

t = 8, ready = {}, active = {6,7}

active is not empty

o = 6: $S(6) + \text{delay}(6) \leq 8$? yes

active := active \ {6}

no data dependent successors

o = 7: $S(7) + \text{delay}(7) \leq 8$? no

ready is empty

t = 9, ready = {}, active = {7}

active is not empty

o = 7: $S(7) + \text{delay}(7) \leq 9$? yes

active := active \ {7}

ready := ready \cup {8,9}


ready is not empty

o := 8 (with priority 2)

ready := {9}, active = {8}

$S(8) = 9$

op	priority
1	9
3	9
2	6
4	6
7	5
5	2
6	2
8	2
9	2
exit	1



Example

■ Applying height-based list scheduling to the previous example

Iteration:

t = 10, ready = {9}, active = {8}
 active is not empty
 o = 8: $S(8) + \text{delay}(8) \leq 10$? yes
 active := active \ {8}
 no data dependent successors

ready is not empty
 o := 9 (with priority 2)
 ready := {}, active = {9}
 $S(9) = 10$

t = 11, ready = {}, active = {9}
 active is not empty
 o = 9: $S(9) + \text{delay}(9) \leq 11$? yes
 active := active \ {9}
 no data dependent successors


ready is empty

done.

op	priority
1	9
3	9
2	6
4	6
7	5
5	2
6	2
8	2
9	2
exit	1

S	
t	op
1	1
2	3
3	5
4	2
5	4
6	7
7	6
8	
9	8
10	9

S	
t	op
1	1
2	3
3	5
4	2
5	4
6	7
7	6
8	
9	8
10	9



Instruction Scheduling & Register Allocation

- Phase ordering between instruction scheduling and register allocation
 - effects of the scheduler on the RA
 - ▶ the scheduler can use renaming to get rid of anti dependences to obtain more freedom in scheduling
 - ▶ the resulting overlap of previously constrained operations may increase register pressure
 - ▶ which, in turn, may force the register allocator to spill one more variable
 - and vice versa (the RA constrains the scheduler in a RA-first compiler)
- Combining scheduling and register allocation
 - has the potential to produce better solutions
 - typically not done due to complexity