


Introduction to Computer Architecture



Computer Architecture

“Computer architecture is a specification detailing how a set of software and hardware technology standards interact to form a computer system or platform. In short, computer architecture refers to how a computer system is designed and what technologies it is compatible with. “

source: Technopedia

- Computer architecture teaches you
 - how a computer is **controlled**
 - how a computer is **built**

After This Course, You Will...

- understand the **functionality and operation of the basic elements of a computer system** including the processor, memory, and input/output
- understand the **hardware/software interface**
- understand and be able to decode programs written in **assembly language**

Credits For The Entire Course

- Slides and material adapted mainly from
 - the csapp textbook
 - H&P's "Computer Architecture: A Quantitative Approach"
 - Tannenbaum's "Structured Computer Organization"
 - slides and projects from CMU
 - slides from David Black-Schaffer (introduction)

Why Study Computer Architecture?

ARM Unveils its Most Energy Efficient Application Processor Ever; Redefines Traditional Power And Performance Relationship With big.LITTLE Processing

19 October 2011

Addresses one of today's industry challenges: extending consumers' always on, always connected mobile experience with both improved performance AND longer battery life


Cambridge, UK – 19th October 2011 – ARM today announced the ARM® Cortex™-A7 MPCore™ processor - the most energy-efficient application class processor ARM has ever developed, and big.LITTLE processing - a flexible approach that redefines the traditional power and performance relationship. The Cortex-A7 processor builds on the low-power leadership established by the Cortex-A9 processor that is at the heart of many of today's most popular smartphones. A

(source: ARM)

- big.LITTLE processing?
- ARM?

Who Owns Device With an ARM Inside?

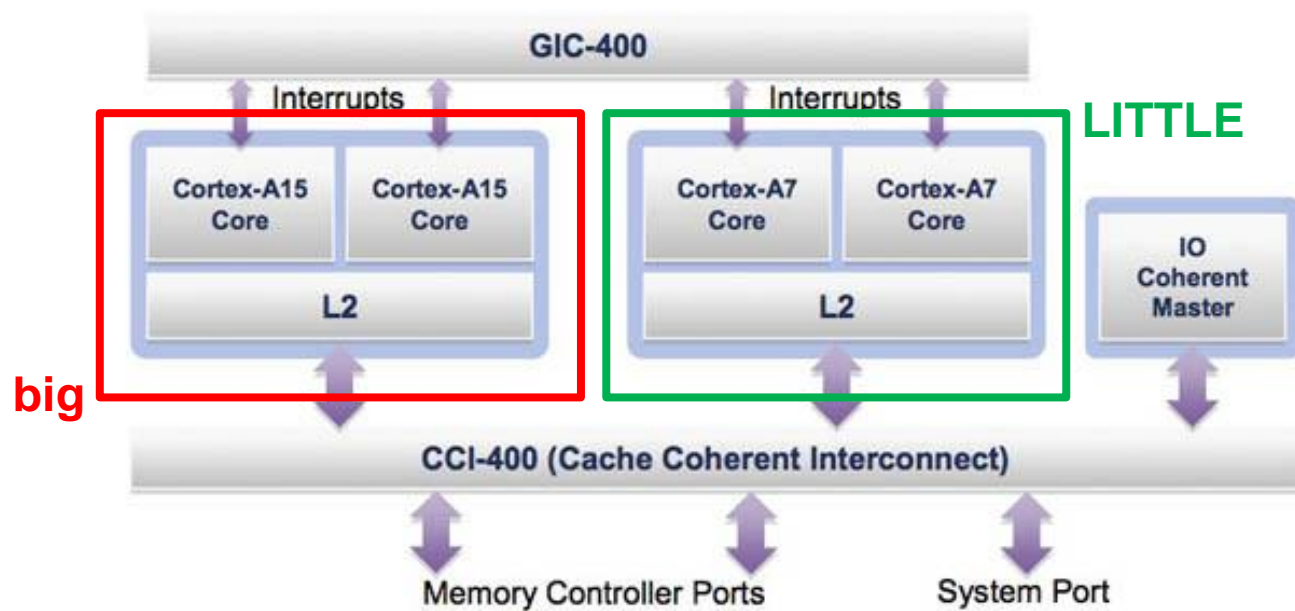
■ All of you:



source: <http://www.gartner.com/newsroom/id/2610015>

What Are They Doing With big.LITTLE?


- big.LITTLE
 - big cores for high performance
 - little cores for low performance



source: http://www.eetimes.com/document.asp?doc_id=1279167

And Why Are They Doing It?

- One reason: **power efficiency = computations / energy**




source: <http://www.mobilegeeks.com/mediatek-outperforms-qualcomm-welcome-true-octa-core-processing/>

The Details

■ LITTLE cores


- simple
- short pipeline



ARM Cortex-A7 Pipeline

■ big cores

- complex
- long pipeline




ARM Cortex-A15 Pipeline

source: http://www.eetimes.com/document.asp?doc_id=1279167


Why Can They Do This?

- Technology scaling: 130nm → 22nm



Why Should They Do This?

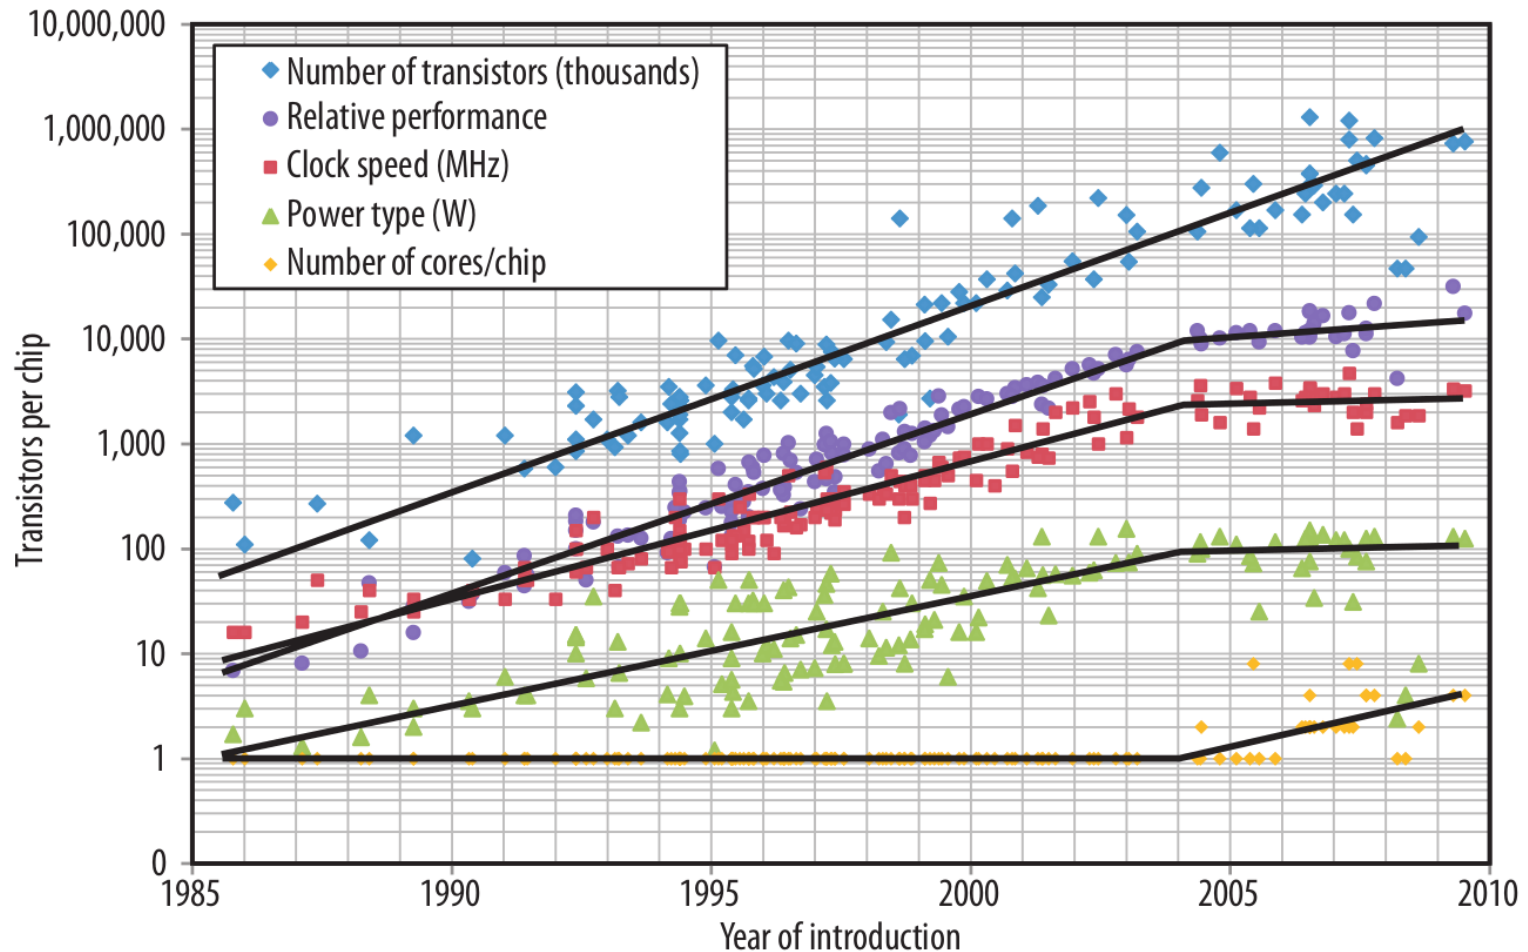
■ Single-core performance



source: Fuller et al. Computing Performance: Game Over or Next Level, IEEE, 2011

The Power Wall

- We can't increase power, need to increase power efficiency

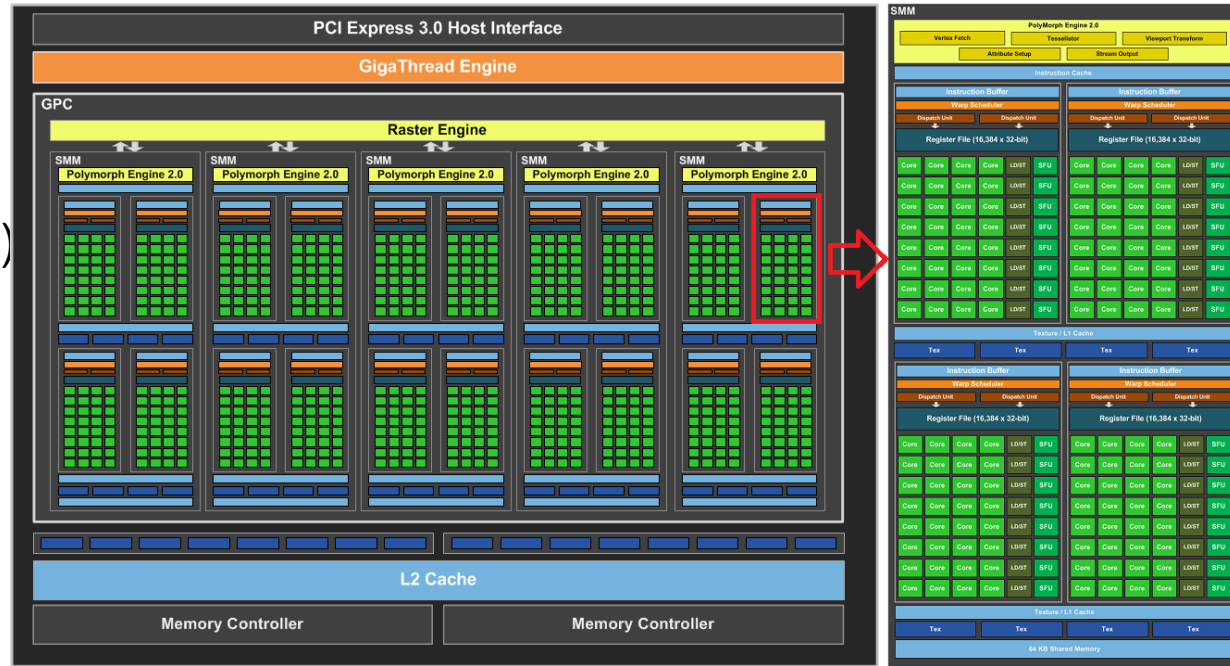


source: Fuller et al. Computing Performance: Game Over or Next Level, IEEE, 2011


What Are Others Doing?

■ NVIDIA

- GPUs (Maxwell architecture)



- mobile CPUs





source: NVIDIA

What Are Others Doing?


■ Intel

- MIC
(Many Integrated Cores)
- desktop/mobile CPUs

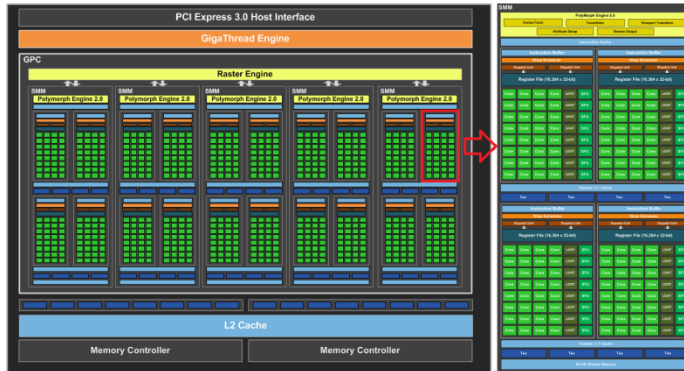
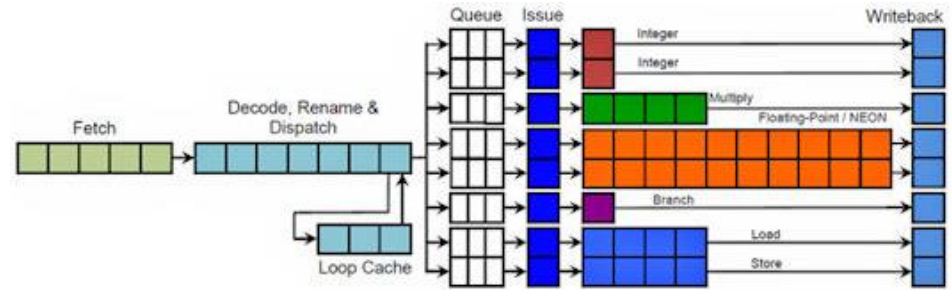


source: Intel


This is Computer Architecture



VS



VS



- Understanding the building blocks of processors and computer systems
- Understanding design tradeoffs such as performance vs efficiency
- Building the hardware
- Making it programmable

Why You Should Care...

- ...apart from this being a mandatory class (at least for CSE students)?
- Understanding computer architecture is important for many other core subjects in computer science (system programming, OS, compilers, programming models and languages, ...)
- Understanding computer architecture will make you a better programmer

```
for (i=0; i<N; i++) {  
    for (j=0; j<N; j++) {  
        C[i,j] = A[i,j] + B[i,j];  
    }  
}
```

vs

```
for (j=0; j<N; j++) {  
    for (i=0; i<N; i++) {  
        C[i,j] = A[i,j] + B[i,j];  
    }  
}
```

- Understanding assembly and a processor's ISA is still an important skill
- It is actually quite fun!

What Are We Going To Do In This Course?

■ We will *not*

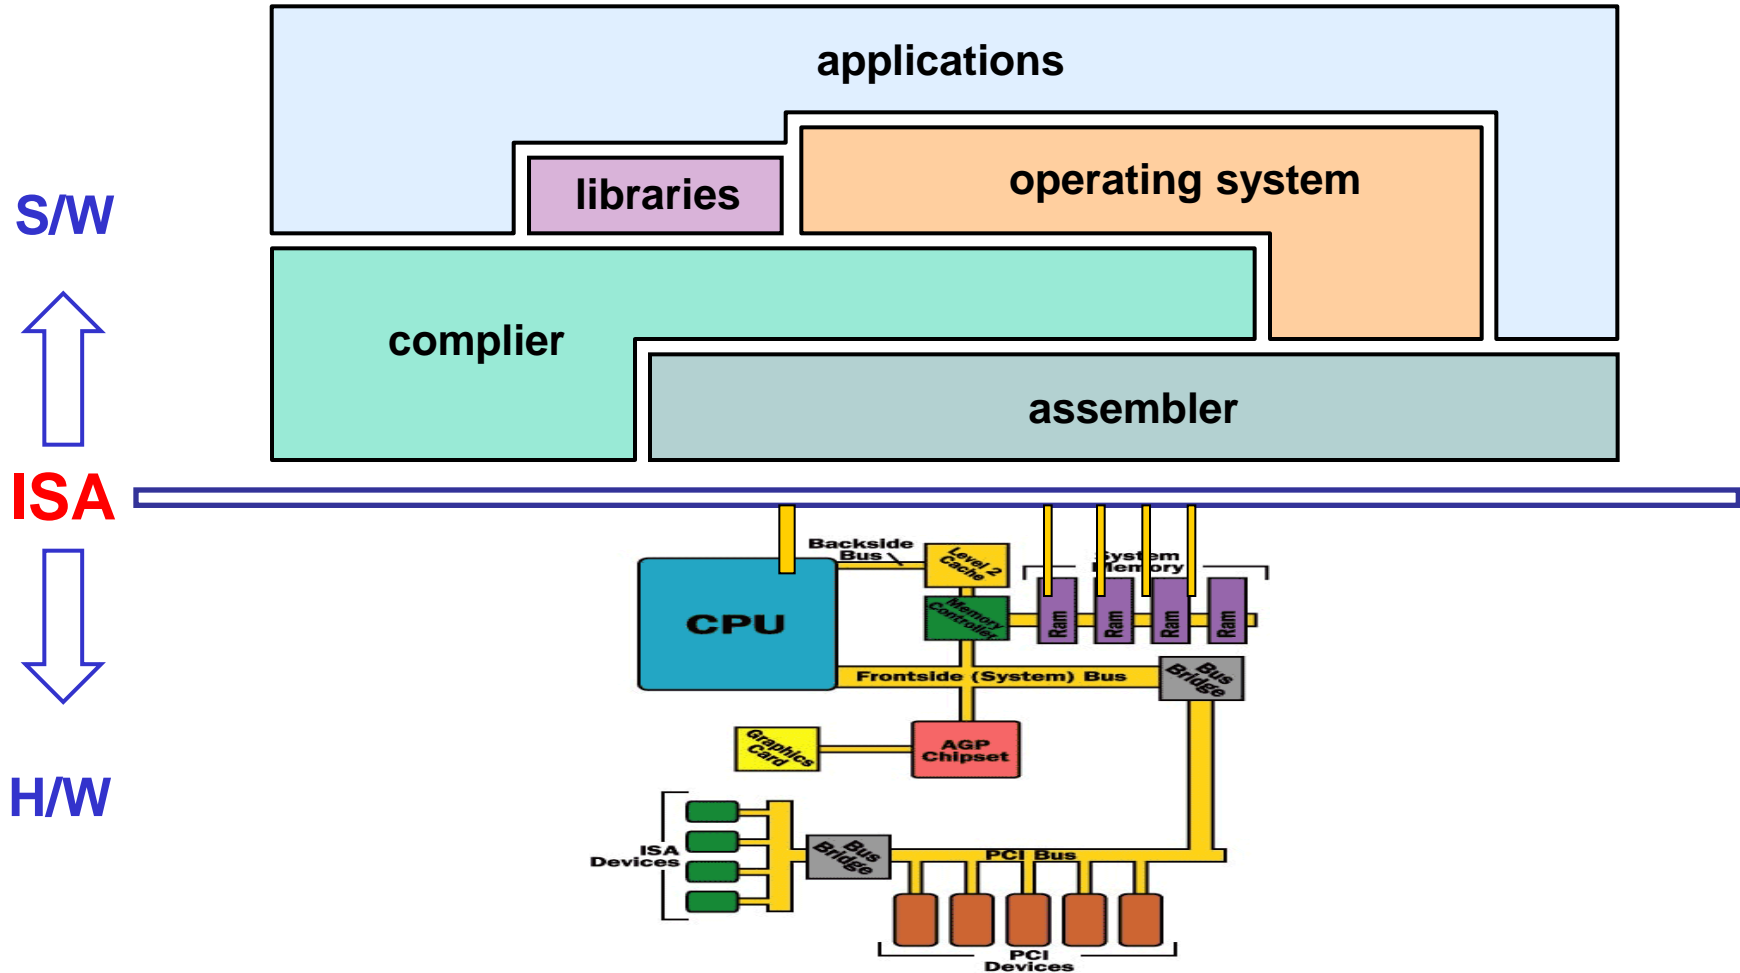
- design the next Core i9 with a 30-stage pipeline
- study how logic gates work (you know that already)
- write entire programs in assembly

■ We will

- learn an ISA (x86) and how to read and understand assembly programs
- learn how a simple pipelined processor is built
- learn about the memory hierarchy in modern computer systems
- have a quick look at modern state-of-the-art processors

What Are We Going To Do In This Course?


- The ISA (Instruction Set Architecture) as the hardware/software interface



©2001 HowStuffWorks

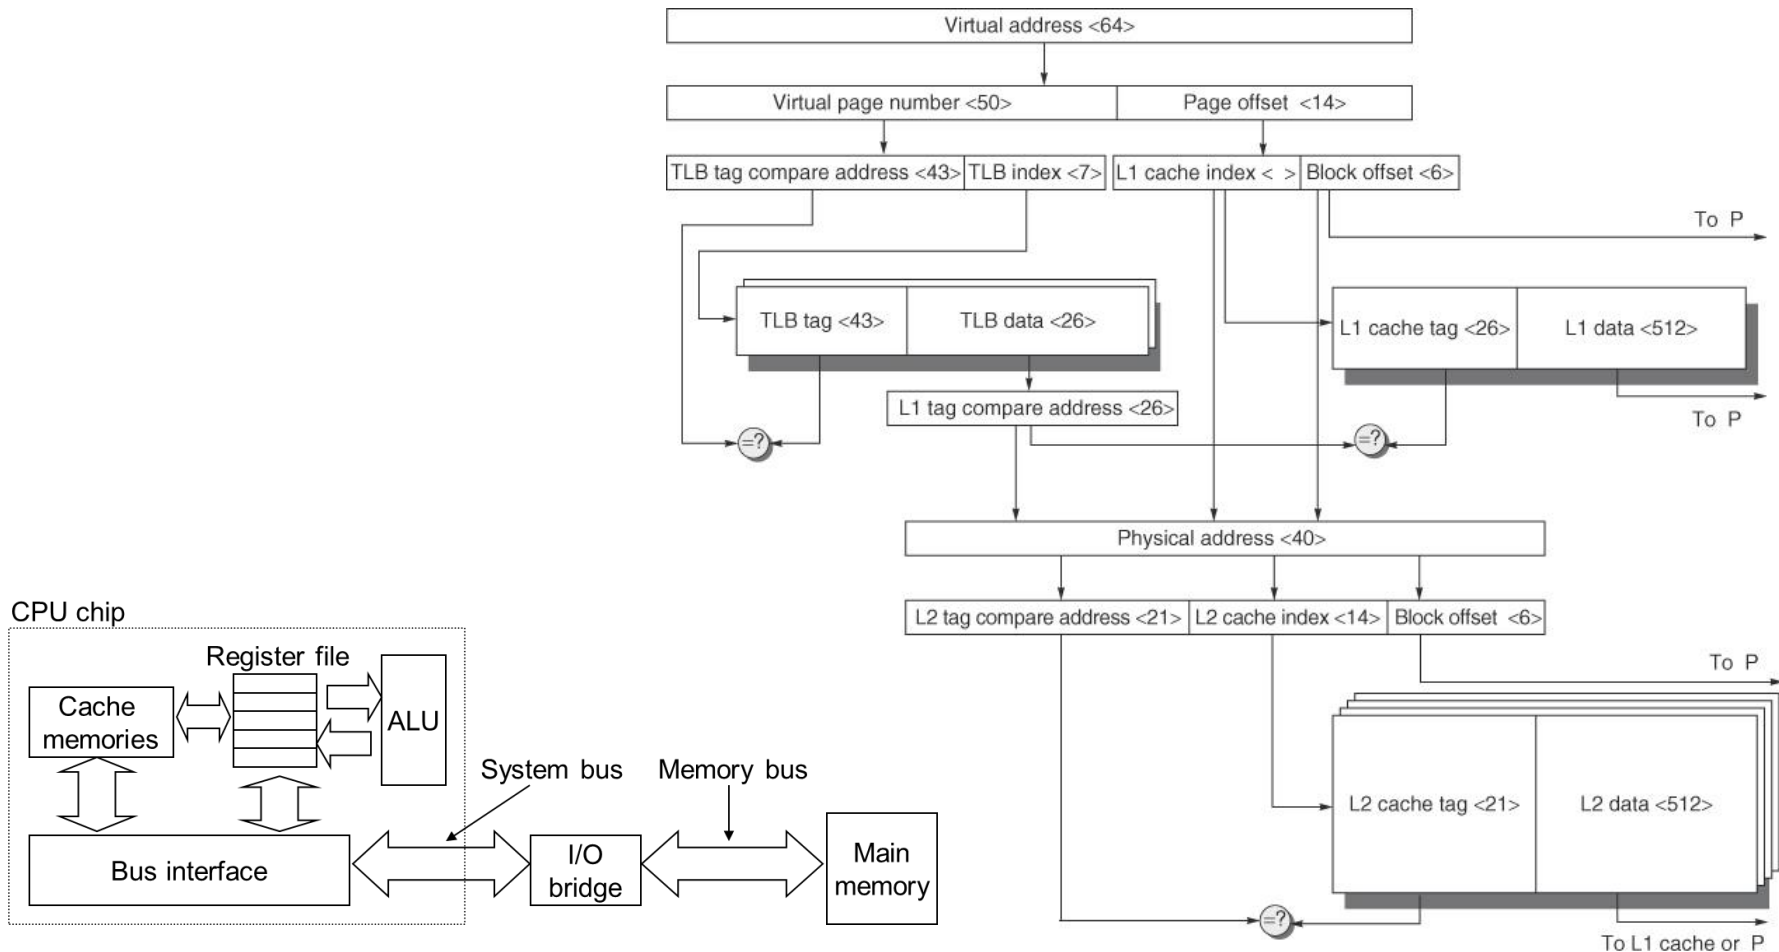
What Are We Going To Do In This Course?

- Study and modify a "simple" pipelined processor



What Are We Going To Do In This Course?

- Study the memory hierarchy in modern computer systems



Great Ideas in Computer Architecture

■ Design for Moore's Law

- anticipate state of technology when the product ships

■ Use Abstraction to Simplify Design


- abstractions are everywhere

■ Make the Common Case Fast

- Amdahl's Law

$$speedup = \frac{1}{(1 - parallel) + \frac{parallel}{\#cores}}$$

Microprocessor Transistor Counts 1971-2011 & Moore's Law



source: Wikipedia

Great Ideas in Computer Architecture

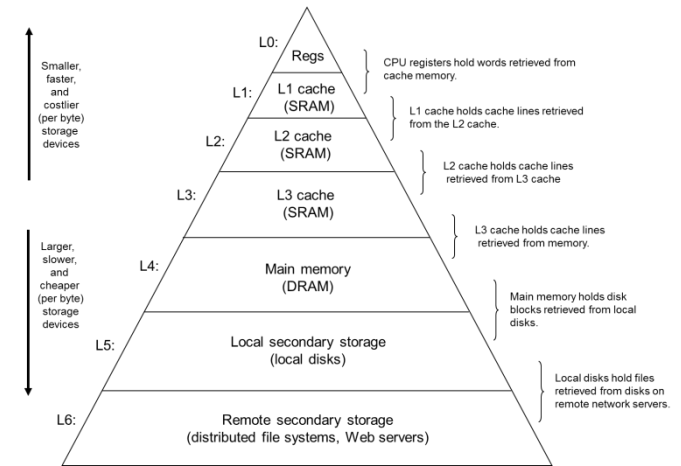
■ Hierarchy of Memories

- exploit spatial and temporal locality (공간/시간 구역성)

■ Performance via Parallelism

■ Performance via Pipelining

■ Performance via Prediction



That's It For Today

- Next class: the ISA and assembly basics
- Late enrollment: students who could not sign up for this class through the online registration system can submit their 초안지 now