

The HW/SW Interface

The x86 ISA: Machine-Level Programming Basics

```
0x08048394 <+0>:  push    %ebp
0x08048395 <+1>:  mov     %esp,%ebp
0x08048397 <+3>:  mov     0xc(%ebp),%eax
0x0804839a <+6>:  add     0x8(%ebp),%eax
0x0804839d <+9>:  pop     %ebp
0x0804839e <+10>: ret
0x8048394 <sum>:  0x55    0x89    0xe5 ...
0x804839c <sum+8>:0x08    0x5d    0xc3
```

Instruction Set Architecture (ISA)

■ ISA classification


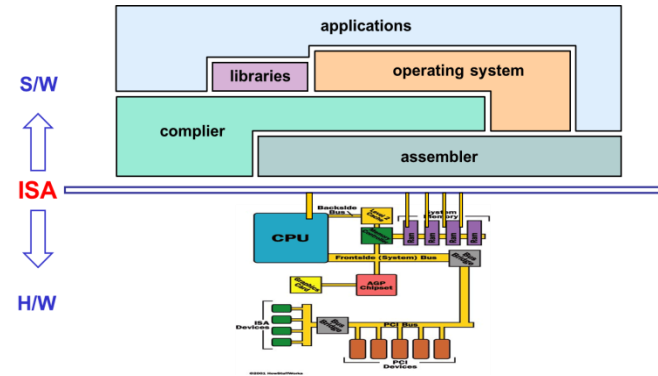
- RISC vs CISC

- general-purpose

- ▶ two or three operands?
- ▶ # of memory operands?
- ▶ classification
 - load-store
 - register-memory
 - memory-memory


- stack

- accumulator



Instruction Set Architecture (ISA)

- The ISA defines the **programmer-visible state** of a processor architecture



- the ISA tells us what we can expect the processor to do when we execute instructions (but not *how* the processor does it)
- memory addressing
 - ▶ interpretation of addresses
 - bytes, half words, words, double words, ...?
 - ▶ byte ordering in memory
 - little-endian vs. big-endian
 - ▶ addressing modes
 - how to access objects in memory

Instruction Set Architecture (ISA)

- The ISA defines the **programmer-visible state** of a processor architecture

- type and size of operands
 - ▶ specifying types
 - byte, half word, word, float, double...

- operations
 - ▶ arithmetic and logical
 - ▶ data transfer
 - ▶ control
 - ▶ system
 - ▶ floating point
 - ▶ string
 - ▶ multimedia


Rank	x86 operation	%
1	load	22
2	conditional branch	20
3	compare	16
4	store	12
5	add	8
6	and	6
7	sub	5
8	move reg-reg	4
9	call	1
10	return	1
	total	96

Instruction Set Architecture (ISA)

- The ISA defines the **programmer-visible state** of a processor architecture
 - instruction encoding
 - ▶ defines binary representation for each operation and operand
 - ▶ fixed size vs. variable size



(a) Variable (e.g., Intel 80x86, VAX)



(b) Fixed (e.g., Alpha, ARM, MIPS, PowerPC, SPARC, SuperH)

x86 Machine-Level Programming Basics

- **History of Intel processors and architectures**
- C, assembly, machine code
- Assembly Basics: Registers, operands, move
- Introduction to x86-64

Acknowledgement: slides based on the cs:app2e material

Intel x86 Processors


- Dominate laptop/desktop/server market
- Evolutionary design
 - Backwards compatible up until 8086, introduced in 1978
 - Added more features as time goes on
- Complex instruction set computer (CISC)
 - Many different instructions with many different formats
 - ▶ But, only small subset encountered with Linux programs
 - Hard to match performance of Reduced Instruction Set Computers (RISC)
 - But, Intel has done just that!
 - ▶ In terms of speed. Less so for low power.

Intel x86 Evolution: Milestones

<i>Name</i>	<i>Date</i>	<i>Transistors</i>	<i>MHz</i>
■ 8086	1978	29K	5-10
<ul style="list-style-type: none">• First 16-bit processor. Basis for IBM PC & DOS• 1MB address space			
■ 386	1985	275K	16-33
<ul style="list-style-type: none">• First 32 bit processor , referred to as IA32• Added “flat addressing”• Capable of running Unix• 32-bit Linux/gcc uses no instructions introduced in later models			
■ Pentium 4F	2004	125M	2800-3800
<ul style="list-style-type: none">• First 64-bit processor, referred to as x86-64			
■ Core i7	2008	731M	2667-3333
<ul style="list-style-type: none">• multiple cores			

Intel x86 Processors: Overview

Architectures	Processors
X86-16	8086
	286
X86-32/IA32	386
	486
	Pentium
	Pentium MMX
	Pentium III
	Pentium 4
X86-64 / EM64t	Pentium 4E
	Pentium 4F
	Core 2 Duo
	Core i7



IA: often redefined as latest Intel architecture

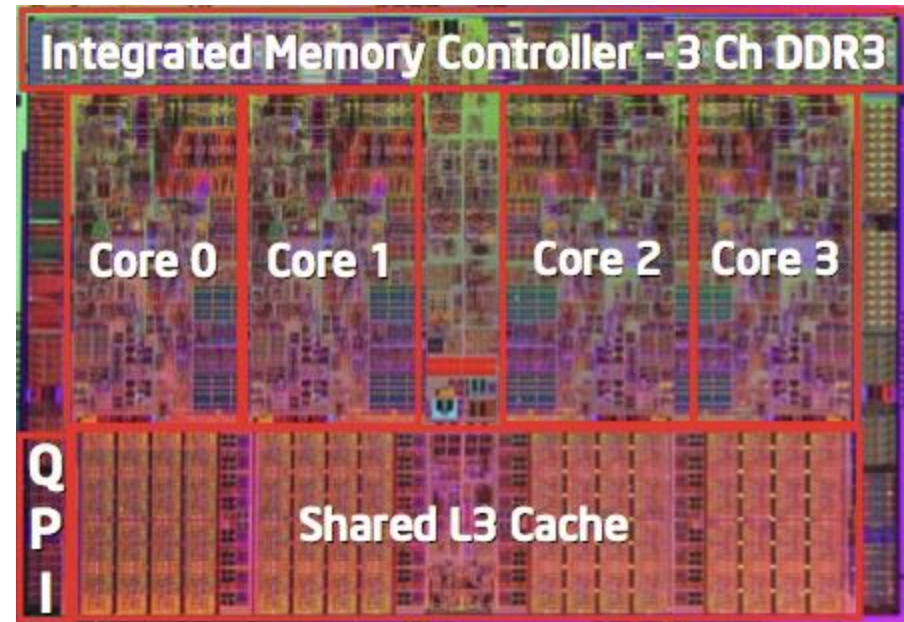
Intel x86 Processors

■ Machine Evolution

● 386	1985	0.3M
● 486	1989	1.9M
● Pentium	1993	3.1M
● Pentium/MMX	1997	4.5M
● PentiumPro	1995	6.5M
● Pentium III	1999	8.2M
● Pentium 4	2001	42M
● Core 2 Duo	2006	291M
● Core i7 (4 cores)	2008	731M
● Core i7 (6 cores)	2011	2'270M
● Xeon E7 v2 (15 c)	2014	4'310M

■ Added Features

- Instructions to support multimedia operations
 - ▶ Parallel operations on 1, 2, and 4-byte data, both integer & FP
- Instructions to enable more efficient conditional operations



Core i7 (45nm)

New Species: ia64, then IPF, then Itanium,...

<i>Name</i>	<i>Date Transistors</i>
■ Itanium	2001 10M
<ul style="list-style-type: none">● First shot at 64-bit architecture: first called IA64● Radically new instruction set designed for high performance● Can run existing IA32 programs<ul style="list-style-type: none">▶ On-board “x86 engine”● Joint project with Hewlett-Packard	
■ Itanium 2	2002 221M
<ul style="list-style-type: none">● Big performance boost	
■ Itanium 2 Dual-Core	2006 1.7B
■ Itanium has not taken off in marketplace	
<ul style="list-style-type: none">● very fast (esp. FP), very hot, and very expensive● Lack of backward compatibility, no good compiler support, Pentium 4 got too good, overtaken by 64-bit x86 designs	

x86 Clones: Advanced Micro Devices (AMD)

■ Historically

- AMD has followed just behind Intel
- A little bit slower, a lot cheaper

■ Then

- Recruited top circuit designers from Digital Equipment Corp. and other downward trending companies
- Built Opteron: tough competitor to Pentium 4
- Developed x86-64, their own extension to 64 bits

■ Recently

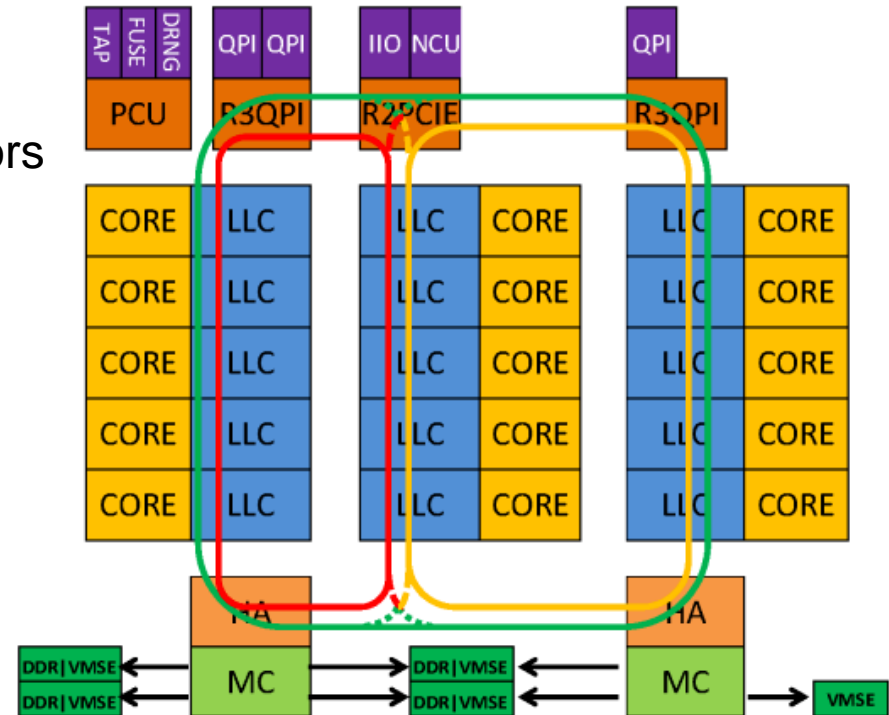

- Intel much quicker with multi-core design
- Intel currently far ahead in performance
- Intel em64t backwards compatible to x86-64

Intel's 64-Bit

- Intel Attempted Radical Shift from IA32 to IA64
 - Totally different architecture (Itanium)
 - Executes IA32 code only as legacy
 - Application performance disappointing
- AMD Stepped in with Evolutionary Solution
 - x86-64 (now called “AMD64”)
- Intel Felt Obligated to Focus on IA64
 - Hard to admit mistake or that AMD is better
- 2004: Intel Announces EM64T extension to IA32
 - Extended Memory 64-bit Technology
 - Almost identical to x86-64!

Intel Ivy Bridge EX

- Current leader
(for the next few months)
 - 15-core desktop: 4.3 billion transistors
 - from 6 up to 15 cores/die
 - 2 threads per core
 - no GPU on chip
 - virtualization support



- 15 cores, 30 threads
- 3.5MB L2, 37.5 MB L3 cache
- 155W TPD

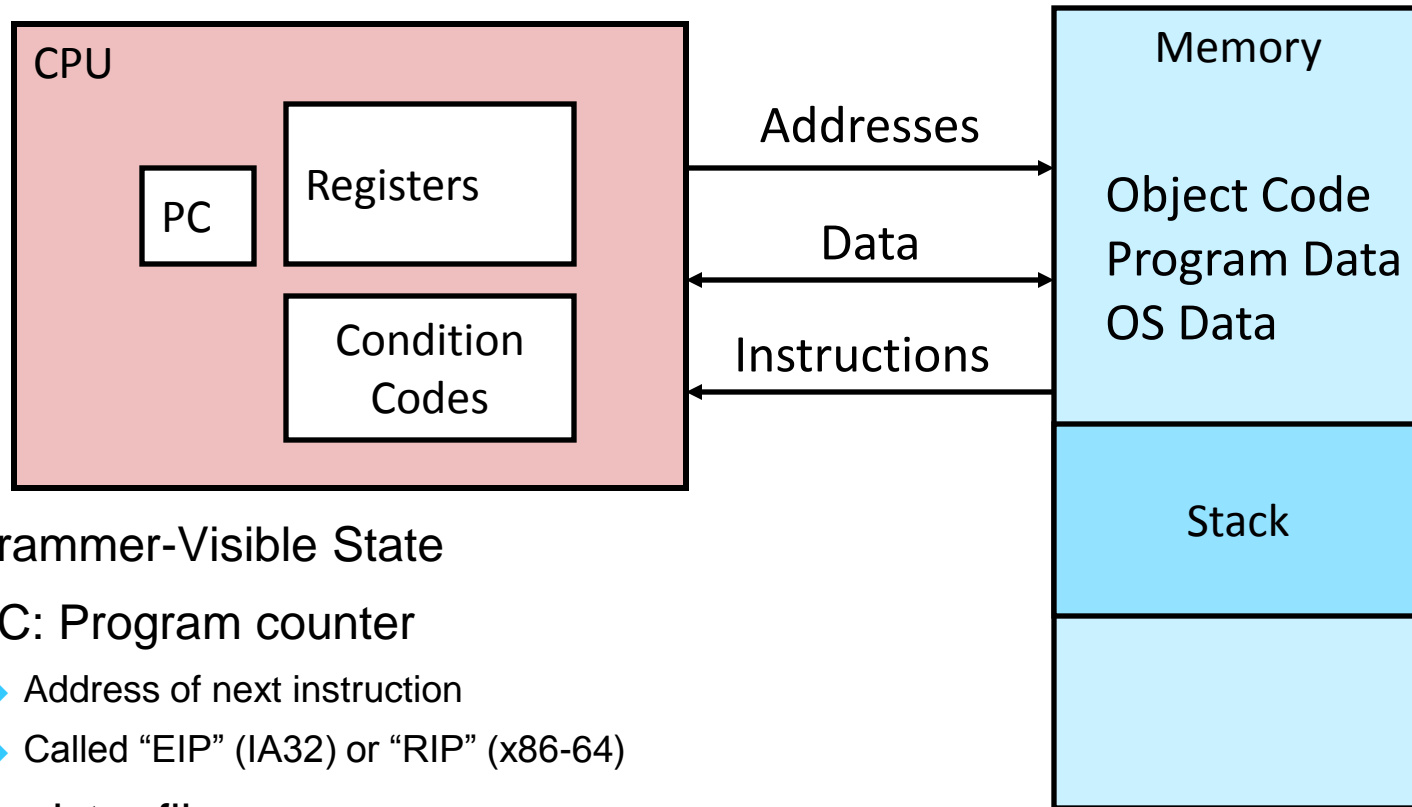
Our Coverage

- IA32
 - The traditional x86
- x86-64/EM64T
 - The emerging standard
- Presentation
 - ia32 in sections 3.1—3.7 of the textbook
 - x86-64 in section 3.13
 - lecture will cover both

x86 Machine-Level Programming Basics

- History of Intel processors and architectures
- **C, assembly, machine code**
- Assembly Basics: Registers, operands, move
- Introduction to x86-64

Assembly Programmer's View



■ Programmer-Visible State

- PC: Program counter
 - ▶ Address of next instruction
 - ▶ Called “EIP” (IA32) or “RIP” (x86-64)
- Register file
 - ▶ Heavily used program data
- Condition codes
 - ▶ Store status information about most recent arithmetic operation
 - ▶ Used for conditional branching


● Memory

- ▶ Byte addressable array
- ▶ Code, user data, (some) OS data
- ▶ Includes stack used to support procedures

x86 ISA

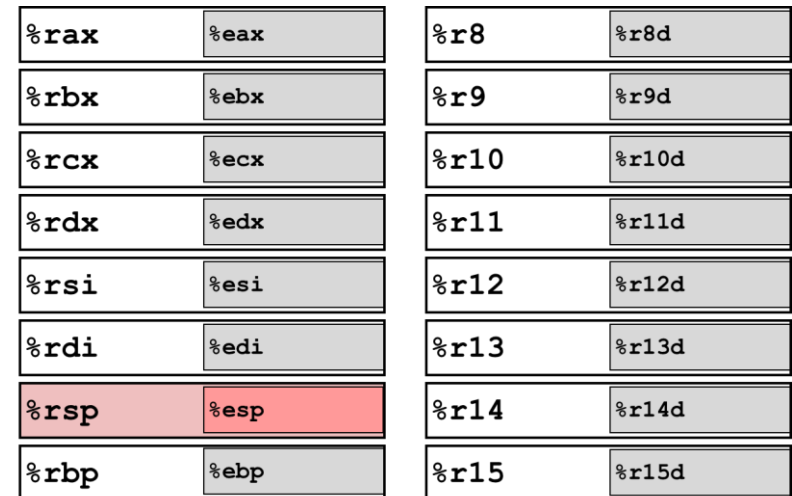
■ x86 (32 bit):

- register-memory architecture
 - ▶ 2 operands, max 1 memory
- registers:
 - ▶ 6(8) general purpose registers, PC (EIP), condition codes, processor status
- byte-addressed memory, little-endian
- 1, 2, 4, 8 byte data types



■ x86_64 (64 bit):

- same as above except
 - ▶ 16 general purpose registers
 - ▶ 16-byte data type



x86 ISA

■ instruction format

- operand size specifier
 - ▶ postfix to the operation
 - ▶ b (1 byte), w (2 bytes), l (4 bytes), q (8 bytes – x86_64 only)
- two-operand format: first source is implicit destination

$a = a + b$

in x86 assembly:

`addl b, a`

- ▶ careful with non-commutative operations:

`subl a, b` \neq `subl b, a`

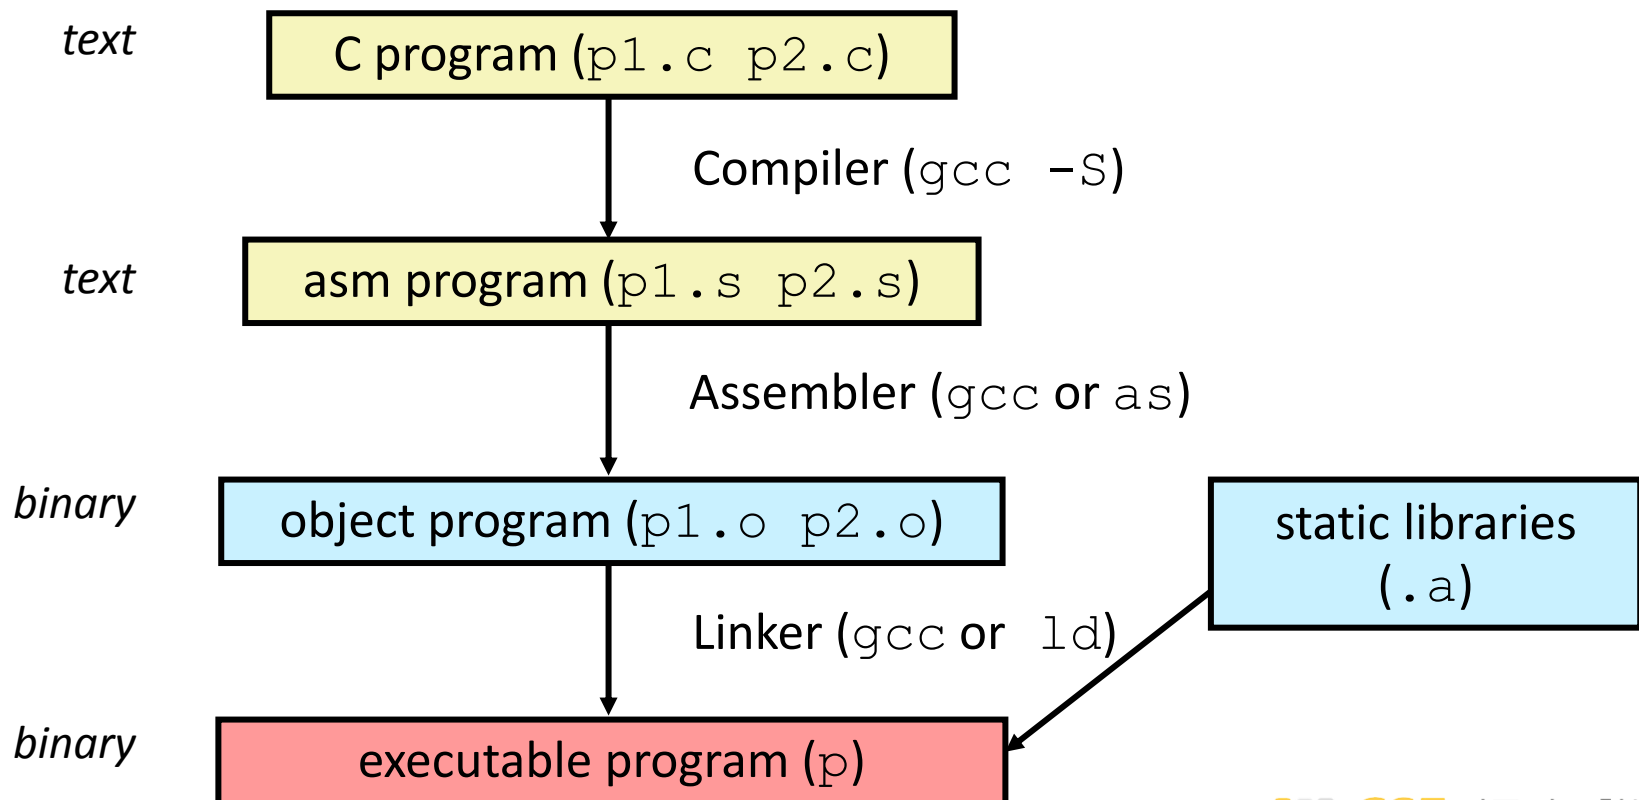
x86 ISA

■ Operand specifiers

Type	Form	Operand value	Name
Immediate	\$Imm	Imm	Immediate
Register	Ea	R[Ea]	Register
Memory	Imm	M[Imm]	Absolute
	(Eb)	M[R[Eb]]	Indirect
	Imm(Eb)	M[R[Eb] + Imm]	Base + displacement
	(Eb, Ei)	M[R[Eb] + R[Ei]]	Indexed
	Imm(Eb, Ei)	M[R[Eb] + R[Ei] + Imm]	Indexed
	(, Ei, s)	M[R[Ei]*s]	Scaled indexed
	Imm(, Ei, s)	M[R[Ei]*s + Imm]	Scaled indexed
	(Eb, Ei, s)	M[R[Eb] + R[Ei]*s]	Scaled indexed
	Imm(Eb, Ei, s)	M[R[Eb] + R[Ei]*s + Imm]	Scaled indexed

From C to Machine Code


- Code in files `p1.c` `p2.c`
- Compile with command: `gcc -o p1.c p2.c -o p`
 - ▶ Use optimizations (`-o`)
 - ▶ Put resulting binary in file `p`



Compiling To IA32 Assembly

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```



\$ gcc -m32 -O -S code.c

produces the file code.s:

Generated IA32 Assembly by gcc-4.4.6
on CentOS 6.3


```
sum:
    pushl    %ebp
    movl     %esp, %ebp
    movl     12(%ebp), %eax
    addl     8(%ebp), %eax
    popl     %ebp
    ret
```

→ compare output generated with `gcc -m32 -O0 -S code.c`

Compiling To IA32 Assembly – GCC Version

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```



\$ gcc -m32 -O -S code.c

produces the file code.s:

Generated IA32 Assembly by gcc-4.6.3
on Gentoo


```
sum:
.LFB0:
    .cfi_startproc
    movl    8(%esp), %eax
    addl    4(%esp), %eax
    ret
    .cfi_endproc
```

→ compare output generated with gcc -m32 -O0 -S code.c

Compiling To x86_64 Assembly

C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```



\$ gcc **-m64** -O -S code.c

produces the file code.s:

Generated x86_64 assembly
(identical for gcc-4.4.6/gcc-4.6.3)

```
sum:
.LFB0:
    .cfi_startproc
    leal    (%rsi,%rdi), %eax
    ret
    .cfi_endproc
```

→ compare output generated with gcc -m64 -O0 -S code.c

Assembly Characteristics: Data Types

- “Integer” data of 1, 2, or 4 bytes
 - Data values
 - Addresses (untyped pointers)
- Floating point data of 4, 8, or 10 bytes
- No aggregate types such as arrays or structures
 - Just contiguously allocated bytes in memory

Assembly Characteristics: Operations

■ Arithmetic/Logic Operations

Perform arithmetic/logic functions on registers or memory data

■ Memory Operations

Transfer data between memory and registers

- Load data from memory into register
- Store register data into memory

■ Control Transfer Operations

Transfer control

- Unconditional jumps to/from procedures
- Conditional branches

■ Special Operations

- processor control, loop operations, ...

Object Code

IA32 Code for sum

0x401040 <sum>:

0x55

0x89

0xe5

0x8b

0x45

0x0c

0x03

0x45

0x08

0x5d

0xc3

- Total of 11 bytes
- Each instruction 1, 2, or 3 bytes
- Starts at address 0x401040

■ Assembler

- Translates `.s` into `.o`
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

■ Linker

- Resolves references between files
- Combines with static run-time libraries
 - ▶ E.g., code for `malloc`, `printf`
- Some libraries are *dynamically linked*
 - ▶ Linking occurs when program begins execution

Machine Instruction Example

```
int t = x+y;
```

```
addl 8(%ebp), %eax
```

Similar to expression:

`x += y`

More precisely:

```
int eax;
```

```
int *ebp;
```

```
eax += ebp[2]
```

```
0x0804809a: 03 45 08
```

■ C Code

- Add two signed integers

■ Assembly

- Add two 4-byte integers
 - ▶ “long” words in GCC parlance
 - ▶ same instruction whether signed or unsigned
- Operands:
 - x:** Register **%eax**
 - y:** Memory **M[%ebp+8]**
 - t:** Register **%eax**
 - return function value in **%eax**

■ Object Code

- 3-byte instruction
- Stored at address **0x0804809a**

Disassembling Object Code

```
$ gcc -m32 -c -O code.c
$ objdump -d code.o
```

```
00000000 <sum>:
   0:  55          push    %ebp
   1:  89 e5       mov     %esp, %ebp
   3:  8b 45 0c    mov     0xc(%ebp), %eax
   6:  03 45 08    add     0x8(%ebp), %eax
   9:  5d          pop     %ebp
  a:  c3          ret
```

■ Disassembler

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces *approximate* rendition of assembly code
- Can be run on either `a.out` (complete executable) or `.o` file

Alternate Disassembly using gdb

■ Within gdb Debugger

```
$ gcc -m32 -O -o p code.c
```

```
$ gdb p
```

```
(gdb) disassemble sum
```

```
Dump of assembler code for function sum:
```

```
0x08048394 <+0>: push    %ebp
0x08048395 <+1>: mov     %esp,%ebp
0x08048397 <+3>: mov     0xc(%ebp),%eax
0x0804839a <+6>: add     0x8(%ebp),%eax
0x0804839d <+9>: pop     %ebp
0x0804839e <+10>: ret
```

```
End of assembler dump.
```

```
(gdb) x/11xb sum
```

```
0x8048394 <sum>:      0x55    0x89    0xe5    0x8b    0x45    ...
0x804839c <sum+8>:    0x08    0x5d    0xc3
```

```
(gdb)
```

What Can be Disassembled?

```
% objdump -d WINWORD.EXE

WINWORD.EXE:      file format pei-i386

No symbols in "WINWORD.EXE".
Disassembly of section .text:

30001000 <.text>:
30001000:  55                      push    %ebp
30001001:  8b ec                  mov     %esp,%ebp
30001003:  6a ff                  push    $0xffffffff
30001005:  68 90 10 00 30 push    $0x30001090
3000100a:  68 91 dc 4c 30 push    $0x304cdc91
```


- Anything that can be interpreted as executable code (not everything makes sense, of course)
- Disassembler examines bytes and reconstructs assembly source

x86 Machine-Level Programming Basics

- History of Intel processors and architectures
- C, assembly, machine code
- **Assembly Basics: Registers, operands, move**
- Introduction to x86-64

Integer Registers (IA32)

Origin
(mostly obsolete)



16-bit virtual registers
(backwards compatibility)

Moving Data: IA32

■ Moving Data

`movx Source, Dest`

with *x* in {b, w, l}

- `movl source, dest`
move 4-byte “long word”
- `movw source, dest`
move 2-byte “word”
- `movb source, dest`
move 1-byte “byte”

- lots of move instructions in typical code

%eax

%ecx

%edx

%ebx

%esi

%edi

%esp

%ebp

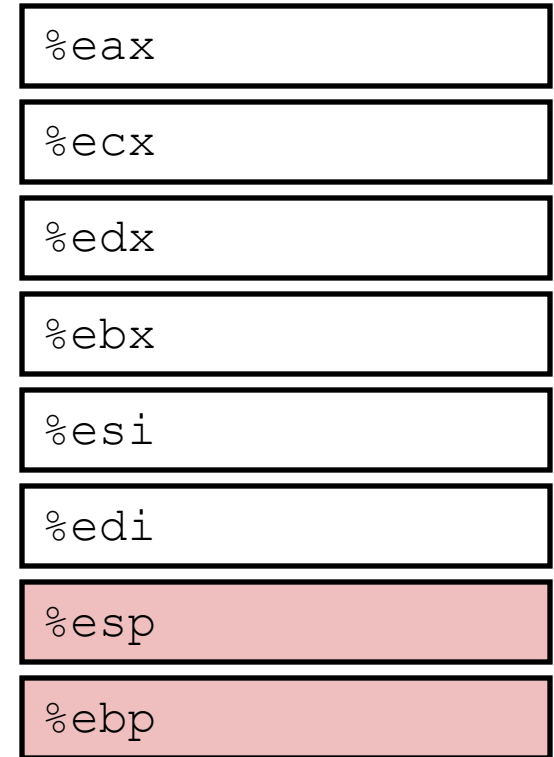
Moving Data: IA32

■ Moving Data

`movl Source, Dest`

■ Operand Types

- **Immediate:** Constant integer data
 - ▶ Example: `$0x400`, `$-533`
 - ▶ Like C constant, but prefixed with ``$'`
 - ▶ Encoded with 1, 2, or 4 bytes
- **Register:** One of 8 integer registers
 - ▶ Example: `%eax`, `%edx`
 - ▶ But `%esp` and `%ebp` reserved for special use
 - ▶ Others have special uses for particular instructions
- **Memory:** 4 consecutive bytes of memory at address given by register
 - ▶ Simplest example: `(%eax)`
 - ▶ Various other “address modes”



movl Operand Combinations

	Source	Dest	Src, Dest	C Analog
movl	Imm	Reg	movl \$0x4, %eax	temp = 0x4;
		Mem	movl \$-147, (%eax)	*p = -147;
	Reg	Reg	movl %eax, %edx	temp2 = temp1;
		Mem	movl %eax, (%edx)	*p = temp;
	Mem	Reg	movl (%eax), %edx	temp = *p;

Cannot do memory-memory transfer with a single instruction

Simple Memory Addressing Modes

■ Normal (R) Mem[Reg[R]]

- Register R specifies memory address

```
movl (%ecx) , %eax
```

■ Displacement D(R) Mem[Reg[R]+D]

- Register R specifies start of memory region
- Constant displacement D specifies offset

```
movl 8(%ebp) , %edx
```

Using Simple Addressing Modes

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
pushl %ebp
movl  %esp,%ebp
pushl %ebx
```

} Set Up

```
movl  8(%ebp), %edx
movl  12(%ebp), %ecx
movl  (%edx), %ebx
movl  (%ecx), %eax
movl  %eax, (%edx)
movl  %ebx, (%ecx)
```


} Body

```
popl  %ebx
popl  %ebp
ret
```

} Finish

Understanding Swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```



Register Value

Register	Value
%edx	xp
%ecx	yp
%ebx	t0
%eax	t1

```
movl 8(%ebp), %edx    # edx = xp
movl 12(%ebp), %ecx   # ecx = yp
movl (%edx), %ebx     # ebx = *xp (t0)
movl (%ecx), %eax     # eax = *yp (t1)
movl %eax, (%edx)     # *xp = t1
movl %ebx, (%ecx)     # *yp = t0
```

Understanding Swap

%eax	
%edx	
%ecx	
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

		Offset	Address
			0x124
			0x120
			0x11c
			0x118
			0x114
yp	12	0x120	0x110
xp	8	0x124	0x10c
	4	return adr	0x108
%ebp	0		0x104
	-4		0x100

```
movl    8(%ebp), %edx    # edx = xp
movl    12(%ebp), %ecx   # ecx = yp
movl    (%edx), %ebx     # ebx = *xp (t0)
movl    (%ecx), %eax     # eax = *yp (t1)
movl    %eax, (%edx)     # *xp = t1
movl    %ebx, (%ecx)     # *yp = t0
```


Understanding Swap

%eax	
%edx	0x124
%ecx	
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		0x124
		0x120
		0x11c
		0x118
		0x114
yp	12	0x110
xp	8	0x10c
	4	0x108
%ebp	0	0x104
	-4	0x100

```

movl 8(%ebp), %edx    # edx = xp
movl 12(%ebp), %ecx   # ecx = yp
movl (%edx), %ebx     # ebx = *xp (t0)
movl (%ecx), %eax     # eax = *yp (t1)
movl %eax, (%edx)     # *xp = t1
movl %ebx, (%ecx)     # *yp = t0
    
```

Understanding Swap

%eax	
%edx	0x124
%ecx	0x120
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		0x124
		0x120
		0x11c
		0x118
		0x114
yp	12	0x120
xp	8	0x124
	4	return adr
%ebp	0	
	-4	
		0x100

```

movl 8(%ebp), %edx    # edx = xp
movl 12(%ebp), %ecx  # ecx = yp
movl (%edx), %ebx     # ebx = *xp (t0)
movl (%ecx), %eax     # eax = *yp (t1)
movl %eax, (%edx)     # *xp = t1
movl %ebx, (%ecx)     # *yp = t0
    
```

Understanding Swap

%eax	
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		0x124
		0x120
		0x11c
		0x118
		0x114
yp	12	0x120
xp	8	0x124
		0x10c
		0x108
		0x104
		0x100

Offset

yp 12

xp 8

4

%ebp → 0

-4

```

movl 8(%ebp), %edx    # edx = xp
movl 12(%ebp), %ecx   # ecx = yp
movl (%edx), %ebx     # ebx = *xp (t0)
movl (%ecx), %eax     # eax = *yp (t1)
movl %eax, (%edx)     # *xp = t1
movl %ebx, (%ecx)     # *yp = t0
    
```

Understanding Swap

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		0x124
		0x120
		0x11c
		0x118
		0x114
yp	12	0x110
xp	8	0x10c
	4	return adr
%ebp	0	0x108
	-4	0x104
		0x100

```

movl 8(%ebp), %edx    # edx = xp
movl 12(%ebp), %ecx   # ecx = yp
movl (%edx), %ebx     # ebx = *xp (t0)
movl (%ecx), %eax     # eax = *yp (t1)
movl %eax, (%edx)     # *xp = t1
movl %ebx, (%ecx)     # *yp = t0
    
```

Understanding Swap

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		0x124
		0x120
		0x11c
		0x118
		0x114
yp	12	0x110
xp	8	0x10c
	4	return adr
%ebp	0	0x108
	-4	0x104
		0x100

```

movl 8(%ebp), %edx    # edx = xp
movl 12(%ebp), %ecx   # ecx = yp
movl (%edx), %ebx     # ebx = *xp (t0)
movl (%ecx), %eax     # eax = *yp (t1)
movl %eax, (%edx)     # *xp = t1
movl %ebx, (%ecx)     # *yp = t0
    
```

Understanding Swap

%eax	456
%edx	0x124
%ecx	0x120
%ebx	123
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		456 0x124
		123 0x120
		0x11c
		0x118
		0x114
yp	12	0x120 0x110
xp	8	0x124 0x10c
	4	return adr 0x108
%ebp	0	0x104
	-4	0x100

```

movl 8(%ebp), %edx    # edx = xp
movl 12(%ebp), %ecx   # ecx = yp
movl (%edx), %ebx     # ebx = *xp (t0)
movl (%ecx), %eax     # eax = *yp (t1)
movl %eax, (%edx)     # *xp = t1
movl %ebx, (%ecx)   # *yp = t0
    
```

Complete Memory Addressing Modes

■ General Form:

$D(Rb, Ri, S)$	$Mem[Reg[Rb] + S * Reg[Ri] + D]$
----------------	----------------------------------

- **D**: Constant “displacement” 1, 2, or 4 bytes
- **Rb**: Base register: Any of 8 integer registers
- **Ri**: Index register: Any, except for `%esp`
 - ▶ Unlikely you’d use `%ebp`, either
- **S**: Scale: 1, 2, 4, or 8 (*why these numbers?*)

■ Special Cases

(Rb, Ri)	$Mem[Reg[Rb] + Reg[Ri]]$
$D(Rb, Ri)$	$Mem[Reg[Rb] + Reg[Ri] + D]$
(Rb, Ri, S)	$Mem[Reg[Rb] + S * Reg[Ri]]$

x86 Machine-Level Programming Basics

- History of Intel processors and architectures
- C, assembly, machine code
- Assembly Basics: Registers, operands, move
- **Introduction to x86-64**

Data Representations: IA32 + x86-64

■ Sizes of C Objects (in Bytes)

C Data Type	Generic 32-bit	Intel IA32	x86-64
▶ unsigned	4	4	4
▶ int	4	4	4
▶ long int	4	4	8
▶ char	1	1	1
▶ short	2	2	2
▶ float	4	4	4
▶ double	8	8	8
▶ long double	8	10/12	16
▶ char *	4	4	8

– *Or any other pointer*

x86-64 Integer Registers

%rax	%eax
%rbx	%ebx
%rcx	%ecx
%rdx	%edx
%rsi	%esi
%rdi	%edi
%rsp	%esp
%rbp	%ebp

%r8	%r8d
%r9	%r9d
%r10	%r10d
%r11	%r11d
%r12	%r12d
%r13	%r13d
%r14	%r14d
%r15	%r15d

- Extend existing registers. Add 8 new ones.
- Make **%ebp/%rbp** general purpose

Instructions

- Long word l (4 Bytes) \leftrightarrow Quad word q (8 Bytes)
- New instructions:
 - `movl` \rightarrow `movq`
 - `addl` \rightarrow `addq`
 - `sall` \rightarrow `salq`
 - etc.
- 32-bit instructions that generate 32-bit results
 - Set higher order bits of destination register to 0
 - Example: `addl`

32-bit code for swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
pushl %ebp
movl  %esp, %ebp
pushl %ebx
```

} Set
Up

```
movl  8(%ebp), %edx
movl  12(%ebp), %ecx
movl  (%edx), %ebx
movl  (%ecx), %eax
movl  %eax, (%edx)
movl  %ebx, (%ecx)
```

} Body

```
popl  %ebx
popl  %ebp
ret
```

} Finish

64-bit code for swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

```
movl    (%rdi), %edx
movl    (%rsi), %eax
movl    %eax, (%rdi)
movl    %edx, (%rsi)
```

ret

} Set Up

} Body

} Finish

- Operands passed in registers (why is that useful?)
 - First (**x**p) in %**r**di, second (**y**p) in %**r**si
 - 64-bit pointers
- No stack operations required
- 32-bit data
 - Data held in registers %**e**ax and %**e**dx
 - **movl** operation

64-bit code for long int swap

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap_1:

```
movq    (%rdi), %rdx
movq    (%rsi), %rax
movq    %rax, (%rdi)
movq    %rdx, (%rsi)
```

ret

} Set Up

} Body

} Finish

■ 64-bit data

- Data held in registers **%rax** and **%rdx**
- **movq** operation
 - ▶ “q” stands for quad-word

Machine-Level Programming Basics: Summary

- History of Intel processors and architectures
 - Evolutionary design leads to many quirks and artifacts
- C, assembly, machine code
 - Compiler must transform statements, expressions, procedures into low-level instruction sequences
- Assembly Basics: Registers, operands, move
 - The x86 move instructions cover wide range of data movement forms
- Intro to x86-64
 - A major departure from the style of code seen in IA32