


# The HW/SW Interface

## The x86 ISA: Procedures



# Recap: Loops

## ■ Do-While


```
do  
    body;  
while (test);
```



```
loop:  
    body;  
    if (test) goto loop;
```

## ■ While

```
while (test)  
    body;
```



```
if (!test) goto done;  
do  
    body;  
while (test);  
done:
```



```
if (!test) goto done;  
loop:  
    body;  
    if (test) goto loop;  
done:
```

## ■ For

```
for(init; test; update)  
    body;
```



```
init;  
while (test) {  
    body;  
    update;  
}
```



```
init;  
if (!test) goto done;  
do {  
    body;  
    update;  
} while (test);  
done:
```



```
init;  
if (!test) goto done;  
loop:  
    body;  
    update;  
    if (test) goto loop;  
done:
```

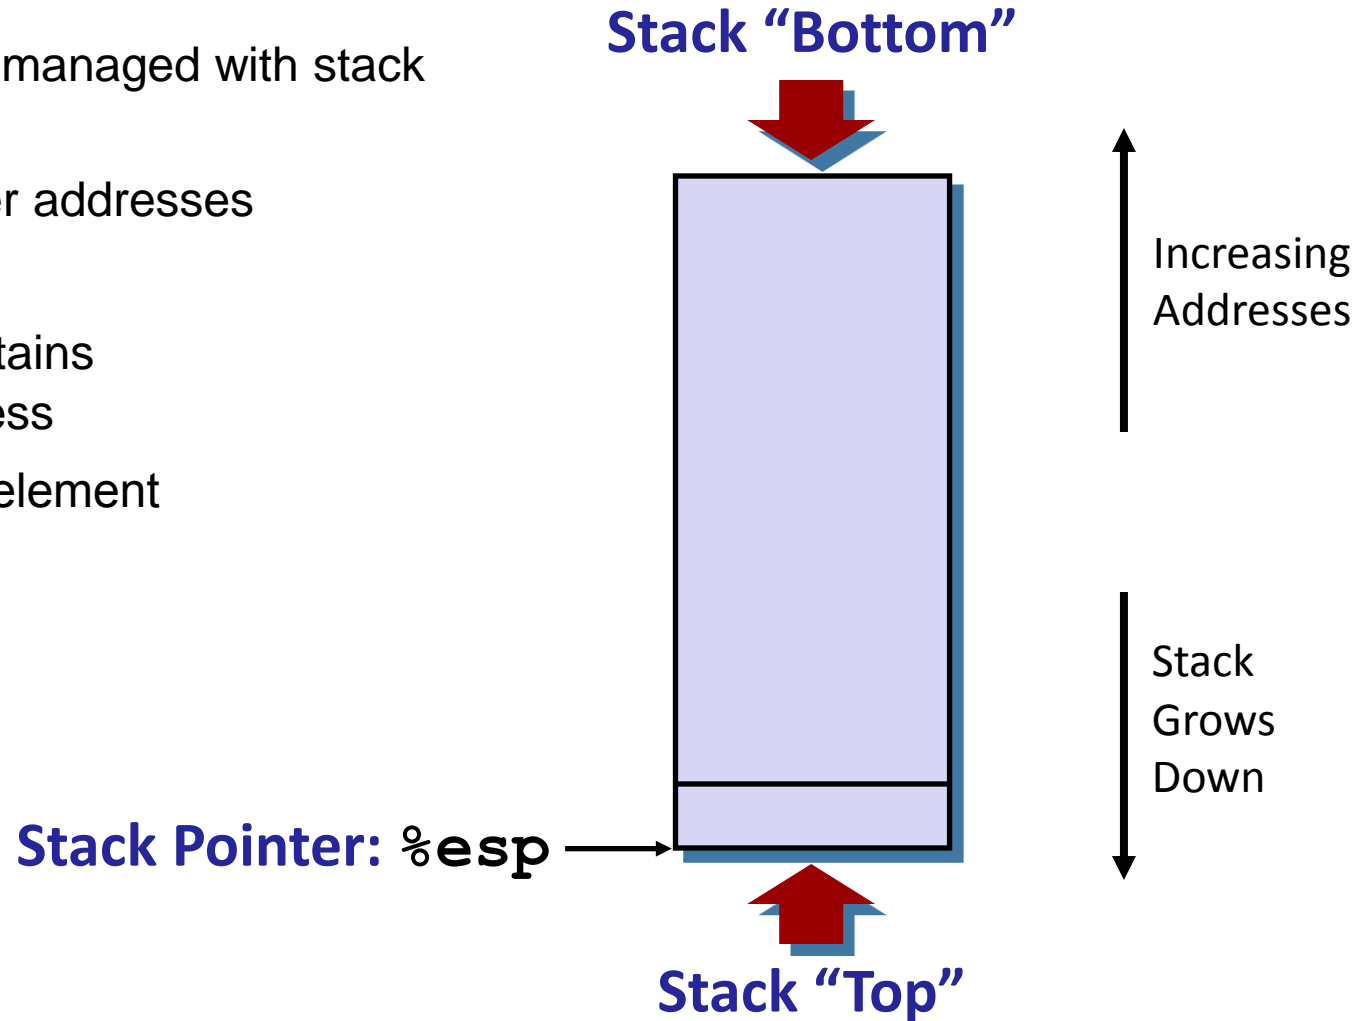
# Procedures

- **IA 32 Procedures**
  - **Stack Structure**
  - Calling Conventions
  - Illustrations of Recursion & Pointers
- x86-64 Procedures

Acknowledgement: slides based on the cs:app2e material


# IA32 Stack

- Region of memory managed with stack discipline
- Grows toward lower addresses
- Register `%esp` contains lowest stack address
  - address of “top” element



# IA32 Stack: Push


- `pushl Src`
  - Fetch operand at *Src*
  - Decrement `%esp` by 4
  - Write operand at address given by `%esp`



# IA32 Stack: Pop

## ■ `popl Dst`

- Fetch operand at `%esp`
- Increment `%esp` by 4
- Write operand to `Dst`



# Procedure Control Flow

- Use stack to support procedure call and return

- **Procedure call:** `call label`

- Push return address on stack
- Jump to *label*

- Return address:

- Address of the next instruction right after call
- Example from disassembly

804854e:	e8 3d 06 00 00	call	8048b90 <main>
8048553:	50	pushl	%eax

- Return address = `0x8048553`


- **Procedure return:** `ret`

- Pop address from stack
- Jump to address

# Procedure Call Example


804854e:	e8 3d 06 00 00	call	8048b90 <main>
8048553:	50	pushl	%eax

**call 8048b90**



`%esp` `0x108`

`%eip` `0x804854e`



`%esp` `0x104`


`%eip` `0x8048b90`

*`%eip`: program counter*



# Procedure Return Example

8048591:          c3                          ret



**%eip**: program counter

# Stack-Based Languages

- Languages that support recursion
  - e.g., C, Pascal, Java
  - Code must be “*reentrant*”
    - ▶ Multiple simultaneous instantiations of single procedure
  - Need some place to store state of each instantiation
    - ▶ Arguments
    - ▶ Local variables
    - ▶ Return pointer
- Stack discipline
  - State for given procedure needed for limited time
    - ▶ From when called to when return
  - Callee returns before caller does
- Stack allocated in **frames**
  - state for single procedure instantiation


# Call Chain Example

```
yoo (...)  
{  
  .  
  .  
  who ();  
  .  
  .  
}
```

```
who (...)  
{  
  . . .  
  amI ();  
  . . .  
  amI ();  
  . . .  
}
```

```
amI (...)  
{  
  .  
  .  
  amI ();  
  .  
  .  
}
```

## Example Call Chain



Procedure **amI ()** is recursive

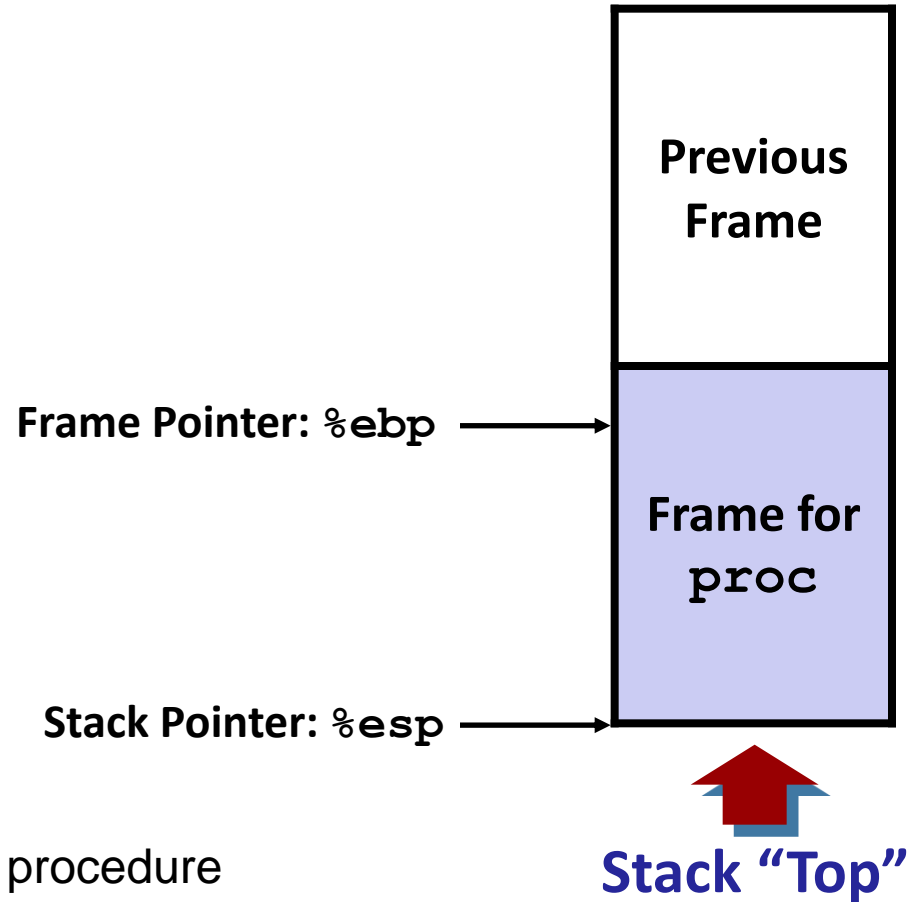
# Stack Frames

## ■ Contents

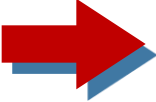
- Local variables
- Return information
- Temporary space

## ■ Management

- Space allocated when entering a procedure
  - ▶ “Set-up” code
- Deallocated when returning to the caller
  - ▶ “Cleanup” code




# Example






```
yoo (...)  
{  
  .  
  .  
  who ( ) ;  
  .  
  .  
}
```




```
yoo  
  ↓  
who  
  ↓  ↘  
amI  amI  
  ↓  
amI  
  ↓  
amI
```






# Example



# Example






# Example


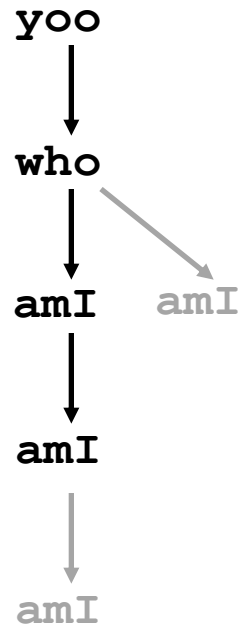








# Example






# Example






# Example




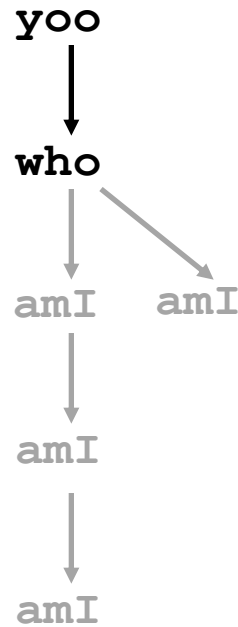

# Example






# Example



# Example




# Example



# IA32/Linux Stack Frame

- Current Stack Frame (Top to Bottom)
  - “Argument build:”  
Parameters for function about to call
  - Local variables  
If can't keep in registers
  - Saved register context
  - Old frame pointer
- Caller Stack Frame
  - Return address
    - ▶ Pushed by `call` instruction
  - Arguments for this call





# Revisiting swap

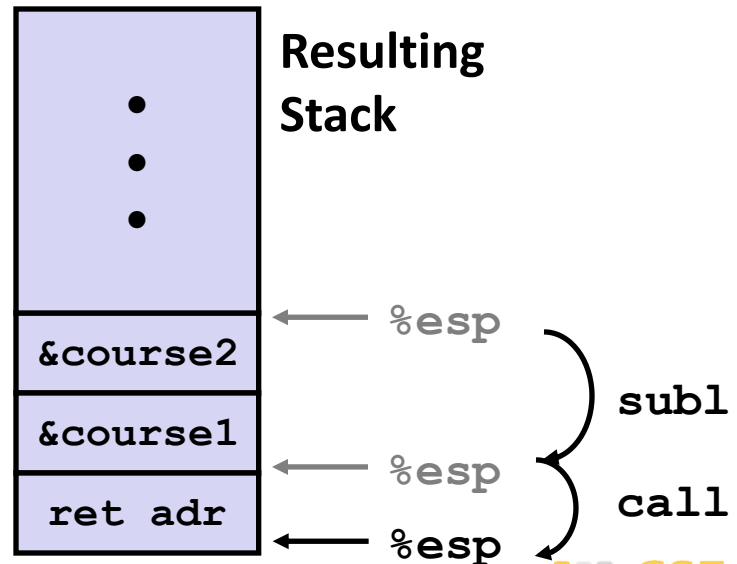
```
int course1 = 15213;
int course2 = 18243;

void call_swap() {
    swap(&course1, &course2);
}
```

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

## Calling swap from call\_swap

```
call_swap:
    . . .
    subl    $8, %esp
    movl    $course2, 4(%esp)
    movl    $course1, (%esp)
    call    swap
    . . .
```



# Revisiting swap


```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:


pushl %ebp	}	Set Up
movl %esp, %ebp		
pushl %ebx		
movl 8(%ebp), %edx	}	Body
movl 12(%ebp), %ecx		
movl (%edx), %ebx		
movl (%ecx), %eax		
movl %eax, (%edx)		
movl %ebx, (%ecx)		
popl %ebx	}	Finish
popl %ebp		
ret		

# swap Setup #1

Entering Stack



Resulting Stack



swap:


**pushl %ebp**

`movl %esp, %ebp`


`pushl %ebx`

# swap Setup #2

Entering Stack



Resulting Stack



`swap:`


`pushl %ebp`

`movl %esp, %ebp`

`pushl %ebx`

# swap Setup #3

Entering Stack




`swap:`

`pushl %ebp`

`movl %esp, %ebp`


`pushl %ebx`

Resulting Stack




# swap Setup #3

Entering Stack



Resulting Stack




swap:

```
...  
    movl 8(%ebp), %edx    # get xp  
    movl 12(%ebp), %ecx   # get yp  
    ...
```


# swap Cleanup

Stack before Cleanup



```
popl    %ebx
popl    %ebp
```

Resulting Stack



## ■ Observations

- Saved and restored register `%ebx`, `%ebp`
- Not so for `%eax`, `%ecx`, `%edx`
- Modified `%esp`, but value after the call is the same as before the call

# Disassembled swap

08048384 <swap>:

8048384:	55	push	%ebp
8048385:	89 e5	mov	%esp, %ebp
8048387:	53	push	%ebx
8048388:	8b 55 08	mov	0x8(%ebp), %edx
804838b:	8b 4d 0c	mov	0xc(%ebp), %ecx
804838e:	8b 1a	mov	(%edx), %ebx
8048390:	8b 01	mov	(%ecx), %eax
8048392:	89 02	mov	%eax, (%edx)
8048394:	89 19	mov	%ebx, (%ecx)
8048396:	5b	pop	%ebx
8048397:	5d	pop	%ebp
8048398:	c3	ret	

leave :=  
movl %ebp, %esp  
popl %ebp

## Calling Code

80483b4:	movl	\$0x8049658, 0x4(%esp)	# Copy &course2
80483bc:	movl	\$0x8049654, (%esp)	# Copy &course1
80483c3:	call	8048384 <swap>	# Call swap
80483c8:	leave		# Prepare to return
80483c9:	ret		# Return



# Procedures

## ■ IA 32 Procedures

- Stack Structure
- **Calling Conventions**
- Illustrations of Recursion & Pointers

## ■ x86-64 Procedures

# Register Saving Conventions

- When procedure `yoo` calls `who`:
  - `yoo` is the *caller*
  - `who` is the *callee*
- Can registers be used for temporary storage?

```
yoo:
    . . .
    movl $15213, %edx
    call who
    addl %edx, %eax
    . . .
    ret
```

```
who:
    . . .
    movl 8(%ebp), %edx
    addl $18243, %edx
    . . .
    ret
```

- Contents of register `%edx` overwritten by `who`
- This could be trouble → something should be done!
  - ▶ Need some coordination

# Register Saving Conventions

- When procedure **yoo** calls **who**:
  - **yoo** is the **caller**
  - **who** is the **callee**
- Can registers be used for temporary storage?
- **Calling Convention**
  - **“Caller Save”**
    - ▶ registers that the callee can overwrite  
(the caller assumes their value is not preserved across procedure calls)
    - ▶ Caller saves temporary values in its frame before the call
  - **“Callee Save”**
    - ▶ registers that the callee must preserve before overwriting with a new value  
(the caller can reuse the value across procedure calls)
    - ▶ Callee saves temporary values in its frame before using

# IA32/Linux+Windows Register Usage

## ■ **%eax, %ecx, %edx**

- caller saved prior to call (if values are used later)
- %eax used to return integer value

## ■ **%ebx, %esi, %edi**

- callee saved (if used)

## ■ **%esp, %ebp**

- used to manage the stack frames
- must restore original values upon exit from procedure (= special form of callee saved)

%eax	Caller saved / Return value
%ecx	Caller saved
%edx	Caller saved
%ebx	Callee saved
%esi	Callee saved
%edi	Callee saved
%esp	Stack pointer
%ebp	Frame pointer

# Procedures

## ■ IA 32 Procedures

- Stack Structure
- Calling Conventions
- **Illustrations of Recursion & Pointers**

## ■ x86-64 Procedures

# Recursive Function

```
/* Recursive popcount */
int pcount_r(unsigned x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

## ■ Registers

- **%eax**, **%edx** used without first saving
- **%ebx** used, but saved at beginning & restored at end

```
pcount_r:
    pushl %ebp
    movl %esp, %ebp
    pushl %ebx
    subl $4, %esp
    movl 8(%ebp), %ebx
    movl $0, %eax
    testl %ebx, %ebx
    je .L3
    movl %ebx, %eax
    shrl %eax
    movl %eax, (%esp)
    call pcount_r
    movl %ebx, %edx
    andl $1, %edx
    leal (%edx,%eax), %eax
.L3:
    addl $4, %esp
    popl %ebx
    popl %ebp
    ret
```


# Recursive Call #1

```
/* Recursive popcount */
int pcount_r(unsigned x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

## ■ Actions

- Save old value of `%ebx` on stack
- Allocate space for argument to (recursive) call
- Store `x` in `%ebx`

```
pcount_r:
    pushl %ebp
    movl  %esp, %ebp
    pushl %ebx
    subl  $4, %esp
    movl  8(%ebp), %ebx
    . . .
```



# Recursive Call #2

```
/* Recursive popcount */
int pcount_r(unsigned x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

```
    • • •
movl  $0, %eax
testl %ebx, %ebx
je   .L3
    • • •
.L3:
    • • •
ret
```

## ■ Actions

- If  $x == 0$ , return
  - ▶ with `%eax` set to 0

`%ebx`

`x`



# Recursive Call #3

```
/* Recursive popcount */
int pcount_r(unsigned x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```


## ■ Actions

- Store  $x \gg 1$  on stack
- Make recursive call

## ■ Effect

- `%eax` set to function result
- `%ebx` still has value of  $x$

```
• • •
movl  %ebx, %eax
shrl  %eax
movl  %eax, (%esp)
call  pcount_r
• • •
```



# Recursive Call #4

```
/* Recursive popcount */
int pcount_r(unsigned x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

```
• • •
movl  %ebx, %edx
andl  $1, %edx
leal  (%edx,%eax), %eax
• • •
```

- Assume
  - **%eax** holds value from recursive call
  - **%ebx** holds x
- Actions
  - Compute (x & 1) + computed value
- Effect
  - **%eax** set to function result

**%ebx**

**x**

# Recursive Call #5


```
/* Recursive popcount */
int pcount_r(unsigned x) {
    if (x == 0)
        return 0;
    else return
        (x & 1) + pcount_r(x >> 1);
}
```

## ■ Actions

- Deallocate space for argument
- Restore values of `%ebx` and `%ebp`
- `ret` will pop the return address into `%eip`

L3:

```
addl    $4, %esp
popl    %ebx
popl    %ebp
ret
```



`%ebx` old %ebx

# Observations About Recursion

- Handled Without Special Consideration
  - Stack frames mean that each function call has private storage
    - ▶ Saved registers & local variables
    - ▶ Saved return pointer
  - Register saving conventions prevent one function call from corrupting another's data
  - Stack discipline follows call / return pattern
    - ▶ If P calls Q, then Q returns before P
    - ▶ Last-In, First-Out
- Also works for mutual recursion
  - P calls Q; Q calls P

# Pointer Code

- add3 creates pointer and passes it to `incrk`

## Generating a Pointer

```
/* Compute x + 3 */
int add3(int x) {
    int localx = x;
    incrk(&localx, 3);
    return localx;
}
```

## Referencing a Pointer

```
/* Increment value by k */
void incrk(int *ip, int k) {
    *ip += k;
}
```


# Creating and Initializing Local Variables

```
int add3(int x) {  
    int localx = x;  
    incrk(&localx, 3);  
    return localx;  
}
```

- variable localx must be stored on the stack
  - the compiler needs to create a pointer to it
- compute pointer as  $-4(\%ebp)$

## First part of add3

```
add3:  
    pushl %ebp  
    movl %esp, %ebp  
    subl $24, %esp      # Alloc. 24 bytes  
    movl 8(%ebp), %eax  
    movl %eax, -4(%ebp) # Set localx to x
```




# Creating Pointer as Argument

```
int add3(int x) {  
    int localx = x;  
    incr(&localx, 3);  
    return localx;  
}
```

- Use `leal` to compute the address of `localx`

## Middle part of `add3`

```
movl $3, 4(%esp)    # 2nd arg = 3  
leal -4(%ebp), %eax # &localx  
movl %eax, (%esp)   # 1st arg = &localx  
call incr
```




# Retrieving local variable

```
int add3(int x) {  
    int localx = x;  
    incrk(&localx, 3);  
    return localx;  
}
```

- Retrieve `localx` from stack as return value

## Final part of add3

```
movl -4(%ebp), %eax # Return val= localx  
leave  
ret
```





# IA 32 Procedure Summary

## ■ Important Points

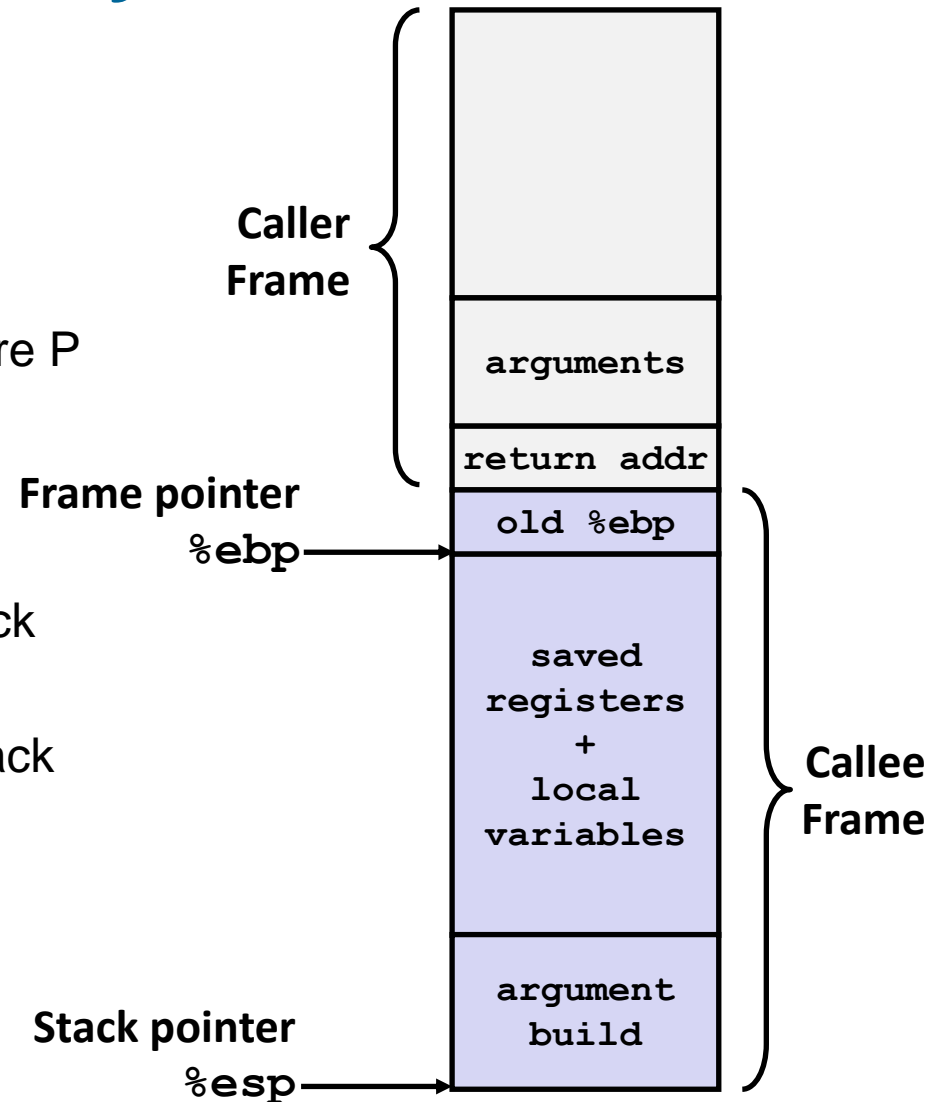
- Stack is the right data structure for procedure call / return
  - ▶ If P calls Q, then Q returns before P

## ■ Recursion (& mutual recursion) handled by normal calling conventions

- Can safely store values in local stack frame and in callee-saved registers
- Put function arguments at top of stack
- Result return in **%eax**

## ■ Pointers are addresses of values

- On stack or global



# Procedures

- IA 32 Procedures
  - Stack Structure
  - Calling Conventions
  - Illustrations of Recursion & Pointers
- **x86-64 Procedures**

# x86-64 Integer Registers: Usage Conventions

<code>%rax</code>	Caller saved / Return value	<code>%r8</code>	Caller saved / Argument #5
<code>%rbx</code>	Callee saved	<code>%r9</code>	Caller saved / Argument #6
<code>%rcx</code>	Caller saved / Argument #4	<code>%r10</code>	Caller saved / Caller saved
<code>%rdx</code>	Caller saved / Argument #3	<code>%r11</code>	Caller Saved
<code>%rsi</code>	Caller saved / Argument #2	<code>%r12</code>	Callee saved
<code>%rdi</code>	Caller saved / Argument #1	<code>%r13</code>	Callee saved
<code>%rsp</code>	Stack pointer	<code>%r14</code>	Callee saved
<code>%rbp</code>	Callee saved	<code>%r15</code>	Callee saved

# x86-64 Registers


- Arguments passed to functions via registers
  - If more than 6 integral parameters, then pass rest on stack
  - These registers can be used as caller-saved as well
- All references to stack frame via stack pointer
  - Eliminates need to update %ebp/%rbp
- Other Registers
  - 6 callee saved
  - 2 caller saved
  - 1 return value (also usable as caller saved)
  - 1 special (stack pointer)

# x86-64 Long Swap

```
void swap_l(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
swap:
    movq    (%rdi), %rdx
    movq    (%rsi), %rax
    movq    %rax, (%rdi)
    movq    %rdx, (%rsi)
    ret
```

- Operands passed in registers
  - First (xp) in %rdi, second (yp) in %rsi
  - 64-bit pointers
- No stack operations required (except ret)
- Avoiding stack
  - Can hold all local information in registers




# x86-64 Locals in the Red Zone

```
/* Swap, using local array */
void swap_a(long *xp, long *yp)
{
    volatile long loc[2];
    loc[0] = *xp;
    loc[1] = *yp;
    *xp = loc[1];
    *yp = loc[0];
}
```

```
swap_a:
    movq    (%rdi), %rax
    movq    %rax, -24(%rsp)
    movq    (%rsi), %rax
    movq    %rax, -16(%rsp)
    movq    -16(%rsp), %rax
    movq    %rax, (%rdi)
    movq    -24(%rsp), %rax
    movq    %rax, (%rsi)
    ret
```

## ■ Avoiding Stack Pointer Change

- Can hold all information within small window beyond stack pointer



# x86-64 NonLeaf without a Stack Frame

```
/* Swap a[i] & a[i+1] */  
void swap_ele(long a[], int i)  
{  
    swap(&a[i], &a[i+1]);  
}
```

- No values held while swap being invoked
- No callee save registers needed
- rep instruction inserted as no-op (recommendation from AMD)

```
swap_ele:  
    movslq %esi,%rsi          # Sign extend i  
    leaq    8(%rdi,%rsi,8), %rax # &a[i+1]  
    leaq    (%rdi,%rsi,8), %rdi  # &a[i] (1st arg)  
    movq    %rax, %rsi          # (2nd arg)  
    call    swap  
    rep  
    ret                       # No-op
```

# x86-64 Stack Frame Example

```
long sum = 0;
/* Swap a[i] & a[i+1] */
void swap_ele_su
    (long a[], int i)
{
    swap(&a[i], &a[i+1]);
    sum += (a[i]*a[i+1]);
}
```

- Keeps values of `&a[i]` and `&a[i+1]` in callee save registers
- Must set up stack frame to save these registers

```
swap_ele_su:
    movq    %rbx, -16(%rsp)
    movq    %rbp, -8(%rsp)
    subq    $16, %rsp
    movslq   %esi, %rax
    leaq    8(%rdi, %rax, 8), %rbx
    leaq    (%rdi, %rax, 8), %rbp
    movq    %rbx, %rsi
    movq    %rbp, %rdi
    call    swap
    movq    (%rbx), %rax
    imulq   (%rbp), %rax
    addq    %rax, sum(%rip)
    movq    (%rsp), %rbx
    movq    8(%rsp), %rbp
    addq    $16, %rsp
    ret
```



# Understanding x86-64 Stack Frame


swap\_ele\_su:

<code>movq</code>	<code>%rbx, -16(%rsp)</code>	# Save %rbx
<code>movq</code>	<code>%rbp, -8(%rsp)</code>	# Save %rbp
<code>subq</code>	<code>\$16, %rsp</code>	# Allocate stack frame
<code>movslq</code>	<code>%esi, %rax</code>	# Extend i
<code>leaq</code>	<code>8(%rdi, %rax, 8), %rbx</code>	# &a[i+1] (callee save)
<code>leaq</code>	<code>(%rdi, %rax, 8), %rbp</code>	# &a[i] (callee save)
<code>movq</code>	<code>%rbx, %rsi</code>	# 2 <sup>nd</sup> argument
<code>movq</code>	<code>%rbp, %rdi</code>	# 1 <sup>st</sup> argument
<code>call</code>	<code>swap</code>	
<code>movq</code>	<code>(%rbx), %rax</code>	# Get a[i+1]
<code>imulq</code>	<code>(%rbp), %rax</code>	# Multiply by a[i]
<code>addq</code>	<code>%rax, sum(%rip)</code>	# Add to sum
<code>movq</code>	<code>(%rsp), %rbx</code>	# Restore %rbx
<code>movq</code>	<code>8(%rsp), %rbp</code>	# Restore %rbp
<code>addq</code>	<code>\$16, %rsp</code>	# Deallocate frame
<code>ret</code>		

# Understanding x86-64 Stack Frame

```
movq    %rbx, -16(%rsp)
movq    %rbp, -8(%rsp)
```


```
# Save %rbx
# Save %rbp
```



```
subq    $16, %rsp
```

```
# Allocate stack frame
```

• • •



```
movq    (%rsp), %rbx
movq    8(%rsp), %rbp
```

```
# Restore %rbx
# Restore %rbp
```

```
addq    $16, %rsp
```

```
# Deallocate frame
```

# Interesting Features of Stack Frames

- Allocate entire frame at once
  - All stack accesses can be relative to %rsp
  - Do by decrementing stack pointer
  - Can delay allocation, since safe to temporarily use red zone
- Simple deallocation
  - Increment stack pointer
  - No base/frame pointer needed

# x86-64 Procedure Summary

- Heavy use of registers
  - Parameter passing
  - More temporaries since more registers
- Minimal use of stack
  - Sometimes none
  - Allocate/deallocate entire block
- Many tricky optimizations
  - What kind of stack frame to use
  - Various allocation techniques