


# The HW/SW Interface

## The x86 ISA: Data Structures



# Recap: Procedures


- Procedure Call and Return  
`call <label>/ * <dest> ... ret`

- Stack Frame

- old frame pointer, saved registers, local variables, temporaries
- during call: arguments, return address
- constructed/destroyed by callee in prologue/epilogue

- Calling Convention

- IA32:
  - ▶ registers: caller saved: eax, ecx, edx; callee saved: ebx, esi, edi, esp, ebp
  - ▶ arguments: on stack above return address, at  $\%ebp + 4 \cdot i$  for the  $i^{\text{th}}$  argument
  - ▶ return value: eax
- x86-64:
  - ▶ registers: caller saved: rax, r10, r11; callee saved: rbx, r12, r13, r14, r15, rsp, rbp
  - ▶ arguments: first 6: rdi, rsi, rdx, rcx, r8, r9, then on stack
  - ▶ return value: rax



# Data Structures

## ■ Arrays

- **One-dimensional**
- Multi-dimensional (nested)
- Multi-level

## ■ Structures

- Allocation
- Access
- Alignment

## ■ Unions

Acknowledgement: slides based on the cs:app2e material

# Basic Data Types

## ■ Integral

- Stored & operated on in general (integer) registers
- Signed vs. unsigned depends on instructions used

▶ Intel	ASM	Bytes	C
▶ byte	b	1	[unsigned] char
▶ word	w	2	[unsigned] short
▶ double word	l	4	[unsigned] int
▶ quad word	q	8	[unsigned] long int (x86-64)

## ■ Floating Point

- Stored & operated on in floating point registers


▶ Intel	ASM	Bytes	C
▶ Single	s	4	float
▶ Double	l	8	double
▶ Extended	t	10/12/16	long double

# Array Allocation


## ■ Basic Principle

- $T \ A[L];$
- Array of data type  $T$  and length  $L$
- Contiguously allocated region of  $L * \text{sizeof}(T)$  bytes


`char string[12];`




`int val[5];`




`double a[3];`



`char *p[3];`



IA32




x86-64

# Array Access

## ■ Basic Principle

- $T \ A[L];$
- Array of data type  $T$  and length  $L$
- Identifier  $A$  can be used as a pointer to array element 0: Type  $T^*$




Reference	Type	Value
<code>val[4]</code>		
<code>val</code>		
<code>val+1</code>		
<code>&amp;val[2]</code>		
<code>val[5]</code>		
<code>*(val+1)</code>		
<code>&amp;val[i] - val</code>		

# Array Example

```
#define ZLEN 5
typedef int zip_dig[ZLEN];


zip_dig cmu = { 1, 5, 2, 1, 3 };
zip_dig mit = { 0, 2, 1, 3, 9 };
zip_dig ucb = { 9, 4, 7, 2, 0 };
```



- Declaration “zip\_dig cmu” equivalent to “int cmu[5]”
- Example arrays were allocated in successive 20 byte blocks
  - Not guaranteed to happen in general

# Array Accessing Example

zip\_dig cmu;



```
int get_digit
(zip_dig z, int dig)
{
    return z[dig];
}
```

IA32

```
# %edx = z
# %eax = dig
movl (%edx,%eax,4),%eax # z[dig]
```

- Register %edx contains starting address of array
- Register %eax contains array index
- Desired digit at  $4 * \%eax + \%edx$
- Use memory reference  $(\%edx, \%eax, 4)$



# Array Loop Example (IA32)

```
void zincr(zip_dig z) {  
    int i;  
    for (i = 0; i < ZLEN; i++)  
        z[i]++;  
}
```

```
# edx = z  
movl    $0, %eax           # %eax = i  
.L4:      # loop:  
addl    $1, (%edx,%eax,4)   # z[i]++  
addl    $1, %eax           # i++  
cmpl    $5, %eax           # i:5  
jne     .L4                # if !=, goto loop
```

# Pointer Loop Example (IA32)

```
void zincr_p(zip_dig z) {  
    int *zend = z+ZLEN;  
    do {  
        (*z)++;  
        z++;  
    } while (z != zend);  
}
```



```
void zincr_v(zip_dig z) {  
    void *vz = z;  
    int i = 0;  
    do {  
        (*(int *) (vz+i))++;  
        i += ISIZE;  
    } while (i != ISIZE*ZLEN);  
}
```

# edx = z = vz	
movl \$0, %eax	# i = 0
.L8:	# loop:
addl \$1, (%edx,%eax)	# Increment vz+i
addl \$4, %eax	# i += 4
cmpl \$20, %eax	# Compare i:20
jne .L8	# if !=, goto loop

# Data Structures

## ■ Arrays

- One-dimensional
- **Multi-dimensional (nested)**
- Multi-level


## ■ Structures

- Allocation
- Access
- Alignment

## ■ Unions

# Nested Array Example

```
#define PCOUNT 4
zip_dig pgh[PCOUNT] =
    {{1, 5, 2, 0, 6},
     {1, 5, 2, 1, 3 },
     {1, 5, 2, 1, 7 },
     {1, 5, 2, 2, 1 }};
```



- “zip\_dig pgh[4]” equivalent to “int pgh[4][5]”
  - Variable pgh: array of 4 elements, allocated contiguously
  - Each element is an array of 5 int’s, allocated contiguously
- “Row-Major” ordering of all elements guaranteed

# Multidimensional (Nested) Arrays

## ■ Declaration


- `T A[R][C];`
- 2D array of data type `T`
- `R` rows, `C` columns
- Type `T` element requires `K` bytes

## ■ Array Size

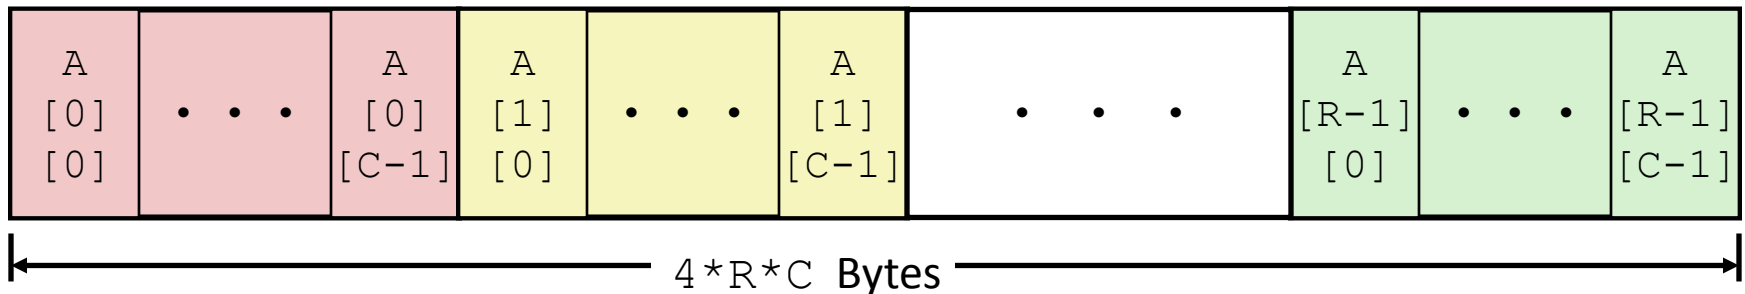
- $R * C * K$  bytes

## ■ Arrangement

- Row-Major Ordering



```
int A[R][C];
```




# Nested Array Row Access

## ■ Row Vectors

- $A[i]$  is array of  $C$  elements
- Each element of type  $T$  requires  $K$  bytes
- Starting address  $A + i * (C * K)$

```
int A[R][C];
```



# Nested Array Row Access Code

```
int *get_pgh_zip(int index)
{
    return pgh[index];
}
```

```
#define PCOUNT 4
zip_dig pgh[PCOUNT] =
    {{1, 5, 2, 0, 6},
     {1, 5, 2, 1, 3 },
     {1, 5, 2, 1, 7 },
     {1, 5, 2, 2, 1 }};
```

```
# %eax = index
leal (%eax,%eax,4),%eax # 5 * index
leal pgh(,%eax,4),%eax  # pgh + (20 * index)
```

## ■ Row Vector

- `pgh[index]` is array of 5 int's
- Starting address `pgh+20*index`

## ■ IA32 Code


- Computes and returns address
- Compute as `pgh + 4*(index+4*index)`

# Nested Array Row Access

## ■ Array Elements

- $A[i][j]$  is element of type  $T$ , which requires  $K$  bytes
- Address  $A + i * (C * K) + j * K = A + (i * C + j) * K$

```
int A[R][C];
```





# Nested Array Element Access Code

```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```

```
movl    8(%ebp), %eax          # index
leal    (%eax,%eax,4), %eax    # 5*index
addl    12(%ebp), %eax         # 5*index+dig
movl    pgh(,%eax,4), %eax     # offset 4*(5*index+dig)
```

## ■ Array Elements

- `pgh[index][dig]` is int
- Address:  $\text{pgh} + 20 \cdot \text{index} + 4 \cdot \text{dig}$ 
  - ▶  $= \text{pgh} + 4 \cdot (5 \cdot \text{index} + \text{dig})$

## ■ IA32 Code

- Computes address  $\text{pgh} + 4 \cdot ((\text{index} + 4 \cdot \text{index}) + \text{dig})$

# Data Structures

## ■ Arrays

- One-dimensional
- Multi-dimensional (nested)
- **Multi-level**

## ■ Structures

- Allocation
- Access
- Alignment


## ■ Unions

# Multi-Level Array Example

```
zip_dig cmu = { 1, 5, 2, 1, 3 };  
zip_dig mit = { 0, 2, 1, 3, 9 };  
zip_dig ucb = { 9, 4, 7, 2, 0 };
```

```
#define UCOUNT 3  
int *univ[UCOUNT] = {mit, cmu, ucb};
```

- Variable univ denotes array of 3 elements
- Each element is a pointer
  - 4 bytes
- Each pointer points to array of int's



# Element Access in Multi-Level Array

```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```

```
movl    8(%ebp), %eax          # index
movl    univ(,%eax,4), %edx     # p = univ[index]
movl    12(%ebp), %eax         # dig
movl    (%edx,%eax,4), %eax     # p[dig]
```


## ■ Computation (IA32)

- Element access  $\text{Mem}[\text{Mem}[\text{univ} + 4 \cdot \text{index}] + 4 \cdot \text{dig}]$
- Must do two memory reads
  - ▶ First get pointer to row array
  - ▶ Then access element within array

# Array Element Accesses


## Nested array

```
int get_pgh_digit
(int index, int dig)
{
    return pgh[index][dig];
}
```



## Multi-level array

```
int get_univ_digit
(int index, int dig)
{
    return univ[index][dig];
}
```



Accesses look similar in C, but addresses very different:

$\text{Mem}[\text{pgh} + 20 * \text{index} + 4 * \text{dig}]$

$\text{Mem}[\text{Mem}[\text{univ} + 4 * \text{index}] + 4 * \text{dig}]$

# N X N Matrix Code

## ■ Fixed dimensions

- Know value of N at compile time
- 1<sup>st</sup> dimension can be left unspecified

## ■ Variable dimensions, explicit indexing

- Traditional way to implement dynamic arrays

## ■ Variable dimensions, implicit indexing

- defined in C99 standard

```
#define N 16
typedef int fix_matrix[N][N];
/* Get element a[i][j] */
int fix_ele
    (fix_matrix a, int i, int j)
{
    return a[i][j];
}
```

```
#define IDX(n, i, j) ((i)*(n)+(j))
/* Get element a[i][j] */
int vec_ele
    (int n, int *a, int i, int j)
{
    return a[IDX(n,i,j)];
}
```

```
/* Get element a[i][j] */
int var_ele
    (int n, int a[n][n], int i, int j)
{
    return a[i][j];
}
```

# 16 X 16 Matrix Access

## ■ Array Elements

- Address  $A + i * (C * K) + j * K$
- $C = 16, K = 4$

```
/* Get element a[i][j] */  
int fix_ele(fix_matrix a, int i, int j) {  
    return a[i][j];  
}
```

```
movl    12(%ebp), %edx    # i  
sall    $6, %edx         # i*64  
movl    16(%ebp), %eax    # j  
sall    $2, %eax         # j*4  
addl    8(%ebp), %eax     # a + j*4  
movl    (%eax,%edx), %eax # *(a + j*4 + i*64)
```

# n X n Matrix Access – Explicit Indexing

## ■ Array Elements

- Address  $A + i * (C * K) + j * K$
- $C = n, K = 4$

```
#define IDX(n, i, j) ((i)*(n)+(j))  
/* Get element a[i][j] */  
int vec_ele(int n, int *a, int i, int j) {  
    return a[IDX(n,i,j)];  
}
```

```
movl    16(%ebp), %edx          # i  
imull   8(%ebp), %edx           # i*n  
addl    20(%ebp), %edx          # i*n + j  
movl    12(%ebp), %eax          # a  
movl    (%eax,%edx,4), %eax     # *(a + (i*n+j)*4)
```



# n X n Matrix Access – Implicit Indexing


## ■ Array Elements

- Address  $A + i * (C * K) + j * K$
- $C = n, K = 4$

```
/* Get element a[i][j] */  
int var_ele(int n, int a[n][n], int i, int j) {  
    return a[i][j];  
}
```

```
movl    8(%ebp), %eax    # n  
sall    $2, %eax        # n*4  
imull   16(%ebp), %eax   # i*n*4  
movl    20(%ebp), %edx   # j  
sall    $2, %edx        # j*4  
addl    12(%ebp), %edx   # a + j*4  
movl    (%edx,%eax), %eax # *(a + j*4 + i*n*4)
```

# Optimizing Fixed Array Access



```
#define N 16  
typedef int fix_matrix[N][N];
```

```
/* Retrieve column j from array */  
void fix_column  
    (fix_matrix a, int j, int *dest)  
{  
    int i;  
    for (i = 0; i < N; i++)  
        dest[i] = a[i][j];  
}
```

- Computation
  - Step through all elements in column j
- Optimization
  - Retrieving successive elements from single column

# Optimizing Fixed Array Access

## ■ Optimization

- Compute  $\text{ajp} = \&a[i][j]$ 
  - ▶ Initially  $= a + 4*j$
  - ▶ Increment by  $4*N$

Register	Value
<b>%ecx</b>	<b>ajp</b>
<b>%ebx</b>	<b>dest</b>
<b>%edx</b>	<b>i</b>

```
/* Retrieve column j from array */
void fix_column
    (fix_matrix a, int j, int *dest)
{
    int i;
    for (i = 0; i < N; i++)
        dest[i] = a[i][j];
}
```

```
.L8:                                # loop:
    movl    (%ecx), %eax             #   Read *ajp
    movl    %eax, (%ebx,%edx,4)      #   Save in dest[i]
    addl    $1, %edx                #   i++
    addl    $64, %ecx                #   ajp += 4*N
    cmpl    $16, %edx               #   i:N
    jne     .L8                     #   if !=, goto loop
```

# Optimizing Variable Array Access

- Compute  $ajp = \&a[i][j]$ 
  - ▶ Initially  $= a + 4*j$
  - ▶ Increment by  $4*n$

Register	Value
%ecx	ajp
%edi	dest
%edx	i
%ebx	4*n
%esi	n

```
/* Retrieve column j from array */
void var_column
(int n, int a[n][n],
 int j, int *dest)
{
    int i;
    for (i = 0; i < n; i++)
        dest[i] = a[i][j];
}
```

```
.L18:                                # loop:
    movl    (%ecx), %eax             #   Read *ajp
    movl    %eax, (%edi,%edx,4)      #   Save in dest[i]
    addl    $1, %edx                #   i++
    addl    $ebx, %ecx               #   ajp += 4*n
    cmpl    $edx, %esi              #   n:i
    jg      .L18                    #   if >, goto loop
```

# Data Structures

## ■ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

## ■ Structures


- **Allocation**
- **Access**
- Alignment

## ■ Unions

# Structure Allocation

```
struct rec {  
    int a[3];  
    int i;  
    struct rec *n;  
};
```

Memory Layout




## ■ Concept

- Contiguously-allocated region of memory
- Refer to members within structure by names
- Members may be of different types

# Structure Access

```
struct rec {  
    int a[3];  
    int i;  
    struct rec *n;  
};
```



## ■ Accessing Structure Member

- Pointer indicates first byte of structure
- Access elements with offsets


```
void  
set_i(struct rec *r,  
      int val)  
{  
    r->i = val;  
}
```

## IA32 Assembly

```
# %edx = val  
# %eax = r  
movl %edx, 12(%eax) # Mem[r+12] = val
```

# Generating Pointer to Structure Member

```
struct rec {  
    int a[3];  
    int i;  
    struct rec *n;  
};
```



## ■ Generating Pointer to Array Element

- Offset of each structure member determined at compile time
- Arguments
  - ▶ `Mem[%ebp+8]: r`
  - ▶ `Mem[%ebp+12]: idx`

```
int *get_ap  
(struct rec *r, int idx)  
{  
    return &r->a[idx];  
}
```

```
movl    12(%ebp), %eax    # Get idx  
sall    $2, %eax         # idx*4  
addl    8(%ebp), %eax     # r+idx*4
```




# Following Linked List

## ■ C Code

```
void set_val
(struct rec *r, int val)
{
    while (r) {
        int i = r->i;
        r->a[i] = val;
        r = r->n;
    }
}
```

```
struct rec {
    int a[3];
    int i;
    struct rec *n;
};
```



Register	Value
<b>%edx</b>	<b>r</b>
<b>%ecx</b>	<b>val</b>

```
.L17:                                # loop:
    movl    12(%edx), %eax            # r->i
    movl    %ecx, (%edx,%eax,4)      # r->a[i] = val
    movl    16(%edx), %edx           # r = r->n
    testl   %edx, %edx               # Test r
    jne     .L17                     # If != 0 goto loop
```

# Data Structures

## ■ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level


## ■ Structures

- Allocation
- Access
- **Alignment**

## ■ Unions

# Structures & Alignment

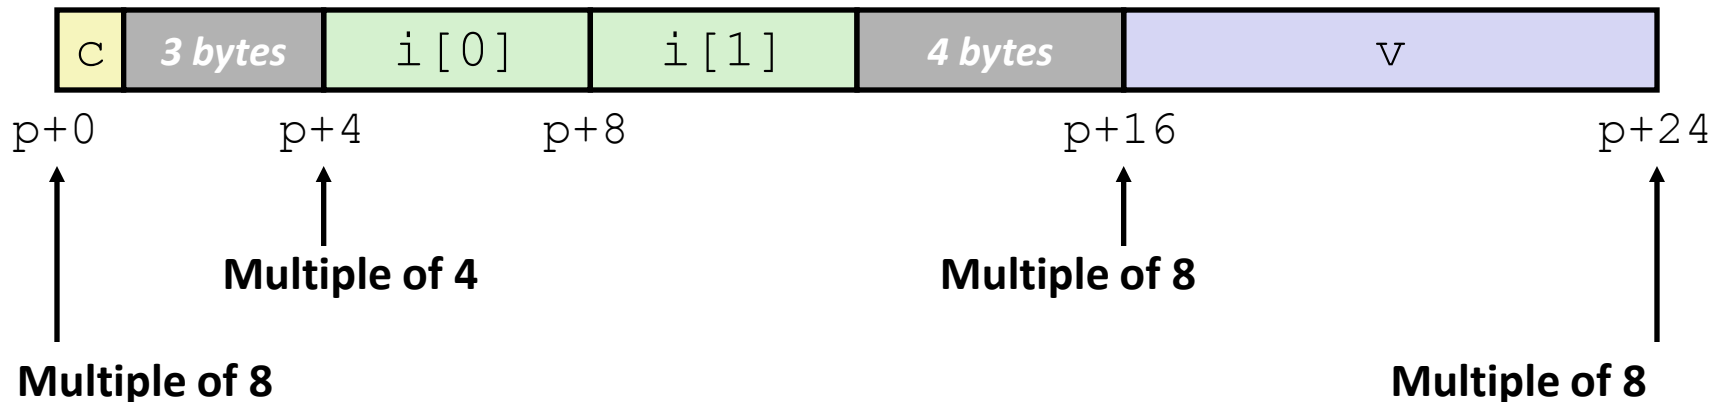
## ■ Unaligned Data



```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```

## ■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K



# Alignment Principles

## ■ Aligned Data

- Primitive data type requires K bytes
- Address must be multiple of K
- Required on some machines; advised on IA32
  - ▶ treated differently by IA32 Linux, x86-64 Linux, and Windows!

## ■ Motivation for Aligning Data

- Memory accessed by (aligned) chunks of 4 or 8 bytes (system dependent)
  - ▶ Inefficient to load or store datum that spans quad word boundaries
  - ▶ Virtual memory very tricky when datum spans 2 pages

## ■ Compiler

- Inserts gaps in structure to ensure correct alignment of fields

# Specific Cases of Alignment (IA32)

- 1 byte: char, ...
  - no restrictions on address
- 2 bytes: short, ...
  - lowest 1 bit of address must be  $0_2$
- 4 bytes: int, float, char \*, ...
  - lowest 2 bits of address must be  $00_2$
- 8 bytes: double, ...
  - Windows (and most other OS's & instruction sets):
    - ▶ lowest 3 bits of address must be  $000_2$
  - Linux:
    - ▶ lowest 2 bits of address must be  $00_2$
    - ▶ i.e., treated the same as a 4-byte primitive data type
- 12 bytes: long double
  - Windows, Linux:
    - ▶ lowest 2 bits of address must be  $00_2$
    - ▶ i.e., treated the same as a 4-byte primitive data type

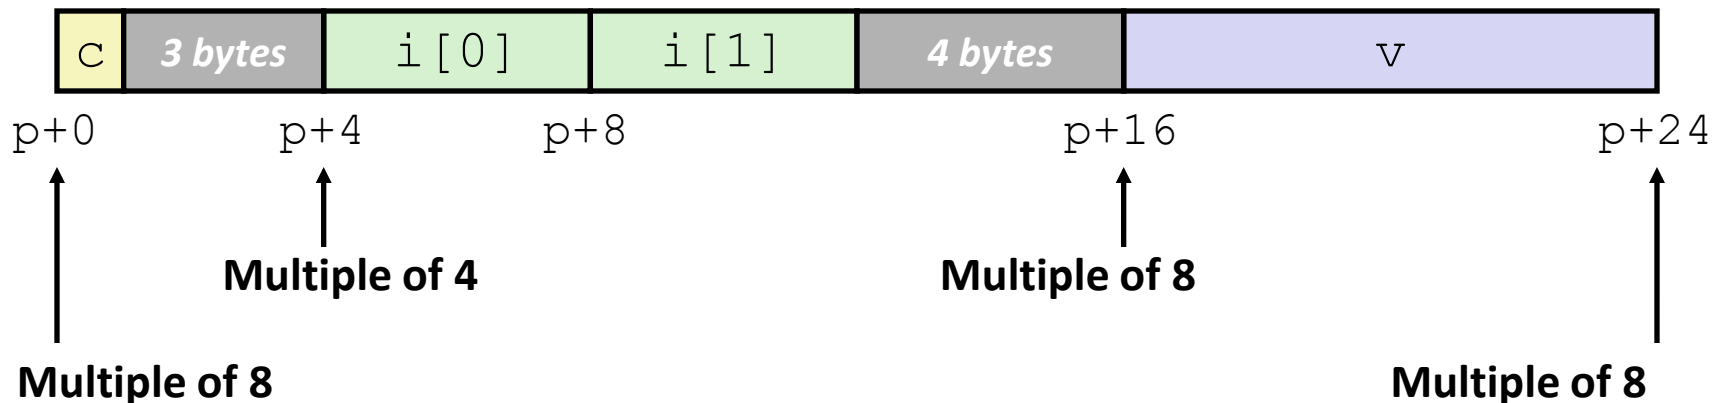
# Specific Cases of Alignment (x86-64)

- 1 byte: char, ...
  - no restrictions on address
- 2 bytes: short, ...
  - lowest 1 bit of address must be  $0_2$
- 4 bytes: int, float, ...
  - lowest 2 bits of address must be  $00_2$
- 8 bytes: double, char \*, ...
  - Windows & Linux:
    - ▶ lowest 3 bits of address must be  $000_2$
- 16 bytes: long double
  - Linux:
    - ▶ lowest 3 bits of address must be  $000_2$
    - ▶ i.e., treated the same as a 8-byte primitive data type

# Satisfying Alignment with Structures

- Within structure:
  - Must satisfy each element's alignment requirement
- Overall structure placement
  - Each structure has alignment requirement K
    - ▶ K = Largest alignment of any element
  - Initial address & structure length must be multiples of K
- Example (under Windows or x86-64):
  - K = 8, due to double element


```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```




# Different Alignment Conventions

- x86-64 or IA32 Windows:
  - $K = 8$ , due to double element

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *p;
```



- IA32 Linux
  - $K = 4$ ; double treated like a 4-byte data type






# Meeting Overall Alignment Requirement

- For largest alignment requirement K
- Overall structure must be multiple of K


```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} *p;
```



# Arrays of Structures

- Overall structure length multiple of K
- Satisfy alignment requirement for every element


```
struct S2 {  
    double v;  
    int i[2];  
    char c;  
} a[10];
```



# Accessing Array Elements

- Compute array offset  $12i$ 
  - `sizeof(S3)`, including alignment spacers
- Element `j` is at offset 8 within structure
- Assembler gives offset `a+8`
  - Resolved during linking

```
struct S3 {  
    short i;  
    float v;  
    short j;  
} a[10];
```



```
short get_j(int idx)  
{  
    return a[idx].j;  
}
```

```
# %eax = idx  
leal (%eax,%eax,2),%eax # 3*idx  
movswl a+8(,%eax,4),%eax
```

# Saving Space



- Put large data types first

```
struct S4 {  
    char c;  
    int i;  
    char d;  
} *p;
```



```
struct S5 {  
    int i;  
    char c;  
    char d;  
} *p;
```

- Effect (K=4)



# Data Structures

## ■ Arrays

- One-dimensional
- Multi-dimensional (nested)
- Multi-level

## ■ Structures

- Allocation
- Access
- Alignment


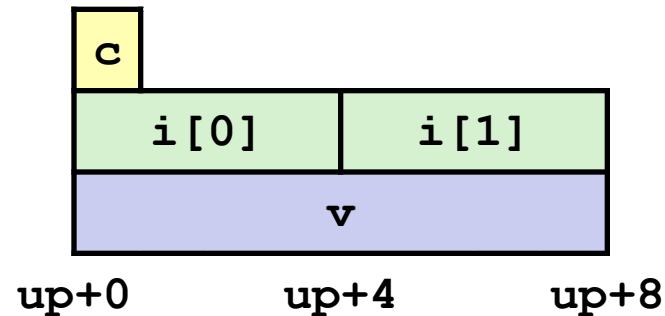
## ■ Unions

# Union Allocation

- Allocate according to largest element
- Can only use one field at a time


```
union U1 {  
    char c;  
    int i[2];  
    double v;  
} *up;
```

```
struct S1 {  
    char c;  
    int i[2];  
    double v;  
} *sp;
```



# Using Union to Access Bit Patterns

```
typedef union {  
    float f;  
    unsigned u;  
} bit_float_t;
```



```
float bit2float(unsigned u)  
{  
    bit_float_t arg;  
    arg.u = u;  
    return arg.f;  
}
```

Same as (float) u ?

```
unsigned float2bit(float f)  
{  
    bit_float_t arg;  
    arg.f = f;  
    return arg.u;  
}
```

Same as (unsigned) f ?

# Byte Ordering Revisited

## ■ Idea

- Short/long/quad words stored in memory as 2/4/8 consecutive bytes
- Which is most (least) significant?
- Can cause problems when exchanging binary data between machines

## ■ Big Endian

- Most significant byte has lowest address
- Sparc

## ■ Little Endian

- Least significant byte has lowest address
- Intel x86



# Byte Ordering Example

```
union {
    unsigned char c[8];
    unsigned short s[4];
    unsigned int i[2];
    unsigned long l[1];
} dw;
```

32-bit	c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
	s[0]		s[1]		s[2]		s[3]	
	i[0]				i[1]			
	l[0]							

64-bit

c[0]	c[1]	c[2]	c[3]	c[4]	c[5]	c[6]	c[7]
s[0]		s[1]		s[2]		s[3]	
i[0]				i[1]			
l[0]							

# Byte Ordering Example (Cont).

```
int j;
for (j = 0; j < 8; j++)
    dw.c[j] = 0xf0 + j;

printf("Characters 0-7 ==
[0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x,0x%x] \n",
    dw.c[0], dw.c[1], dw.c[2], dw.c[3],
    dw.c[4], dw.c[5], dw.c[6], dw.c[7]);

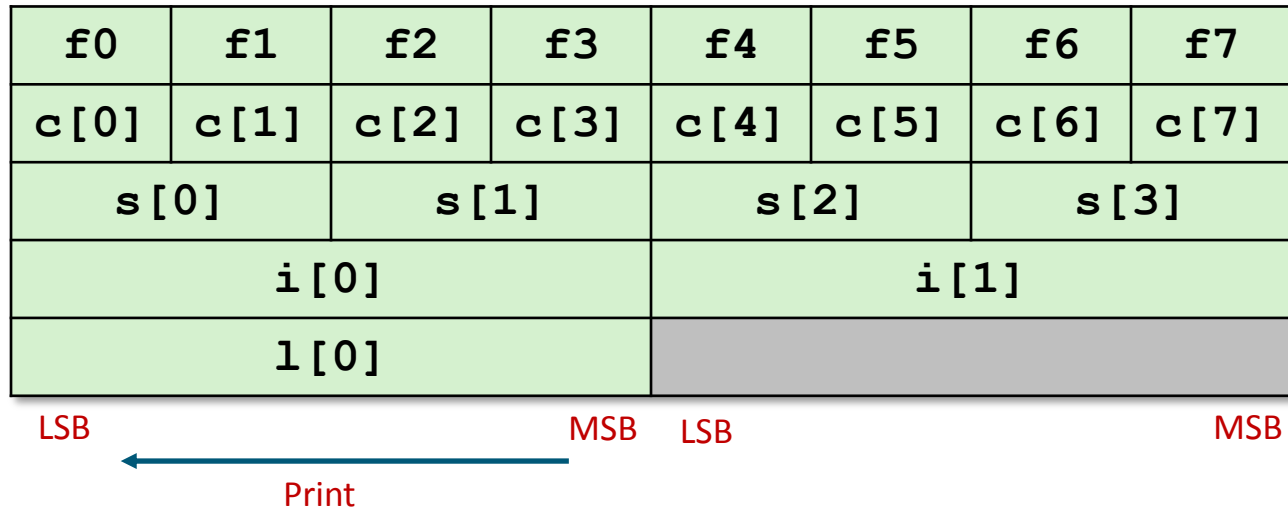
printf("Shorts 0-3 == [0x%x,0x%x,0x%x,0x%x] \n",
    dw.s[0], dw.s[1], dw.s[2], dw.s[3]);

printf("Ints 0-1 == [0x%x,0x%x] \n",
    dw.i[0], dw.i[1]);

printf("Long 0 == [0x%lx] \n",
    dw.l[0]);
```

# Byte Ordering on IA32

## Little Endian



## Output:

Characters 0-7 == [0xf0, 0xf1, 0xf2, 0xf3, 0xf4, 0xf5, 0xf6, 0xf7]

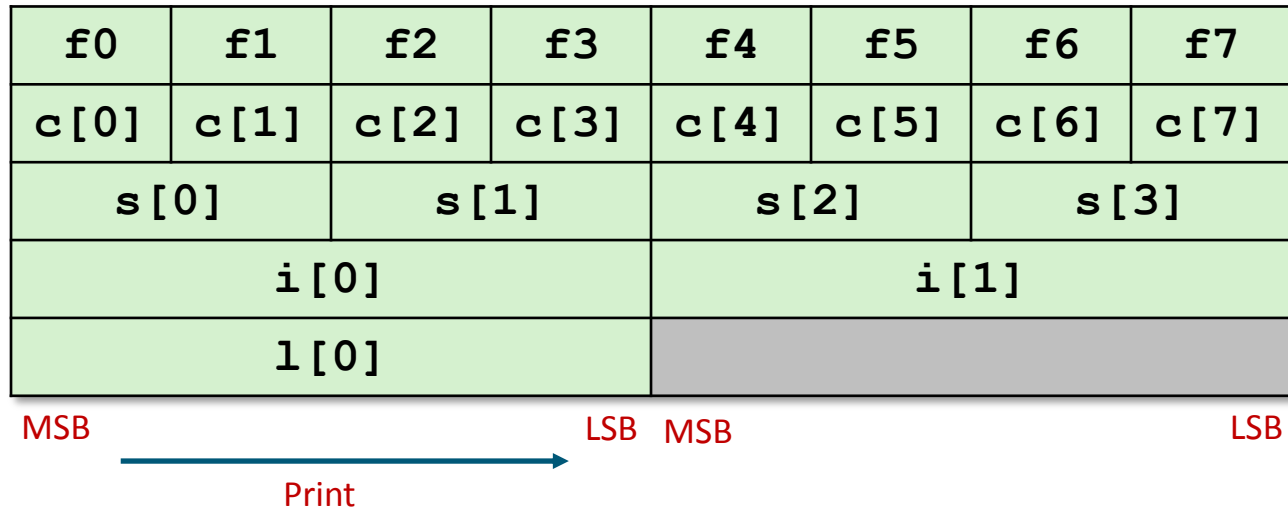
Shorts 0-3 == [0xf1f0, 0xf3f2, 0xf5f4, 0xf7f6]

Ints 0-1 == [0xf3f2f1f0, 0xf7f6f5f4]

Long 0 == [0xf3f2f1f0]

# Byte Ordering on Sun

## Big Endian



## Output on Sun:

Characters 0-7 == [0xf0, 0xf1, 0xf2, 0xf3, 0xf4, 0xf5, 0xf6, 0xf7]

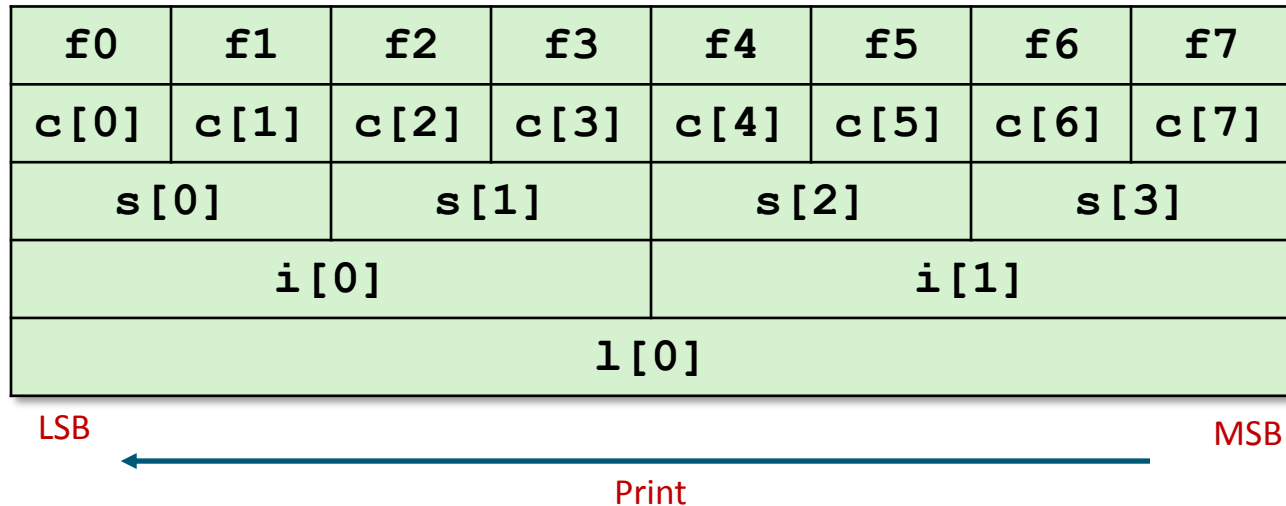
Shorts 0-3 == [0xf0f1, 0xf2f3, 0xf4f5, 0xf6f7]

Ints 0-1 == [0xf0f1f2f3, 0xf4f5f6f7]

Long 0 == [0xf0f1f2f3]

# Byte Ordering on x86-64

## Little Endian



## Output on x86-64:

Characters 0-7 == [0xf0, 0xf1, 0xf2, 0xf3, 0xf4, 0xf5, 0xf6, 0xf7]

Shorts 0-3 == [0xf1f0, 0xf3f2, 0xf5f4, 0xf7f6]

Ints 0-1 == [0xf3f2f1f0, 0xf7f6f5f4]

Long 0 == [0xf7f6f5f4f3f2f1f0]

# Summary

## ■ Arrays in C

- Contiguous allocation of memory
- Aligned to satisfy every element's alignment requirement
- Pointer to first element
- No bounds checking

## ■ Structures

- Allocate bytes in order declared
- Pad in middle and at end to satisfy alignment

## ■ Unions

- Overlay declarations
- Way to circumvent type system