
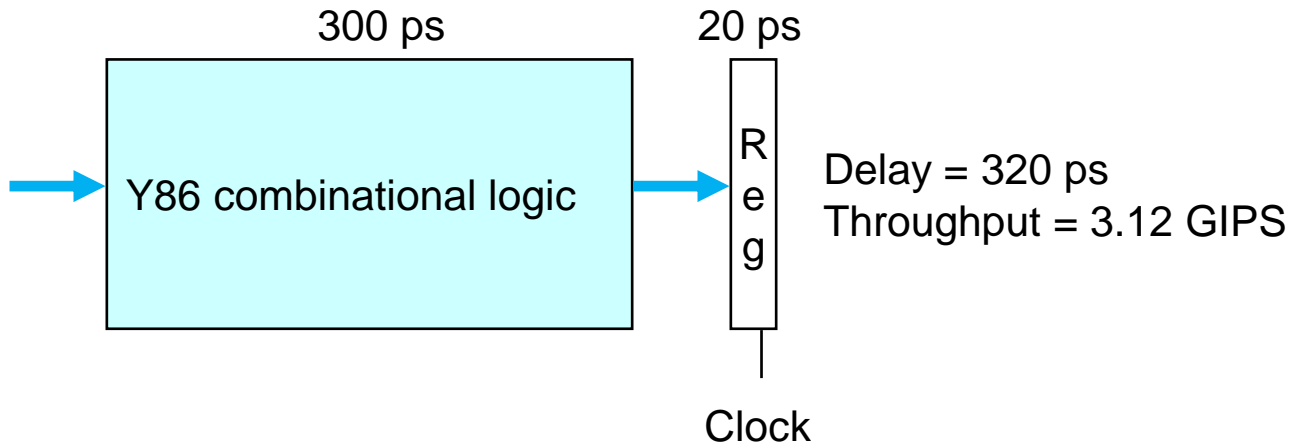


Processor Architecture

Principles of Pipelining



Limitations of Sequential Circuits



■ Sequential Implementation

- Computation requires total of 300 picoseconds
- Additional 20 picoseconds to save result in register
- Must have clock cycle of at least 320 ps

Sequential vs. Parallel vs. Pipelining

■ Car Wash – Sequential



Sequential vs. Parallel vs. Pipelining

- Car Wash – Parallel (Coarse-grained)




Sequential vs. Parallel vs. Pipelining

- Car Wash – Parallel (Fine-grained)



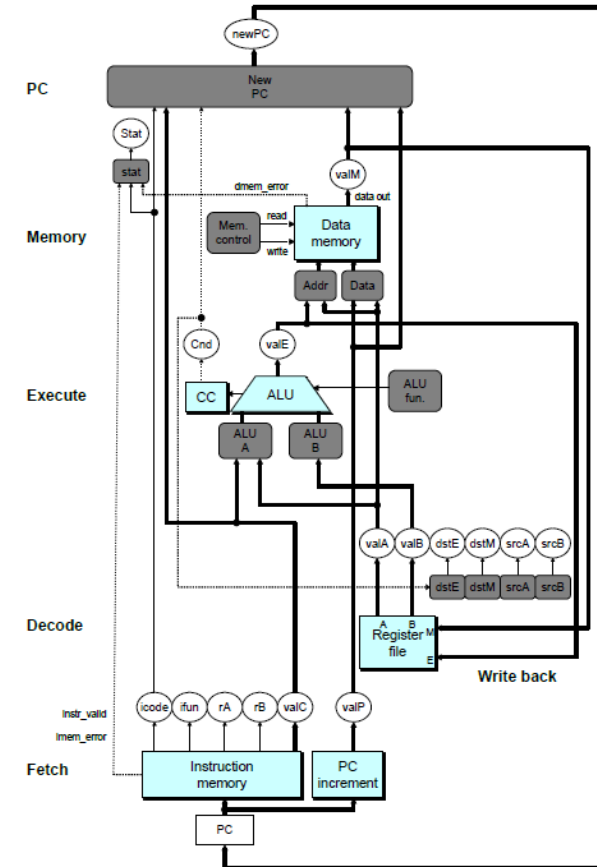

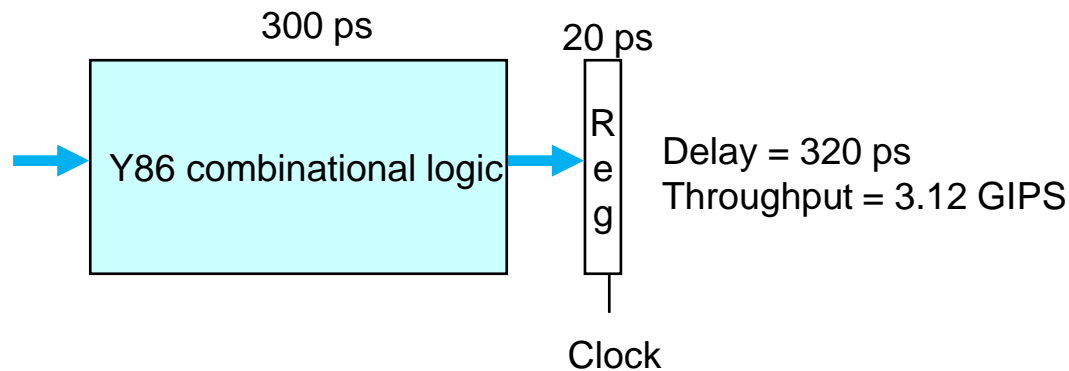
Sequential vs. Parallel vs. Pipelining

■ Car Wash – Pipelined



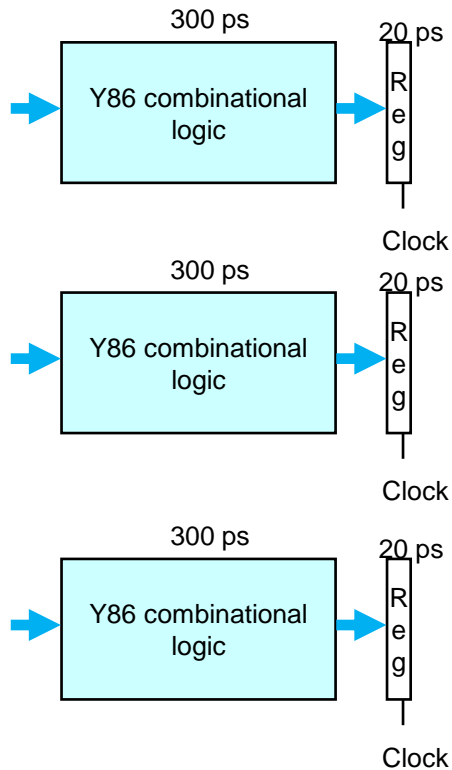
Sequential Architectures

Sequential Logic in Parallel




Parallel Architectures

■ Coarse-grained Parallelism – Multiprocessor, Multicore




Delay = 320 ps
Throughput = $3 * 3.12 \text{ GIPS}$
= 9.36 GIPS




Parallel Architectures

■ Fine-grained Parallelism – Superscalar, VLIW



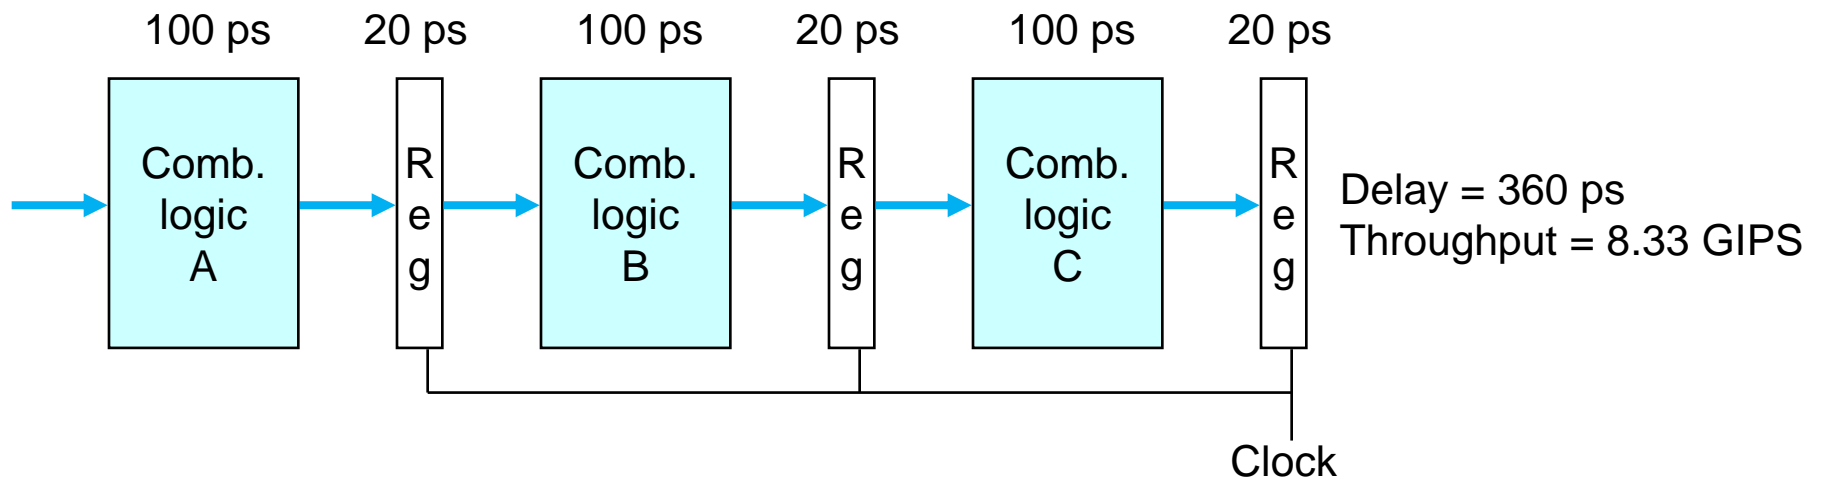
Delay = 320 ps
Throughput $\leq 3 * 3.12$ GIPS
 ≤ 9.36 GIPS



Pipelined Architectures

■ Pipelining – (almost) all modern processors


- idea:
 - ▶ split work up into several independent phases
 - ▶ run independent phases in parallel




- can begin new operation as soon as previous has passed through stage A
- overall latency (for one operation) *increases*

Pipeline Stages

■ 3-way pipelined

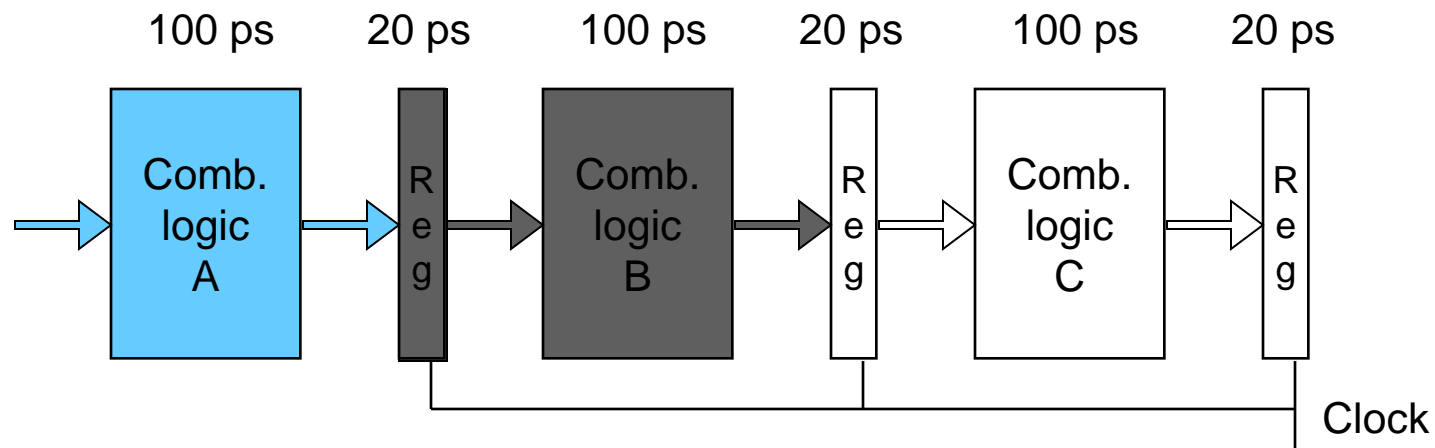



■ 5-way pipelined



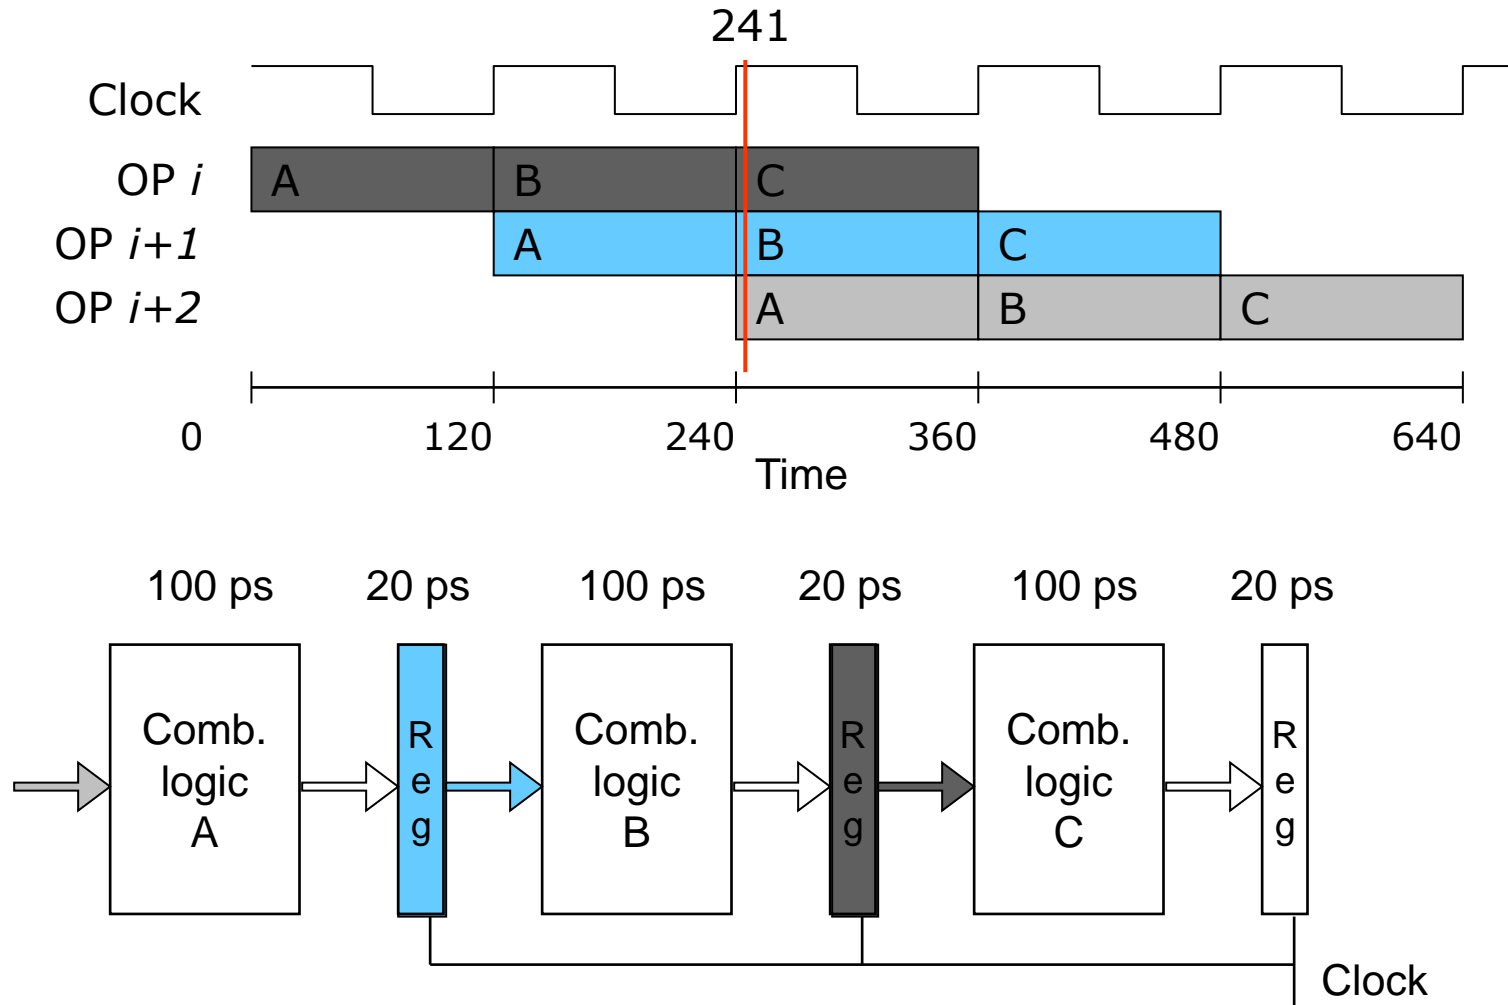
Pipeline Operation

- Immediately before a new clock cycle starts




Pipeline Operation

- Immediately after a new clock cycle has started




Pipeline Operation

- In the middle of a clock cycle



Pipeline Operation

- Immediately before a clock cycle ends



Basic Pipeline

■ Pipeline Stages

1. Instruction fetch step (F)
2. Instruction decode/register fetch step (D)
3. Execution/effective address step (E)
4. Memory access (M)
5. Register write-back step (W)

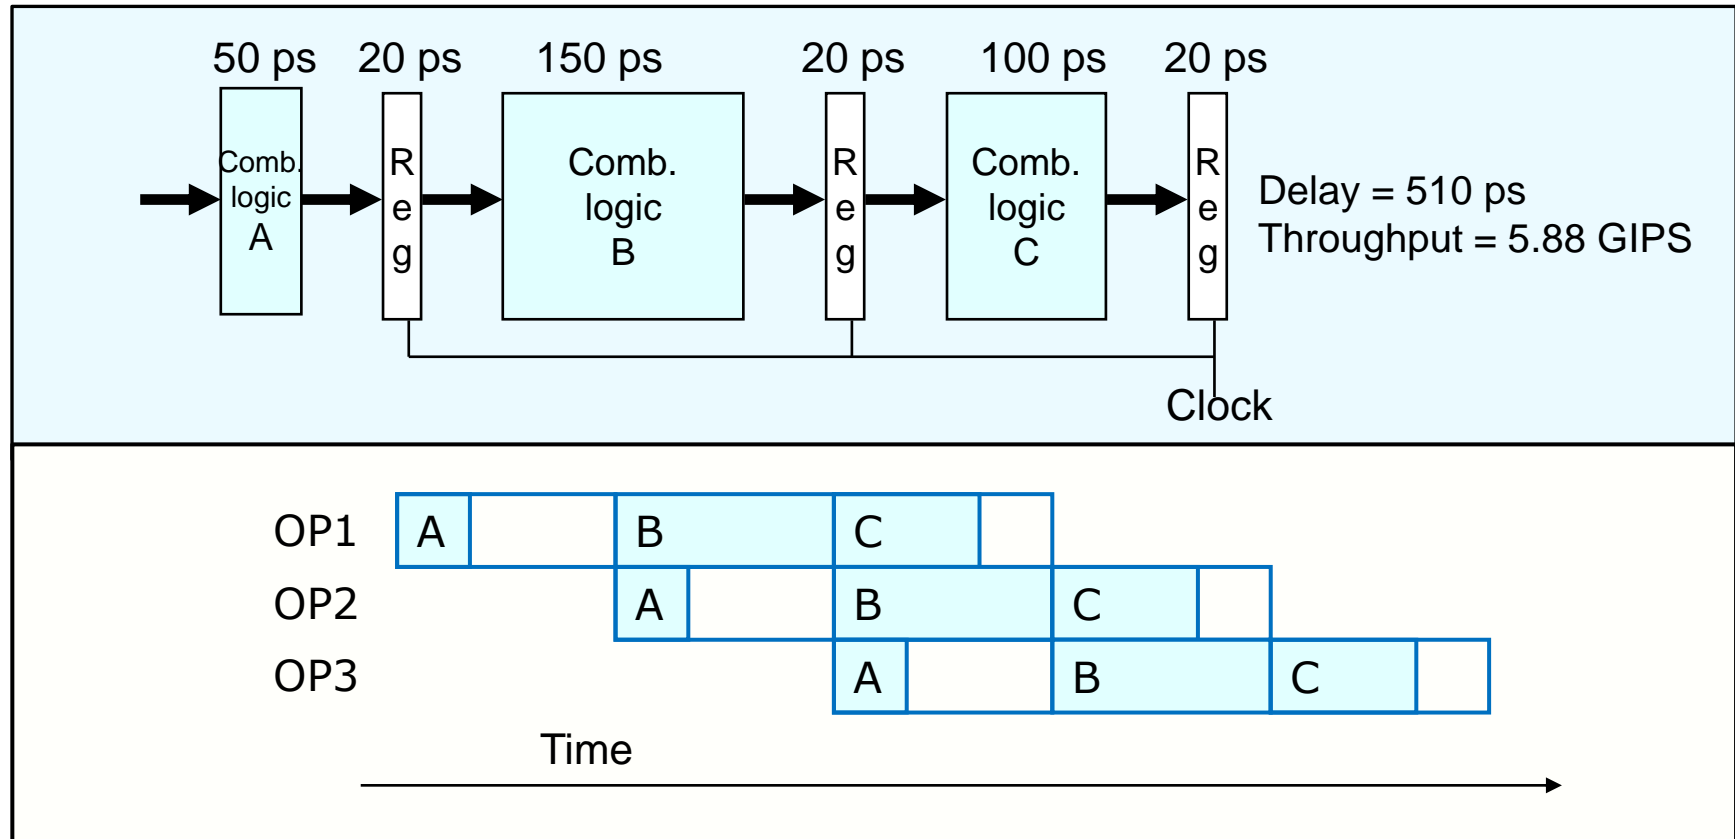
■ Pipeline Operation

	Clock number								
Instruction number	1	2	3	4	5	6	7	8	9
Instruction i	F	D	E	M	W				
Instruction $i + 1$		F	D	E	M	W			
Instruction $i + 2$			F	D	E	M	W		
Instruction $i + 3$				F	D	E	M	W	
Instruction $i + 4$					F	D	E	M	W

Limitations and Major Hurdles of Pipelining

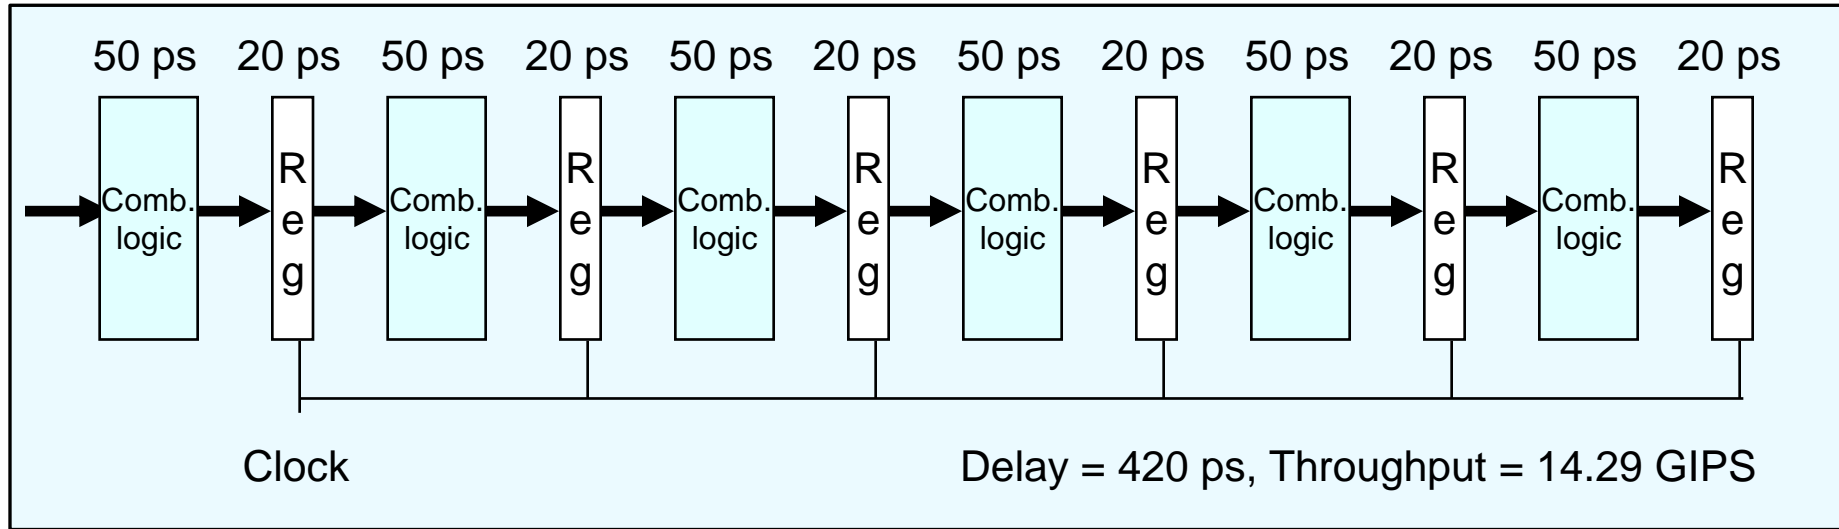
- Non-uniform delays
- Overhead
- Pipeline hazards

Limitations: Nonuniform Delays



- Throughput limited by slowest stage
- Other stages sit idle for much of the time
- Challenging to partition system into balanced stages

Limitations: Register Overhead




- As we try to deepen the pipeline, the overhead of loading registers becomes more significant
- Percentage of clock cycle spent loading register:
 - 1-stage pipeline: 6.25%
 - 3-stage pipeline: 16.67%
 - 6-stage pipeline: 28.57%
- High speeds of modern processor designs obtained through very deep pipelining

Pipeline Hazards

- Three major hurdles of pipelining
 - Structural Hazard
 - Data Hazard
 - Control Hazard

Structural Hazard



	Clock number								
Instruction number	1	2	3	4	5	6	7	8	9
Load Instruction	F	D	E	M	W				
Instruction $i + 1$		F	D	E	M	W			
Instruction $i + 2$			F	D	E	M	W		
Instruction $i + 3$				F	D	E	M	W	
Instruction $i + 4$					F	D	E	M	W

- and any other form of hardware resource contention

Solutions to Structural Hazards


■ Resource Duplication

- example

- ▶ Separate I and D caches for memory access conflict
- ▶ Time-multiplexed or multi-port register file for register file access conflict

Data Hazard (RAW hazard)

OP i `addl %edx, %eax`
OP $i+1$ `subl %eax, %ecx`




Solutions to Data Hazards


- Freezing the pipeline
- (Internal) Forwarding
- Compiler scheduling

Freezing The Pipeline

- ALU result to next instruction




- Load result to next instruction




(Internal) Forwarding

- ALU result to next instruction (Stall X)
 - requires forwarding path



- Load result to next instruction (Stall 1)
 - requires forwarding path




Load interlock

Compiler Scheduling

- Rearranging instructions in order to avoid data hazards

```
...  
addl %eax, %ebx  
rmmovl 0(%edx), %ebx  
rmmovl 4(%edx), %eax  
addl %eax, %ebx  
irmovl $4, %ecx  
...
```



```
...  
addl %eax, %ebx  
rmmovl 4(%edx), %eax  
rmmovl 0(%edx), %ebx  
irmovl $4, %ecx  
addl %eax, %ebx  
...
```

- requires a careful data flow analysis and knowledge about the inner workings of the pipeline

Control Hazard

- Caused by PC-changing instructions
 - (conditional/unconditional) jumps, function call/return

Branch Instruction	F	D	E	M	W					
Branch successor		F	<i>stall</i>	<i>stall</i>	F	D	E	M	W	
Branch successor + 1						F	D	E	M	W
Branch successor + 2							F	D	E	M
Branch successor + 3								F	D	E
Branch successor + 4									F	D
Branch successor + 5										F


**For 5-stage pipeline, 3 cycle penalty
15% branch frequency. CPI = 1.45**

Solutions to Control Hazards

- Branch delay slots
- Branch prediction

Branch Delay Slots

- Execute the next n instructions following a branch no matter what the result of the branch
 - up to the compiler to make good use of the branch delay slots
 - exposes the pipeline structure to the user



branch taken: stall 2 cycles
not taken: no stall

addl, mrmovl executed no
matter whether the jump is
taken or not (no stall)

Branch Prediction

■ Predict-taken

Taken branch instruction	F	D	E	M	W					
Branch target	F		D	E	M	W				
Branch target + 1	F			D	E	M	W			
Branch target + 2	F				D	E	M	W		
Branch target + 3	F					D	E	M	W	

Untaken branch instruction	F	D	E	M	W				
Branch target	F		<i>idle</i>	<i>idle</i>	<i>idle</i>	<i>idle</i>			
Untaken branch instruction + 1	F			D	E	M	W		
Untaken branch instruction + 2	F				D	E	M	W	
Untaken branch instruction + 3	F					D	E	M	W

~60% success rate

■ other techniques:

- predict backwards-taken, forward-untaken
- hardware branch prediction