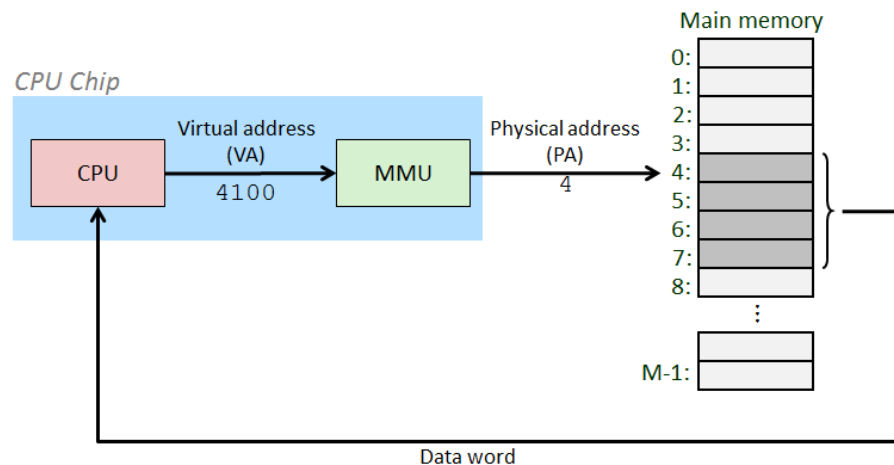


The Memory Hierarchy

Virtual Memory Concepts




Virtual Memory: Theory

- **Virtual Memory(VM): Overview and motivation**
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation
- Allocation

Acknowledgement: slides based on the cs:app2e material


Memory Up Until Now

- Programs refer to virtual memory addresses
 - `mov (%ebx), %eax`
 - conceptually a very large array of bytes
 - implemented with a hierarchy of different memory types
 - system provides private address space to each process
- Allocation: Compiler and run-time system
 - where different program objects should be stored
 - all objects allocated within a single virtual address space
- Why would we need virtual memory?




Problem 1: How does everything fit?

virtual memory:
64-bit addresses,
16 exabytes



physical memory:
a few gigabytes



...and there are many processes

Problem 2: Memory Management

physical main memory

process 1
process 2
process 3
...
process n


X

.text
.data
heap
stack


what goes
where?




Problem 3: Protection



Problem 4: Sharing



Solution: Indirection

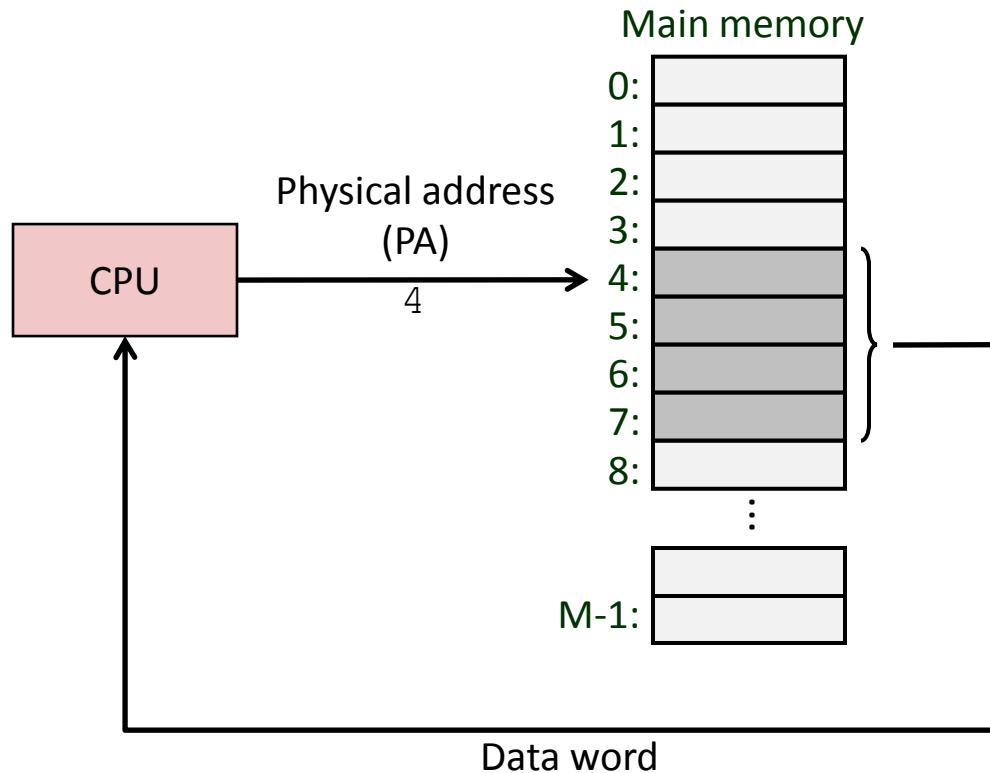


- Each process gets its own private memory space
- Solves all previous problems

Address Spaces

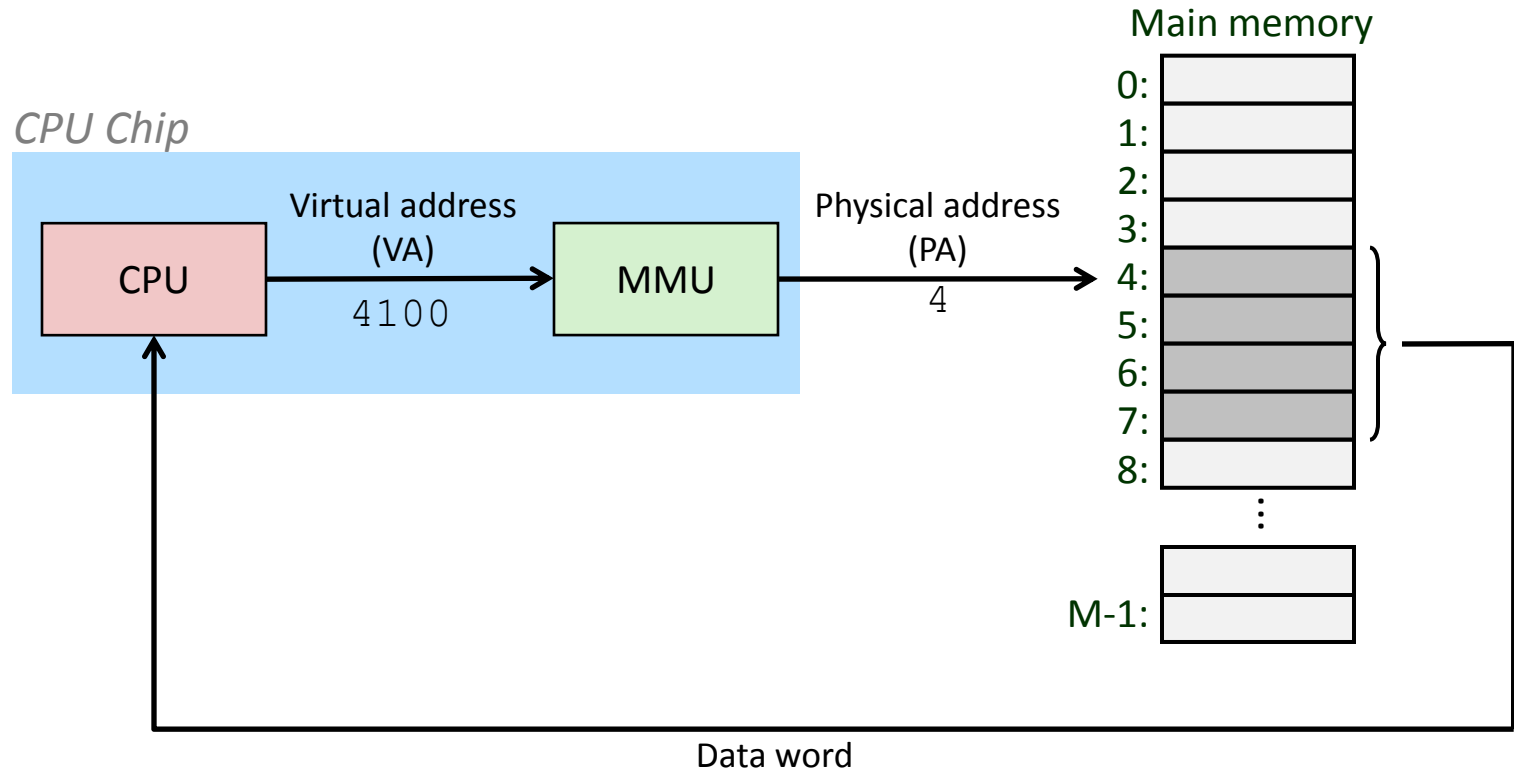
- Linear address space: Ordered set of contiguous non-negative integer addresses:
 $\{0, 1, 2, 3 \dots\}$
- Virtual address space: Set of $N = 2^n$ virtual addresses
 $\{0, 1, 2, 3, \dots, N-1\}$
- Physical address space: Set of $M = 2^m$ physical addresses
 $\{0, 1, 2, 3, \dots, M-1\}$
- Clean distinction between data (bytes) and their attributes (addresses)
- Each object can now have multiple addresses
- Every byte in main memory:
one physical address, one (or more) virtual addresses

A System Using Physical Addressing



- Still used in “simple” systems like embedded microcontrollers in devices like cars, elevators, and digital picture frames

A System Using Virtual Addressing



- Used in all modern servers, desktops, and laptops
- One of the great ideas in computer science

Why Virtual Memory (VM)?


- Efficient use of limited main memory
 - uses DRAM as a cache for the parts of a virtual address space
 - ▶ some non-cached parts stored on disk
 - ▶ some (unallocated) and non-cached parts stored nowhere
 - keep only active areas of virtual address space in memory
 - ▶ transfer data back and forth as needed
- Simplified memory management for programmers
 - Each process gets the same full, uniform linear address space
- Isolated address spaces
 - One process can't interfere with another's memory
 - ▶ because they operate in different address spaces
 - User program cannot access privileged kernel information
 - ▶ different sections of address spaces have different permissions

Virtual Memory: Theory

- Virtual Memory(VM): Overview and motivation
- **VM as a tool for caching**
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation
- Allocation

VM as a Tool for Caching

- Virtual memory is an array of N contiguous bytes stored on disk
 - conceptually: stored somewhere on disk
- Physical memory = cache for allocated virtual memory
- Cache blocks are called pages (size is $P = 2^p$ bytes)

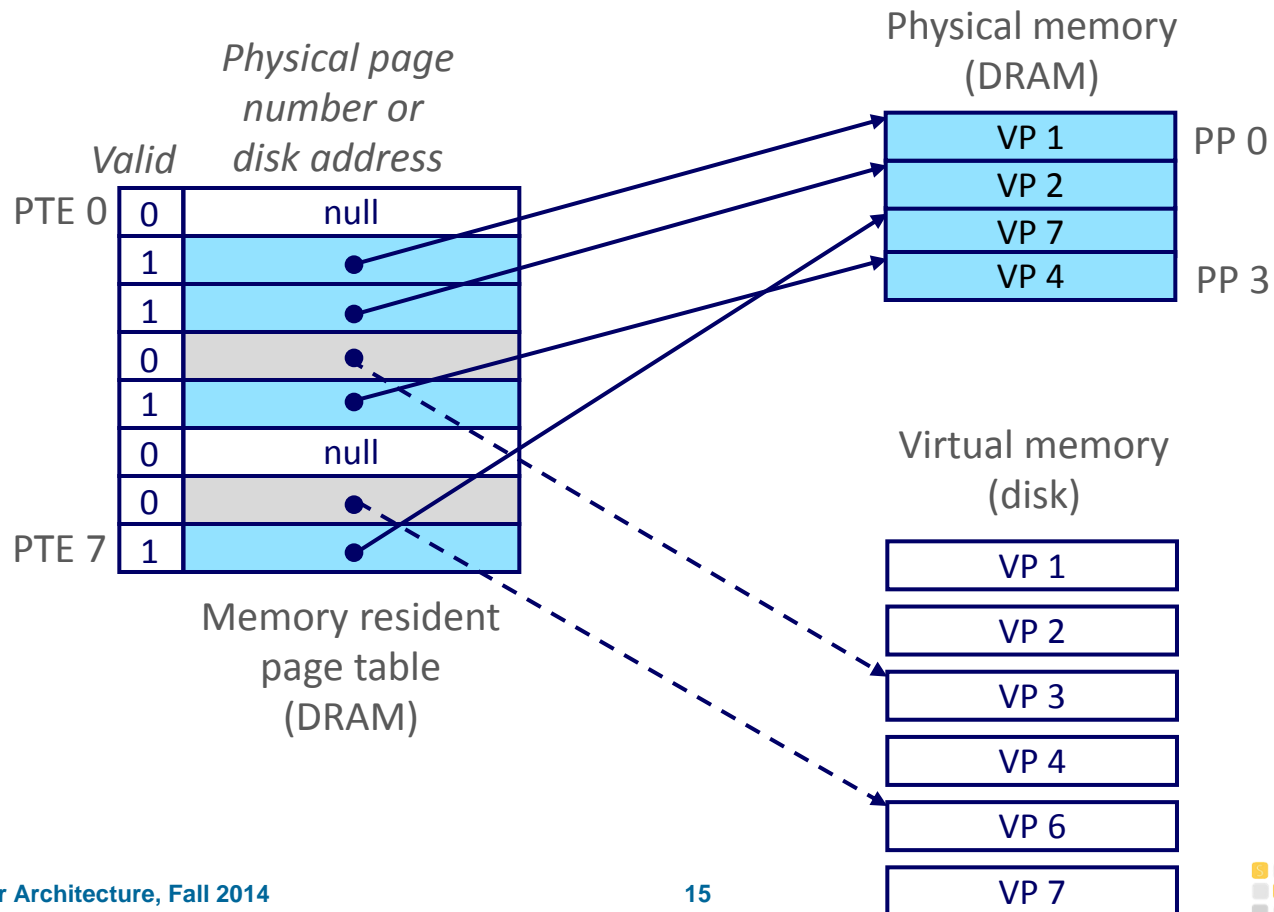


DRAM Cache Organization

- DRAM cache organization driven by the enormous miss penalty
 - DRAM is about 10x slower than SRAM
 - Disk is about 10,000x slower than DRAM
- Consequences
 - Large page (block) size: typically 4-8 KB, sometimes 4 MB
 - Fully associative
 - ▶ Any VP can be placed in any PP
 - ▶ Requires a “large” mapping function – different from CPU caches
 - Highly sophisticated, expensive replacement algorithms
 - ▶ Too complicated and open-ended to be implemented in hardware
 - Write-back rather than write-through

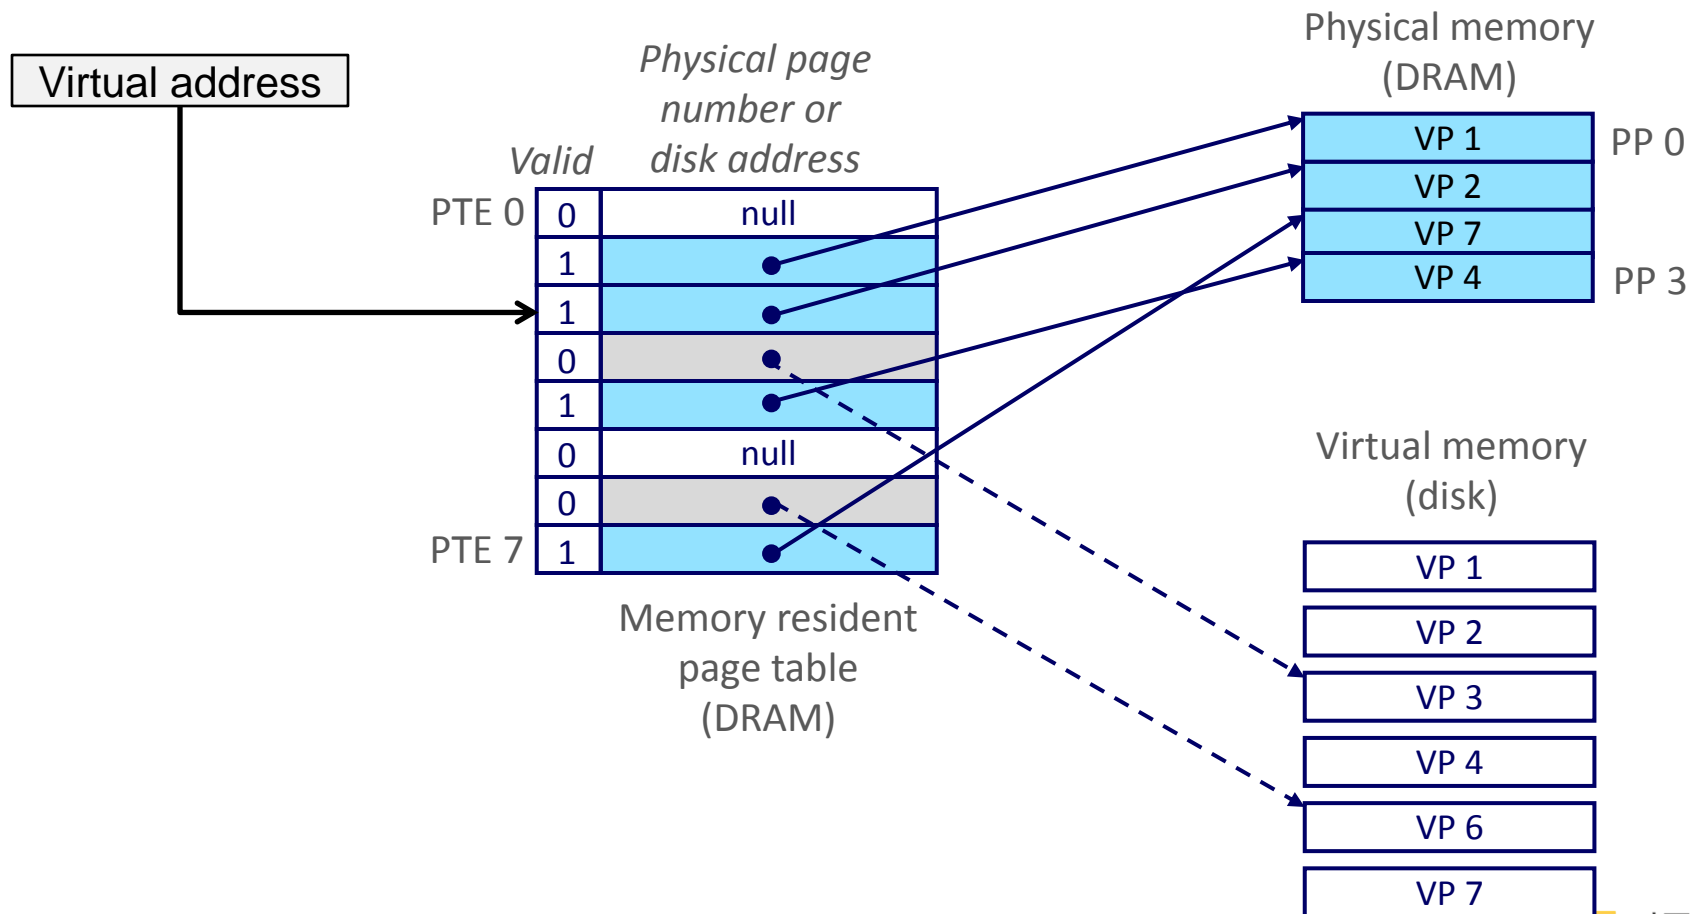
Address Translation: Page Tables

- A **page table** is an array of page table entries (PTEs) that maps virtual pages to physical pages.
 - Per-process kernel data structure in DRAM



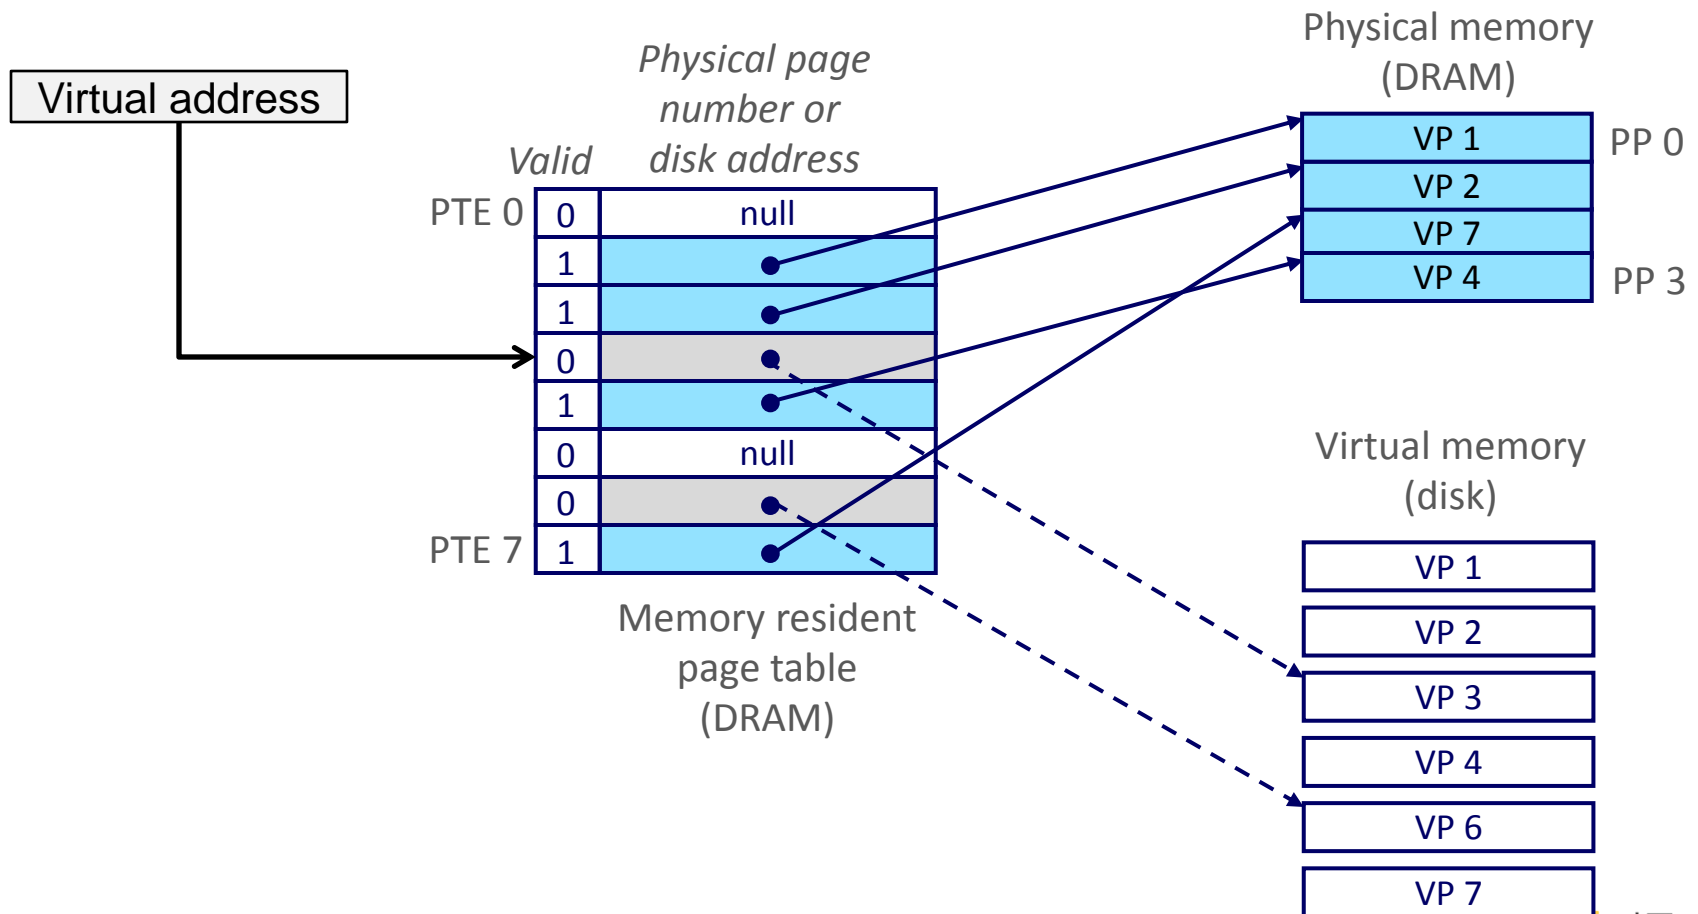
Page Hit

- Reference to VM word that is in physical memory (DRAM cache hit)



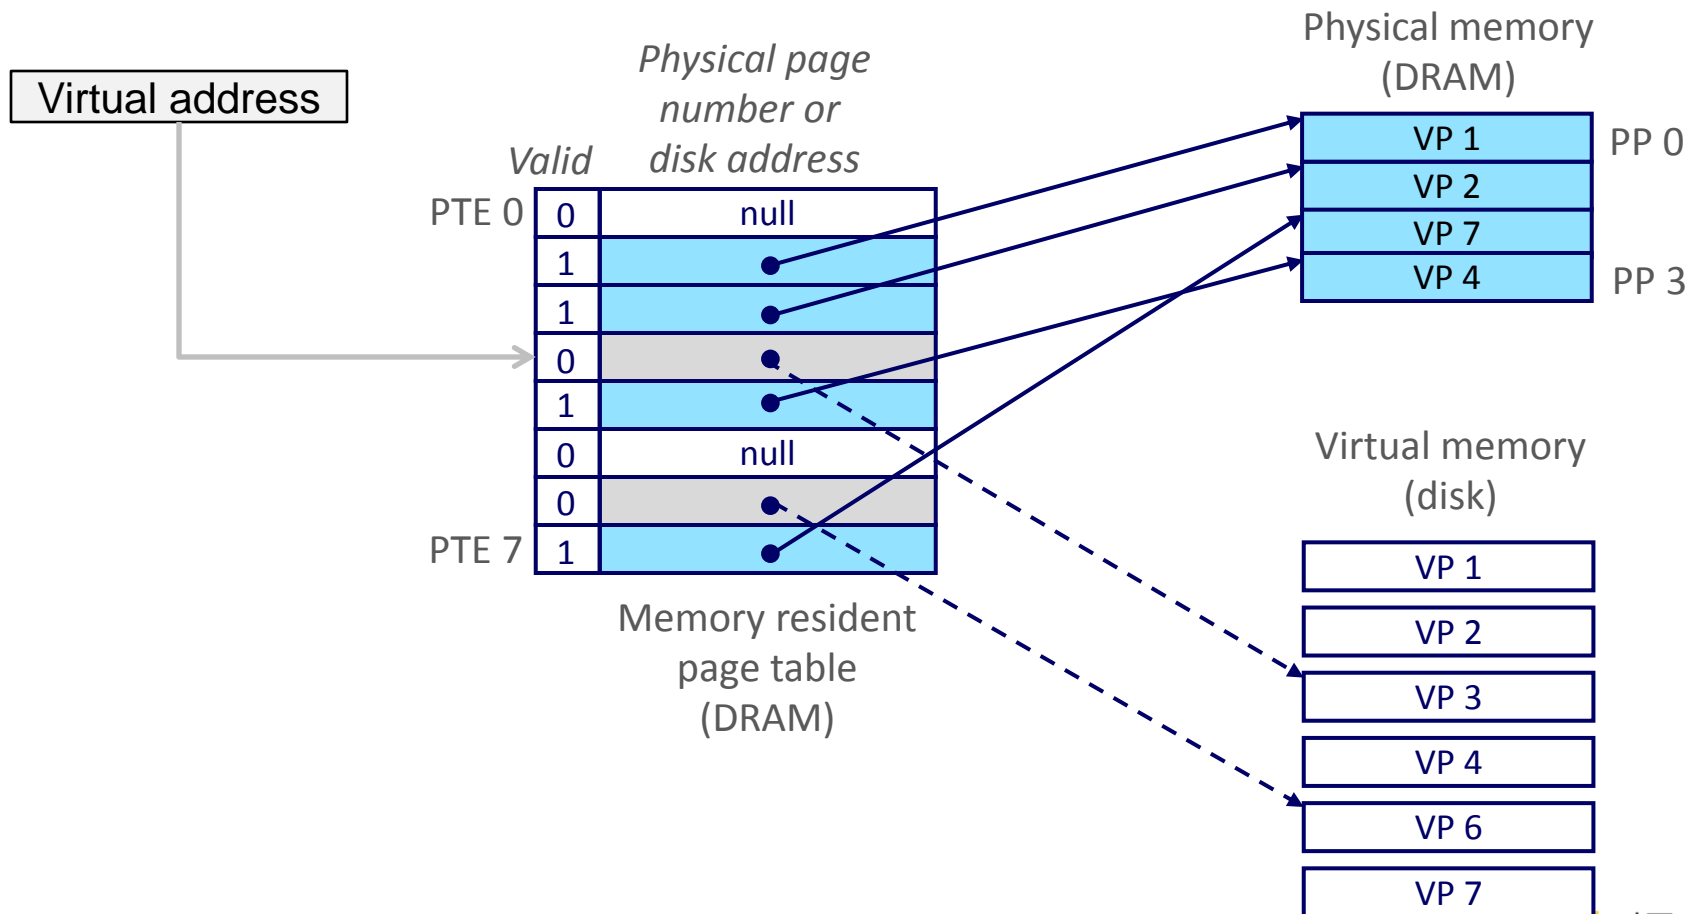
Page Miss

- Reference to VM word that is not in physical memory
 - in the context of VM called “page fault”



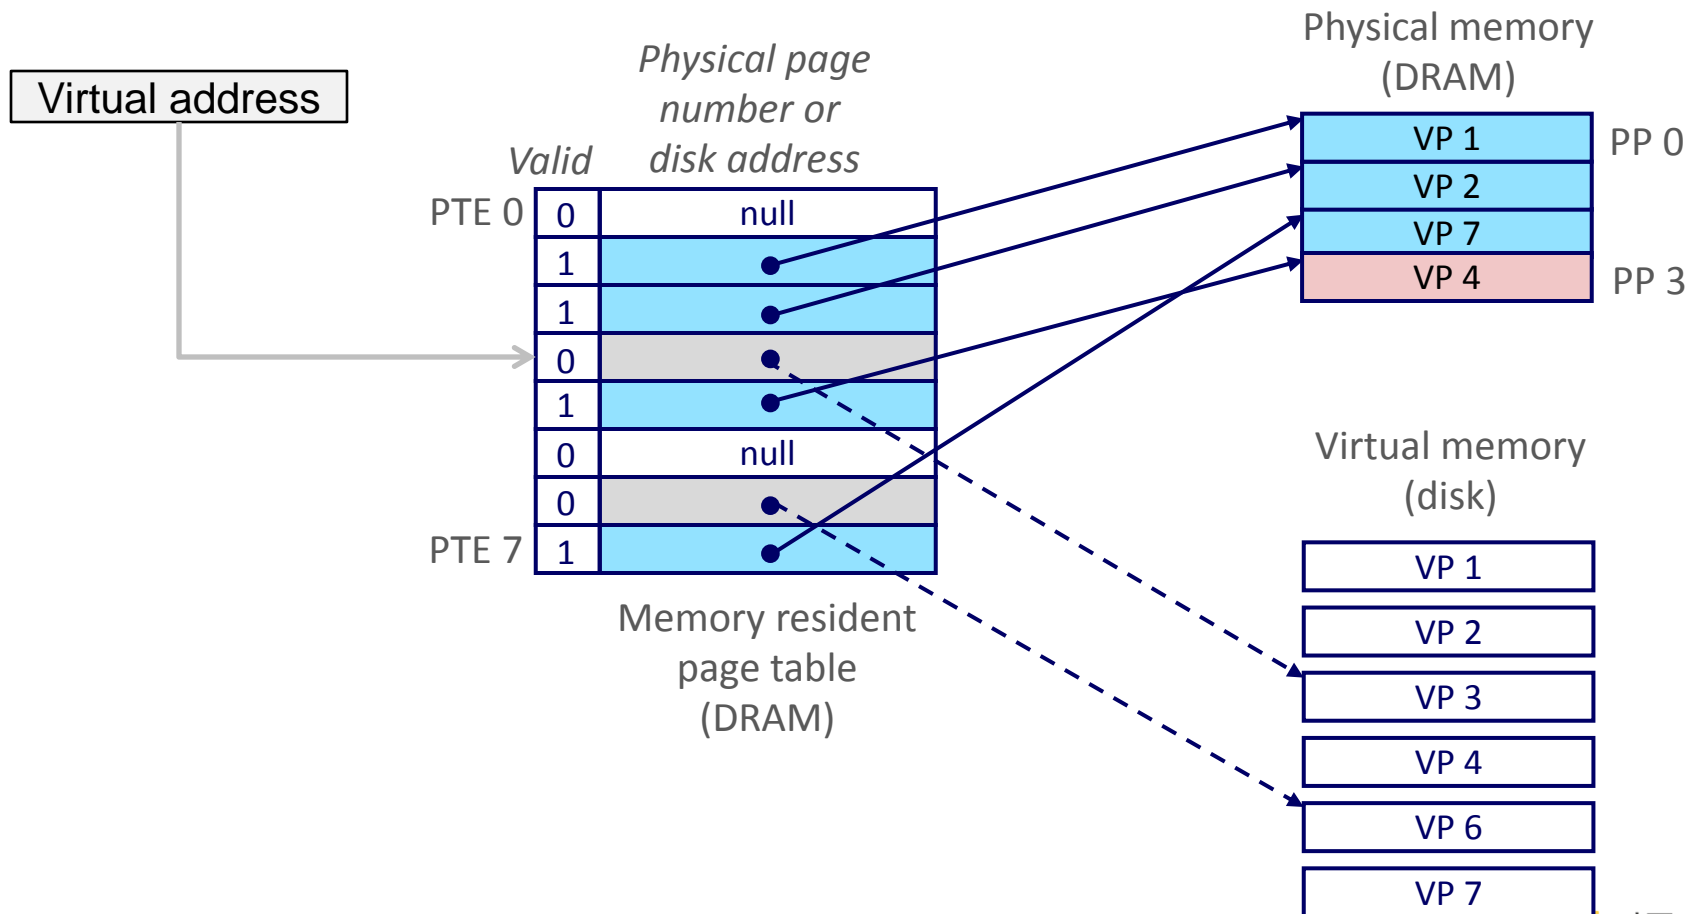
Handling a Page Fault

- Page miss causes page fault (an exception)



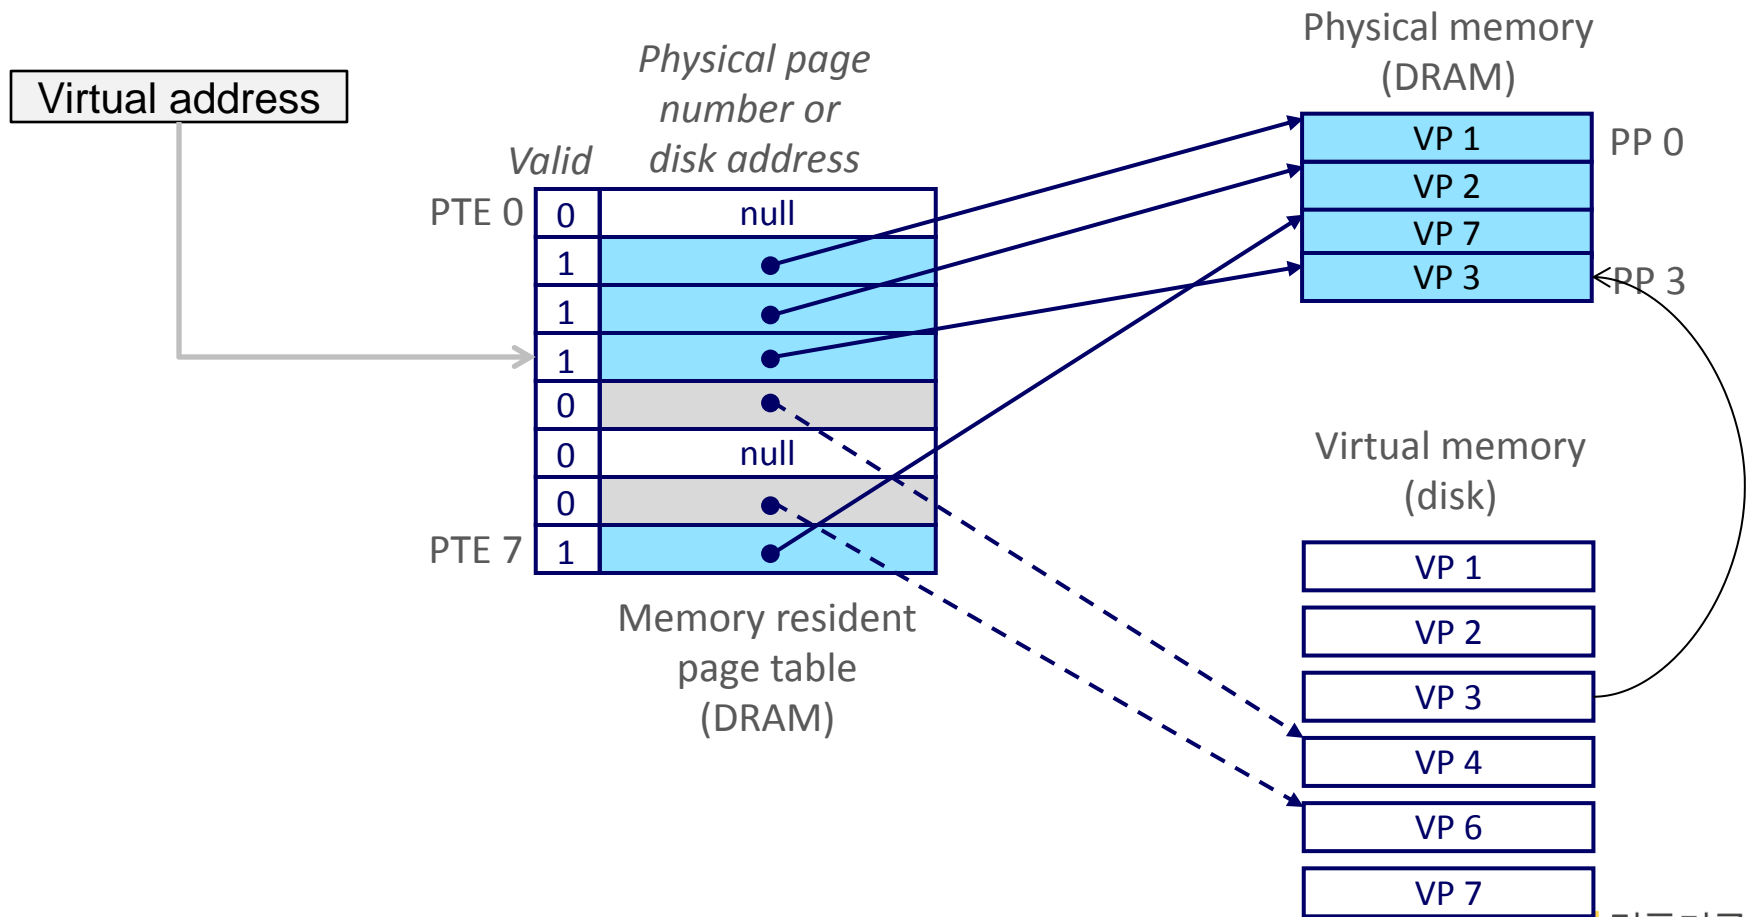
Handling a Page Fault

- Page fault handler selects a victim to be evicted (here VP 4)



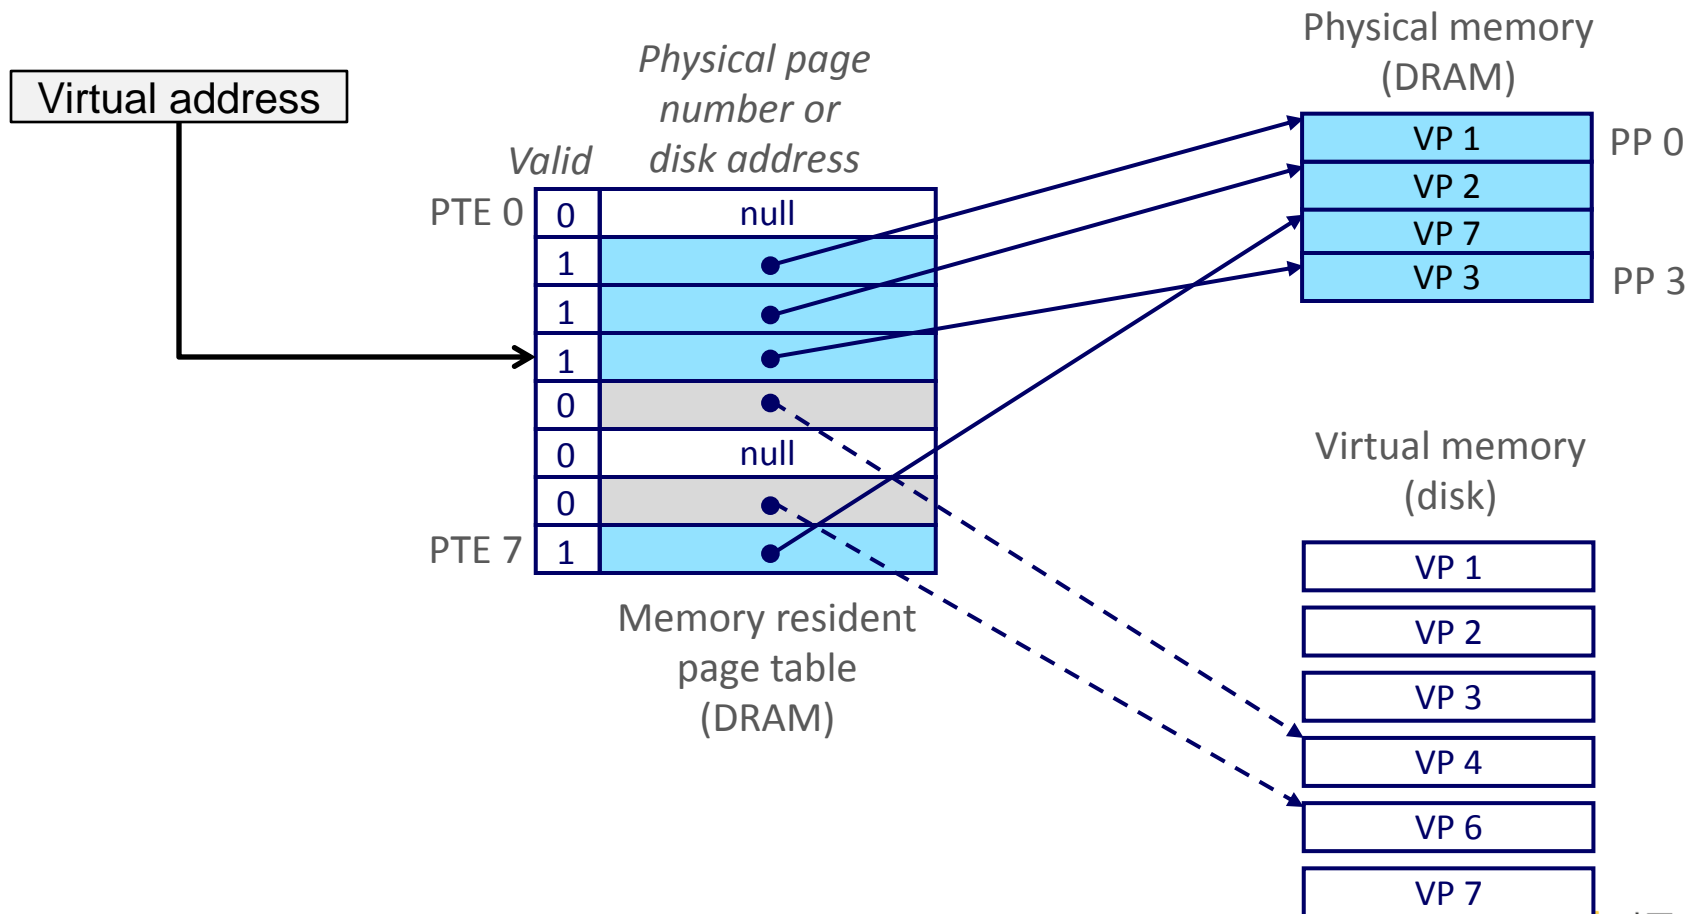
Handling a Page Fault

- Page fault handler selects a victim to be evicted (here VP 4) and loads the requested page into physical memory



Handling Page Fault

- Offending instruction is restarted: page hit!



Why does it work? Locality


- Virtual memory works because of locality
- At any point in time, programs tend to access a set of active virtual pages called the **working set**
 - Programs with better temporal locality have smaller working sets
- If (working set size < main memory size)
 - Good performance for one process after compulsory misses
- If (SUM(working set sizes) > main memory size)
 - **Thrashing**: Performance meltdown where pages are swapped (copied) in and out continuously

Virtual Memory: Theory

- Virtual Memory(VM): Overview and motivation
- VM as a tool for caching
- **VM as a tool for memory management**
- VM as a tool for memory protection
- Address translation
- Allocation

VM as a Tool for Memory Management

- Key idea: each process has its own virtual address space
 - It can view memory as a simple linear array
 - Mapping function scatters addresses through physical memory
 - ▶ Well chosen mappings simplify memory allocation and management




VM as a Tool for Memory Management

■ Memory allocation

- Each virtual page can be mapped to *any* physical page
- A virtual page can be stored in different physical pages at different times

■ Sharing code and data among processes

- Map virtual pages to the same physical page (here: PP 6)



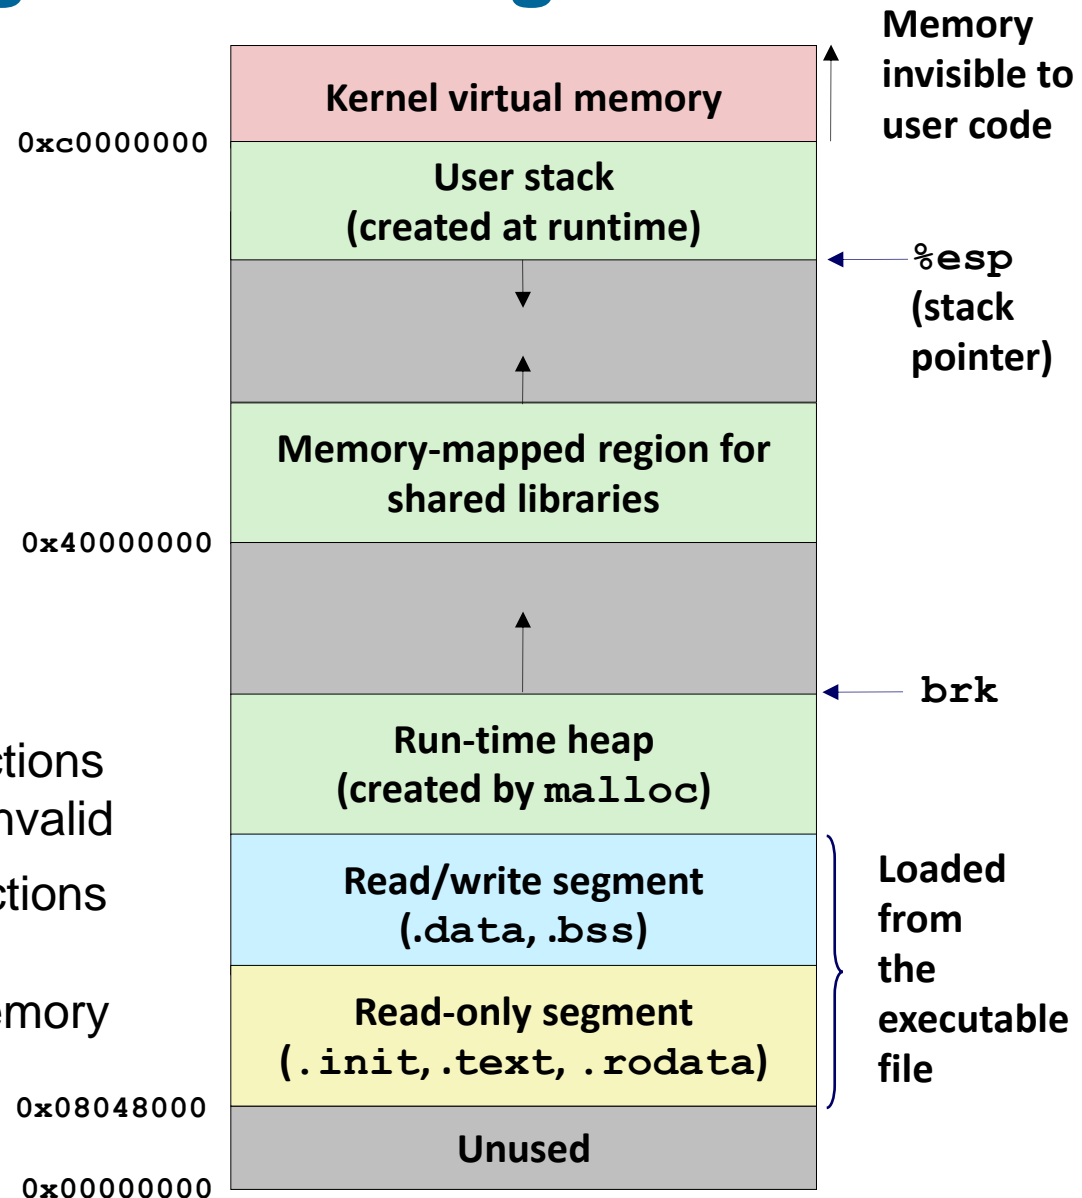
Simplifying Linking and Loading

■ Linking

- Each program has similar virtual address space
- Code, stack, and shared libraries always start at the same address

■ Loading

- `execve()` allocates virtual pages for `.text` and `.data` sections = creates PTEs marked as invalid
- The `.text` and `.data` sections are copied, page by page, on demand by the virtual memory system

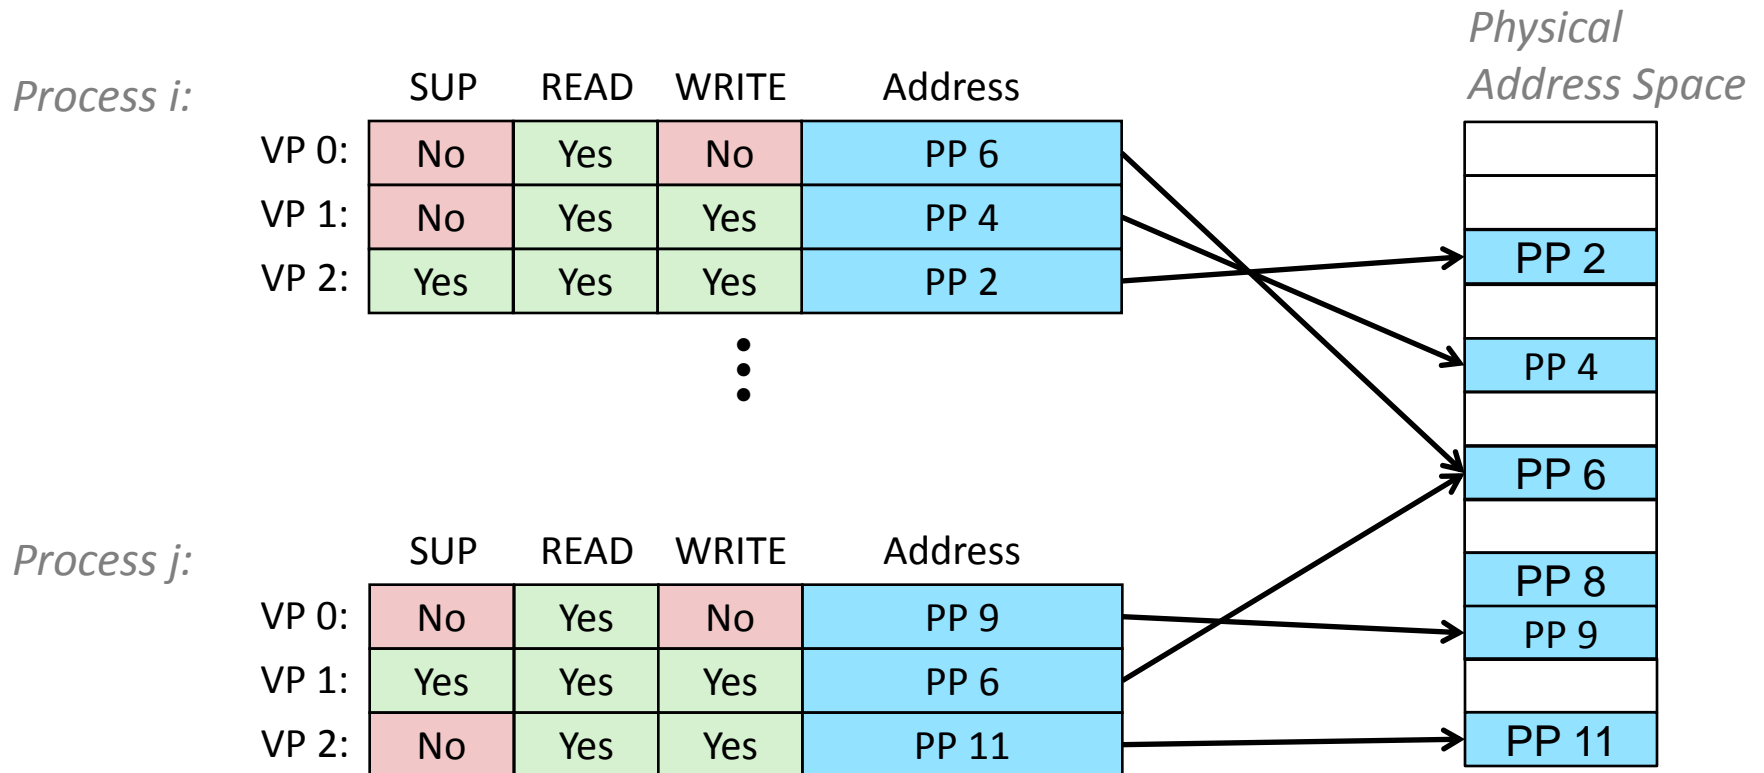


Virtual Memory: Theory

- Virtual Memory(VM): Overview and motivation
- VM as a tool for caching
- VM as a tool for memory management
- **VM as a tool for memory protection**
- Address translation
- Allocation

VM as a Tool for Memory Protection

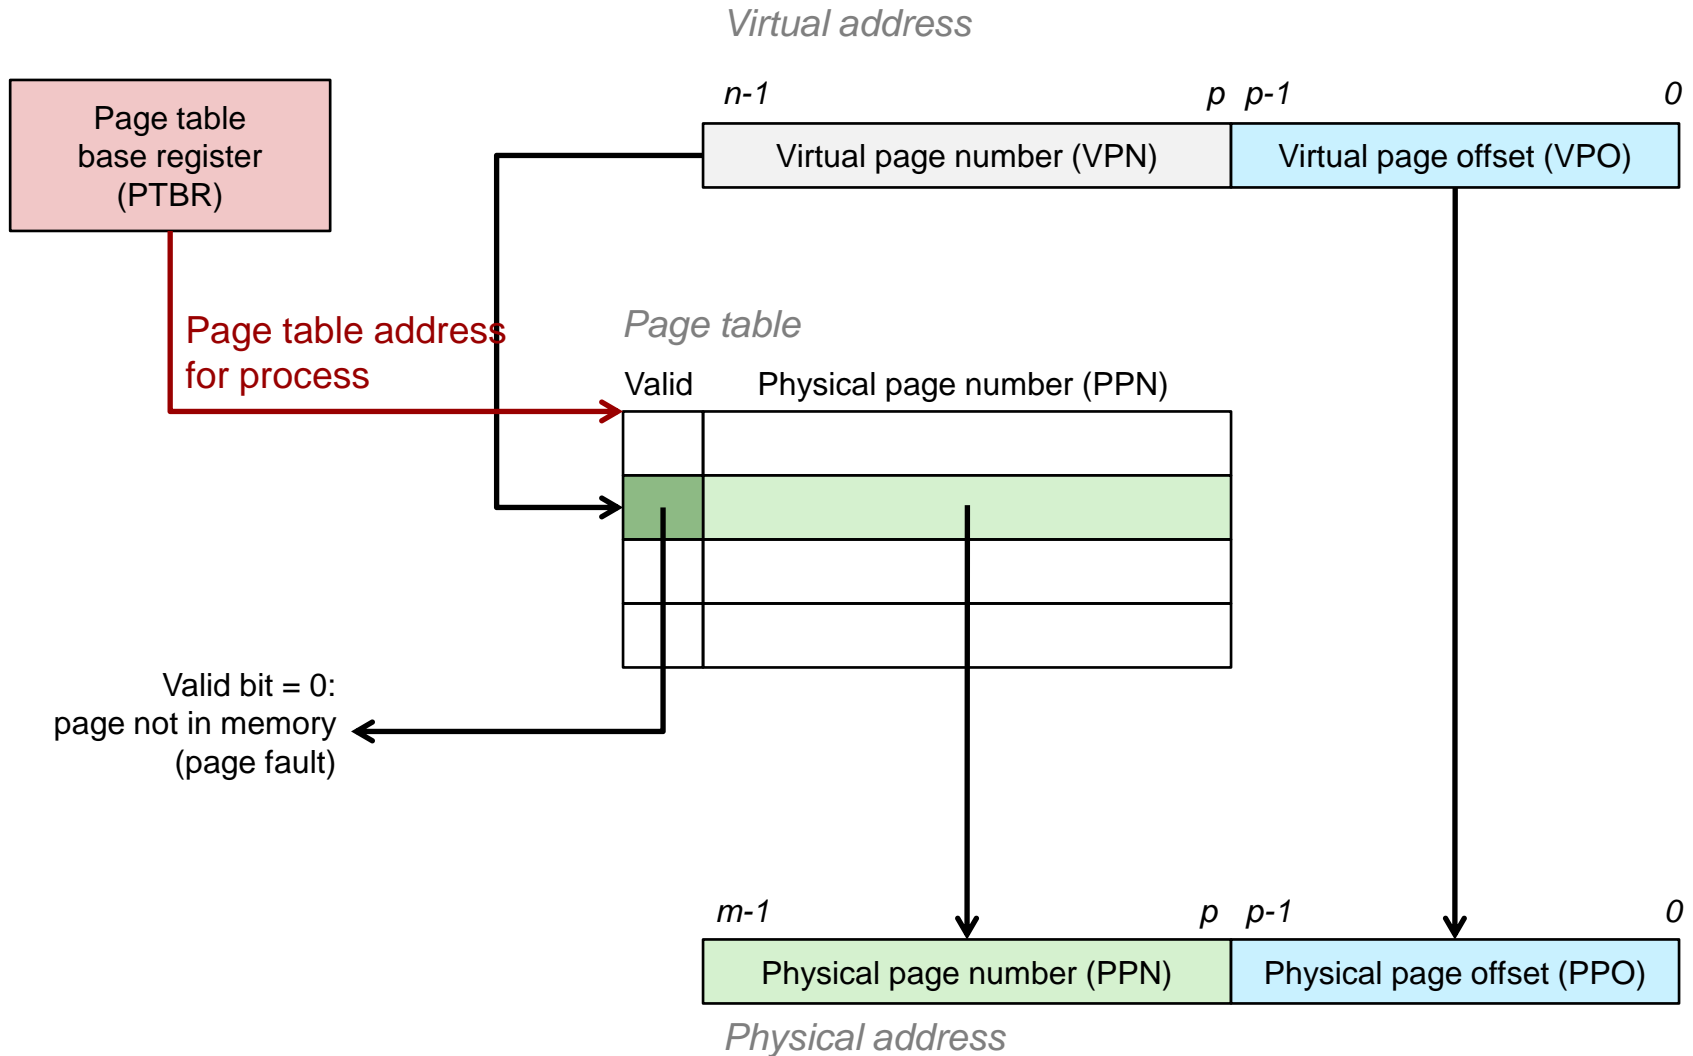
- Extend PTEs with permission bits
- Page fault handler checks these before remapping
 - If violated, send process SIGSEGV (segmentation fault)



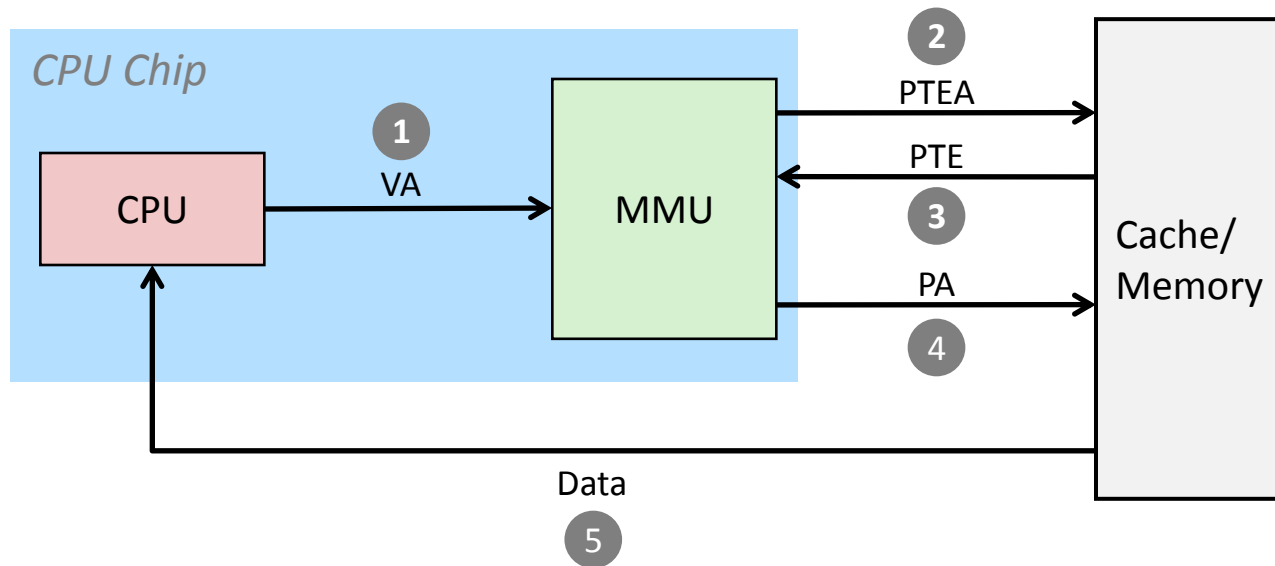
Virtual Memory: Theory

- Virtual Memory(VM): Overview and motivation
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- **Address translation**
- Allocation

Address Translation With a Page Table

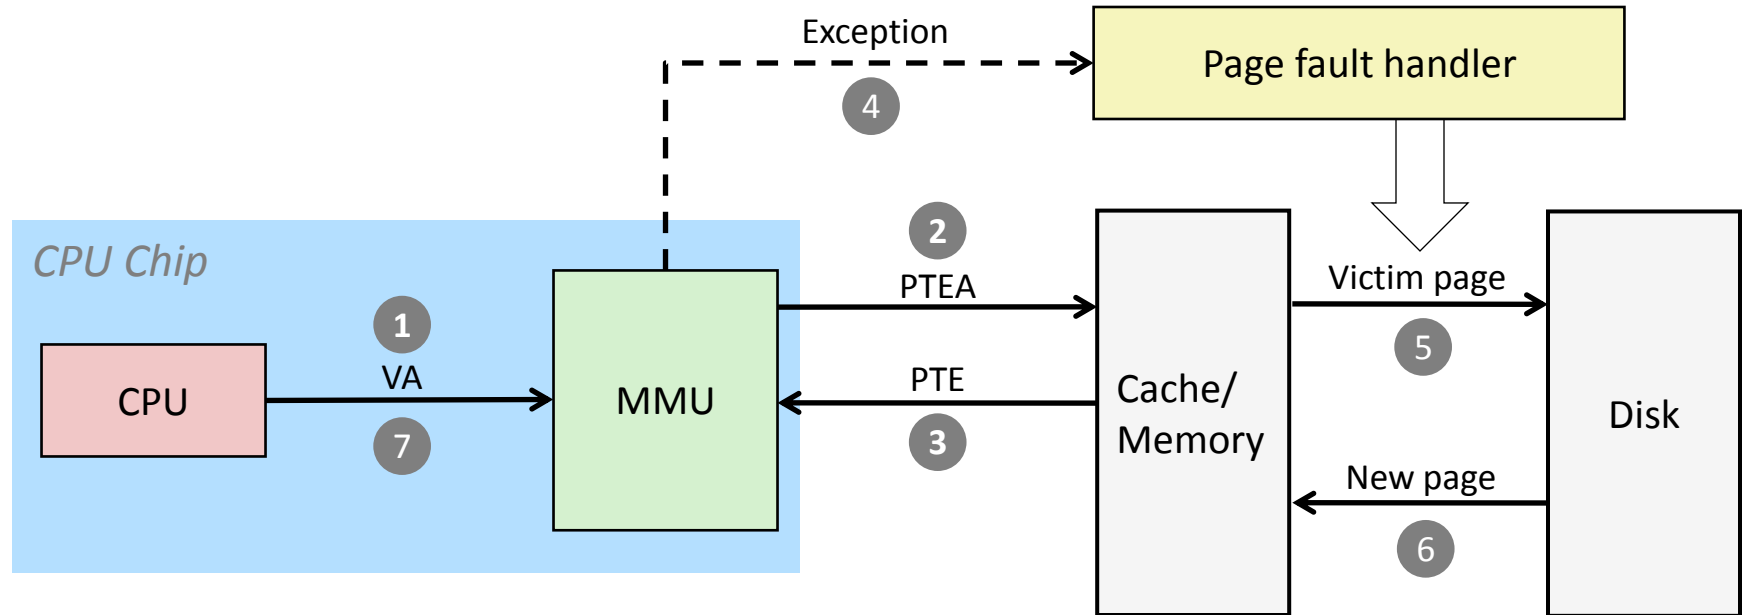


Address Translation: Page Hit



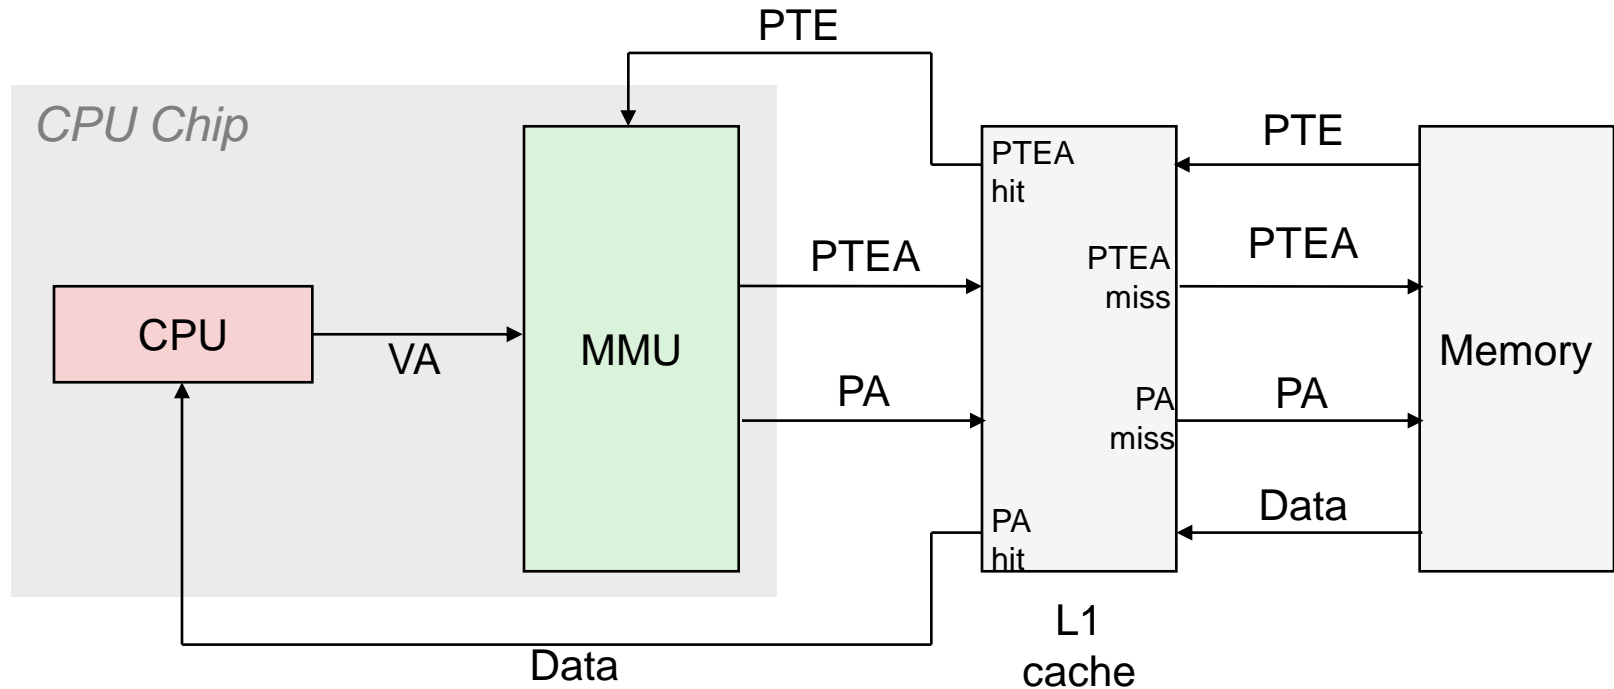
- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) MMU sends physical address to cache/memory
- 5) Cache/memory sends data word to processor

Address Translation: Page Fault



- 1) Processor sends virtual address to MMU
- 2-3) MMU fetches PTE from page table in memory
- 4) Valid bit is zero, so MMU triggers page fault exception
- 5) Handler identifies victim (and, if dirty, pages it out to disk)
- 6) Handler pages in new page and updates PTE in memory
- 7) Handler returns to original process, restarting faulting instruction

Integrating VM and Cache

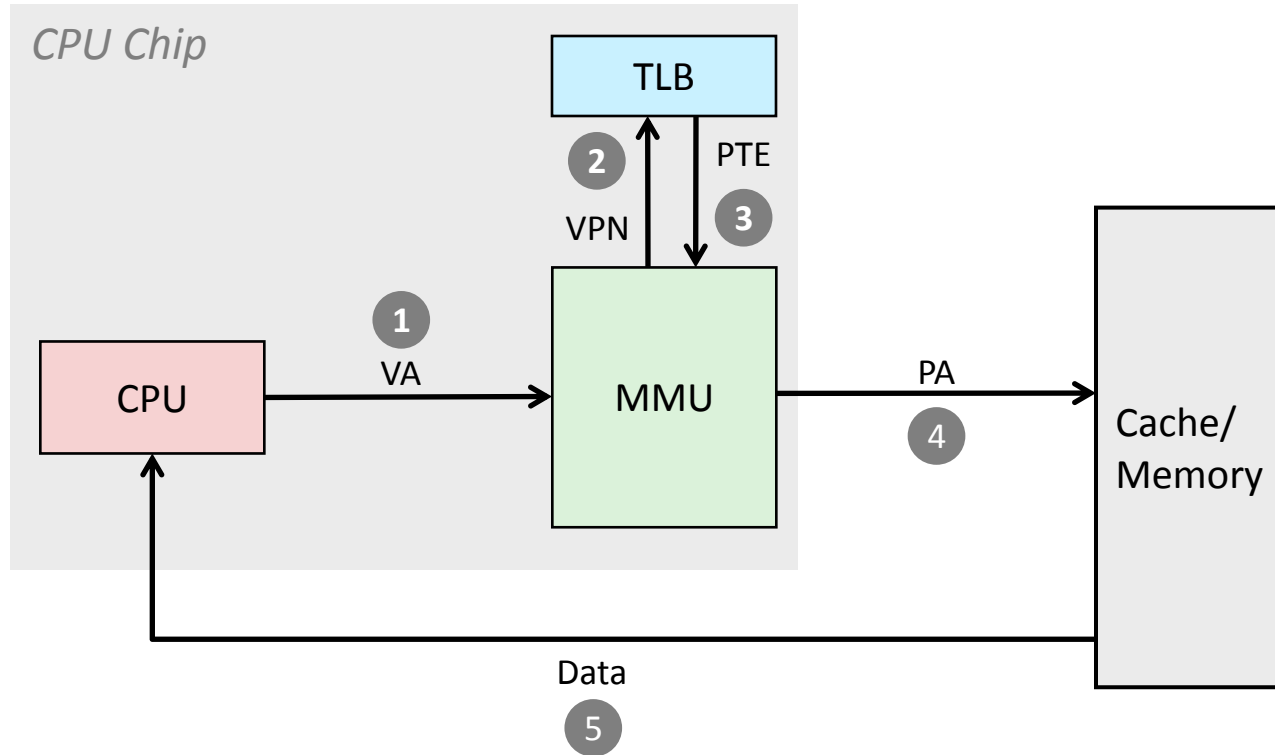


VA: virtual address, PA: physical address, PTE: page table entry, PTEA = PTE address

Speeding up Translation with a TLB

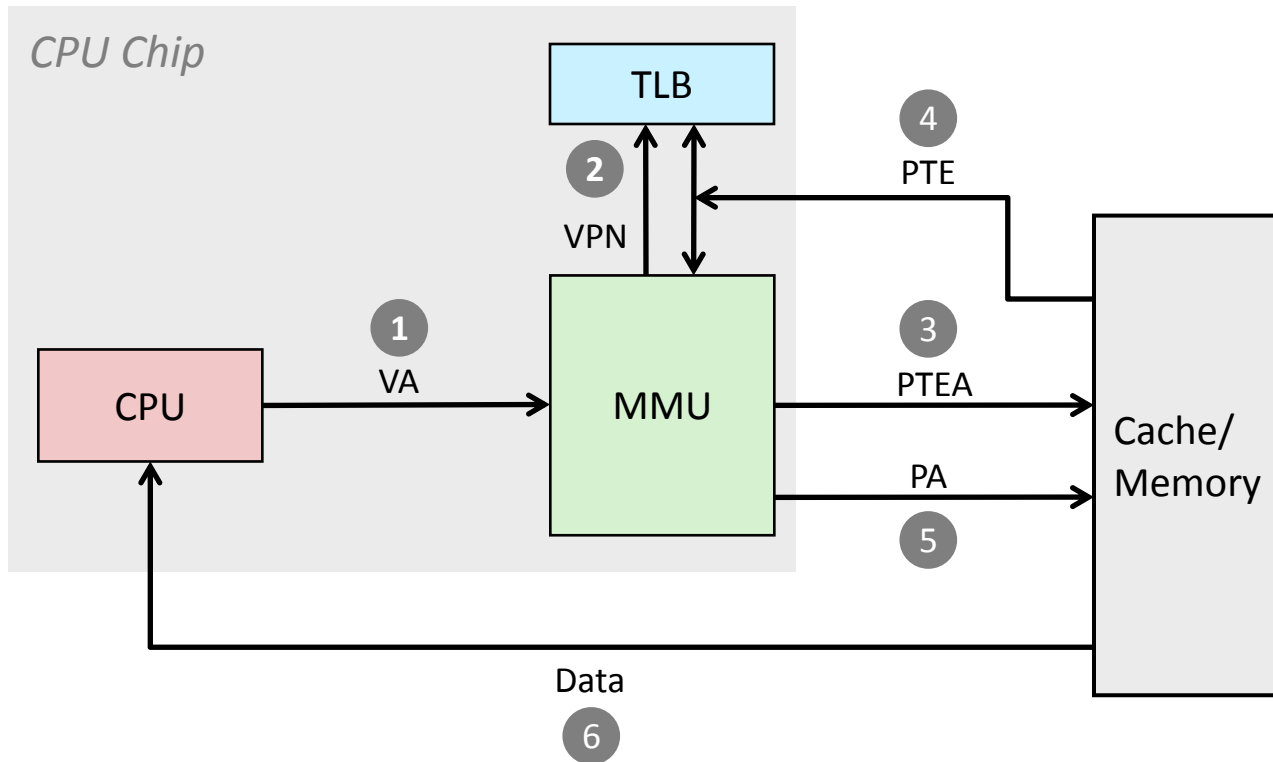
- Page table entries (PTEs) are cached in L1 like any other memory word
 - PTEs may be evicted by other data references
 - PTE hit still requires a small L1 delay
- Solution: **Translation Lookaside Buffer (TLB)**
 - Small hardware cache in MMU
 - Maps virtual page numbers to physical page numbers
 - Contains complete page table entries for small number of pages

TLB Hit



- A TLB hit eliminates a memory access


TLB Miss



- A TLB miss incurs one (or more) additional memory access(es) (the PTE)
 - fortunately, TLB misses are rare.

Virtual Addresses and TLB Index/Tag

- The TLB is a small cache with a high associativity
 - $T = 2^t$ sets
 - TLBI (TLB Index) = t least significant bits of VPN
 - TLBT (TLB Tag) = remaining bits of VPN



Set	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid	Tag	PPN	Valid
0	03	–	0	09	0D	1	00	–	0	07	02	1
1	03	2D	1	02	–	0	04	–	0	0A	–	0
2	02	–	0	08	–	0	06	–	0	03	–	0
3	07	–	0	03	0D	1	0A	34	1	02	–	0

Summary

- Programmer's view of virtual memory
 - Each process has its own private linear address space
 - Cannot be corrupted by other processes
- System view of virtual memory
 - Uses memory efficiently by caching virtual memory pages
 - ▶ Efficient only because of locality
 - Simplifies memory management and programming
 - Simplifies protection by providing a convenient interpositioning point to check permissions

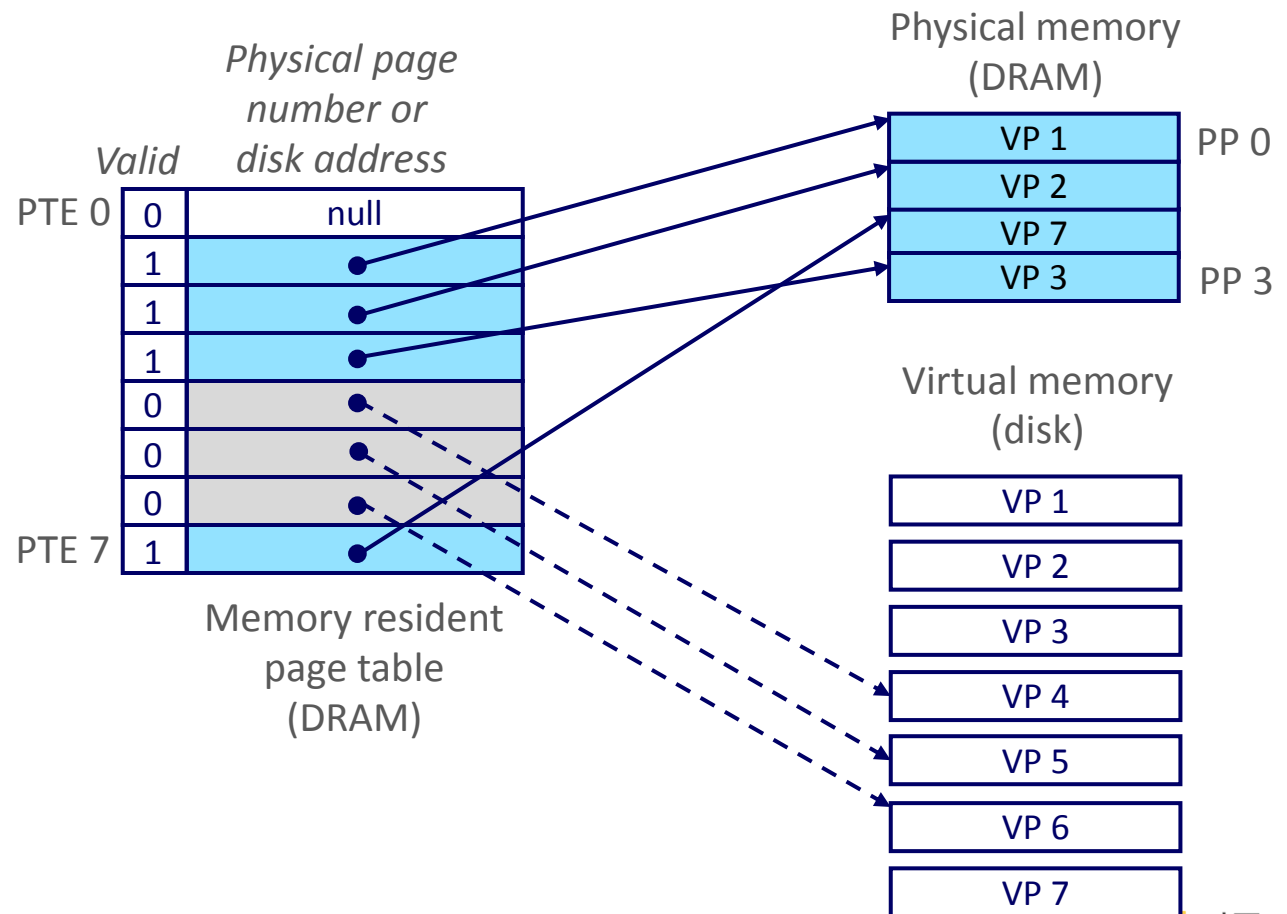
Virtual Memory: Theory

- Virtual Memory(VM): Overview and motivation
- VM as a tool for caching
- VM as a tool for memory management
- VM as a tool for memory protection
- Address translation
- **Allocation**

Allocating Virtual Pages

■ Example: Allocating VP 5

- kernel allocates VP 5 on disk and points PTE 5 to it



Memory System Summary

■ L1/L2 Memory Cache

- purely a speed-up technique
- behavior invisible to application programmer and (mostly) OS
- implemented entirely in hardware

■ Virtual Memory

- supports many OS-related functions
 - ▶ process creation, task switching, protection, sharing
- software
 - ▶ allocates/shares physical memory among processes
 - ▶ maintains high-level tables tracking memory type, source, sharing
 - ▶ handles exceptions, fills in hardware-defined mapping tables
- hardware
 - ▶ translates virtual addresses via mapping tables, enforcing permissions
 - ▶ accelerates mapping via translation cache (TLB)