

4190.308 Computer Architecture, Fall 2014  
Optimizing the Performance of a Pipelined Processor  
Assigned: Tue, Nov. 18, Due: Sun, Nov. 30, 23:59

## 1 Introduction

In this lab, you will learn about the design and implementation of a pipelined Y86 processor, optimizing both it and a benchmark program to maximize performance. You are allowed to make any semantics preserving transformations to the benchmark program, or to make enhancements to the pipelined processor, or both. When you have completed the lab, you will have a keen appreciation for the interactions between code and hardware that affect the performance of your programs.

The lab is organized into two parts, each with its own handin. In Part A, you will extend the SEQ simulator with one new instruction. This part will prepare you for Part B, the heart of the lab, where you will optimize the Y86 benchmark program and the processor design.

## 2 Logistics

You will work on this lab alone.

Any clarifications and revisions to the assignment will be posted on the course Web page.

## 3 Handout Instructions

**Download the `archlab-handout.tar` file from eTL**

1. Start by copying the file `archlab-handout.tar` to a directory in which you plan to do your work.
2. Then give the command: `tar xvf archlab-handout.tar`. This will cause the following file to be unpacked into the directory: `README`, `Makefile`, `sim.tar`, `archlab.ps`, `archlab.pdf`, and `simguide.pdf`.
3. Next, give the command `tar xvf sim.tar`. This will create the directory `sim`, which contains your personal copy of the Y86 tools. You will be doing all of your work inside this directory.
4. Finally, change to the `sim` directory and build the Y86 tools:

```
unix> cd sim
unix> make clean; make
```

## 4 Part A

You will be working in directory `sim/seq` in this part.

Your task in Part A is to extend the SEQ processor to support one new instruction: `iaddl` (described in CS:APP textbook, practice problems 4.48 and 4.50). To add this instruction, you will modify the file `seq-full.hcl`, which implements the version of SEQ described in the textbook. In addition, it contains declarations of some constants that you will need for your solution.

### Building and Testing Your Solution

Once you have finished modifying the `seq-full.hcl` file, then you will need to build a new instance of the SEQ simulator (`ssim`) based on this HCL file, and then test it:

- *Building a new simulator.* You can use `make` to build a new SEQ simulator:

```
unix> make VERSION=full
```

This builds a version of `ssim` that uses the control logic you specified in `seq-full.hcl`. To save typing, you can assign `VERSION=full` in the Makefile.

- *Testing your solution on a simple Y86 program.* For your initial testing, we recommend running a simple program such as `asumi.yo` (testing `iaddl`) in TTY mode, comparing the results against the ISA simulation:

```
unix> ./ssim -t ../y86-code/asumi.yo
```

- *Retesting your solution using the benchmark programs.* Once your simulator is able to correctly execute small programs, then you can automatically test it on the Y86 benchmark programs in `../y86-code`:

```
unix> (cd ../y86-code; make testssim)
```

This will run `ssim` on the benchmark programs and check for correctness by comparing the resulting processor state with the state from a high-level ISA simulation. Note that none of these programs test the added instruction. You are simply making sure that your solution did not inject errors for the original instructions. See file `../y86-code/README` file for more details.

- *Performing regression tests.* Once you can execute the benchmark programs correctly, then you should run the extensive set of regression tests in `../ptest`. To test everything except `iaddl`:

```
unix> (cd ../ptest; make SIM=../seq/ssim)
```

To test your implementation of `iaddl`:

```
unix> (cd ../ptest; make SIM=../seq/ssim TFLAGS=-i)
```

For more information on the SEQ simulator refer to the handout *CS:APP Guide to Y86 Processor Simulators* (`simguide.pdf`).

```

1  /*
2   * ncopy - copy src to dst, returning number of positive ints
3   * contained in src array.
4   */
5 int ncopy(int *src, int *dst, int len)
6 {
7     int count = 0;
8     int val;
9
10    while (len > 0) {
11        val = *src++;
12        *dst++ = val;
13        if (val > 0)
14            count++;
15        len--;
16    }
17    return count;
18 }
```

Figure 1: **C version of the ncopy function.** See sim/pipe/ncopy.c.

## 5 Part B

You will be working in directory sim/pipe in this part.

The ncopy function in Figure 1 copies a len-element integer array `src` to a non-overlapping `dst`, returning a count of the number of positive integers contained in `src`. Figure 2 shows the baseline Y86 version of ncopy. The file `pipe-full.hcl` contains a copy of the HCL code for PIPE, along with a declaration of the constant value `IIADDL`.

Your task in Part C is to modify `ncpy.ys` and `pipe-full.hcl` with the goal of making `ncpy.ys` run as fast as possible.

### Coding Rules

You are free to make any modifications you wish, with the following constraints:

- Your `ncpy.ys` function must work for arbitrary array sizes. You might be tempted to hardwire your solution for 64-element arrays by simply coding 64 copy instructions, but this would be a bad idea because we will be grading your solution based on its performance on arbitrary arrays.
- Your `ncpy.ys` function must run correctly with YIS. By correctly, we mean that it must correctly copy the `src` block *and* return (in `%eax`) the correct number of positive integers.
- Your `pipe-full.hcl` implementation must pass the regression tests in `../y86-code` and `../ptest` (without the `-i1flags` that test `iaddl`).

Other than that, you are free to implement the `iaddl` instruction if you think that will help. You are free to alter the branch prediction behavior or to implement techniques such as load bypassing. You may make any semantics preserving transformations to the `ncpy.ys` function, such as swapping instructions, replacing groups of instructions with single instructions, deleting some instructions, and adding other instructions.

## **Building and Running Your Solution**

In order to test your solution, you will need to build a driver program that calls your `ncopy` function. We have provided you with the `gen-driver.pl` program that generates a driver program for arbitrary sized input arrays. For example, typing

```
unix> make drivers
```

will construct the following two useful driver programs:

- `sdriver.yo`: A *small driver program* that tests an `ncopy` function on small arrays with 4 elements. If your solution is correct, then this program will halt with a value of 2 in register `%eax` after copying the `src` array.
- `ldriver.yo`: A *large driver program* that tests an `ncopy` function on larger arrays with 63 elements. If your solution is correct, then this program will halt with a value of 31 (`0x1f`) in register `%eax` after copying the `src` array.

Each time you modify your `ncopy.ys` program, you can rebuild the driver programs by typing

```
unix> make drivers
```

Each time you modify your `pipe-full.hcl` file you can rebuild the simulator by typing

```
unix> make psim VERSION=full
```

If you want to rebuild the simulator and the driver programs, type

```
unix> make VERSION=full
```

To test your solution on a small 4-element array, type

```
unix> ./psim sdriver.yo
```

To test your solution on a larger 63-element array, type

```
unix> ./psim ldriver.yo
```

```

1 ##########
2 # ncopy.ys - Copy a src block of len ints to dst.
3 # Return the number of positive ints (>0) contained in src.
4 #
5 # Include your name and ID here.
6 #
7 # Describe how and why you modified the baseline code.
8 #
9 #########
10          # Function prologue. Do not modify.
11 ncopy: pushl %ebp           # Save old frame pointer
12         rrmovl %esp,%ebp      # Set up new frame pointer
13         pushl %esi           # Save callee-save regs
14         pushl %ebx
15         mrmovl 8(%ebp),%ebx   # src
16         mrmovl 12(%ebp),%ecx   # dst
17         mrmovl 16(%ebp),%edx   # len
18
19         # Loop header
20         xorl %esi,%esi       # count = 0;
21         andl %edx,%edx       # len <= 0?
22         jle Done              # if so, goto Done:
23
24         # Loop body.
25 Loop:    mrmovl (%ebx), %eax      # read val from src...
26         rmmovl %eax, (%ecx)     # ...and store it to dst
27         andl %eax, %eax       # val <= 0?
28         jle Npos              # if so, goto Npos:
29         irmovl $1, %edi
30         addl %edi, %esi       # count++
31 Npos:    irmovl $1, %edi
32         subl %edi, %edx       # len--
33         irmovl $4, %edi
34         addl %edi, %ebx       # src++
35         addl %edi, %ecx       # dst++
36         andl %edx,%edx       # len > 0?
37         jg Loop               # if so, goto Loop:
38
39         # Function epilogue. Do not modify.
40 Done:   rrmovl %esi, %eax
41         popl %ebx
42         popl %esi
43         rrmovl %ebp, %esp
44         popl %ebp
45         ret

```

Figure 2: **Baseline Y86 version of the ncopy function.** See sim/pipe/ncopy.ys.

Once your simulator correctly runs your version of `ncopy.ys` on these two block lengths, you will want to perform the following additional tests:

- *Testing your driver files on the ISA simulator.* Make sure that your `ncopy.ys` function works properly with YIS:

```
unix> make drivers
unix> ./misc/yis sdriver.yo
```

- *Testing your code on a range of block lengths with the ISA simulator.* The Perl script `correctness.pl` generates driver files with block lengths from 0 up to some limit (default 65), plus some larger sizes. It simulates them (by default with YIS), and checks the results. It generates a report showing the status for each block length:

```
unix> ./correctness.pl
```

If you get incorrect results for some length  $K$ , you can generate a driver file for that length that includes checking code, and where the result varies randomly:

```
unix> ./gen-driver.pl -f ncopy.ys K -rc > driver.ys
unix> make driver.yo
unix> ./misc/yis driver.yo
```

The program will end with register `%eax` having value:

**0xaaaa** : All tests pass

**0xbbbb** : Incorrect count

**0xcccc** : Function `ncpy` is more than 1000 bytes long.

**0xdddd** : Some of the source data was not copied to its destination.

**0xeeee** : Some word just before or just after the destination region was corrupted.

- *Testing your simulator on the benchmark programs.* Once your simulator is able to correctly execute `sdriver.ys` and `ldriver.ys`, you should test it against the Y86 benchmark programs in `../y86-code`:

```
unix> (cd ../y86-code; make testpsim)
```

This will run `psim` on the benchmark programs and compare results with YIS.

- *Testing your pipeline simulator with extensive regression tests.* Once you can execute the benchmark programs correctly, then you should check it with the regression tests in `../ptest`. For example, if your solution implements the `iaddl` instruction, then

```
unix> (cd ../ptest; make SIM=../pipe/psim TFLAGS=-i)
```

## 6 Evaluation

### Part A

This part of the lab is worth 35 points:

- 10 points for your description of the computations required for the `iaddl` instruction.
- 10 points for passing the benchmark regression tests in `y86-code`, to verify that your simulator still correctly executes the benchmark suite.
- 15 points for passing the regression tests in `ptest` for `iaddl`.

### Part B

This part of the Lab is worth 80 points:

- 20 points each for your descriptions in the headers of `ncopy.ys` and `pipe-full.hcl`.
- 60 points for performance. To receive credit here, your solution must be correct, as defined earlier. That is, `ncopy` runs correctly with `YIS`, and `pipe-full.hcl` passes all tests in `y86-code` and `ptest`.

We will express the performance of your function in units of *cycles per element* (CPE). That is, if the simulated code requires  $C$  cycles to copy a block of  $N$  elements, then the CPE is  $C/N$ . The PIPE simulator displays the total number of cycles required to complete the program. The baseline version of the `ncopy` function running on the standard PIPE simulator with a large 63-element array requires 1037 cycles to copy 63 elements, for a CPE of  $1037/63 = 16.46$ .

Since some cycles are used to set up the call to `ncopy` and to set up the loop within `ncopy`, you will find that you will get different values of the CPE for different block lengths (generally the CPE will drop as  $N$  increases). We will therefore evaluate the performance of your function by computing the average of the CPEs for blocks ranging from 1 to 64 elements. You can use the Perl script

`benchmark.pl` in the `pipe` directory to run simulations of your `ncopy.ys` code over a range of block lengths and compute the average CPE. Simply run the command

```
unix> ./benchmark.pl
```

to see what happens. For example, the baseline version of the `ncopy` function has CPE values ranging between 45.0 and 16.45, with an average of 18.15. Note that this Perl script does not check for the correctness of the answer. Use the script `correctness.pl` for this.

You should be able to achieve an average CPE of less than 12.0. Our best version averages 7.43.

By default, `benchmark.pl` and `correctness.pl` compile and test `ncopy.ys`. Use the `-f` argument to specify a different file name. The `-h` flag gives a complete list of the command line arguments.

## 7 Handin Instructions

- You will be handing in three files
  - Part A: seq-full.hcl.
  - Part B: ncopy.ys and pipe-full.hcl.
- Make a tarball file containing the above files
- Send the tar file in an email ([comparch-ta@csap.snu.ac.kr](mailto:comparch-ta@csap.snu.ac.kr))  
with the Subject Line **[Archlab] #student ID #name (ex. [Archlab] 2012-20798 Daeyong Shin)**

## 8 Hints

refer to archlab.hints.pdf (on eTL) and README files for each directory (in a archlab-handout.tar file)