

Intro to DB

CHAPTER 4

INTERMEDIATE SQL

Chapter 4: Intermediate SQL

- Join Expressions
- Views
- Transactions
- Integrity Constraints
- SQL Data Types and Schemas
- Authorization

Joined Relations

- **Join operation**
 - takes two relations and return as a result another relation.
 - a Cartesian product which requires that tuples in the two relations match (under some condition).
 - also specifies the attributes that are present in the result of the join
- The join operations are typically used as subquery expressions in the **from** clause

Join Operation

- Join matches tuples with the same values specified in the on condition.

```
select *
from instructor join teaches on (instructor.ID=teaches.ID);
```

ID	name	dept_name	salary	ID	course_id	sec_id	semester	year
10101	Srinivasan	Comp. Sci.	65000	10101	CS-101	1	Fall	2009
12121	Wu	Finance	90000	10101	CS-315	1	Spring	2010
15151	Mozart	Music	40000	10101	CS-347	1	Fall	2009
22222	Einstein	Physics	95000	12121	FIN-201	1	Spring	2010
32343	El Said	History	60000	15151	MU-199	1	Spring	2010
45565	Katz	Comp. Sci.	75000	22222	PHY-101	1	Fall	2009
76766	Crick	Biology	72000	32343	HIS-351	1	Spring	2010
76766	Crick	Biology	72000	45565	CS-101	1	Spring	2010
				45565	CS-319	1	Spring	2010
				76766	BIO-101	1	Summer	2009
				76766	BIO-301	1	Summer	2010

Join Operation (cont.)

- Let $r_1(A_1, \dots, A_k), r_2(B_1, \dots, B_n)$
select *
from r_1 **join** r_2 **on** *condition*;

is equivalent to



- You can use expressions (that evaluate to relations) in **from** clauses
from E_1, \dots, E_n

select distinct *name, title*
from *instructor* **join** *teaches* **on** (*instructor.ID=teaches.ID*), *course*
where *teaches.course_id=course.course_id*

Joined Relations

- take two relations and return as a result another relation.
- typically used as subquery expressions in the **from** clause
- **Join condition** – defines which tuples in the two relations match, and what attributes are present in the result of the join.
- **Join type** – defines how tuples in each relation that do not match any tuple in the other relation (based on the join condition) are treated.

<i>Join types</i>	<i>Join Conditions</i>
inner join left outer join right outer join full outer join	natural on <predicate> using (A_1, A_1, \dots, A_n)

Inner join

- Inner join = join (default)

select * from *instructor* join *teaches* on (*instructor.ID*=*teaches.ID*);

= select * from *instructor* inner join *teaches* on (*instructor.ID*=*teaches.ID*);

- ON clause

- an explicit join clause (acts like a **where** clause)

- USING clause

- specifies which columns to test for *equality*

-

- Above example is equivalent to

select * from *instructor* inner join *teaches* using (*ID*);

Natural Join


- (Section 3.3.3)
- *Natural join* matches tuples with the same values for all common attributes, and retains only one copy of each common column

select *
from *instructor* **natural join** *teaches*;

- Let $r_1(A_1, \dots, A_k, C_1, \dots, C_m)$, $r_2(B_1, \dots, B_n, C_1, \dots, C_m)$
select *
from r_1 **natural join** r_2

is equivalent to

select $A_1, \dots, A_k, r_1.C_1, \dots, r_1.C_m, B_1, \dots, B_n$




Natural Join (cont.)

- **Danger in natural join:** beware of unrelated attributes with same name which get equated incorrectly
- List the names of instructors along with the titles of courses that they teach
 - Incorrect version (makes `course.dept_name = instructor.dept_name`)
 - **select distinct** *name, title*
from *instructor* **natural join** *teaches* **natural join** *course*;
 - Correct version
 - **select distinct** *name, title*
from *instructor* **natural join** *teaches, course*
where *teaches.course_id = course.course_id*;
 - Another correct version
 - **select distinct** *name, title*
from (*instructor* **natural join** *teaches*)
join *course* **using** (*course_id*);

Natural Join (cont.)

- Find the IDs of all students who were taught by an instructor named Einstein in building 301 (remove duplicates in the result).

```
select    distinct takes.ID
from      takes, section, teaches, instructor
where      takes.course_id=section.course_id and takes.sec_id=section.sec_id and
             takes.semester=section.semester and takes.year=section.year and
             section.course_id=teaches.course_id and section.sec_id=teaches.sec_id and
             section.year=teaches.year and section.year=teaches.year and
             teaches.ID=instructor.ID and
             instructor.name='Einstein' and building='301'
```



Join operation – Example

course

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

■ *course* **natural join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

Note: prereq information missing for CS-315 and
course information missing for CS-437.

Outer Join

- An extension of the join operation that avoids loss of information.
- Computes the join and then adds tuples from one relation that does not match tuples in the other relation to the result of the join.
- Uses *null* values.

Left Outer Join

- *course* **natural left outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>

= *course* **left outer join** *prereq* **on** *course.course_id* = *prereq.course_id*

= *course* **left outer join** *prereq* **using** (*course_id*)

course

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

Right Outer Join

- *course* **natural right outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

course

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

Full Outer Join

- *course* **natural full outer join** *prereq*

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>
BIO-301	Genetics	Biology	4	BIO-101
CS-190	Game Design	Comp. Sci.	4	CS-101
CS-315	Robotics	Comp. Sci.	3	<i>null</i>
CS-347	<i>null</i>	<i>null</i>	<i>null</i>	CS-101

course

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>
BIO-301	Genetics	Biology	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3

prereq

<i>course_id</i>	<i>prereq_id</i>
BIO-301	BIO-101
CS-190	CS-101
CS-347	CS-101

Joined Relations – Examples

- **course inner join prereq on**

course.course_id = prereq.course_id

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190

- **course left outer join prereq on**

course.course_id = prereq.course_id

<i>course_id</i>	<i>title</i>	<i>dept_name</i>	<i>credits</i>	<i>prere_id</i>	<i>course_id</i>
BIO-301	Genetics	Biology	4	BIO-101	BIO-301
CS-190	Game Design	Comp. Sci.	4	CS-101	CS-190
CS-315	Robotics	Comp. Sci.	3	<i>null</i>	<i>null</i>

Integrity Constraints

- Integrity constraints guard against accidental damage to the database

-

- Examples

- A checking account must have a balance greater than \$10,000.00
 - A salary of a bank employee must be at least \$4.00 an hour
 - A customer must have a (non-null) phone number

- Constraints on a single relation

- **not null**
 - **unique**
 - **primary key**
 - **check** (P), where P is a predicate

- Referential integrity

Not Null Constraint

- Declare *name* for *budget* to be **not null**

name **varchar(20) not null**

budget **numeric(12,2) not null**

- Example



The Unique & Primary Key Constraints

- **unique** (A_1, A_2, \dots, A_m)

- **primary key** (A_1, A_2, \dots, A_m)

- A_1, A_2, \dots, A_m form the primary key for the table

```
create table instructor (  
    ID          char(5),  
    name       varchar(20) not null,  
    dept_name varchar(20),  
    salary    numeric(8,2)  
    primary key (ID) );
```

The check clause

- **check** (P), where P is a predicate
- Example:
 - ensure that semester is one of fall, winter, spring or summer

```
create table section (  
    course_id varchar (8),  
    sec_id varchar (8),  
    semester varchar (6),  
    year numeric (4,0),  
    building varchar (15),  
    room_number varchar (7),  
    time slot id varchar (4),  
    primary key (course_id, sec_id, semester, year),  
    check (semester in ('Fall', 'Winter', 'Spring', 'Summer'))  
)
```

Referential Integrity

- “A value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.”
 - Example:
 - If “Biology” is a department name appearing in one of the tuples in the *course* relation,
 - then there exists a tuple in the *department* relation for “Biology”.
- specified as part of **create table** statement
- **Foreign Key**
 - the attributes that comprise the foreign key, and
 - the name of the relation referenced by the foreign key

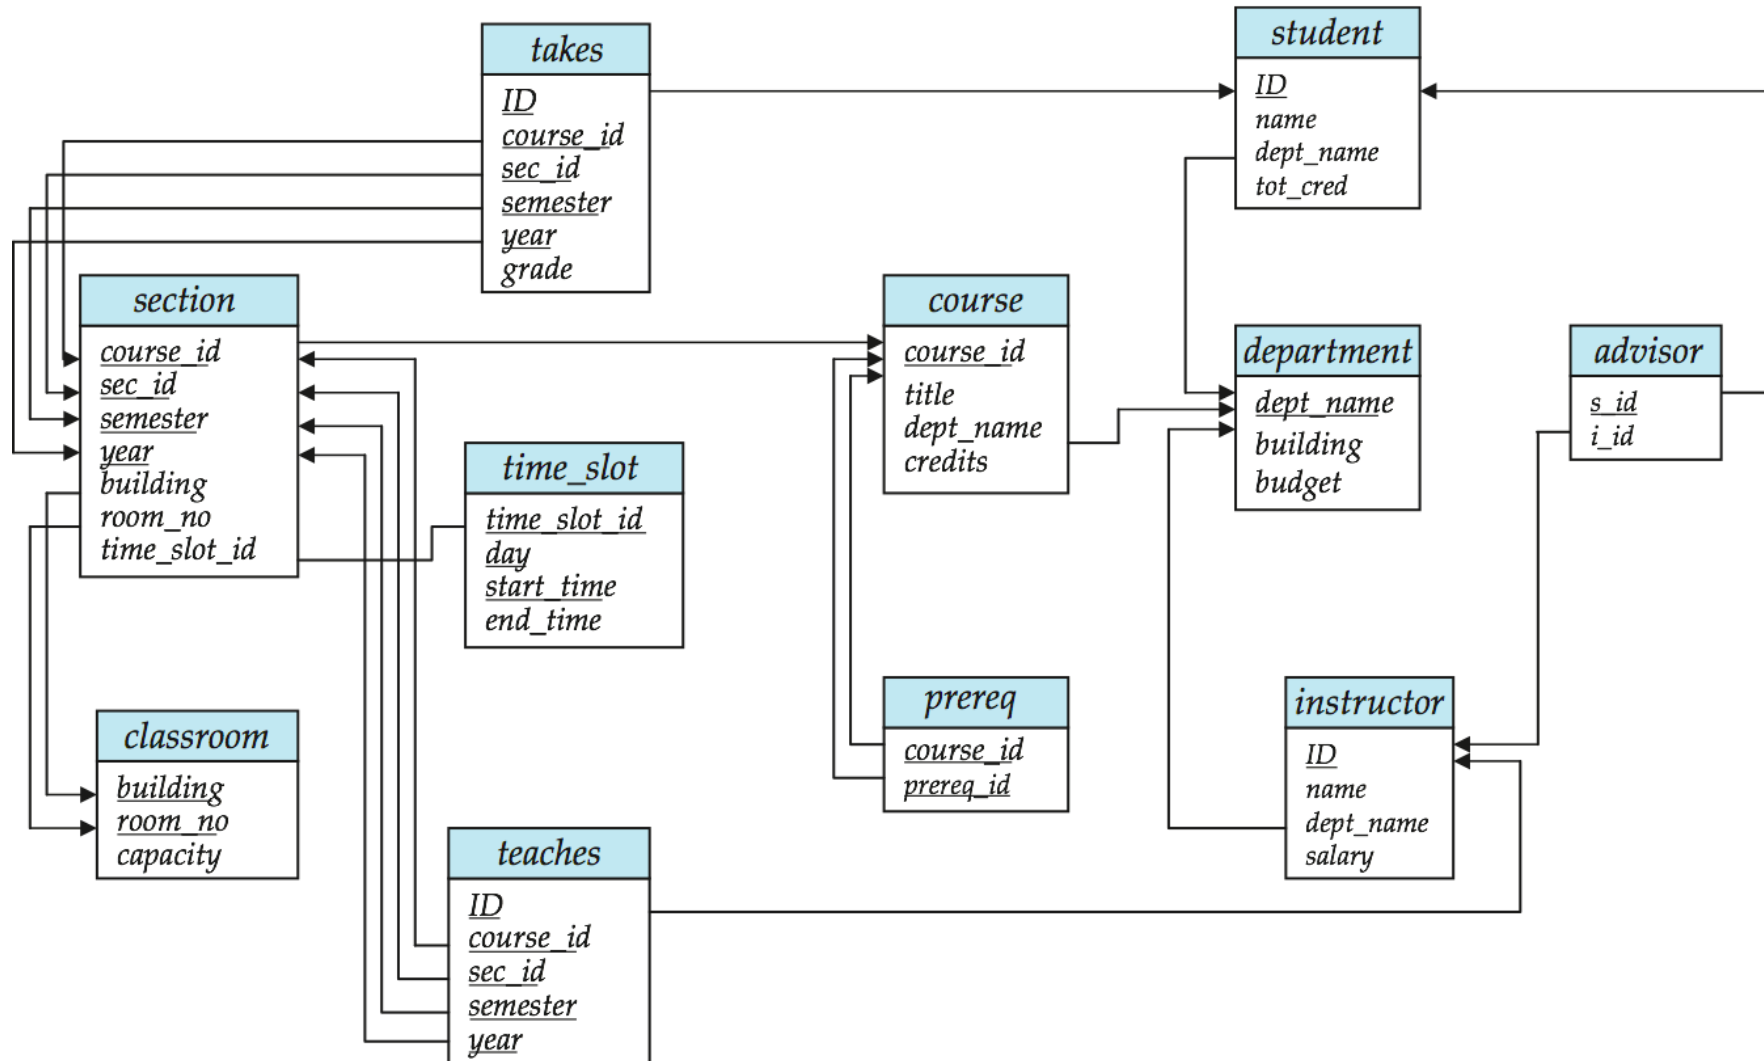
course_id	title	dept_name	credits
BIO-101	Intro. to Biology	Biology	4
BIO-301	Genetics	Biology	4
BIO-399	Computational Biology	Biology	3
CS-101	Intro. to Computer Science	Comp. Sci.	4
CS-190	Game Design	Comp. Sci.	4
CS-315	Robotics	Comp. Sci.	3
CS-319	Image Processing	Comp. Sci.	3
CS-347	Database System Concepts	Comp. Sci.	3
EE-181	Intro. to Digital Systems	Elec. Eng.	3
FIN-201	Investment Banking	Finance	3

dept_name	building	budget
Biology	Watson	90000
Comp. Sci.	Taylor	100000
Elec. Eng.	Taylor	85000
Finance	Painter	120000
History	Painter	50000
Music	Packard	80000
Physics	Watson	70000

Cascading Actions in Referential Integrity

- **create table** *course* (
 course_id **char**(5) **primary key**,
 title **varchar**(20),
 dept_name **varchar**(20) **references** *department*
)
- **create table** *course* (
 ...
 dept_name **varchar**(20),
 foreign key (*dept_name*) **references** *department*
 on delete cascade
 on update cascade,
 ...
)
-


Primary & Foreign Keys in Schema Diagram



Authorization Specification in SQL


- The **grant** statement is used to confer authorization
 grant <privilege list>
 on <relation name or view name>
 to <user list>
- <user list> is:
 - a user-id
 - **public**, which allows all valid users the privilege granted
 - a role (explained later)

▪



Privileges in SQL

- **select**: allows read access to relation
 - (or the ability to query using the view)
 - Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *branch* relation:



- **insert**: the ability to insert tuples
- **update**: the ability to update using the SQL update statement
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges
- **index**: allows creation and deletion of indices.
- **resources**: allows creation of new relations.
- **alteration**: allows addition or deletion of attributes in a relation.
- **drop**: allows deletion of relations.

Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.

revoke <privilege list>

on <relation name or view name>


from <user list>

- Example:

revoke select on *branch* from U_1, U_2, U_3

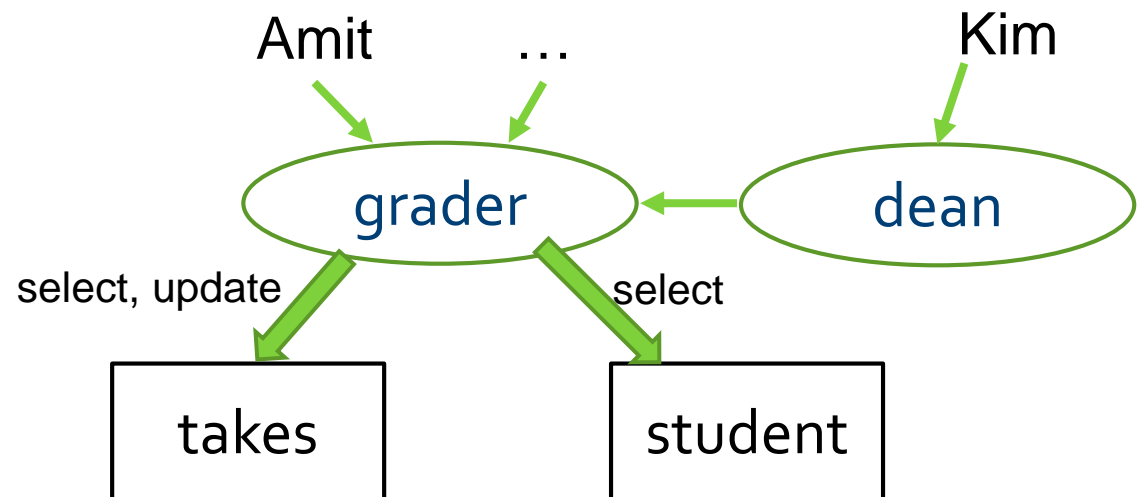
- <privilege-list> may be **all** to revoke all privileges the revokee may hold.
- All privileges that depend on the privilege being revoked are also revoked.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.

Authorization-Grant Graph



Roles

- Roles are used to represent a group of users and their privileges
 - **create role** *grader*;
- Privileges can be granted to roles:
 - **grant select on** *student* **to** *grader*
 - **grant select, update on** *takes* **to** *grader*
- Roles can be granted to users, as well as to other roles
 - **grant** *grader* **to** Amit;
 - **create role** *dean*;
 - **grant** *grader* **to** *dean*;
 - **grant** *dean* **to** Kim;



END OF CHAPTER 4