

Intro to DB

# CHAPTER 12

# QUERY PROCESSING

# Chapter 12: Query Processing

- Overview
- Measures of Query Cost
- Selection Operation
- Sorting
- Join Operation
- Other Operations
- Evaluation of Expressions
- Cost Estimation of Expressions (Chap. 13)

# Basic Steps in Query Processing

## 1. Parsing and translation

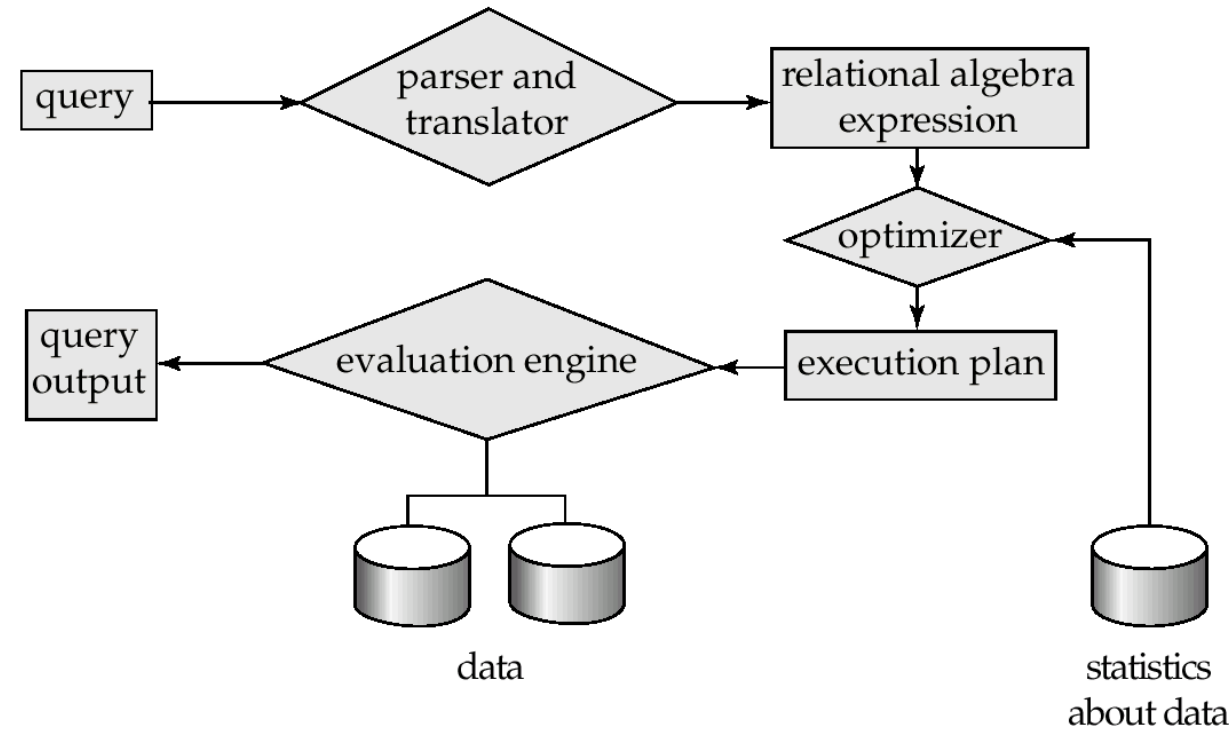
- translate the query into an internal form (eg., relational algebra)
- Parser checks syntax, verifies relations

## 2. Optimization

- More than one way to evaluate a query


## 3. Evaluation

- The query-execution engine takes a query-evaluation plan, executes that plan, and returns the answers to the query.



# Query Plan

- Evaluation primitive
  - a relational algebra expression annotated with instructions on how to evaluate it
    - $\sigma_{balance>2500}(account)$ : use index 1
    - $\sigma_{balance>2500}(account)$ : use table scan
- - a sequence of primitive operations that can be used to evaluate a query
- Example:
  - SELECT *balance* FROM *account* WHERE *balance*>2500



# Query Optimization

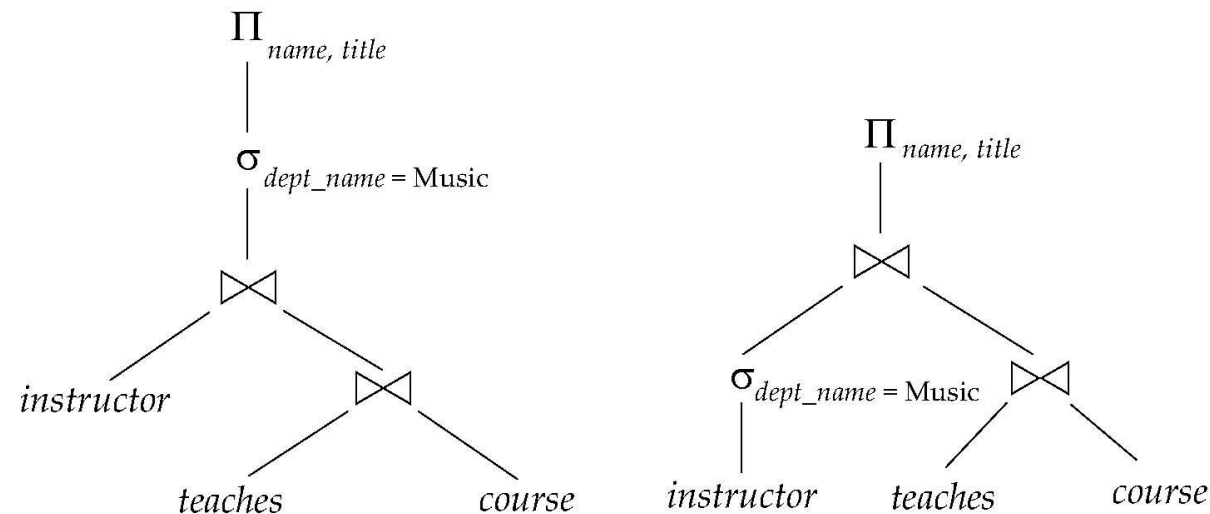
- More than one way to evaluate a query

- 

Given a DB schema  $S$ , a query  $Q$  on  $S$  is equivalent to another query  $Q'$  on  $S$ , if the answer sets of  $Q$  and  $Q'$  are the same in *any* instances of the DB.

$\Pi_{name, title}(\sigma_{dept="Music"}(instructor \bowtie (teaches \bowtie course)))$  vs

$\Pi_{name, title}((\sigma_{dept="Music"}(instructor)) \bowtie (teaches \bowtie course))$



# Query Optimization

- Query optimization is the process of *selecting the most efficient query evaluation plan* for a given query
- Generation of Evaluation Plan
  1. Generating logically equivalent expressions
    - Use *equivalence rules* to transform an expression into an equivalent one.
  2. Annotating resulting expressions to get alternative query plans
  3. Choosing the cheapest plan based on *estimated cost*
- The overall process is called

# Measures of Query Cost

- Cost is generally measured as total elapsed time for answering query
  - Many factors contribute to time cost
    - *disk accesses, CPU, or even network communication*
- - also relatively easy to estimate.
  - # of seeks  $\times$  average-seek-cost
  - # of blocks read  $\times$  average-block-read-cost
  - # of blocks written  $\times$  average-block-write-cost

# Measures of Query Cost (Cont.)

- For simplicity, we only use
  - **number of block transfers** *from disk*
    - $t_T$  – time to transfer one block: 10~40 ms
  - **number of seeks**
    - $t_S$  – time for one seek: 8~20 ms (disk seek + rotational delay)
  - Cost for  $b$  block transfers plus  $s$  seeks
- We ignore CPU costs
- We do not include cost of writing final output to disk
  - output of an operation may be sent to the parent operation without being written to disk



# Selection Operation

- File scan
  - locate and retrieve records that fulfill a selection condition
  - Full file scan: retrieve all records of a file (relation)
- A1: *linear search*
  - Scan each file block and test all records on the selection condition
  - Cost estimate (number of disk blocks scanned) = 
    - $b_r$ : # of blocks containing records from relation  $r$
  - If selection is on a key attribute, cost = 
    - stop on finding record
  - Linear search can be applied regardless of
    - selection condition or
    - ordering of records in the file, or
    - availability of indices

# Selections Using Indices

- **Index scan** – search algorithms that use an index
  - selection condition must be on search-key of index.
- **A2 (primary index, equality on key)**. Retrieve a single record that satisfies the corresponding equality condition
  - $Cost = (h_i + 1) * (t_T + t_S)$
- **A3 (primary index, equality on nonkey)** Retrieve multiple records.
  - Records will be on consecutive blocks
    - Let  $b$  = number of blocks containing matching records
  - $Cost = h_i * (t_T + t_S) + t_S + t_T * b$

# Selections Using Indices

- **A4 (secondary index, equality).**
  - Retrieve a single record if the search-key is a **candidate key**
    - $Cost =$
  - Retrieve multiple records if search-key is **not a candidate key**
    - each of  $n$  matching records may be on a different block
    - $Cost =$  
      - Can be very expensive!

# Selections Involving Comparisons

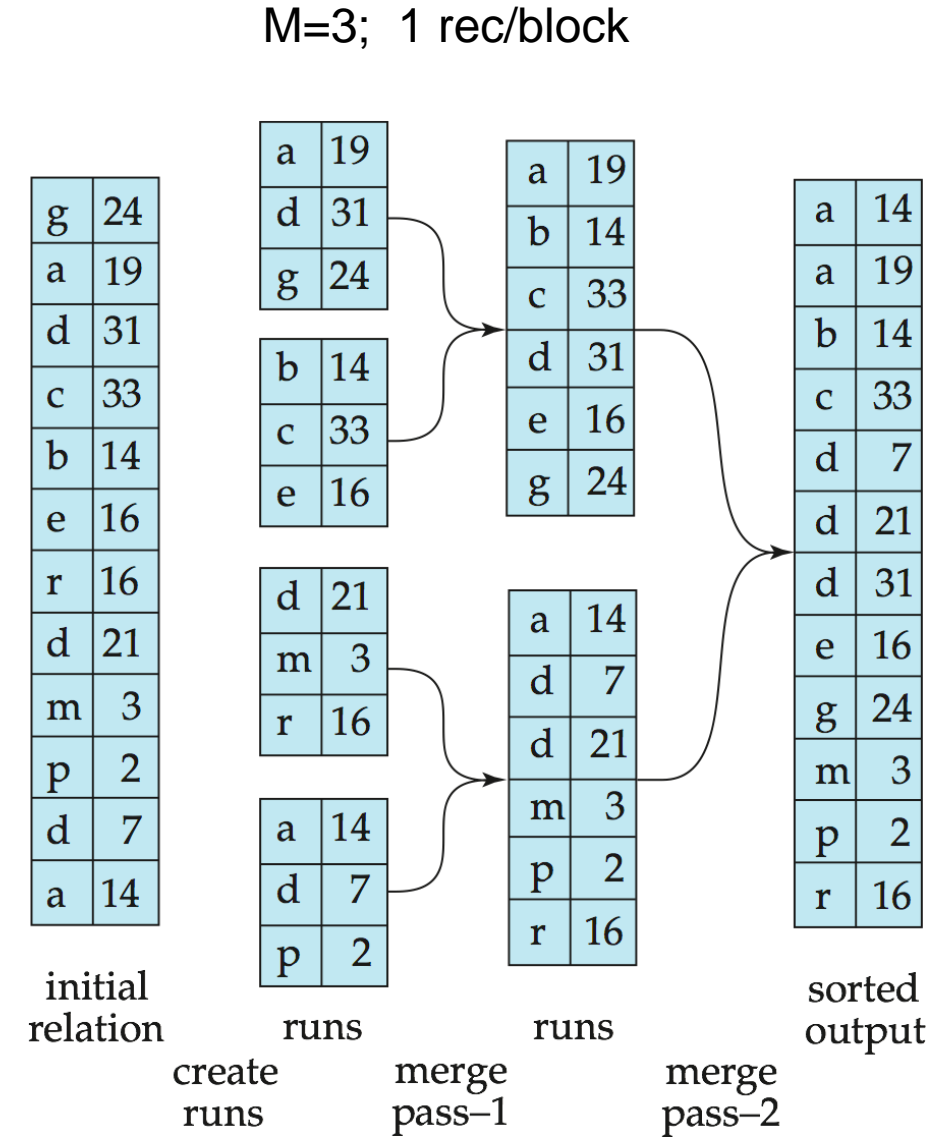
- Can implement selections of the form  $\sigma_{A \leq V}(r)$  or  $\sigma_{A \geq V}(r)$  by using
  - a linear file scan,
  - or by using indices in the following ways:
- **A5 (primary index, comparison)**. (Relation is sorted on A)
  - For  $\sigma_{A \geq V}(r)$  use index to find first tuple  $\geq v$  and scan relation sequentially from there
  - For  $\sigma_{A \leq V}(r)$  just scan relation sequentially till first tuple  $> v$ ; do not use index

# Selections Involving Comparisons (Cont.)

- Can implement selections of the form  $\sigma_{A \leq V}(r)$  or  $\sigma_{A \geq V}(r)$  by using
  - a linear file scan,
  - or by using indices in the following ways:
- **A6 (secondary index, comparison).**
  - For  $\sigma_{A \geq V}(r)$  use index to find first index entry  $\geq v$  and scan index sequentially from there, to find pointers to records.
  - For  $\sigma_{A \leq V}(r)$  just scan leaf pages of index finding pointers to records, till first entry  $> v$
  - In either case, retrieve records that are pointed to
    - requires an I/O for each record
    - Linear file scan may be cheaper

# External Sort-Merge

- **Sorting**
  - Important core operation: order by, group by, join, ...
  - One option: build an index, and use the index to read the relation in sorted order. May lead to one disk block access for each tuple.
- **External Sort-Merge**
  - Good choice for relations that don't fit in memory



# External Sort-Merge

- Let  $M$  denote memory size (in pages).

## 1. Create sorted **runs**. Let $i$ be 0 initially.

Repeatedly do the following till the end of the relation:

- Read  $M$  blocks of relation into memory
- Sort the in-memory blocks
- Write sorted data to run  $R_i$ ; increment  $i$

Let the final value of  $i$  be  $N$

$M=3$ ; 1 rec/block

g	24
a	19
d	31
c	33
b	14
e	16
r	16
d	21
m	3
p	2
d	7
a	14

initial  
relation

a	19
d	31
g	24

b	14
c	33
e	16

d	21
m	3
r	16

a	14
d	7
p	2

runs

create  
runs

# External Sort-Merge (Cont.)

## 2. Merge the runs (N-way merge). (if $N < M$ )

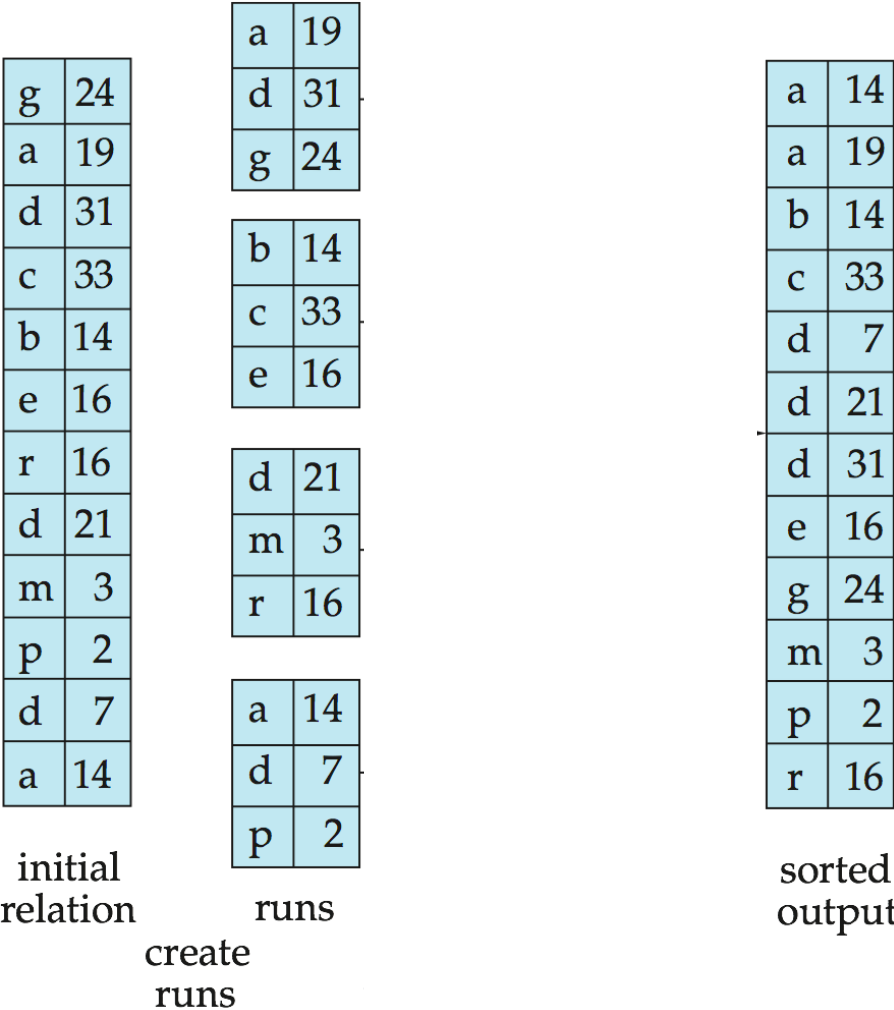
- 1. Use  $N$  buffer blocks for input runs,  
1 buffer block for output.

Read the first block of each run into its buffer page

### 2. repeat

- 1. Select the first record (in sort order) among all buffer pages
- 2. Write the record to the output buffer. If the output buffer is full write it to disk.
- 3. Delete the record from its input buffer page.  
**If** the buffer page becomes empty **then**  
read next block (if any) of the run into the buffer.

- 3. **until** all input buffer pages are empty:



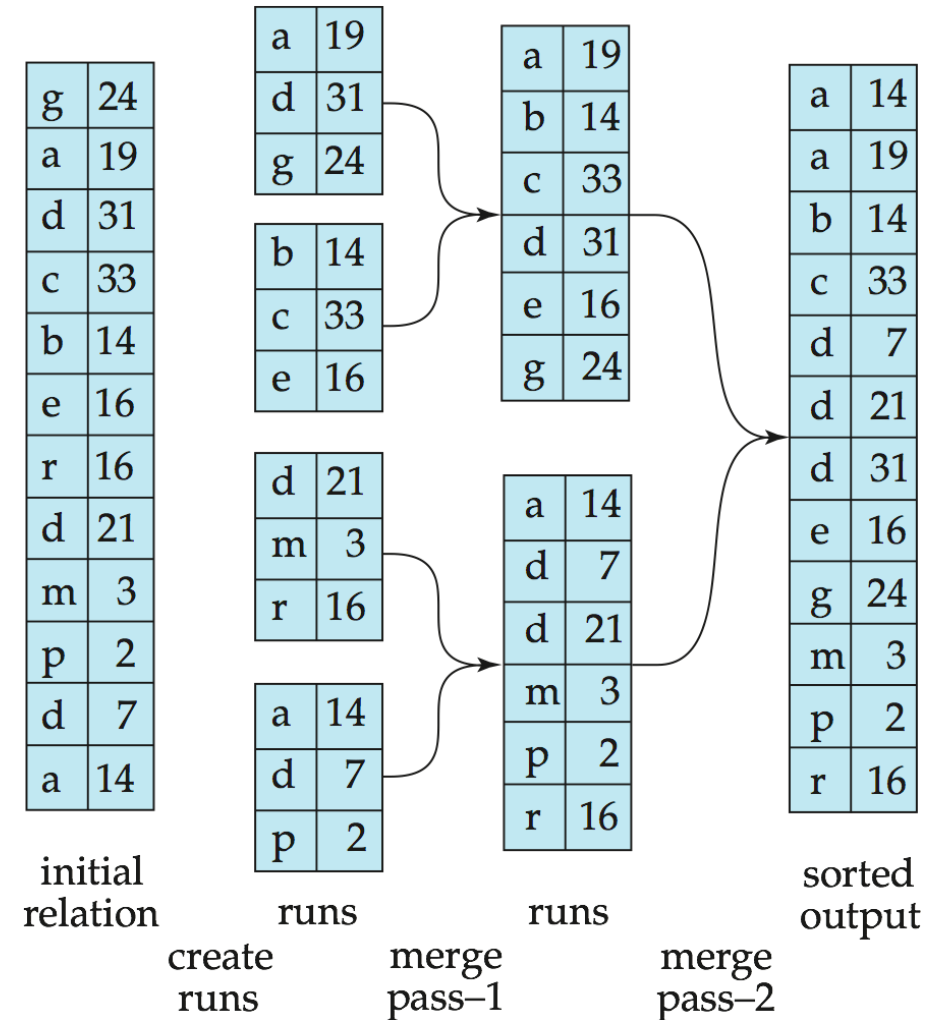


# External Sort-Merge (Cont.)

## 2. Merge the runs (if $N \geq M$ )

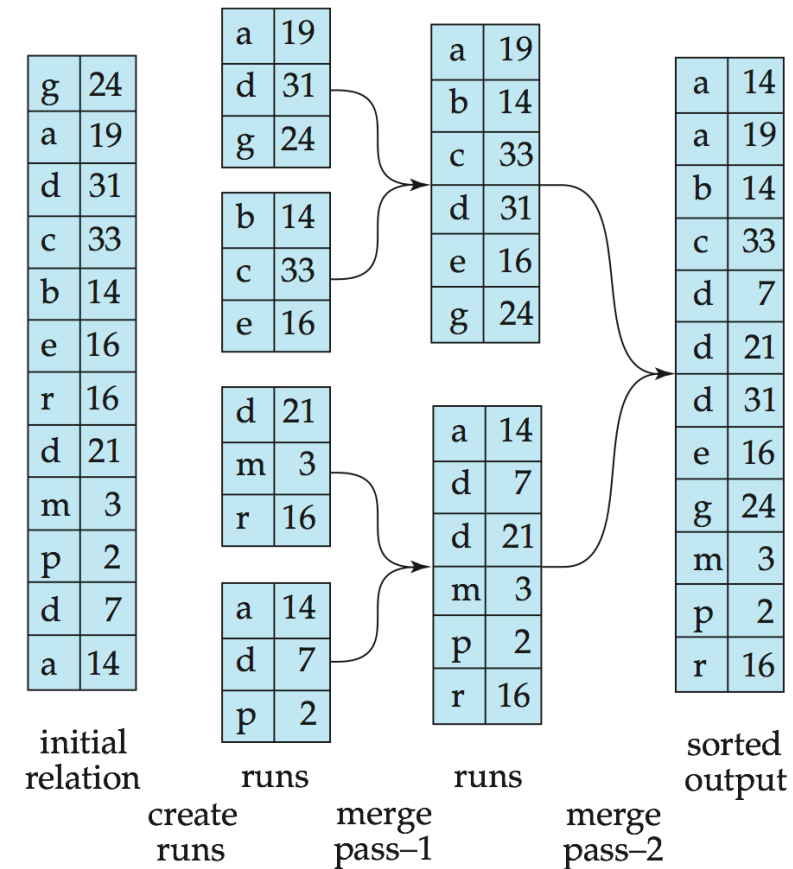
- If  $N \geq M$ , several merge passes are required.
  - In each pass, contiguous groups of  $M - 1$  runs are merged.
  - A pass reduces the number of runs by a factor of  $M - 1$ , and creates runs longer by the same factor.
    - E.g. If  $M=11$ , and there are 90 runs, one pass reduces the number of runs to 9, each 10 times the size of the initial runs
  - Repeated passes are performed till all runs have been merged into one.

$M=3$ ; 1 rec/block



# External Sort-Merge – Cost Analysis

- Total number of merge passes required:  $\lceil \log_{M-1}(b_r/M) \rceil$ .
- Block transfers
  - for initial run creation as well as in each pass:  $2b_r$
  - Thus total number of block transfers for external sorting:  (we don't count final write cost)
- Seeks
  - Run generation: 1 seek to read and 1 seek to write each run
    - $2 \lceil b_r/M \rceil$
  - During the merge phase
    - Buffer size per run:  $b_b$  (read/write  $b_b$  blocks at a time)
    - Need  $2 \lceil b_r/b_b \rceil$  seeks for each merge pass
      - except the final one which does not require a write
  - Total number of seeks:



# Join Operation

- Several different algorithms to implement joins
  - Nested-loop join
  - Block nested-loop join
  - Indexed nested-loop join
  - Merge-join
  - Hash-join
- Choice based on cost estimate
- Running Example
  - Number of records of *student*: 5,000    *takes*: 10,000
  - Number of blocks of *student*: 100    *takes*: 400

# Nested-Loop Join

- To compute the theta join  $r \bowtie_{\theta} s$

```
for each tuple  $t_r$  in  $r$  do
  for each tuple  $t_s$  in  $s$  do
    if pair  $(t_r, t_s)$  satisfy the join condition  $\theta$ 
      add  $t_r \cdot t_s$  to the result.
```

- $r$  is called the
- $s$  the
- Requires no indices and can be used with any kind of join condition.
- Expensive since it examines every pair of tuples in the two relations.

# Nested-Loop Join (Cont.)

- Worst case

- there is memory only to hold one block of each relation

$$n_r * b_s + b_r \text{ block transfers} + n_r + b_r \text{ seeks}$$

- Best case

- the smaller relation fits entirely in memory: use it as the inner relation



- Example

- with *student* as outer relation:



- with *takes* as the outer relation

- $10000 * 100 + 400 = 1,000,400$  block transfers and 10,400 seeks

- If smaller relation (*student*) fits entirely in memory

- $100 + 400 = 500$

# Block Nested-Loop Join

- Variant of nested-loop join
    - every block of inner relation is paired with every block of outer relation
- for each block  $B_r$  of  $r$  do  
for each block  $B_s$  of  $s$  do  
for each tuple  $t_r$  in  $B_r$  do  
for each tuple  $t_s$  in  $B_s$  do  
if  $(t_r, t_s)$  satisfy the join condition  
then add  $t_r \cdot t_s$  to the result
- Worst case estimate
    - Each block in the inner relation  $s$  is read once for each *block* in the outer relation
    - block transfers +  seeks

# Block Nested-Loop Join (Cont.)

- Best case


- block transfers + 2 seeks.

- Improvements

- Use  $M-2$  disk blocks as blocking unit for outer relations, and remaining two blocks to buffer inner relation and output
  - Cost =
- If equi-join attribute forms a key of inner relation, stop inner loop on first match
- Scan inner loop forward and backward alternately, to make use of the blocks remaining in buffer (with LRU replacement)

# Hash-Join

- Applicable for equi-joins and natural joins.
- A hash function  $h$  is used to partition tuples of both relations
- $h$  maps  $JAttrs$  values to  $\{0, 1, \dots, n-1\}$ 
  - $JAttrs$ : attributes of  $r$  and  $s$  used in the equi-join.
  - $r_0, r_1, \dots, r_{n-1}$  : partitions of  $r$ 
    - $t_r \in r$  is put in partition  $r_i$  where  $i = h(t_r[JAttrs])$ .
  - $s_0, s_1, \dots, s_{n-1}$  : partitions of  $s$
- $r$  tuples in  $r_i$  need to be compared with  $s$  tuples in  $s_i$  only



(Note: In the textbook,  $r_i$  is denoted as  $H_{ri}$ ,  $s_i$  is denoted as  $H_{si}$ , and  $n$  is denoted as  $n_h$ .)



# Hash-Join Algorithm

1. Partition the relation  $s$  using hashing function  $h$ .
    - When partitioning a relation, one block of memory is reserved as the output buffer for each partition.
  2. Partition  $r$  similarly.
  3. For each  $i$ :
    - (a) Load  $s_i$  into memory
      - build an in-memory hash index on it using the join attribute (using a different hash function than the earlier one  $h$ )
    - (b) Read the tuples in  $r_i$  from the disk one by one
      - For each tuple  $t_r$  locate each matching tuple  $t_s$  in  $s_i$  using the in-memory hash index. Output the concatenation of their attributes.
- Relation  $s$  is called the  and  
 $r$  is called the

# Hash-Join algorithm (Cont.)

- The value  $n$  and the hash function  $h$  is chosen such that each  $s_i$  should fit in memory.
- *Hash-table overflow* occurs in partition  $s_i$  if  $s_i$  does not fit in memory.
  - Many tuples in  $s$  with same value for join attributes or bad hash function
  - Overflow resolution can be done in build phase
    - Partition  $s_i$  is further partitioned using different hash function.
    - Partition  $r_i$  must be similarly partitioned.
  - Fails with large numbers of duplicates
    - Fallback option: use block nested loops join on overflowed partitions

# Cost of Hash-Join

- Cost of hash join (without recursive partitioning)


 block transfers ( $4n_h$  can be ignored) + seeks

- If the entire build input can be kept in main memory,  $n$  can be set to 0 and the algorithm does not partition the relations into temporary files. Cost estimate goes down to  $b_r + b_s$ .
- Example:
  - $student \bowtie takes$
  - memory size: 20 blocks;  $b_{stud} = 100$  and  $b_{takes} = 400$ .
  - $student$  is build input
  - Partition it into 5 partitions, each of size less than 20 blocks
  - Similarly, partition  $takes$  into 5 partitions each of size about 80
  - Therefore total cost:


$$\begin{aligned} 3(100 + 400) &= 1500 && \text{block transfers} \\ 2(\lceil 100/3 \rceil + \lceil 400/3 \rceil) &= 336 && \text{seeks} \end{aligned}$$

# Evaluation of Expressions

- So far, we have seen algorithms for individual operations
- How do you evaluate an entire expression tree?
- - generate results of an expression whose inputs are relations or are already computed, **materialize** (store) it on disk. Repeat.
- - pass on tuples to parent operations even as an operation is being executed

# Materialization

- Evaluate one operation at a time, starting at the lowest-level.
- Use intermediate results *materialized into temporary relations* to evaluate next-level operations.
- E.g.,
  - compute and store  $\sigma_{building="Watson"}(department)$
  - then compute and store its join with *instructor*
  - and finally compute the projections on *name*.




# Materialization (Cont.)


- Materialized evaluation is always applicable
- Cost of writing results to disk and reading them back can be quite high
  - Our cost formulas for operations ignore cost of writing results to disk
  - Overall cost = Sum of costs of individual operations +

# Pipelining

- Evaluate several operations simultaneously
  - passing the results of one operation on to the next.
- E.g., in expression tree
  - don't store result of selection
  - instead, pass tuples directly to the join
  - Similarly, don't store result of join but pass tuples directly to projection




# Pipelining (cont.)

- Much cheaper than materialization
  - no need to store a temporary relation to disk.
- Pipelining may not always be possible
  - e.g., sort, hash-join: 
  - Very difficult to achieve a lengthy chain of pipeline
- For pipelining to be effective
  - use evaluation algorithms that generate output tuples even as tuples are received for inputs to the operation



# Cost Estimation of Expressions

- Cost of computing each operator is as described
  - Need statistics of input relations
    - E.g. number of tuples, sizes of tuples
- Inputs can be results of sub-expressions
  - 
  - Additional statistics are needed
    - number of distinct values for an attribute, histograms, ...

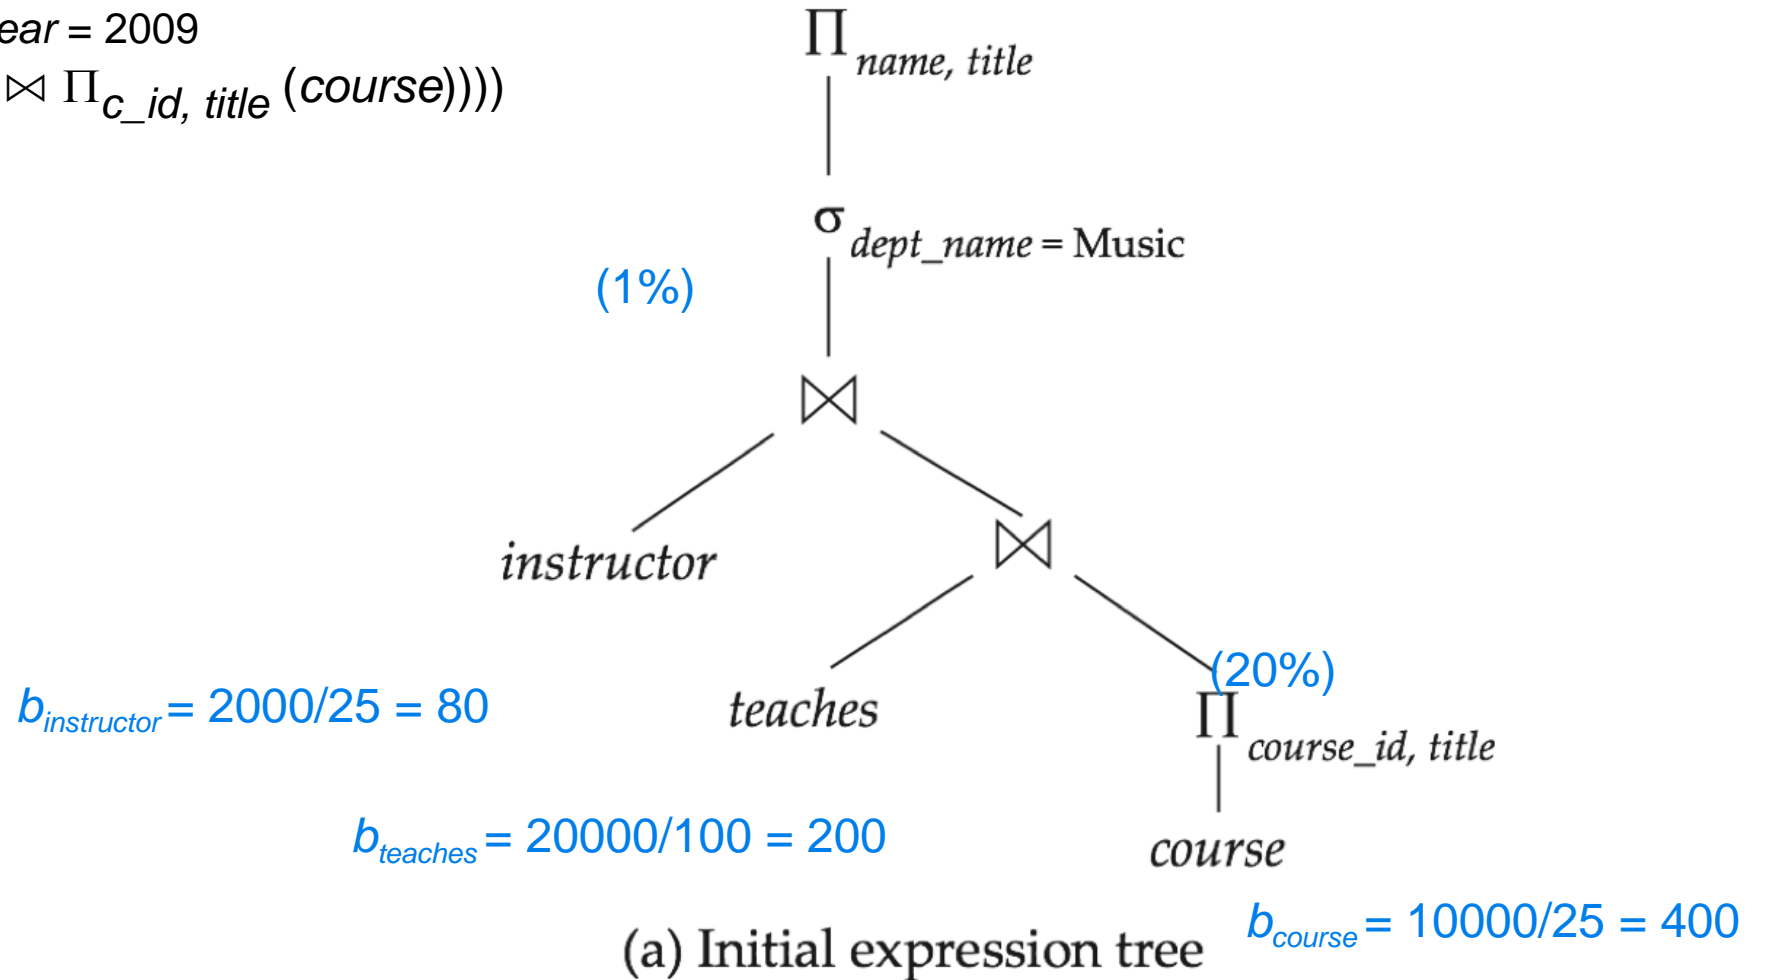
# Cost Estimation of Expressions

$\Pi_{name, title}(\sigma_{dept="Music" \wedge year = 2009}$   
 $(instructor \bowtie (teaches \bowtie \Pi_{c\_id, title}(course))))$

$$M = 22$$

$$N: \lceil b_r / (M-2) \rceil * b_s + b_r$$

$$H: 3(b_r + b_s) + 4 * n_h$$



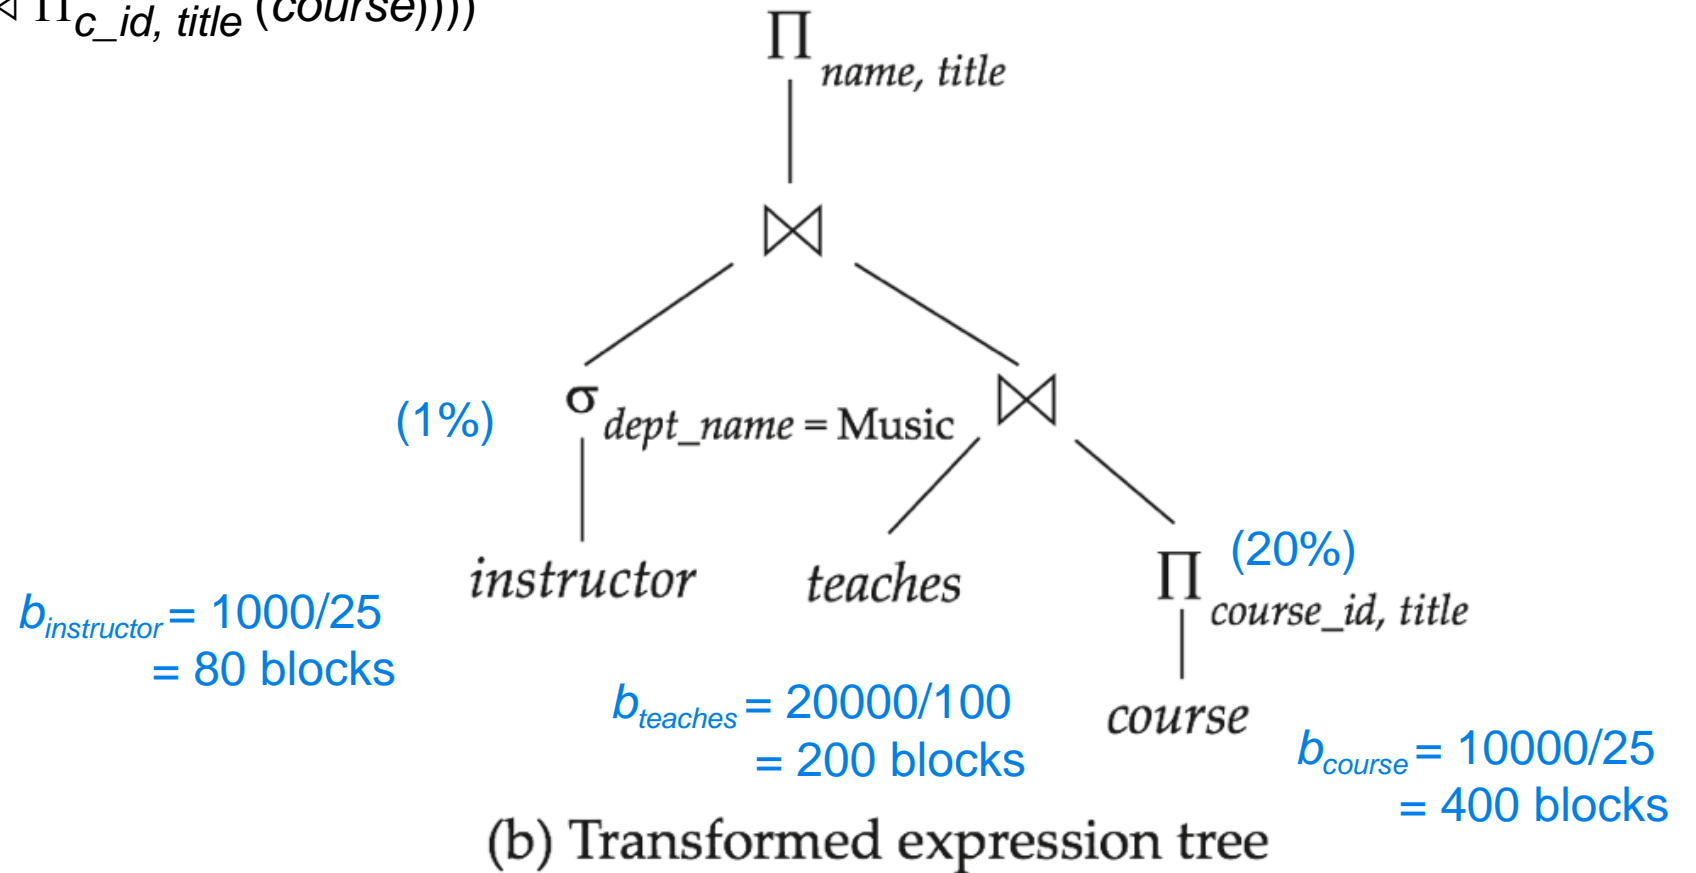
# Cost Estimation of Expressions

$\Pi_{name, title}(\sigma_{dept="Music" \wedge year = 2009}$   
 $((instructor \bowtie teaches) \bowtie \Pi_{c\_id, title}(course)))$

$$M = 22$$

$$N: \lceil b_r / (M-2) \rceil * b_s + b_r$$

$$H: 3(b_r + b_s) + 4 * n_h$$



**END OF CHAPTER 12**