

Intro to DB

# CHAPTER 14

# TRANSACTIONS

# Chapter 14: Transactions

- Transaction Concept
- Transaction Model
- Storage Structure
- Transaction States
- Transaction Isolation & Schedules
- Serializability
- Recoverability
- Isolation Levels
- Transactions as SQL Statements

# Transactions

- What is a transaction?
  - (user view)
  - A program unit that accesses and possibly updates various data items (system view)
  - DBMS must guarantee certain properties (ACID properties) for units of works declared as a DB transaction

# Examples of Transactions

- Banks
  - Withdraw \$100 to account A.
  - Transfer \$50 from account A to B.
- Schools
  - Register course #409.433 for student #4321.
- Airlines
  - Check if two seats are available on flight #453.
  - Reserve the two seats on flight #453.
- Companies
  - Increase every employee's salary by 5%.

# Dangers for Transactions

- Various types of failure
  - system crash
  - disk failure
  - system error
- - disk access is performed in chunks: page (block)
  - i.e., write operation performed after the right amount of data has been gathered
  - buffer manager may pin a page
- Concurrent execution of multiple transactions

# Properties of a Transaction

## ACID properties

- **A**tomicity

- **C**onsistency (Correctness)

- **I**solation

- **D**urability

# Atomicity

- **All or nothing**

*“Transfer \$50 from account A to account B”*

Begin transaction

read(A, a)

a = a-50


write(A, a)

read(B, b)

b = b+50

write(B, b)

End Transaction



# Consistency

- Move from a **consistent state to another consistent state**

*“Withdraw \$100 from account A”*

Begin transaction

read(A, a)

a = a-100

write(A, a)

End Transaction

*What if A only had \$20?*

▪



# Isolation

- Should **not be interfered** by other transactions (concurrency)

*“Transaction T1”*

Begin transaction

read(A, a1)

a1 = a1-50

write(A, a1)

read(B, b1)

b1 = b1+50

write(B, b1)

End Transaction

*“Transaction T2”*


Begin transaction

read(A, a2)

a2 = a2-100

write(A, a2)

End Transaction



■

# Durability

- The effect of a completed transaction should be **durable & public**

*“Withdraw \$100 from account A”*

Begin transaction

read (A,a)


a = a-100

write(A,a)


End Transaction

...

\*buffer flush\*



# Transaction States



# Transaction States (Cont.)

- **Active**: the initial state, transaction stays in this state while it is executing
- **Partially committed**: the final statement has been executed
- **Failed**: normal execution can no longer proceed
- **Aborted**: transaction has been rolled back and the database restored to its state prior to the start of the transaction.
  - Two options after it has been aborted:
    - restart the transaction – only if no internal logical error
    - kill the transaction
- **Committed**: after *successful completion*.

# Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system
  - increased processor and disk utilization
  - reduced average response time
- – mechanisms to achieve isolation
  - to control the interaction among the concurrent transactions in order to prevent them from destroying the consistency of the database

# Schedules

- Simplified view of transactions
  - Assume that transactions may perform arbitrary computations on data in local buffers in between reads and writes.
  - Ignore operations other than **read** and **write** instructions  
=> simplified schedules consist of only **read** and **write** instructions.
- *Schedules*
  - 
  - schedule for a set of transactions:
    - must consist of all instructions of those transactions
    - must preserve the order in which the instructions appear in each individual transaction

# Example Schedules

- $T_1$ : transfer \$50 from  $A$  to  $B$ ;  $T_2$ : transfer 10% of the balance from  $A$  to  $B$ .

- 

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

Schedule 1 (fig 14.2)

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$ write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ ) $B := B + temp$ write ( $B$ ) commit

Schedule 2 (fig 14.3)

# Example Schedule (Cont.)

- Let  $T_1$  and  $T_2$  be the transactions defined previously.
- Schedule 3 is **not** a serial schedule, but it
- In both Schedule 1 and 3, the sum  $A+B$  is *preserved*.

$T_1$	$T_2$
read (A) $A := A - 50$ write (A)	
	read (A) $temp := A * 0.1$ $A := A - temp$ write (A)
read (B) $B := B + 50$ write (B) commit	
	read (B) $B := B + temp$ write (B) commit

Schedule 3 (fig 14.4)



## Example Schedules (Cont.)

- Schedule 4 **does not** preserve the value of the the sum  $A + B$ .

Schedule 4 is **not** equivalent to Schedule 1.

$T_1$	$T_2$
read ( $A$ ) $A := A - 50$	read ( $A$ ) $temp := A * 0.1$ $A := A - temp$ write ( $A$ ) read ( $B$ )
write ( $A$ ) read ( $B$ ) $B := B + 50$ write ( $B$ ) commit	$B := B + temp$ write ( $B$ ) commit

Schedule 4 (fig 14.5)

# Serializability

- Basic Assumption – Each transaction preserves database consistency.
  - Thus, **serial execution** of a set of transactions preserves database consistency.
- A (possibly concurrent) schedule is serializable if it is **equivalent** to a serial schedule.
- $S1$  is *equivalent* to  $S2$  if, for every possible instance of the database,

# Serializable Schedules

$T_1$	$T_2$
read( $A$ ) write( $A$ )	
	read( $A$ ) write( $A$ )
read( $B$ ) write( $B$ )	
	read( $B$ ) write( $B$ )

$T_8$	$T_9$
read( $A$ ) write( $A$ )	
	read( $A$ )
read( $B$ )	

# Conflict Serializability

- Instructions  $I_i$  and  $I_j$  of transactions  $T_i$  and  $T_j$  respectively, **conflict** if and only if
  - there exists some item  $Q$  accessed by both  $I_i$  and  $I_j$ , and
  - at least one of these instructions is a write( $Q$ )
    1.  $I_i = \text{read}(Q)$ ,  $I_j = \text{read}(Q)$ . not conflict.
    2.  $I_i = \text{read}(Q)$ ,  $I_j = \text{write}(Q)$ . conflict.
    3.  $I_i = \text{write}(Q)$ ,  $I_j = \text{read}(Q)$ . conflict
    4.  $I_i = \text{write}(Q)$ ,  $I_j = \text{write}(Q)$ . conflict
- Intuitively, a conflict between  $I_i$  and  $I_j$  forces a (logical) temporal order between them.


# Conflict Serializability (Cont.)

- Schedules  $S$  and  $S'$  are *conflict equivalent* if
  - $S$  can be transformed into a schedule  $S'$
  - by a series of swaps of non-conflicting instructions
- A schedule  $S$  is *conflict serializable* if it is conflict equivalent to a serial schedule
- Example of a schedule that is **not** conflict serializable:

$T_3$	$T_4$
read ( $Q$ )	write ( $Q$ )
write ( $Q$ )	




# Testing for Serializability

- Consider some schedule of a set of transactions  $T_1, T_2, \dots, T_n$
- **Precedence graph** — a direct graph where the vertices are the transactions
  - draw an arc from  $T_i$  to  $T_j$  if the two transaction conflict, and  $T_i$  accessed the data item on which the conflict arose earlier.
  - We may label the arc by the item that was accessed.
- Example



# Test for Conflict Serializability

- A schedule is conflict serializable if and only if its precedence graph is acyclic.
- Cycle-detection algorithms exist which take order  $n^2$  time, where  $n$  is the number of vertices in the graph.
  - (Better algorithms take order  $n + e$  where  $e$  is the number of edges.)
- If precedence graph is acyclic, the serializability order can be obtained by a *topological sorting* of the graph.
  - a linear order consistent with the partial order of the graph.



# Recoverability

- Need to address the effect of transaction failures on concurrently running transactions
- Example: Schedule 9 is not recoverable
- *Recoverable schedule*
  - if a transaction  $T_j$  reads a data item previously written by a transaction  $T_i$ ,
  -
- DBMS must ensure that schedules are recoverable.

$T_8$	$T_9$
read (A)	
write (A)	
	read (A)
read (B)	commit

Schedule 9 (fig 14.14)



# Cascading Rollback

- 
- Example: If  $T_{10}$  fails,  $T_{11}$  and  $T_{12}$  must also be rolled back.

$T_{10}$	$T_{11}$	$T_{12}$
read (A) read (B) write (A)	read (A) write (A)	
abort		read (A)

Schedule 10 (fig 14.15)

Can lead to the undoing of a significant amount of work

# Cascadeless Schedules

- Cascading rollbacks cannot occur if
  - for each pair of transactions  $T_i$  and  $T_j$
  - such that  $T_j$  reads a data item previously written by  $T_i$ ,
  -
- 
- It is *desirable* to restrict the schedules to those that are cascadeless

**END OF CHAPTER 14**