


Chapter 1: Introduction

WHAT'S AHEAD:

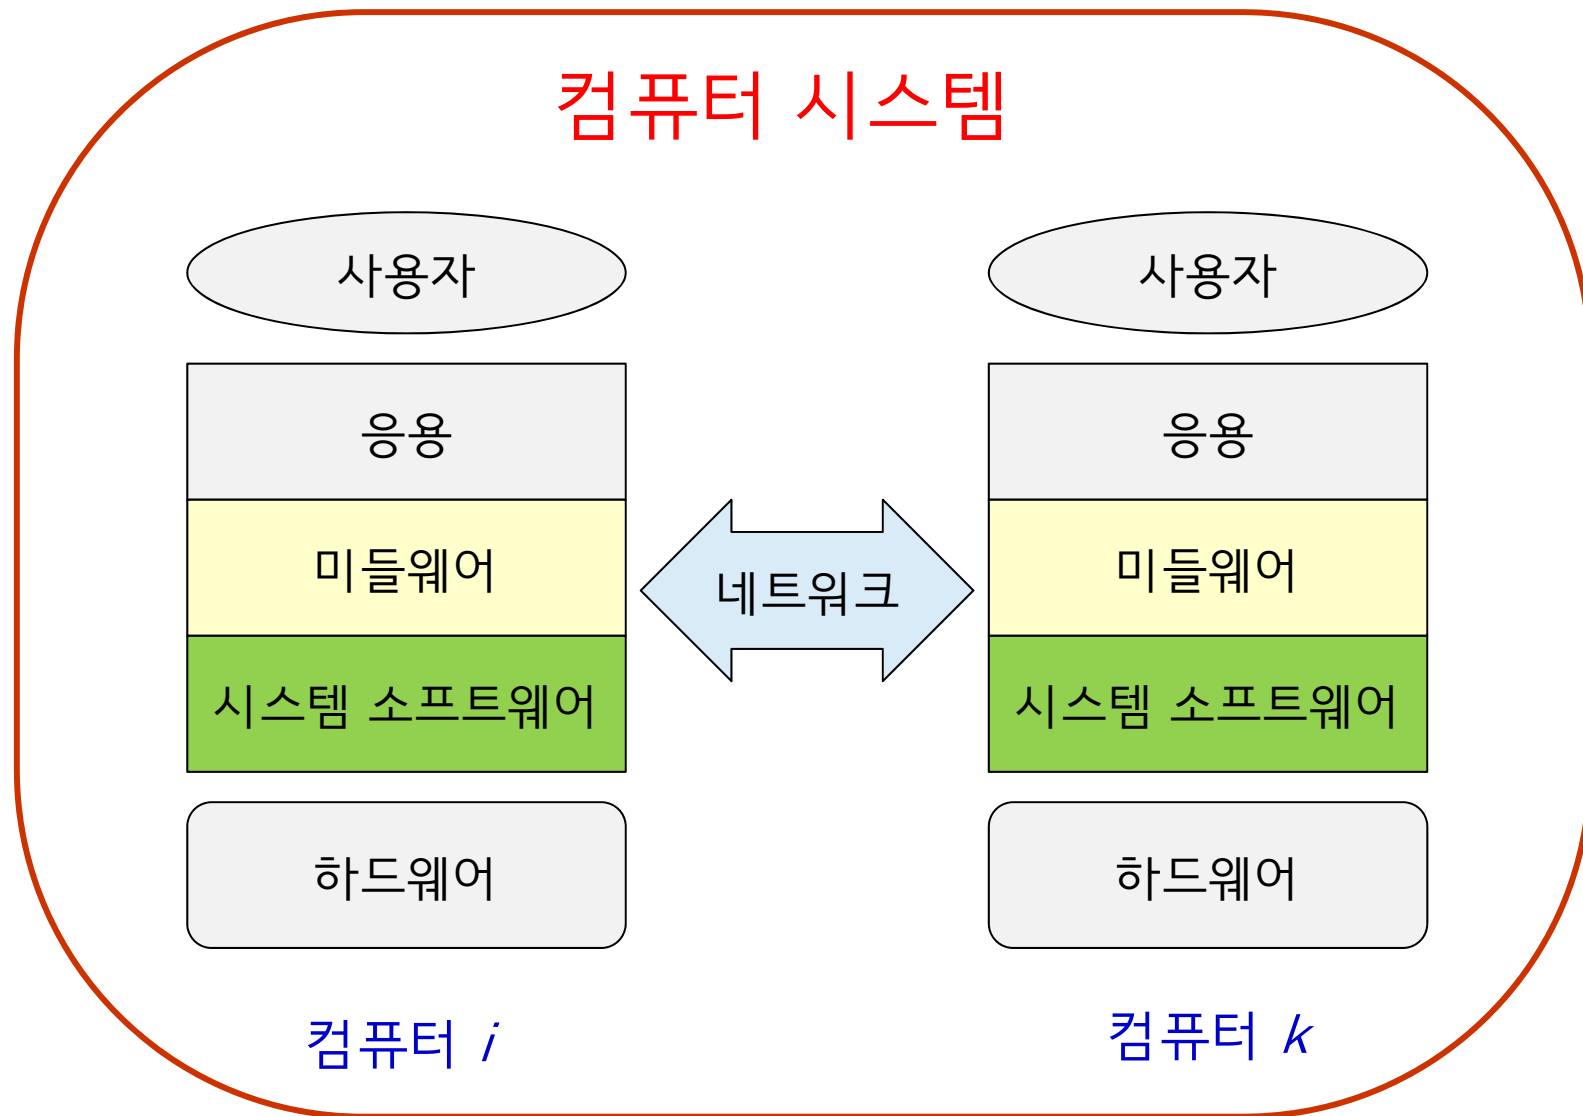
- What Operating Systems Do
- Computer System Organization
 - Computer System Architecture
- Operating System Structure
- Operating System Operations
 - Major functions of OS
 - Kernel Data Structures
 - Computing Environments

WE AIM:

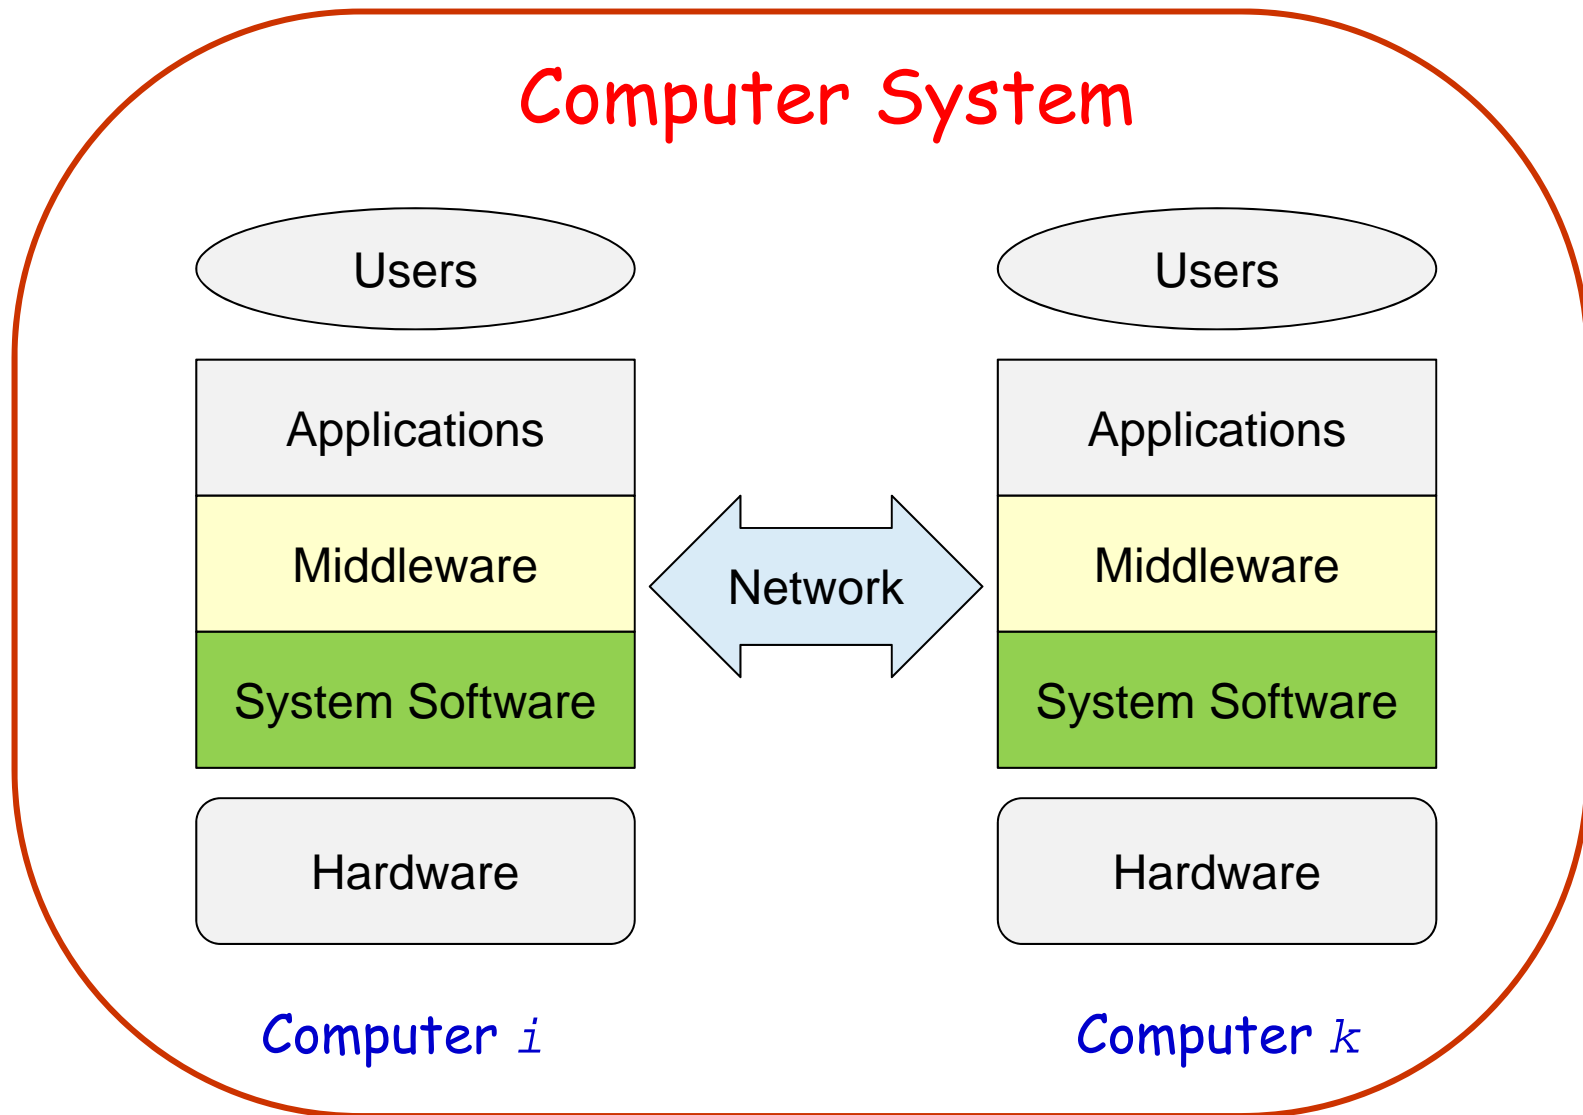
- To describe the structural aspect of computer systems
- To preview the major components of OSs
- To overview various types of computing environments

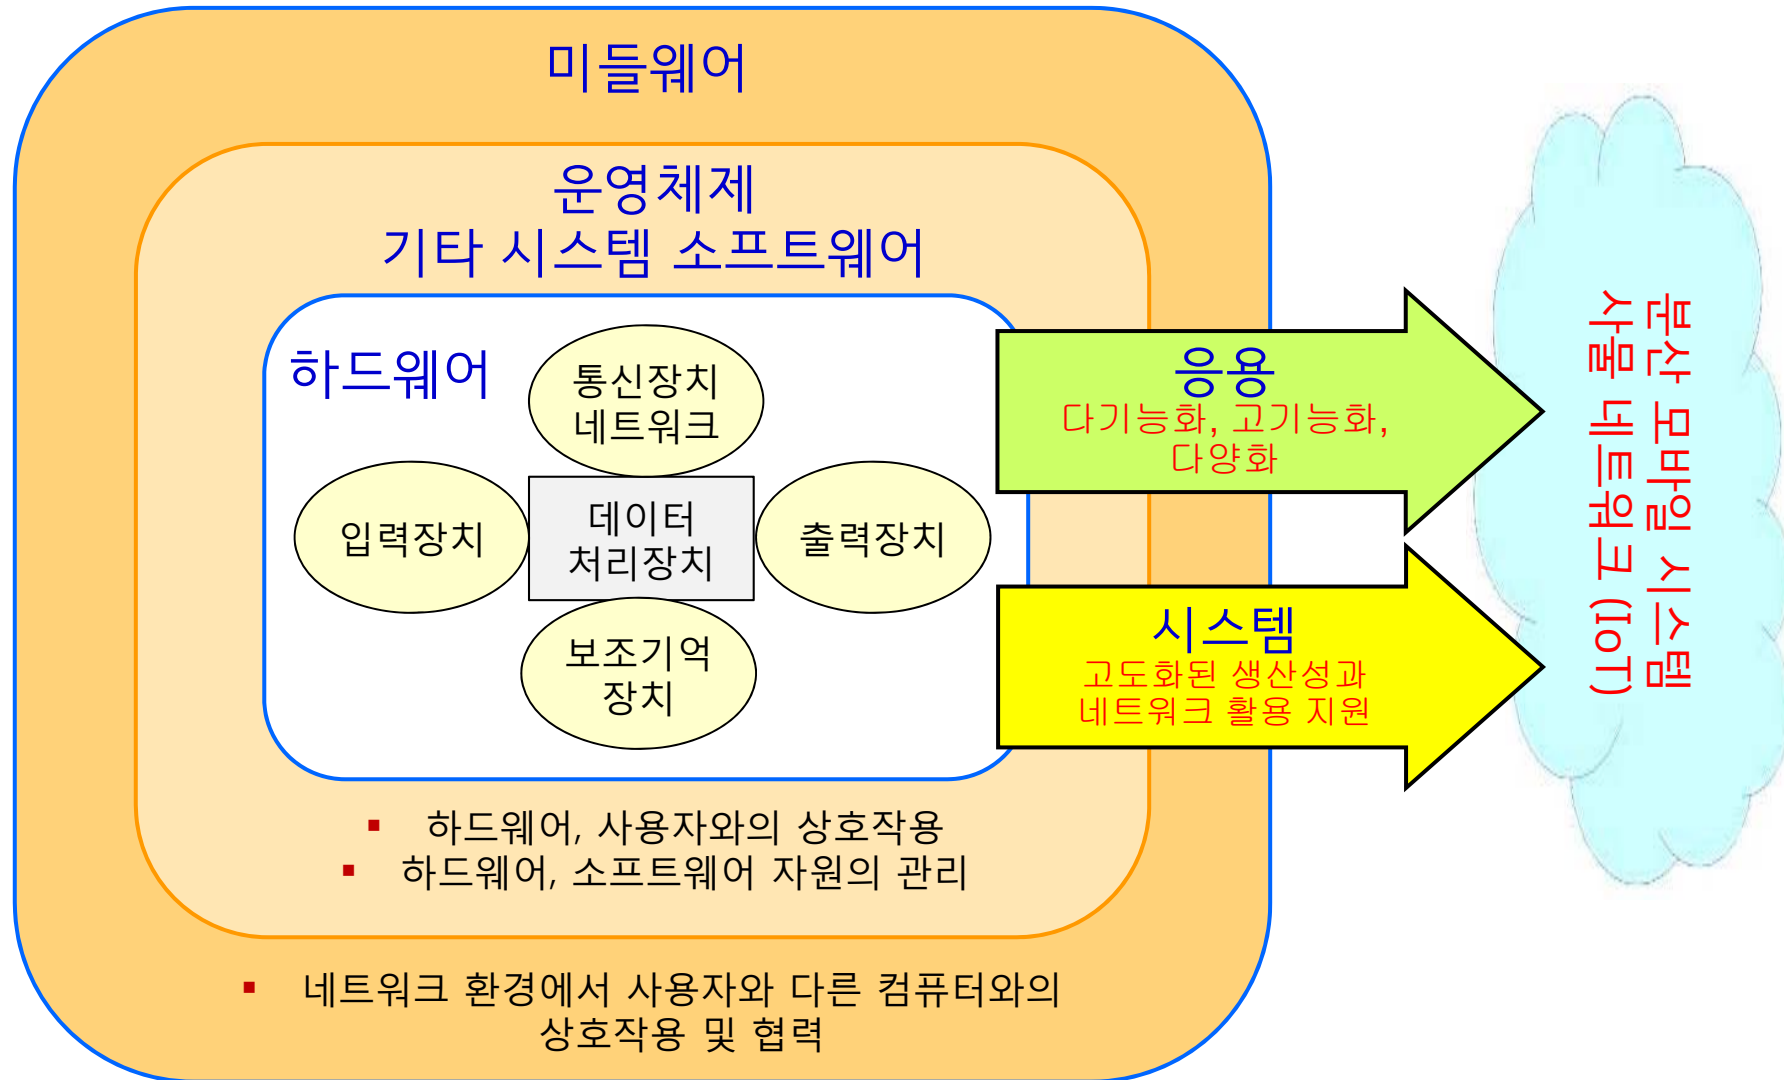


Note: These lecture materials are based on the lecture notes prepared by the authors of the book titled *Operating System Concepts*, 9e (Wiley)

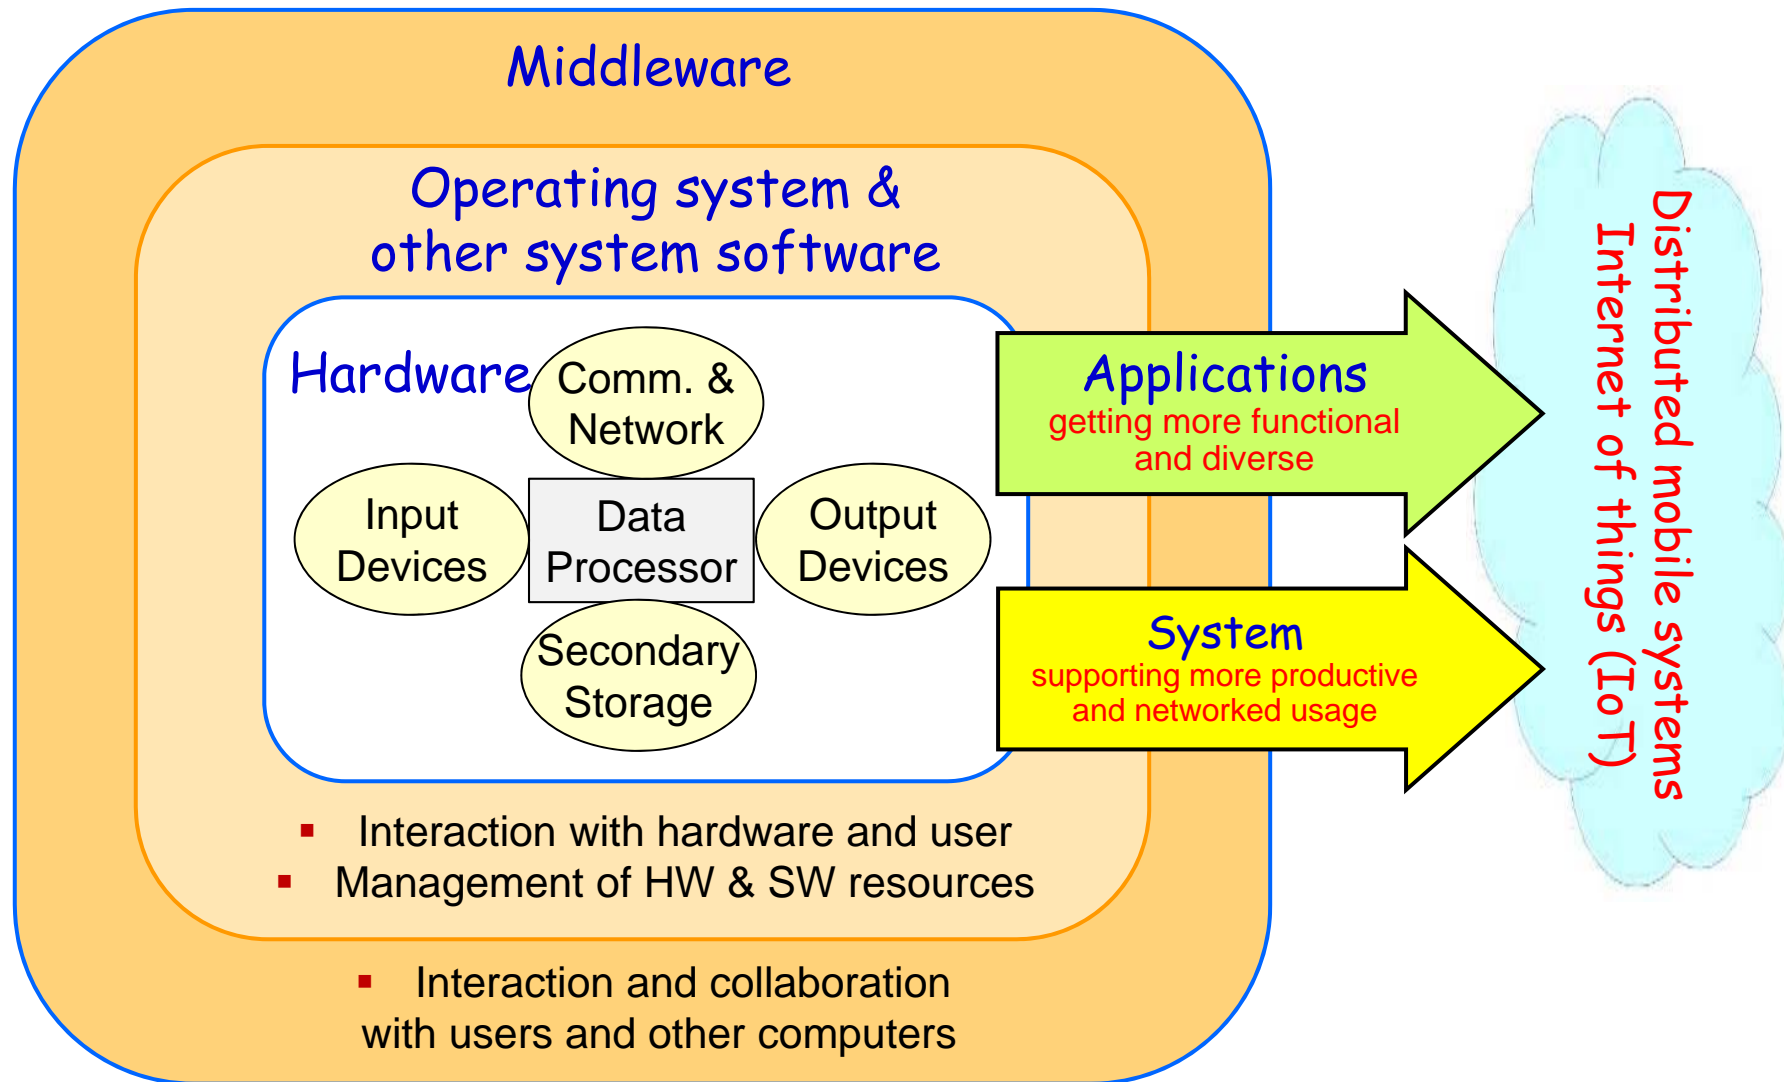


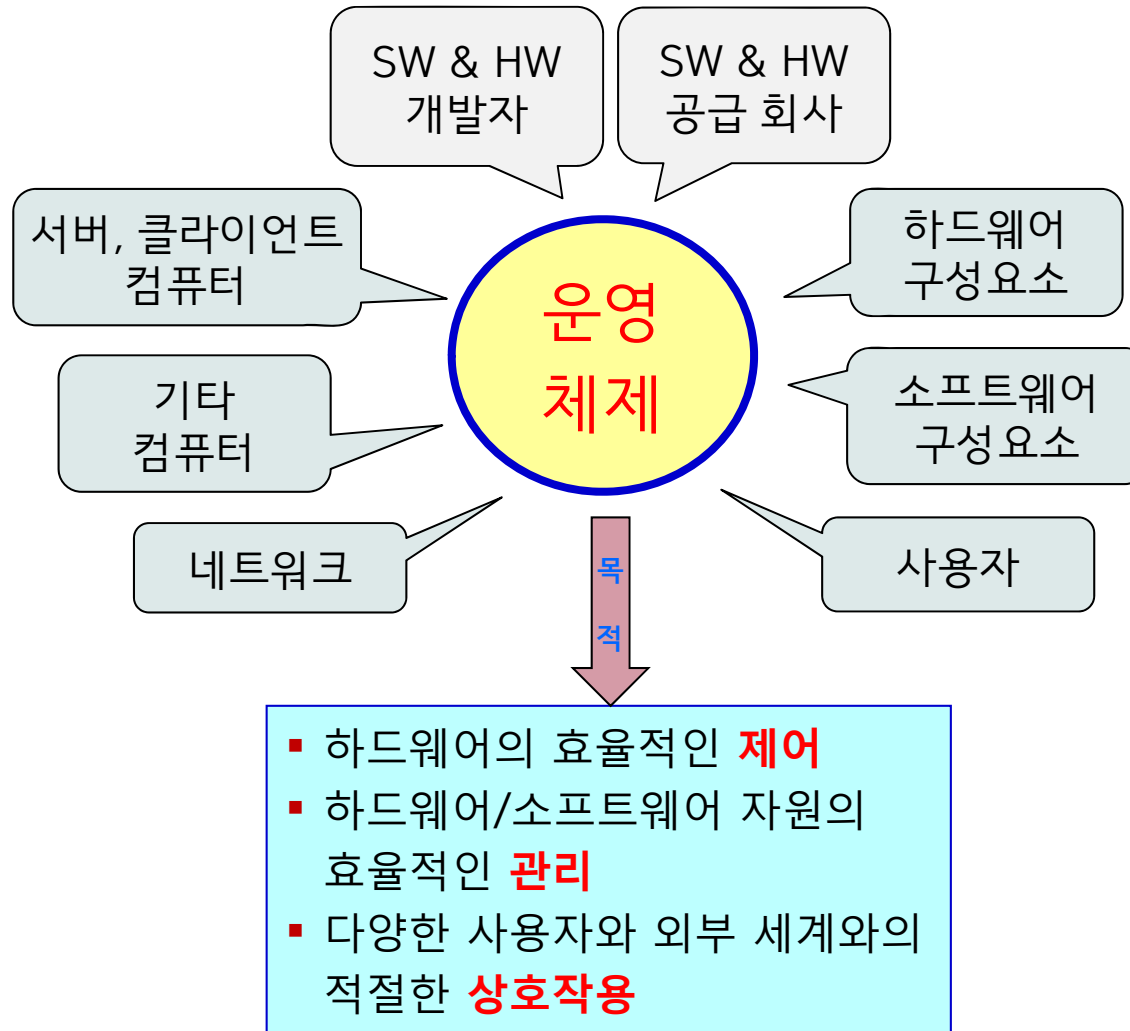
Core Ideas How Computer Systems Look



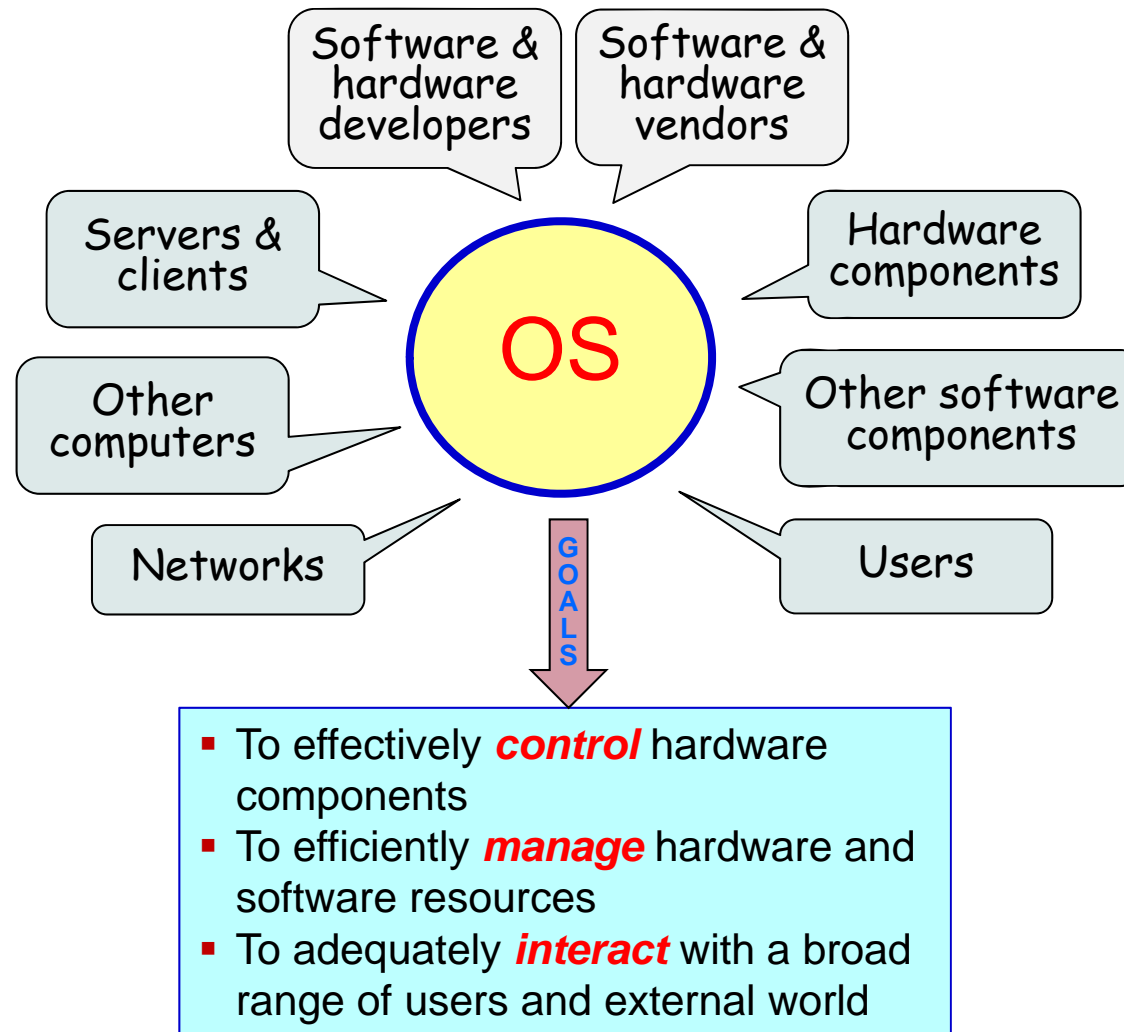


Core Ideas Evolution of Computer Systems





Core Ideas OS Goals & Its Ecosystem



What is an Operating System?




- A program that acts as an intermediary between a user of a computer and the computer hardware
- Operating system goals:
 - Execute user programs and make solving user problems easier
 - Make the computer system convenient to use
 - Use the computer hardware in an efficient manner
- Also, OS aims to
 - Please end users, as many as possible
 - Please system managers
 - Please software developers



Computer System Structure

- Computer system can be divided into four components:
 - **Hardware** - provides basic computing resources
 - CPU, memory, I/O devices
 - **Operating system**
 - Controls and coordinates use of hardware among various applications and users
 - In many cases, includes middleware functionalities
 - **System and application programs** - define the ways in which the system resources are used to solve the computing problems of the users
 - Word processors, video games, compilers, web browsers, database systems
 - **Users**
 - People, *machines*, other computers

Four Components of a Computer System



Hierarchical in nature



What Operating Systems Do

- Depends on the point of view
- Users want convenience, ease of use
 - Don't care about resource utilization
- But **shared computers** such as mainframes or minicomputers must keep all users happy
- Users of **dedicated systems** such as workstations have dedicated resources but frequently use shared resources from servers
- **Handheld computers** are resource poor, optimized for usability and battery life
- Some computers have little or no user interface, such as **embedded computers** in devices and automobiles



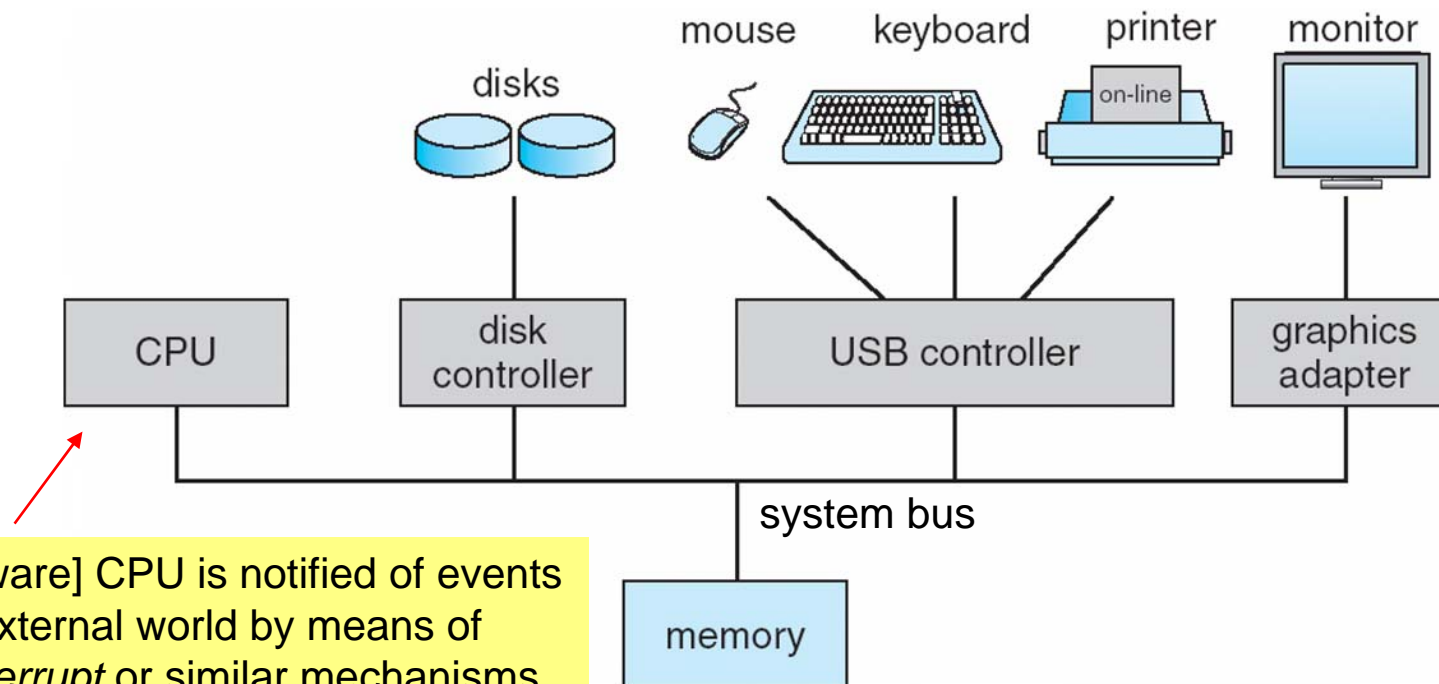
Operating System Definition

- OS is a **resource manager**
 - Allocates and manages all hardware and software resources
 - Decides between conflicting requests for efficient and/or fair resource use
 - Allows users and application programs to make use of resources
- OS is a **control program**
 - Controls execution of programs to prevent errors and improper use of the computer
- In computer business, no universally accepted definition
 - Usually, include *middleware* functionalities, such as networking, multimedia, GUI, etc.

Computer System Organization



- Computer system operations
 - One or more CPUs, device controllers connect through common bus providing access to shared memory
 - Parallel execution of CPUs and devices competing for memory cycles



[Hardware] CPU is notified of events from external world by means of the *interrupt* or similar mechanisms.



Interrupts

- (Hardware) Interrupt
 - The (hardware) action that an I/O device performs onto CPU to notify an event
 - The primary mechanism by which OS interacts with computer hardware
 - Invisible to users, but drives OS and program executions
 - Most important mechanism for OS implementation
- Interrupt service routine (ISR)
 - Upon occurrence of interrupt, CPU automatically execute ISR to service that interrupt
 - Each interrupt usually has its own ISR
- Software interrupt
 - Invoked by a machine instruction or by the result of execution of an instruction
 - Serviced similarly to the hardware interrupt
 - Caused either by an error or a user request
 - Examples: trap, exception

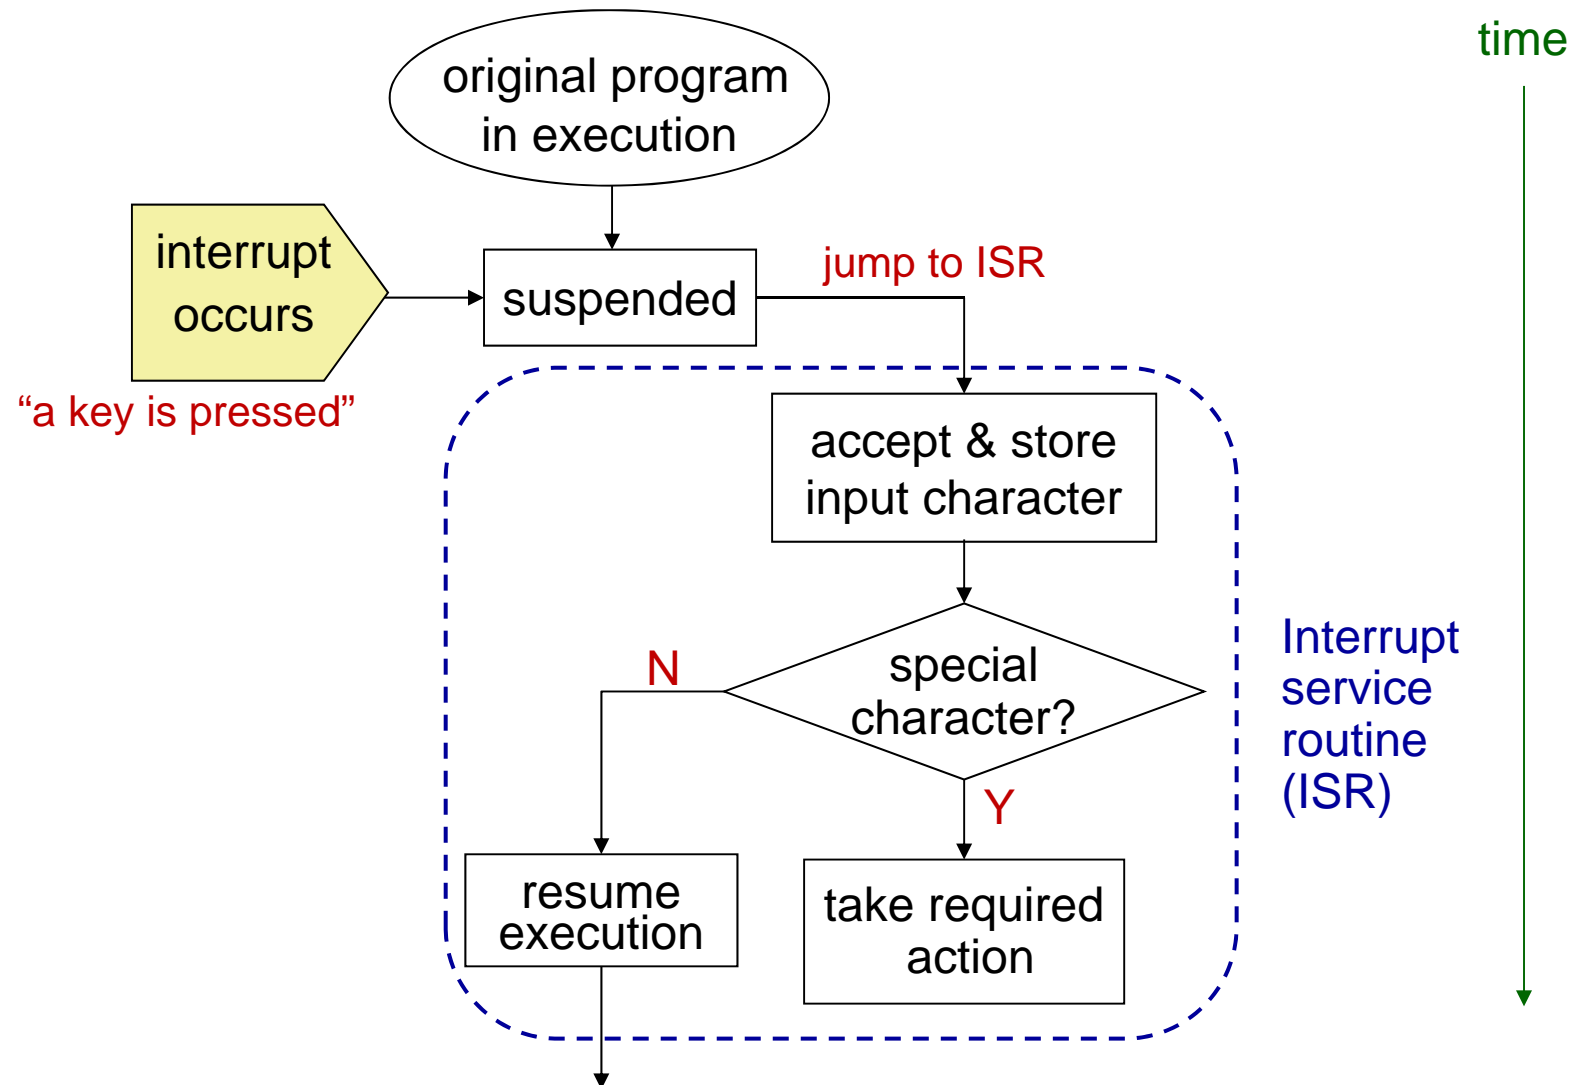


Interrupt Processing

- If an interrupt occurs
 - Saves the program counter (that is, return address) of the interrupted process
 - Transfers control to the ISR, i.e., jump to the ISR
- Upon completion of ISR
 - Resume the execution of interrupted process by restoring the PC
 - Or, return to the location or a program that OS designates
- Examples
 - An event occurs at an I/O device, such as keyboard, mouse, disk, timer, etc.
 - Abnormal, critical, urgent event occurs, such as power outage, component failure, etc.
 - Illegal instructions, such as non-existing instruction, divide-by-zero, etc.



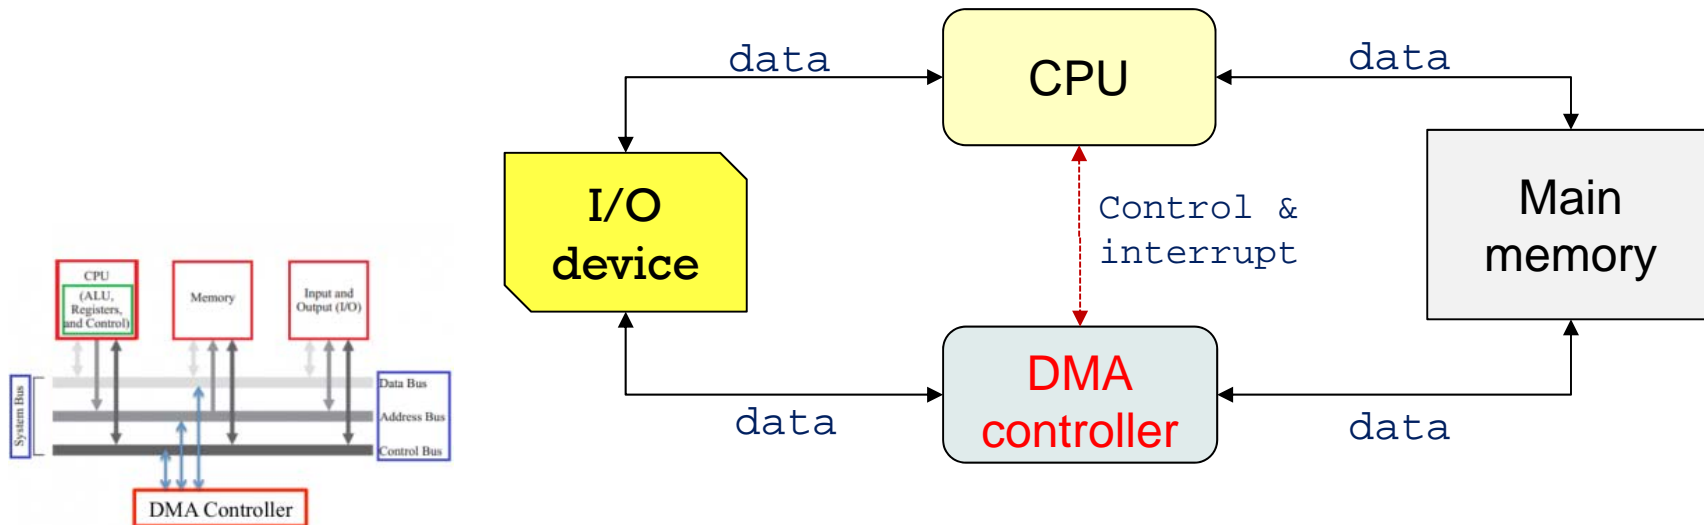
Interrupts – *Keyboard Example*





Direct Memory Access (DMA)

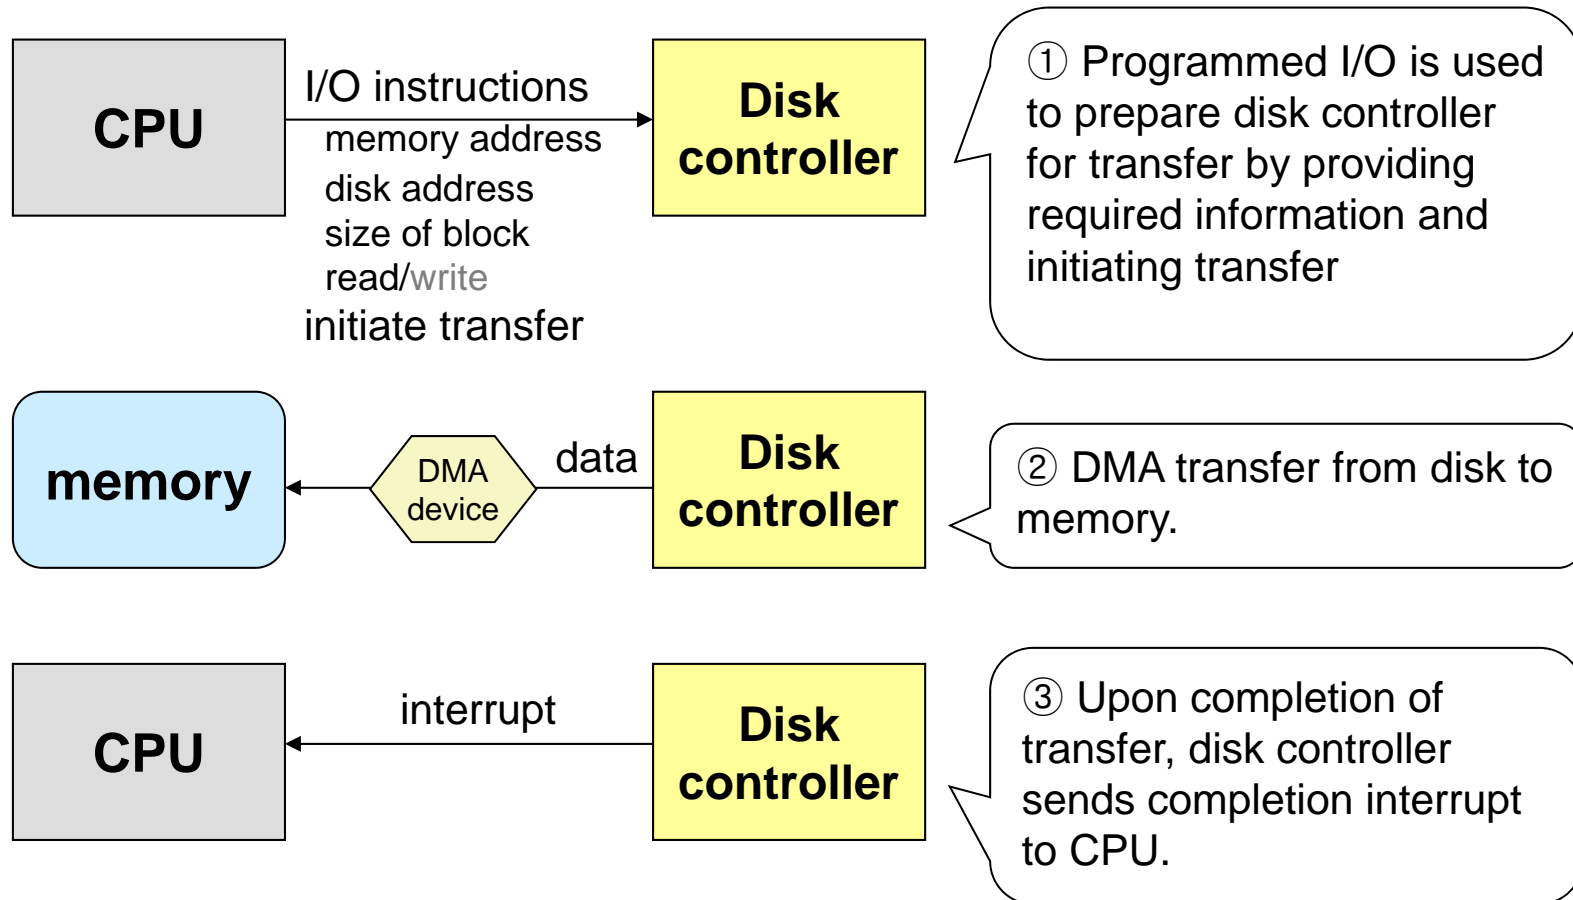
- Used for high-speed I/O devices able to transmit information at close to memory speeds
- DMA controller transfers blocks of data from buffer storage **directly** to main memory without CPU intervention
 - Only one interrupt is generated per block, rather than the one interrupt per byte





DMA Operations – *An Example*

- Moving data from disk to the main memory





I/O Structure

After I/O starts

1. Control returns to user program only upon I/O completion
(Synchronous)
 - Wait instruction idles the CPU until the next interrupt
 - At most one I/O request is outstanding at a time, no simultaneous I/O processing
2. Control returns to user program without waiting for I/O completion (Asynchronous)
 - System call - request to the OS to allow user to wait for I/O completion (The user is forced to wait in the wait queue)
 - Device-status table contains entry for each I/O device indicating its type, address, and state
 - OS indexes into I/O device table to determine device status and to modify table entry to include interrupt



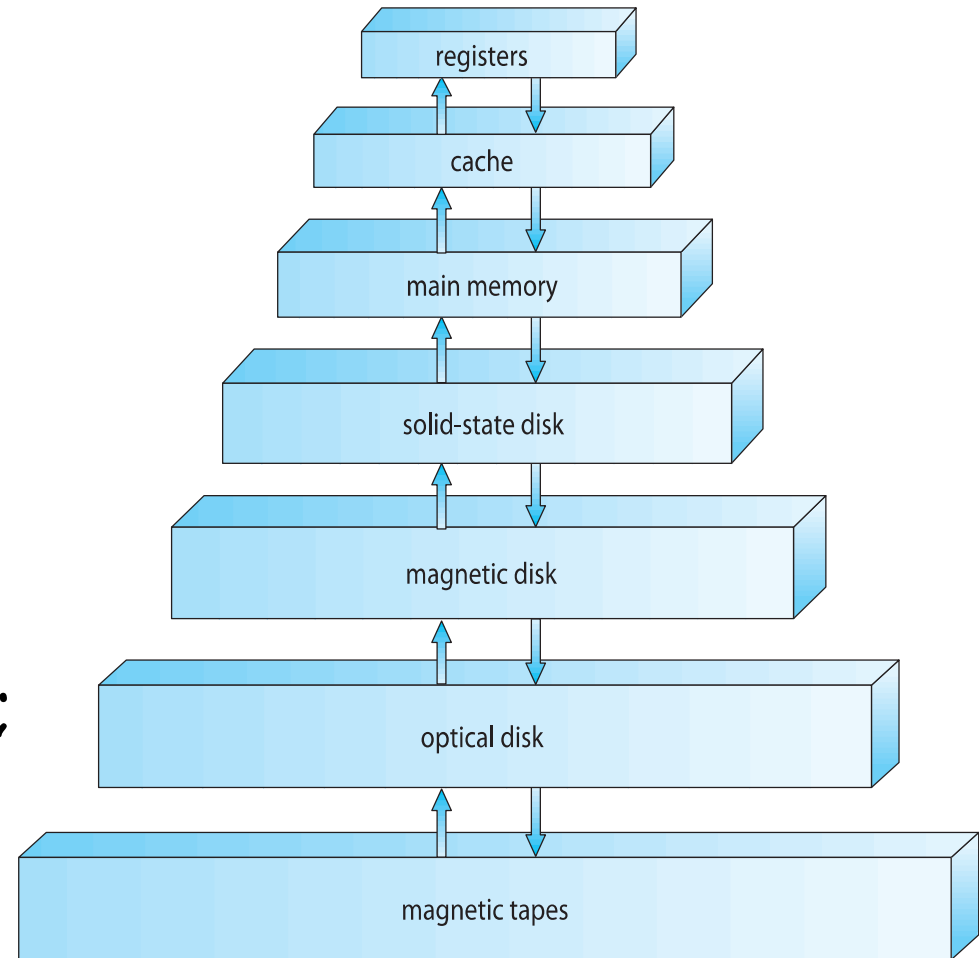
Storage Structure

- Main memory - only large storage media that the CPU can access directly
 - Random access
 - Typically volatile
- Secondary storage - extension of main memory that provides large nonvolatile storage capacity
 - Typically realized by magnetic disks (or hard disks) and SSDs
- Magnetic disks - large capacity, nonvolatile
 - Disk surface is logically divided into tracks, which are subdivided into sectors
 - The disk controller determines the logical interaction between the device and the computer
- Solid-state disks (SSDs) - faster than magnetic disks, nonvolatile
 - Currently, based on flash memories, especially NAND flash
 - Becoming more popular



Storage Hierarchy

- Storage systems organized in hierarchy based on
 - Speed
 - Cost
 - Volatility
- Caching - copying information into faster storage system; e.g., main memory can be viewed as a cache for secondary storage





Caching

- Important principle, performed at many levels in a computer (in hardware, operating system, software)
 - Information in use copied from slower to faster storage temporarily
 - Faster storage (cache) checked first to determine if information is there
 - If it is, information used directly from the cache (fast)
 - If not, data copied to cache and used there
 - Cache smaller than storage being cached
 - Cache management important design problem
 - Cache size and replacement policy
- *You must understand the difference between caching and buffering.*



Units of Digital Information

■ Basic units

- bit
- byte: 8 bits on most computers. may be of different size
- word: the number of bytes that a computer processes physically or logically at a time, typically 4 bytes

■ More units: K, M, G, T, P, E, Z, Y

- powers of 1000 (SI unit) or powers of $2^{10}=1024$ (binary unit)?
- base 10: kilobyte (KB, 10^3) - megabyte (MB, 10^6) - gigabyte (GB, 10^9) - terabyte (TB, 10^{12}) - petabyte (PB, 10^{15}) - exabyte (EB, 10^{18}) - zettabyte (ZB, 10^{21}) - yottabyte (YB, 10^{24})
- base 2: (use binary prefix) kibibyte (KiB, 2^{10}), mebibyte (MiB, 2^{20}), gibibyte (GiB, 2^{30}), tebibyte (TiB, 2^{40}), pebibyte (PiB, 2^{50}), exbibyte (EiB, 2^{60}), zebibyte (ZiB, 2^{70}), yobibyte (YiB, 2^{80})

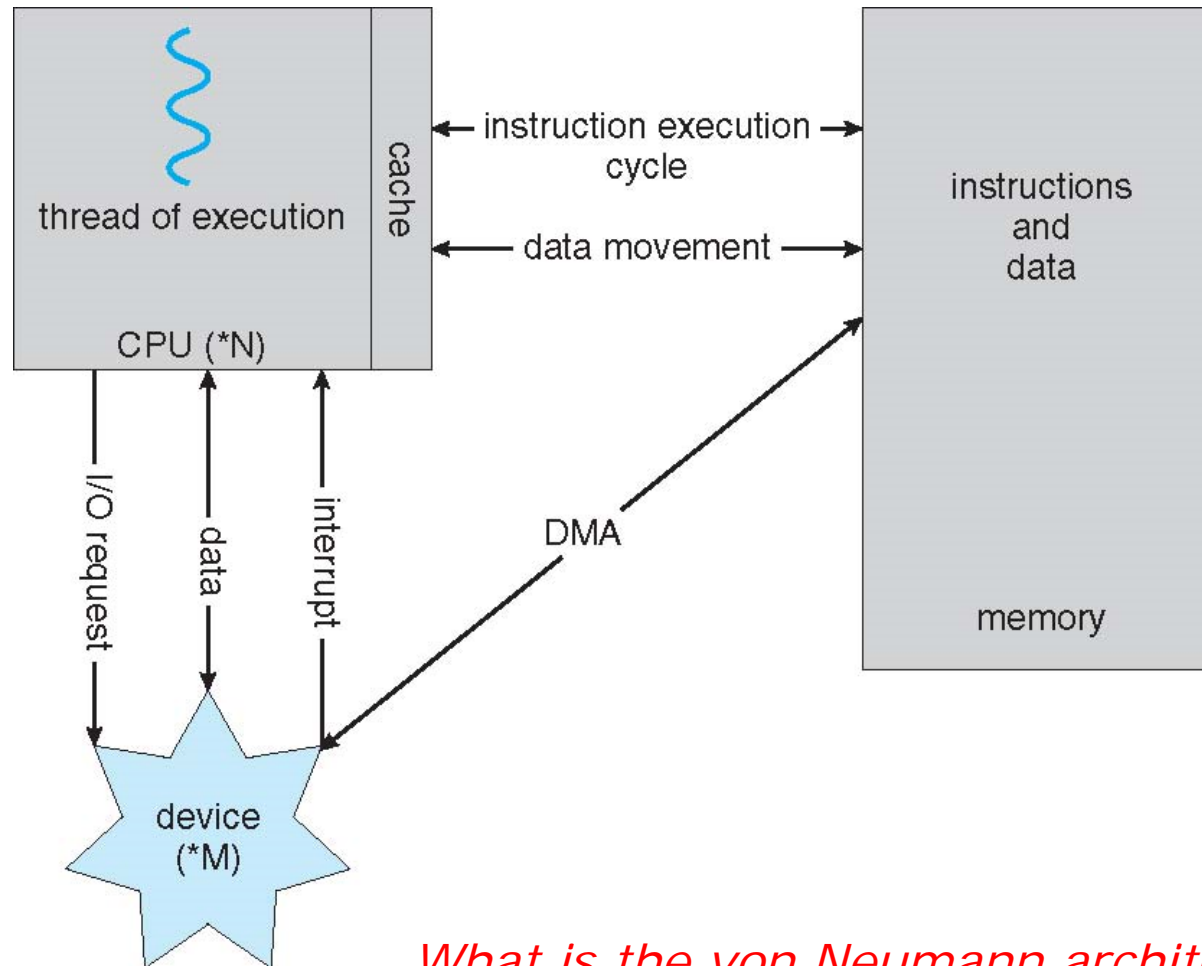
• http://en.wikipedia.org/wiki/Binary_prefix

■ Usage

- in computer hardware and software: interpret on a case-by-case basis, either base 10 or base 2
- networking and communication: usually SI prefix (powers of 10^3)



How a Modern Computer Works



What is the von Neumann architecture?

- memory locations – sequential addresses
- code + data in memory - mixed

Computer System Architecture




- Most systems use a single **general-purpose processor** (PDAs through mainframes)
 - Most systems have special-purpose processors as well
- **Multiprocessor systems**
 - Also known as parallel systems, tightly-coupled systems sharing the main memory
 - Advantages include:
 - Increased throughput
 - Economy of scale
 - Increased reliability - graceful degradation or fault tolerance
 - Two types:
 - Asymmetric Multiprocessing
 - Symmetric Multiprocessing
- **Clustered systems**
 - Multiple autonomic computers connected via a high-speed network
 - Loosely-coupled system usually sharing the secondary storage




Multiprocessor Systems

- Characteristics
 - CPUs communicate through **shared data** in the (common) memory
 - OS may run either on a single CPU or all CPUs
- Asymmetric multiprocessing (ASMP, AMP)
 - CPUs may be different in functionality or structure
 - Usually, only a single processor runs the OS and others run as slaves: master-slave relationship
- **Symmetric multiprocessing (SMP)**
 - All CPUs are treated equally
 - OS may run on any processors
- **Multicore**
 - A microprocessor containing multiple CPUs
 - Either SMP or ASMP



Symmetric multiprocessing architecture




Dual-core microprocessor



Clustered Systems

- Multiple autonomous systems working together
 - Usually sharing storage via a **storage area network (SAN)**
 - Provides a high-availability service which survives failures
 - Asymmetric clustering has one machine in hot-standby mode
 - Symmetric clustering has multiple nodes running applications, monitoring each other
 - Some clusters are for high-performance computing (HPC)
 - Applications must be written to use parallelization
 - Useful for Internet servers, DB servers, graphics servers, etc.



- ★ **autonomous system:**
a computer system with its own OS
- ★ **storage area network (SAN):**
a dedicated network of storage devices
such as disk arrays and tape drives

Operating System Structure & Operations




■ Multiprogramming needed for efficiency

- When it has to wait (for I/O for example), OS switches to another job
- Multiprogramming organizes jobs (code and data) so CPU always has one to execute
- A subset of total jobs in system is kept in memory
- One job selected and run via job scheduling
- Multitasking
 - Multiple tasks running on CPU during the same period of time
 - Based on the multiprogramming technique

■ Timesharing systems

- CPU switches jobs regularly while each job is given a fixed time slot (say, 20 msec) in round robin fashion
- Based on multiprogramming and multitasking methods
- Provide the basis for interactive computing in multi-user environment
 - Response time should be < 1 second
- Each user has at least one program executing in memory \Rightarrow process
- If several jobs ready to run at the same time \Rightarrow CPU scheduling

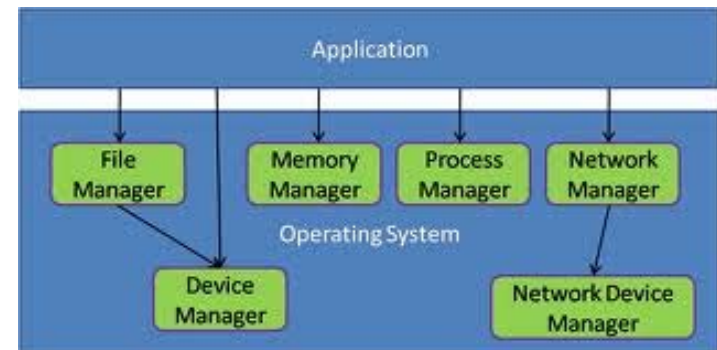


Memory layout for a multiprogrammed system



Operating System Operations

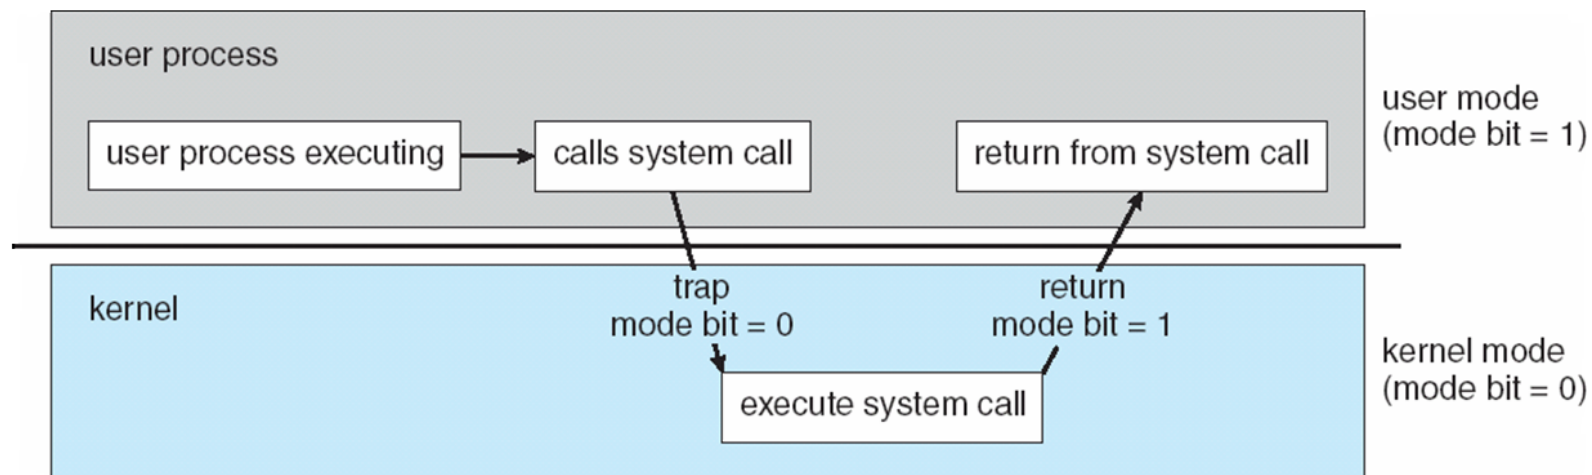
- Often driven by hardware **interrupt** mechanism
- Software error or request creates **exception** or **trap**, considered "**software interrupt**"
 - Division by zero, request for OS service, etc.
 - May be caused by other process problems including an infinite loop, and processes modifying each other or the OS
- Dual-mode operation allows OS to protect itself and other system components
 - Dual modes: user mode and kernel mode
 - Mode bit provided by hardware
 - Provides ability to distinguish when system is running user code or kernel code
 - Some instructions designated as privileged, only executable in kernel mode
 - System call changes mode to *kernel*, return from call resets it to *user*





Transition from User to Kernel Mode

- Example: Timer to prevent infinite loop / process hogging resources
 - [User mode] Call a system service for checking infinite loop
 - System service: Set interrupt after specific period
 - [Kernel mode] (OS) sets the timer to a desired value
 - Hardware timer continues decrementing the timer value
 - When the value becomes zero, generate an interrupt
 - [Kernel mode] Set up the calling process to regain control, or terminate program that exceeds allotted time
 - [User mode] (*If allotted time is not exceeded*) Return from system call → resume the execution



Major Functions of OS

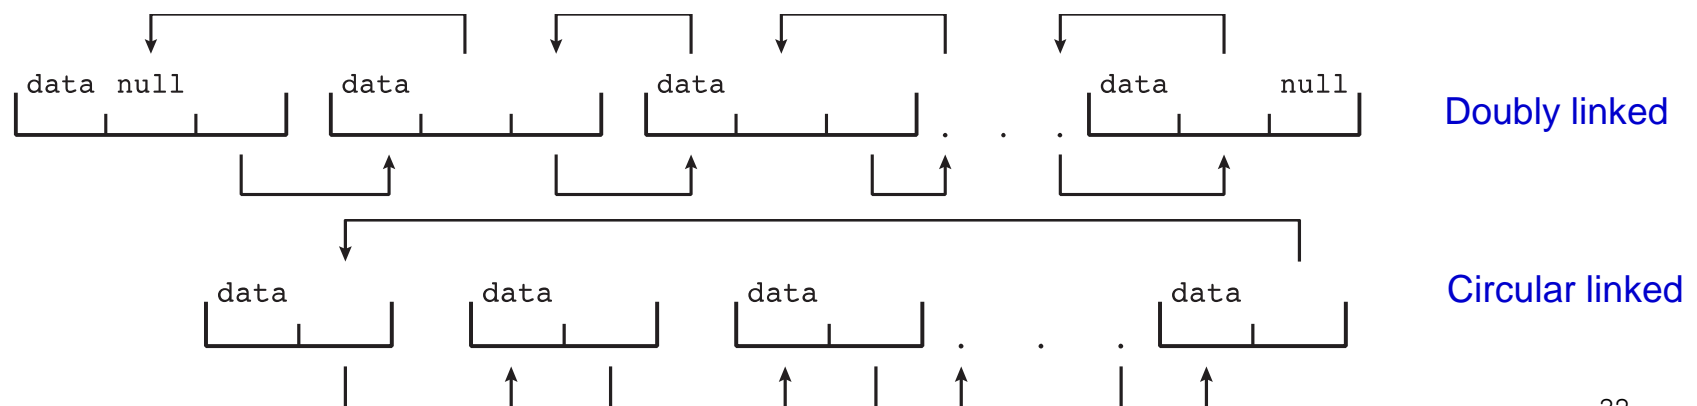


- Process and thread management
 - A process is the object that executes a program
 - OS creates, runs, suspends and terminates a process
- Memory management
 - In von Neumann architecture, all instructions and data are located at main memory
 - Memory space are always short of the space needed for program execution → allocation and deallocation of memory
- Storage and I/O management
 - Abstracts physical storage properties to logical storage unit - file
 - General interface to I/O devices using device drivers
 - I/O memory management through buffering, caching or spooling
- System protection and security
 - Protection - controls access of processes or users to resources defined by the OS
 - Security - defenses the system against internal and external attacks

Kernel Data Structures



- Many similar to standard programming data structures
- Kernel characteristics that affect data structures
 - Dynamic - must accommodate current varying status and operations
 - Intensive search and scanning operations
 - Mostly reside in main memory
 - Consists of data and pointers (addresses)
- Stacks, queues, and lists
 - Stacks and queues implement LIFO and FIFO, respectively
 - Lists
 - Singly linked list, doubly linked list, circular linked list






Kernel Data Structures

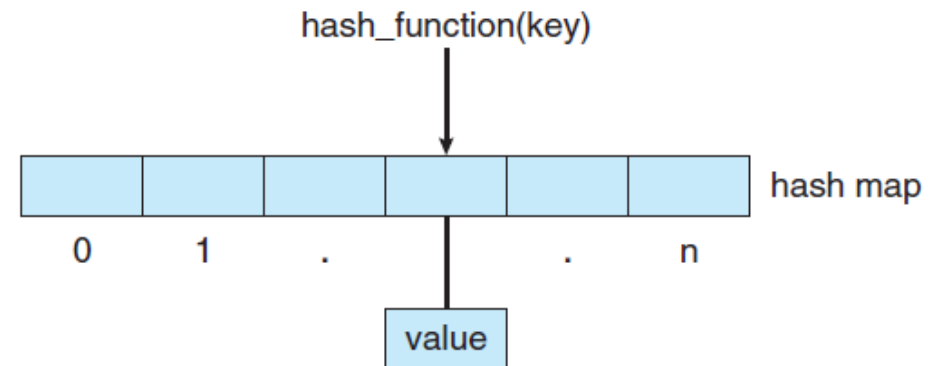
■ Trees

- General tree
- Binary search tree
left \leq right
 - Search performance is $O(n)$
 - Balanced binary search tree is $O(\log n)$



■ Hash function and maps

- can create a hash map



■ Bitmap

- String of n binary digits representing the status of n items

■ Linux data structures

- Defined in include files
<linux/list.h>, <linux/kfifo.h>,
<linux/rbtree.h>

Computing Environments



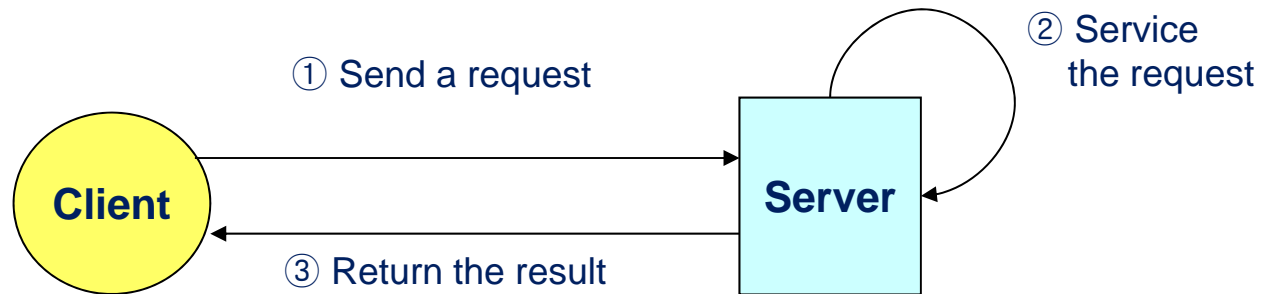
- Traditional: Stand-alone general purpose machines
 - "Centralized computing"
 - But blurred as most systems interconnect with others (i.e. the Internet)
- Distributed computing paradigm
 - Based physically on a collection of separate, possibly heterogeneous, autonomous systems networked together
 - Based logically on a client-server model
 - Computers communicate with each other using messages
 - Diverse types of approaches and implementations, such as Internet/Web computing, mobile computing, cloud computing, etc.

Computing Env. – Distributed System



■ Client-server model

- Basic model for distributed systems



- Implementations

- Client: PCs, laptop computers, smartphones, tablets, etc.
- Server: compute server, file server, Web server, etc.
- Network: LAN(e.g., Ethernet, WiFi), WAN(e.g., Internet), etc.

■ Services involving multiple servers

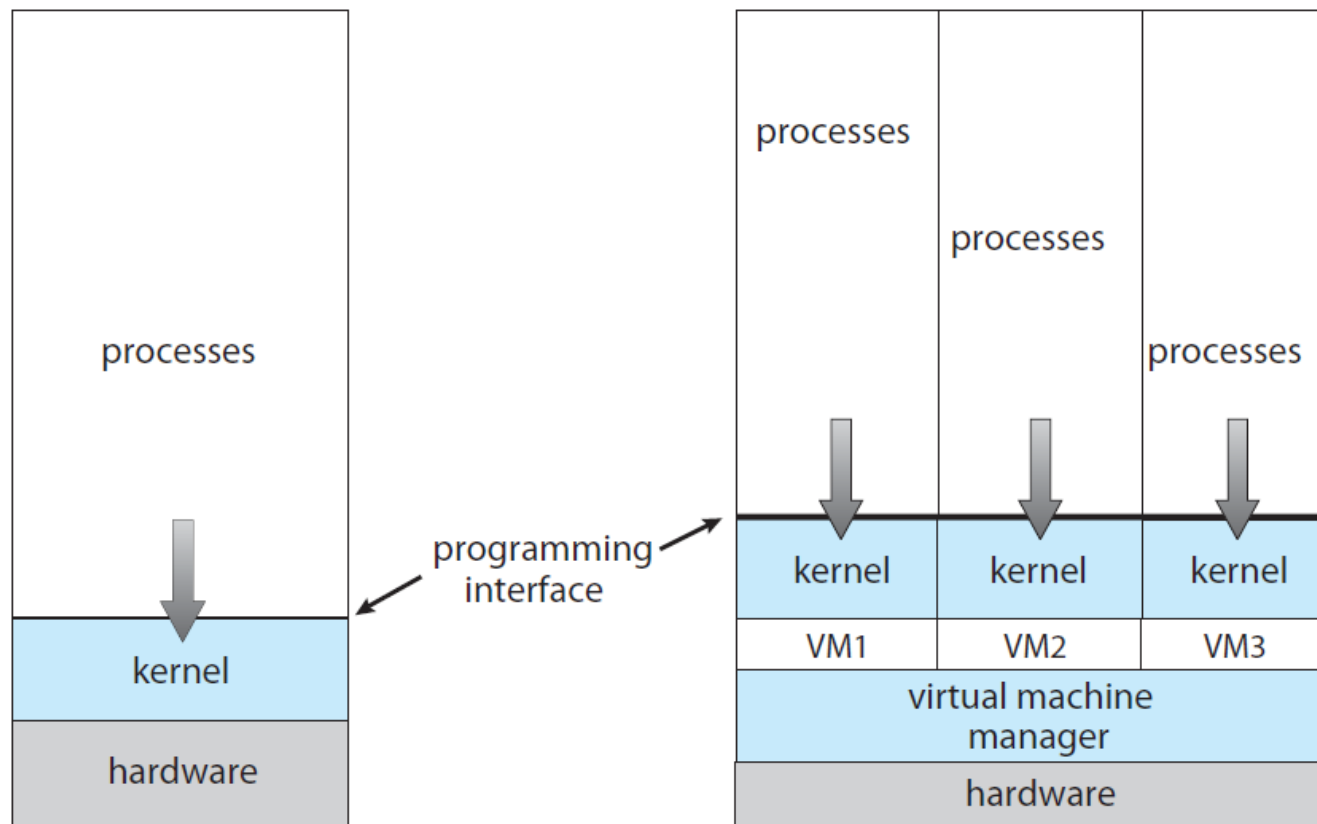
- e.g. sets of data that are distributed across many servers



Computing Env. - Virtualization

- Allows OSs to run on another OS platforms
 - Vast and growing industry
 - Emulation used when source CPU type different from target type (i.e. PowerPC to Intel x86) - very slow, though
- [Hardware] **Virtualization**
 - Creation of a virtual machine that operates like a real computer with an OS
 - OS natively compiled for CPU, running guest OSs also natively compiled
 - Virtual machine manager [or monitor] (**VMM**) provides virtualization services
- Use cases involve laptops and desktops running multiple OSs for exploration or compatibility
 - Apple laptop running Mac OS X host, Windows as a guest
 - Developing apps for multiple OSs without having to have multiple systems
 - Executing and managing compute environments within data centers or cloud servers

Computing Env. – Virtualization Model



(a)

Without virtualization

(b)

VMware: Virtual machine manager(VMM) may run with or without the host OS

Computing Env. – Real-Time Embedded Systems



- **Real-time embedded systems** most prevalent form of computers
 - Applications: diverse, special purpose, limited purpose
 - Examples: robots, mobile phones, guided missiles, spacecrafts, all sorts of machines, and many other special computing environments as well
 - Some have OSs called real-time OS or RTOS, some perform tasks without a full-fledged OS
- Real-time OS must be designed to meet time constraints, such as deadline and period
 - Temporal correctness in addition to logical correctness
 - Correct operation only if time constraints are met

Summary



- 운영체제(OS)
 - 정의: 자원의 관리, 프로그램 실행의 제어
 - 기능: 편리한 사용자 인터페이스, 효율적인 자원의 관리
- 컴퓨터시스템 구조
 - 인터럽트(interrupt)
 - DMA
 - 저장장치의 계층적 구성
 - 저장장치의 계층 구조에서 캐싱(caching)의 역할
- 컴퓨터시스템 구성- CPU 개수
 - 프로세서 개수: =1 or >1 ? (multiprocessors)
 - 컴퓨터 개수: =1 or >1 ? (클러스터형, 분산형)
 - 마이크로프로세서 칩 속의 CPU 개수: =1 or >1 (멀티코어)
- 운영체제(OS) 구조
 - 멀티프로그래밍: why?, how?, 운영체제에서의 의미
 - 시분할시스템
- Operating systems (OSs)
 - Definition: resource manager, control program execution
 - Functions: user-friendly interface, efficient resource management
- Computer system organization
 - Interrupt
 - DMA
 - Storage hierarchy
 - Role of caching in the storage hierarchy
- Computer system - # CPU
 - # processors: =1 or >1? (multiprocessors)
 - # computers: =1 or >1? (clustered, distributed)
 - # CPU in a microprocessor chip: =1 or >1 (multicore)
- OS structure
 - multiprogramming: why?, how?, significance in OS
 - timesharing system



Summary (Cont.)

- 운영체제의 동작
 - 소프트웨어 실행에 따른 예외적인 상황에 대한 처리
 - 사용자 모드 vs. 시스템 모드 (혹은 커널 모드)
- 운영체제의 기능적/관리적 구성 요소
 - 프로세스, 메모리, 저장장치, 입출력(I/O), utility, 등
- 시스템 보호 vs. 시스템 보안
- 커널 자료구조
- 컴퓨팅 환경: 분산 컴퓨팅 시스템
 - 클라이언트-서버 컴퓨팅 모델
 - peer-to-peer (P2P) 모델
 - 클라우드 컴퓨팅, 모바일 컴퓨팅, 인터넷/웹 컴퓨팅, 등
- 컴퓨팅 환경: 가상화
 - 가상기계(virtual machine)
- 컴퓨팅 환경: 실시간 내장 시스템
- OS operations
 - Processing exceptions caused by software
 - User mode vs system mode (or kernel mode)
- Functional/managerial components of OS
 - Process, memory, storage, I/O, utility, etc.
- System protection vs. system security
- Kernel data structures
- Computing environments: distributed computing systems
 - Client/server computing model
 - Peer-to-peer computing (P2P) model
 - Cloud computing, mobile computing, Internet/Web computing, etc.
- Computing environment: virtualization
 - Virtual machines
- Computing environment: real-time embedded systems