


Chapter 3: Process Concept

WHAT'S AHEAD:

- Process Concept
- Process Scheduling
 - Operations on Processes
 - Interprocess Communication
- Examples of IPC Systems

WE AIM:

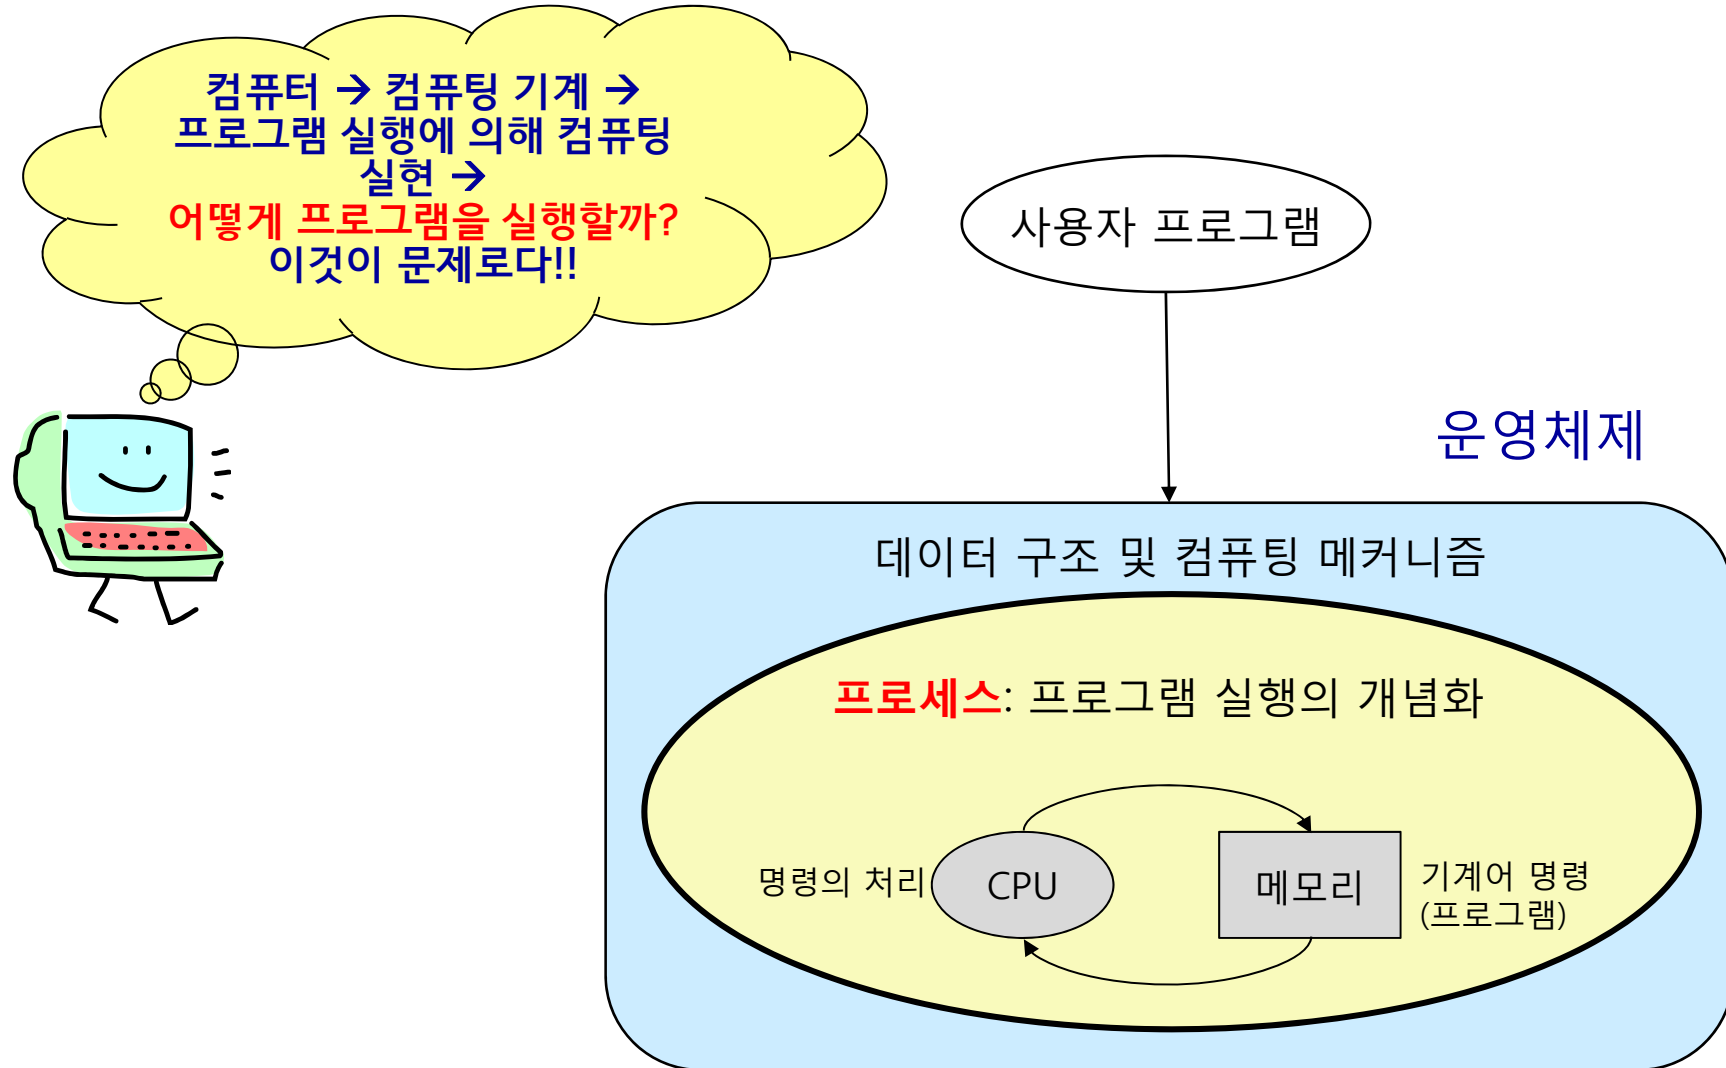
- To introduce the notion of a process
- To describe the various features of processes
- To explore interprocess comm.



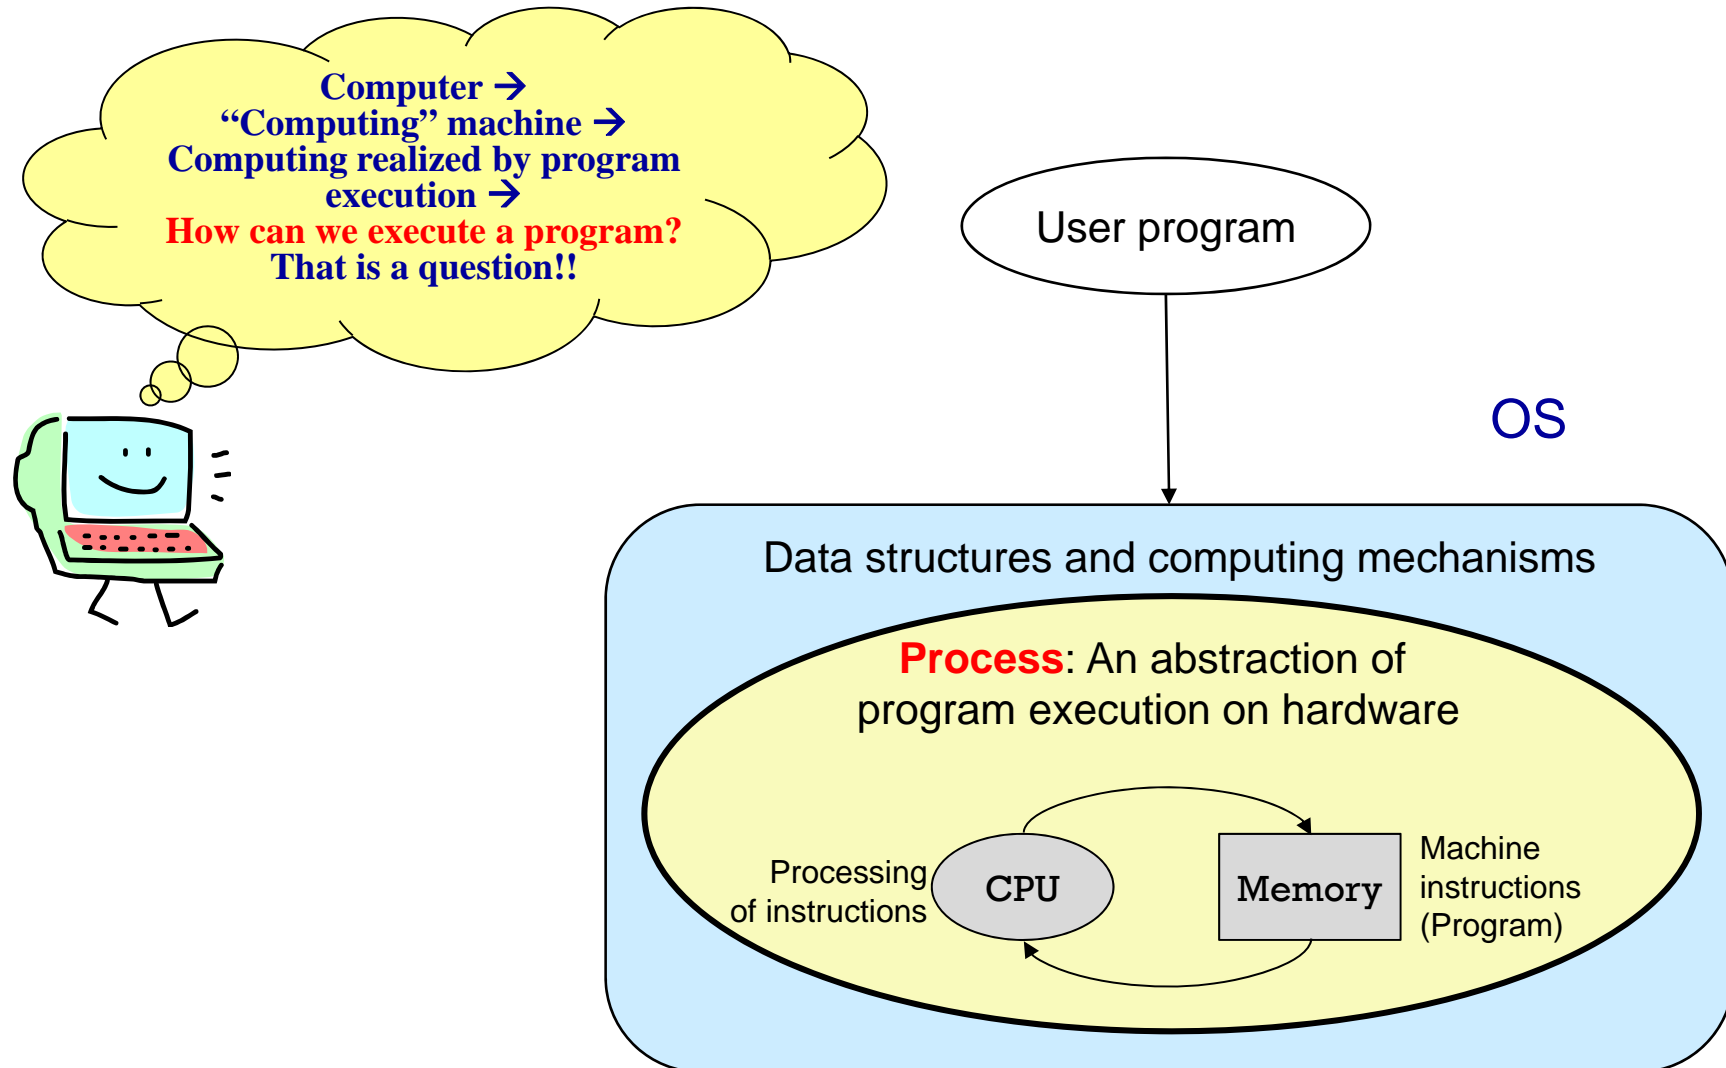
Note: These lecture materials are based on the lecture notes prepared by the authors of the book titled *Operating System Concepts*, 9e (Wiley)

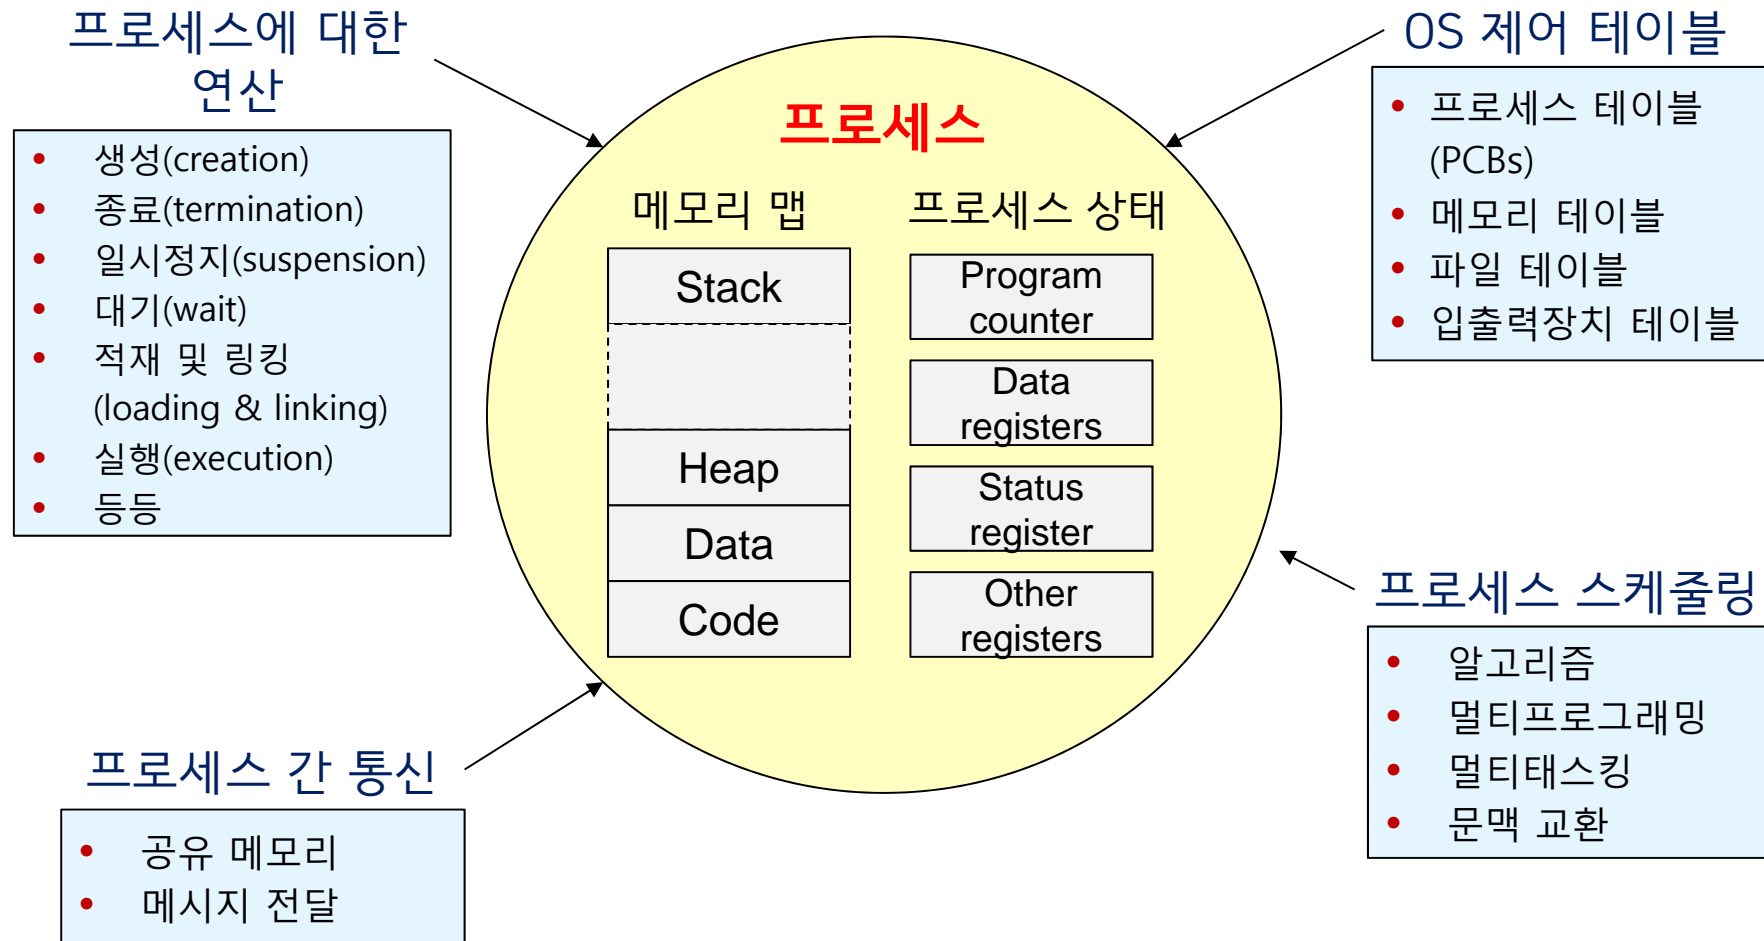
핵심·요점

프로세스란 무엇인가? 사용 동기는?

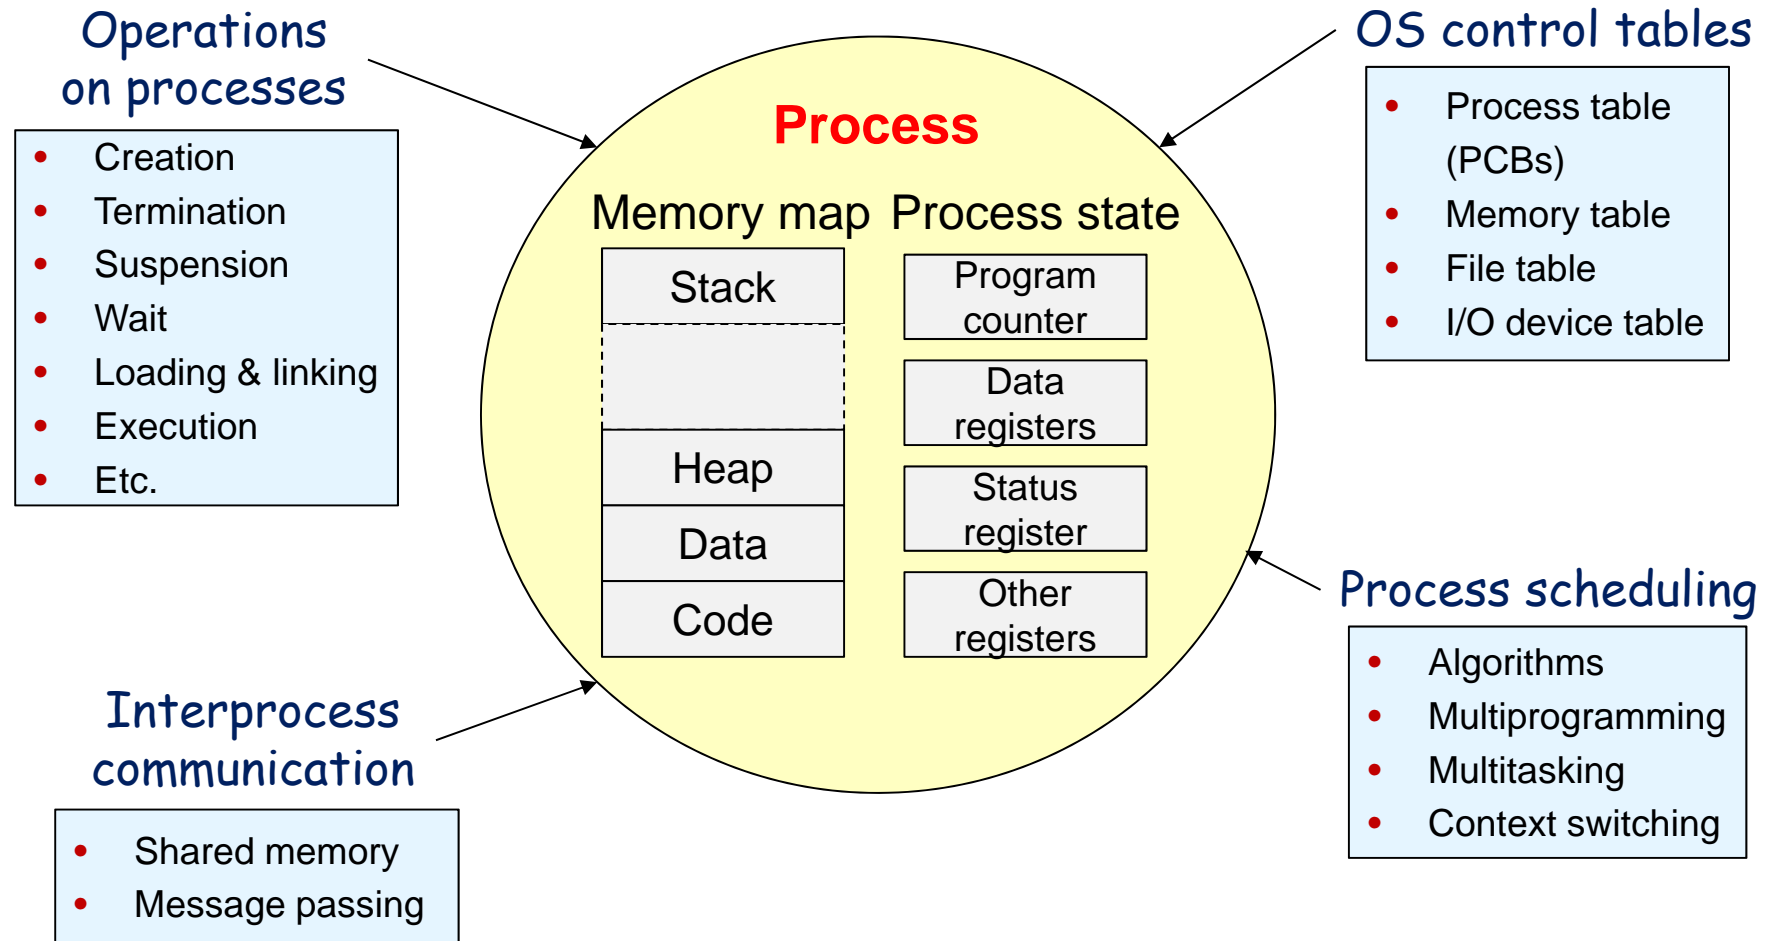


Core Ideas What and Why a Process?





Core Ideas Process and OS Support



Process Concept




- An OS executes a variety of programs:
 - Batch system - **jobs**
 - Time-shared systems - **user programs** or **tasks**
- Textbook uses the terms *job*, *task* and *process* almost interchangeably
- **Process** - Definition
 - A program in execution
 - A series of computational operations performed on the basis of instructions in a program
 - Process execution must progress in sequential fashion
- Program is *passive* entity stored on disk (**executable file**), process is *active*
 - Program becomes process when executable file loaded into memory



Process Concept (Cont.)

- Multiple parts comprising a process
 - The program code, also called **text section**
 - Current activity including **program counter**, processor registers
 - **Stack** containing temporary data
 - Function parameters, return addresses, local variables
 - **Data section** containing global variables
 - **Heap** containing memory dynamically allocated during run time
- Execution of program started via GUI mouse clicks, command line entry of its name, etc
- One program can be several processes
 - Consider multiple users executing the same program

Process in Memory



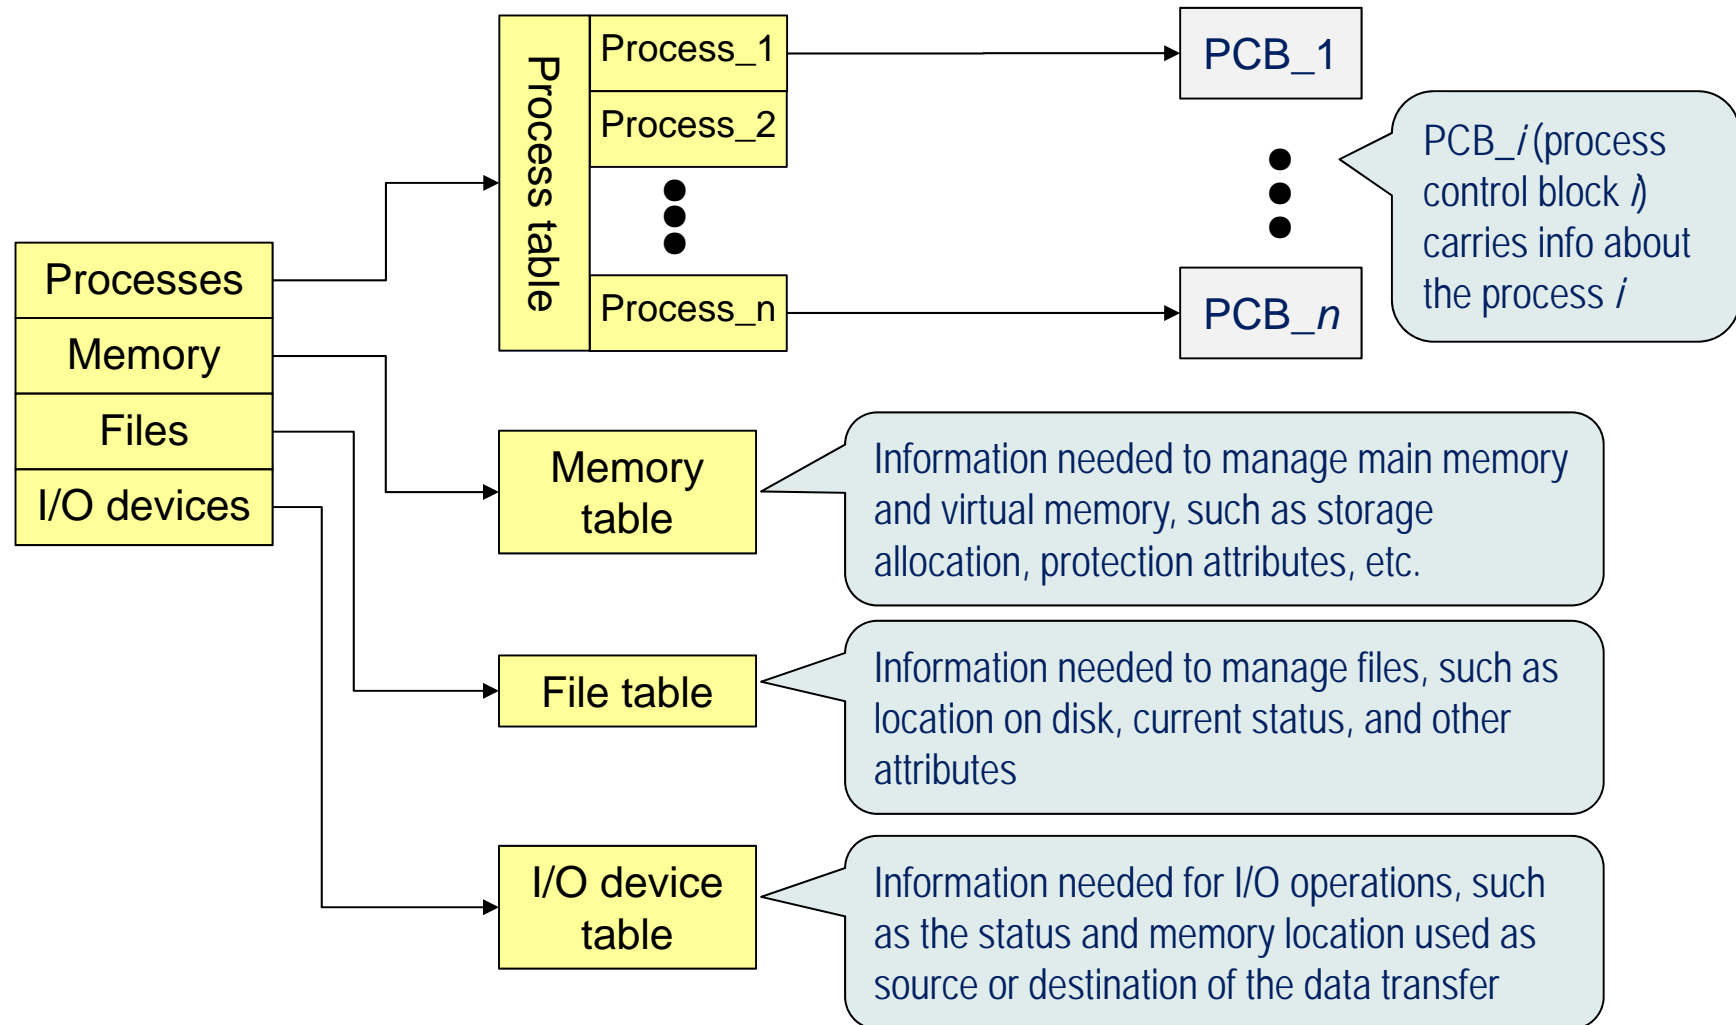
- Part of process image
- Defined in the virtual address space



Process Description

- OS controls physical/logical resources
 - Physical resources: CPU, main memory, I/O devices, etc.
 - Logical resources: processes, virtual memory, data structures, files, etc.
- OS controls resources using “control tables”
 - See next slide
- Process control structures
 - Information on executable program(s) and associated data structures
 - Process image: the collection of program, data, stack, and process control block (PCB)
 - Information on the usage of physical/logical resources
 - PCB contains process attributes that are used by the OS for process control
 - Processor state information: data registers, condition code register, stack pointer, program counter

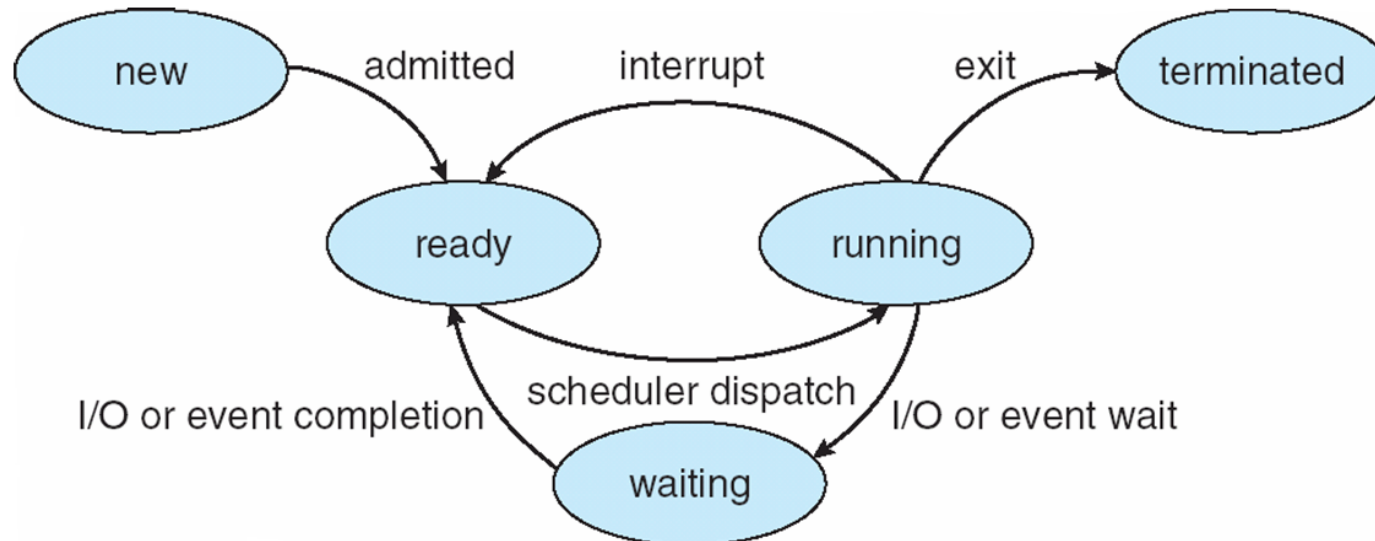
OS Control Tables





Process State

- As a process executes, it changes **state**
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a processor
 - **terminated**: The process has finished execution





Process Control Block (PCB)

Information associated with each process
(also called **task control block**)

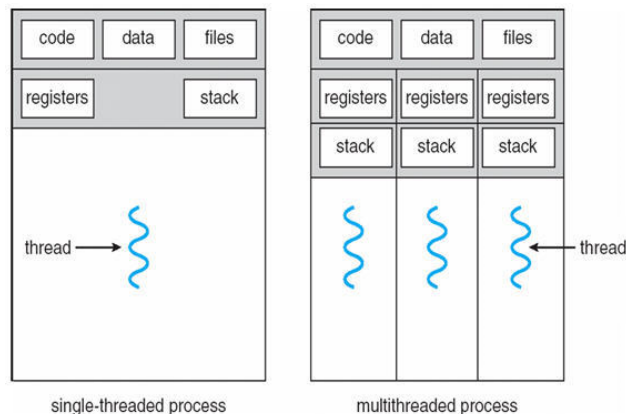
- Process state - running, waiting, etc
- Program counter - location of instruction to next execute
- CPU registers - contents of all process-centric registers
- CPU scheduling information- priorities, scheduling queue pointers
- Memory-management information - memory allocated to the process
- Accounting information - CPU used, clock time elapsed since start, time limits
- I/O status information - I/O devices allocated to process, list of open files





Threads

- So far, process has a single thread of execution
- Consider having multiple program counters per process
 - Multiple locations can execute at once
 - Multiple threads of control -> **threads**
- Must then have storage for thread details, multiple program counters in PCB
- See next chapter



What is a thread (of execution)?

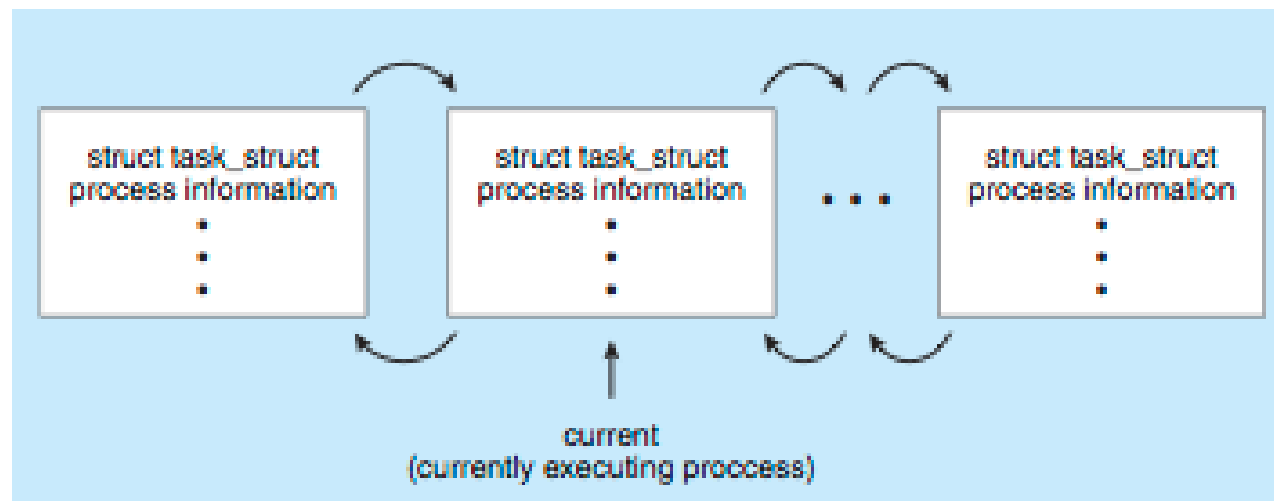
- a sequence of instructions in execution
- an execution trace of instructions
- schedulable as a separate entity
- sometimes, called “lightweight process”



Process Representation in Linux

- Represented by the C structure **task_struct**

```
pid_t pid; /* process identifier */
long state; /* state of the process */
unsigned int time_slice /* scheduling information */
struct task_struct *parent; /* this process's parent */
struct list_head children; /* this process's children */
struct files_struct *files; /* list of open files */
struct mm_struct *mm; /* address space of this process */
```



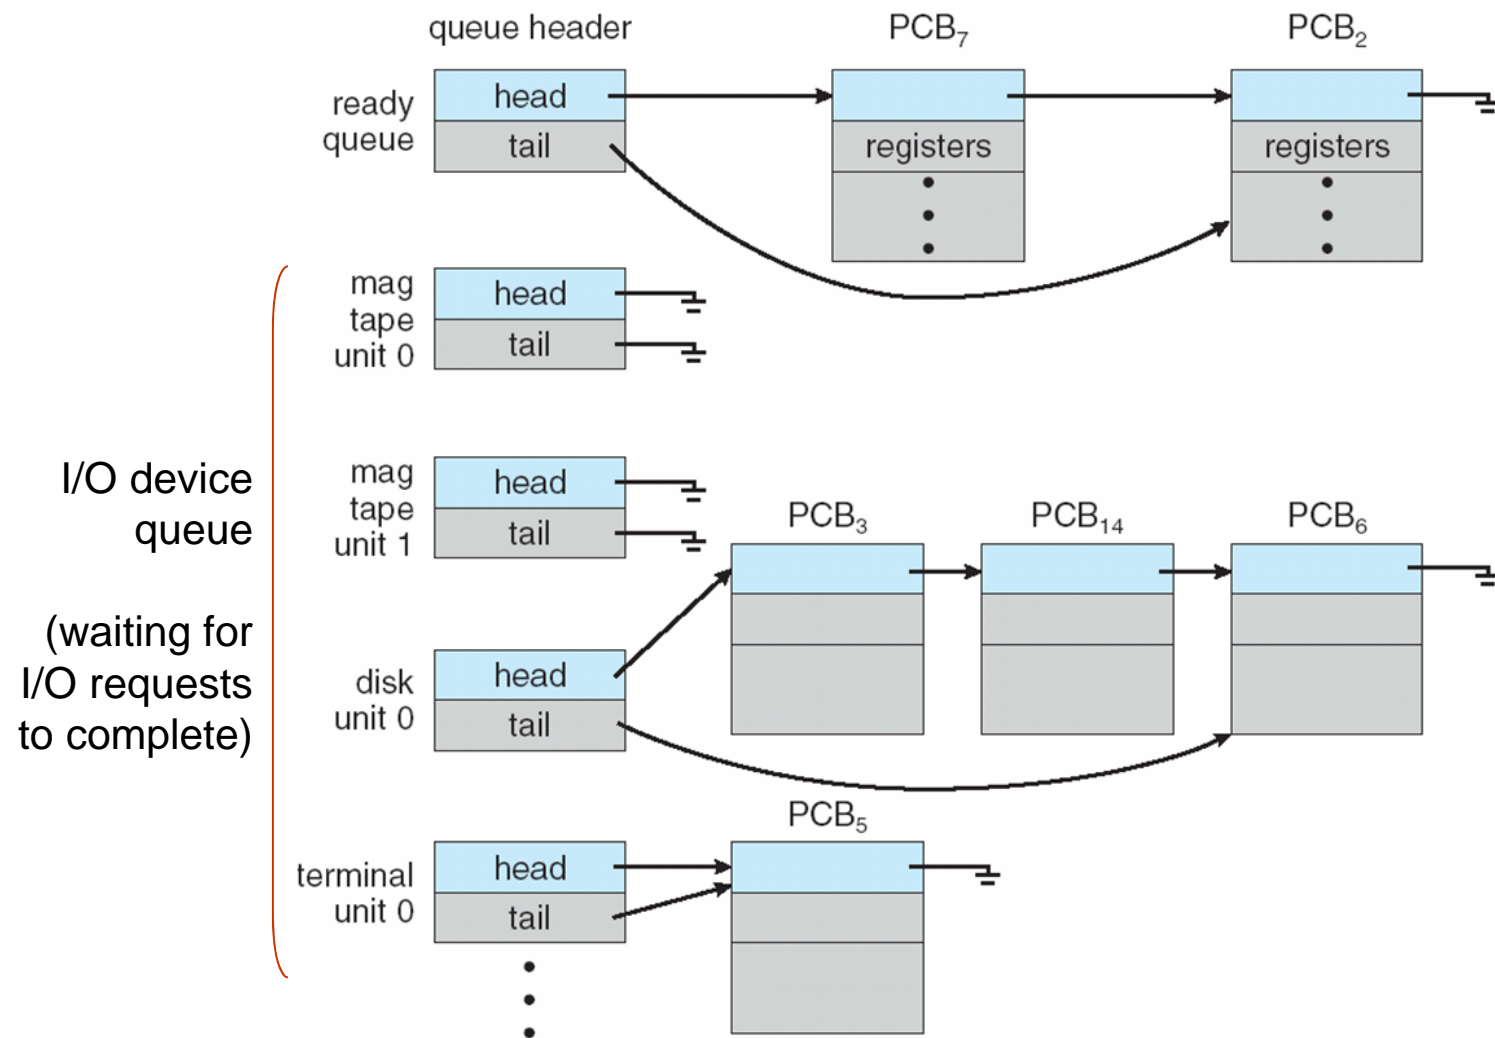
Process Scheduling



- Maximize CPU use, quickly switch processes onto CPU for time sharing
- **Process scheduler** selects one among available processes for next execution on CPU
- Maintains **scheduling queues** of processes
 - **Job queue** - set of all processes in the system
 - **Ready queue** - set of all processes residing in main memory, ready and waiting to execute
 - **Device queues** - set of processes waiting for an I/O device
 - Processes migrate among the various queues



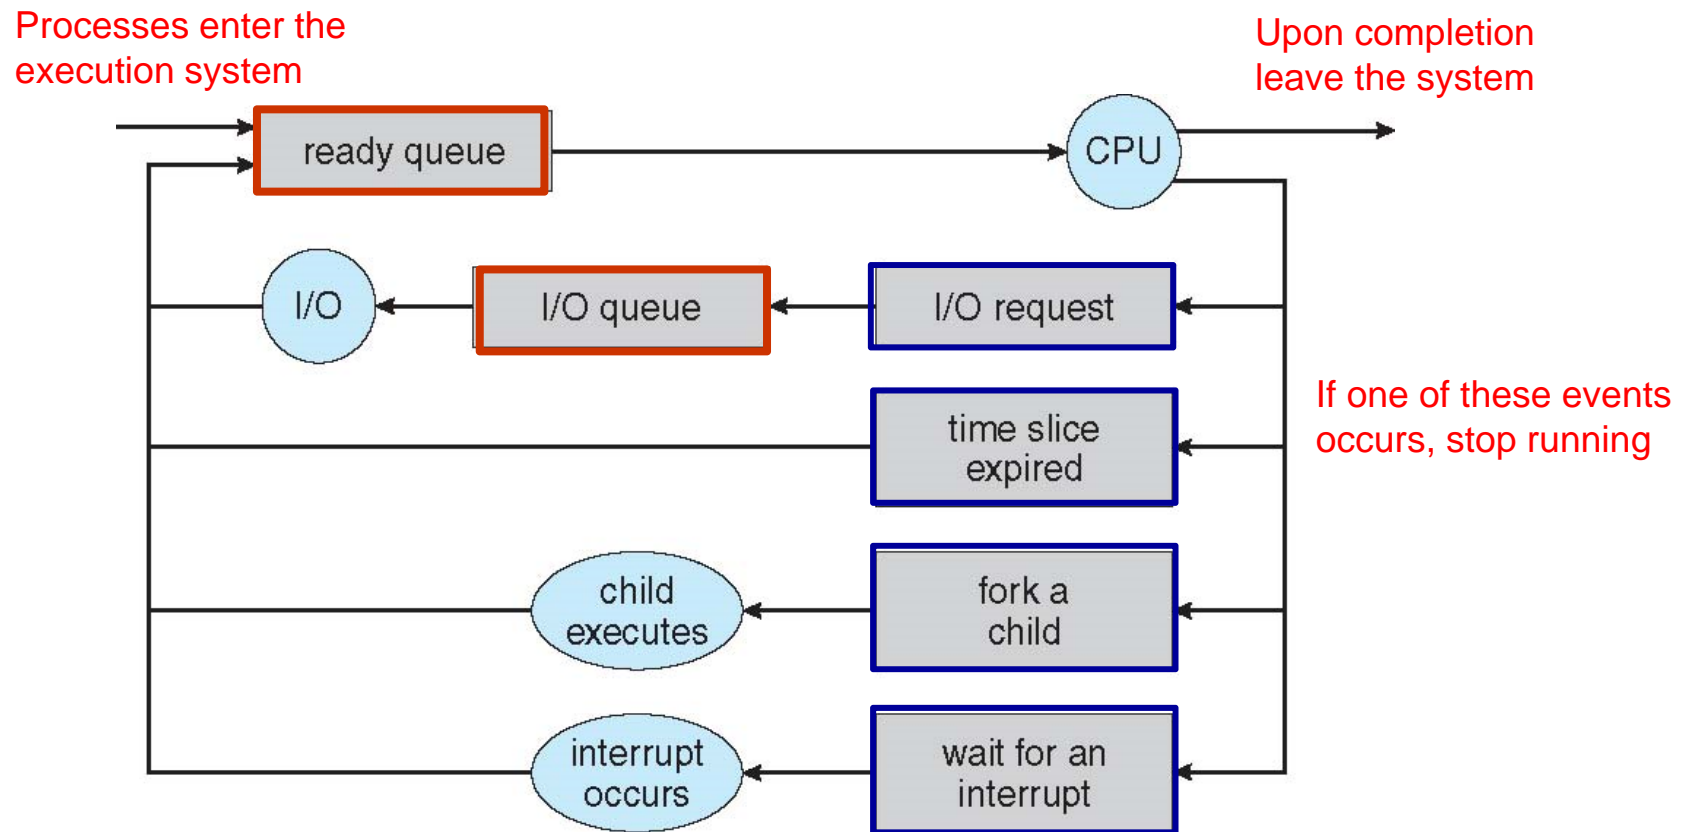
Ready Queue And Various I/O Device Queues



Representation of Process Scheduling



- Queuing diagram represents queues, resources, flows



- Circles: servers, or the resources that serve the queues
- Arrows: the flow of processes in the system
- Red boxes: queues
- Blue boxes: description of events



Schedulers

- **Long-term scheduler** (or **job scheduler**) - selects which processes should be brought into the ready queue
- **Short-term scheduler** (or **CPU scheduler**) - selects which process should be executed next and allocates CPU
 - Sometimes the only scheduler in a system
- Short-term scheduler is invoked very frequently (milliseconds) ⇒ (must be fast)
- Long-term scheduler is invoked very infrequently (seconds, minutes) ⇒ (may be slow)
- The long-term scheduler controls the **degree of multiprogramming**
- Processes can be described as either:
 - **I/O-bound process** - spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** - spends more time doing computations; few very long CPU bursts
- Long-term scheduler strives for good *process mix*



Multitasking in Mobile Systems

- Some systems / early systems allow only one process to run, others suspended
- Apple iOS probably limits multitasking due to battery life and memory use concerns
 - Single **foreground** process - controlled via user interface
 - Multiple **background** processes - in memory, running, but not on the display, and with limits
 - Limits include single, short task, receiving notification of events, specific long-running tasks like audio playback
- Android runs foreground and background, with fewer limits
 - Background process uses a **service** to perform tasks
 - Service can keep running even if background process is suspended
 - Service has no user interface, small memory use



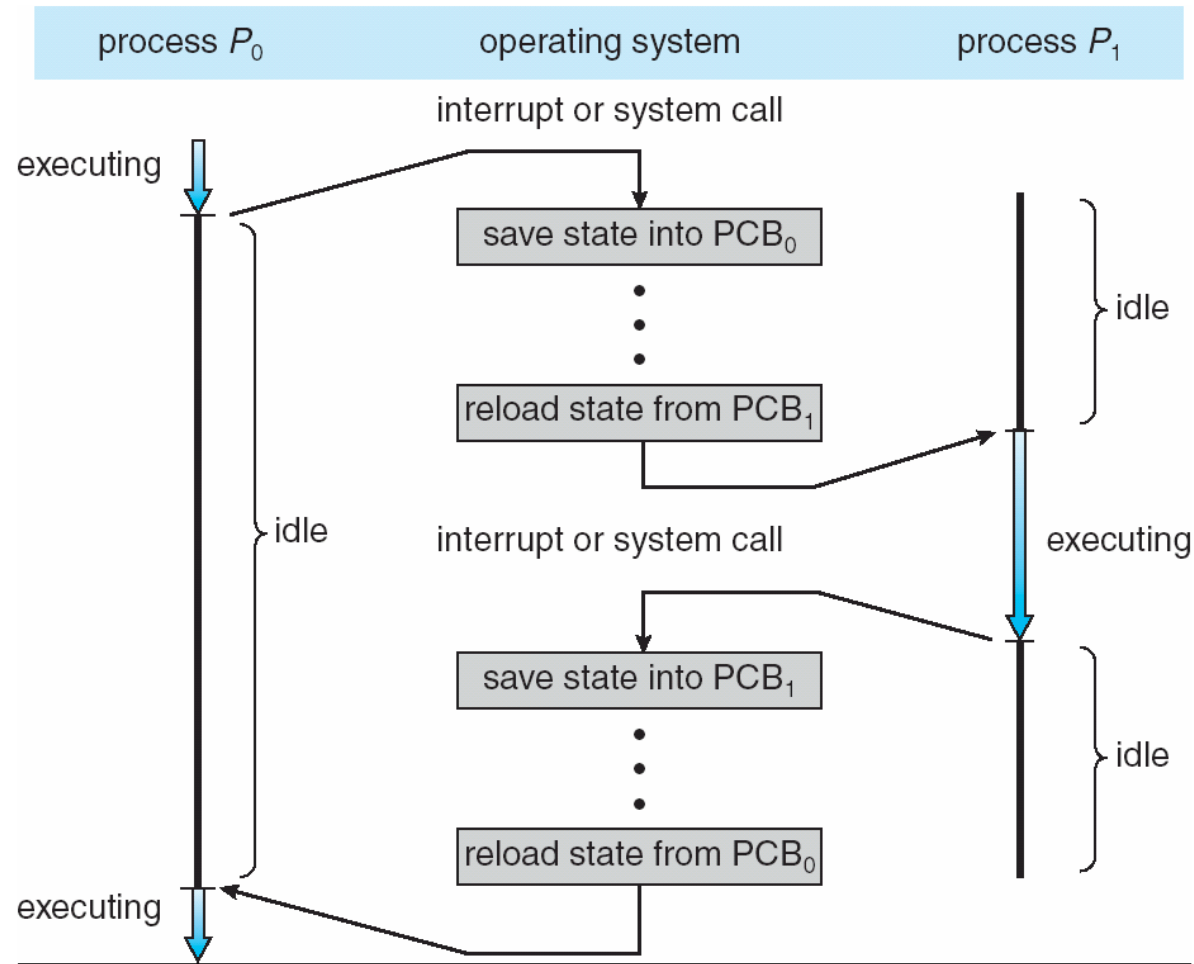
Context Switch

- When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- **Context** of a process is represented in the PCB
 - context: the state of a process
- Context-switch time is overhead; the system does no useful work while switching
 - The more complex the OS and the PCB -> longer the context switch
- Context-switch time: dependent on hardware support
 - Some hardware provides multiple sets of registers per CPU -> multiple contexts loaded at once

Context Switch (Cont.)



CPU switch from process to process



Operations on Processes



- System must provide mechanisms for process management
 - Process creation
 - Process termination
 - Process loading and execution
 - Process suspension
 - and so on

- Such mechanisms are implemented by system services
 - `fork()`, `exec()`, `wait()`, `exit()`, etc.

(You are supposed to be familiar with these systems services.)



Process Creation

- **Parent** process create **children** processes, which, in turn create other processes, forming a **tree** of processes
- Generally, process identified and managed via a **process identifier (pid)**
- Resource sharing options
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution options
 - Parent and children execute concurrently
 - Parent waits until children terminate

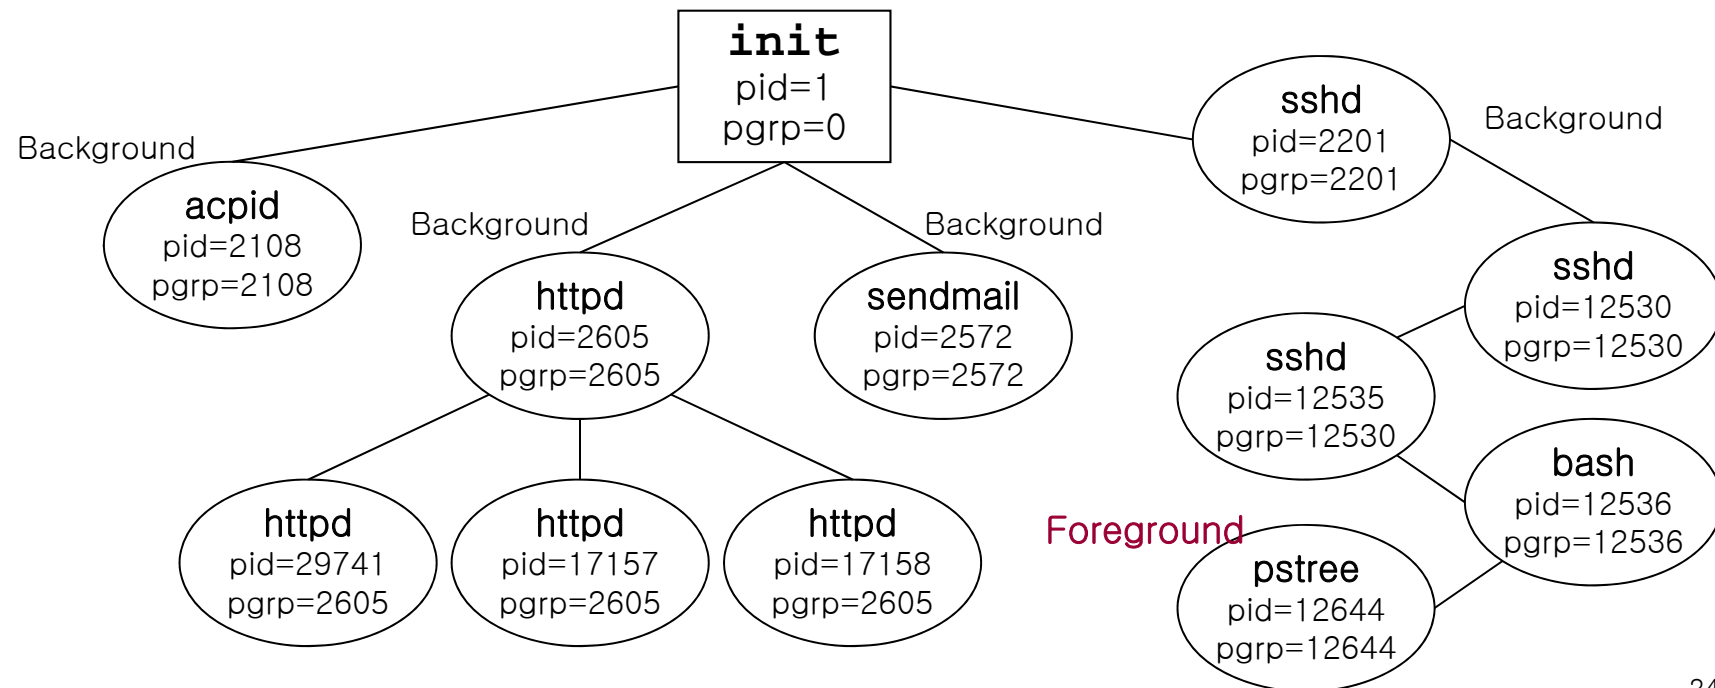


Process Creation (Cont.)

- Example: A tree of processes in Linux

linux> pstree -p // -p: for pid

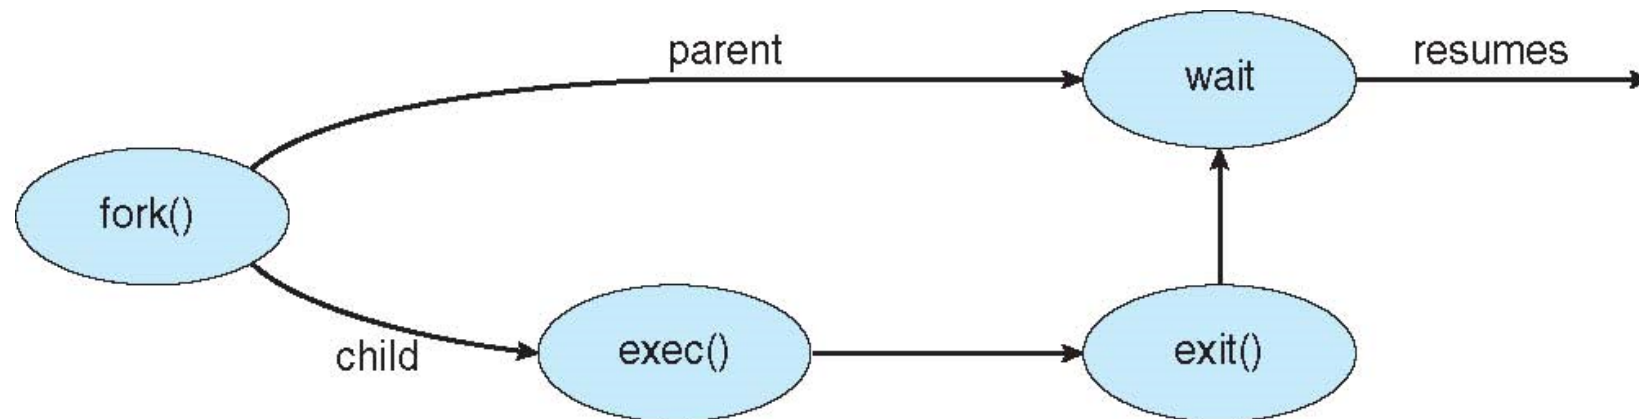
```
init(1)--acpid(2108)
      |-httpd(2605)--httpd(29741)
      |               |-httpd(17157)
      |               |-httpd(17158)
      |               \-httpd(23987)
      |-sendmail(2572)
      \-sshd(2201)---sshd(12530)---sshd(12535)---bash(12536)---pstree(12644)
```





Process Creation (Cont.)

- Address space
 - Child is given a duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - `fork()` system call creates new process
 - `exec()` system call used after a `fork()` to replace the process' memory space with a new program





Examples

Creating a separate process via Unix `fork()` system call (left) and Windows API (right)

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
```

```
int main()
{
    pid_t pid;
```

pid: a local variable, whose value is assigned by `fork()`

```
    /* fork a child process */
    pid = fork();
```

```
    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
```

```
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
```

```
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }
```

```
    return 0;
}
```

The `getpid()` function returns the process ID of the calling process.

```
#include <stdio.h>
#include <windows.h>
```

```
int main(VOID)
{
    STARTUPINFO si;
    PROCESS_INFORMATION pi;
```

```
    /* allocate memory */
    ZeroMemory(&si, sizeof(si));
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));
```

```
    /* create child process */
    if (!CreateProcess(NULL, /* use command line */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* command */
        NULL, /* don't inherit process handle */
        NULL, /* don't inherit thread handle */
        FALSE, /* disable handle inheritance */
        0, /* no creation flags */
        NULL, /* use parent's environment block */
        NULL, /* use parent's existing directory */
        &si,
        &pi))
```

```
    {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
```

```
    /* parent will wait for the child to complete */
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");
```

```
    /* close handles */
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```



Process Termination

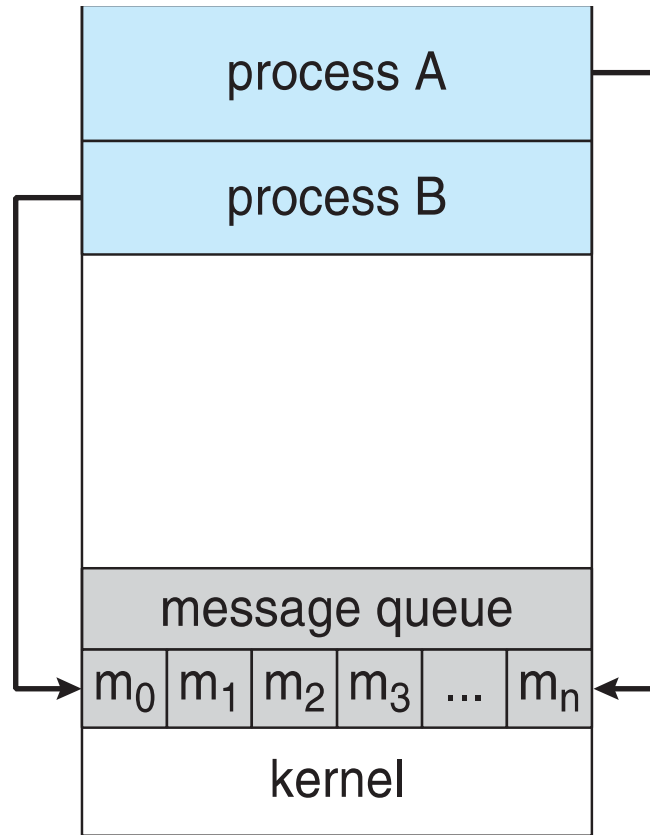
- Process executes last statement and asks the operating system to delete it (`exit()`)
 - Output data from child to parent (via `wait()`)
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (`abort()`)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - Some operating systems do not allow child to continue if its parent terminates
 - All children terminated - **cascading termination**
- Wait for termination, returning the pid:
`pid_t pid; int status;`
`pid = wait(&status);`
- If no parent waiting, then terminated process is a **zombie**
- If parent terminated, processes are **orphans**

Interprocess Communication




- Processes within a system may be *independent* or *cooperating*
- ***Independent*** process
 - cannot affect or be affected by the execution of another process
- ***Cooperating*** process
 - can affect or be affected by the execution of another process
 - Reasons for cooperating processes:
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience
 - Cooperating processes need **interprocess communication (IPC)**
- Two models of IPC
 - **Shared memory**
 - **Message passing**

Communications Models



(a)

Message passing




(b)

Shared memory



Producer-Consumer Problem

- Paradigm for cooperating processes: a *producer* process produces information that is consumed by a *consumer* process
 - **bounded-buffer** assumes that there is a fixed buffer size
 - **unbounded-buffer** places no practical limit on the size of the buffer



Source: <http://www.dcs.warwick.ac.uk/~sgm/open/pc.html>

Interprocess Communication – Shared Memory



- Bounded-Buffer - Shared-Memory Solution
 - Shared data

```
#define BUFFER_SIZE 10
typedef struct {
    . . .
} item;

item buffer[BUFFER_SIZE];
int in = 0;
int out = 0;
```

- Solution is correct, but can only use BUFFER_SIZE-1 elements (next slide)

Bounded-Buffer – Producer & Consumer



Case of shared memory

■ Producer

```
item next_produced;  
while (true) {  
    /* produce an item in next_produced */  
    while (((in + 1) % BUFFER_SIZE) == out)  
        ; /* do nothing */  
    buffer[in] = next_produced;  
    in = (in + 1) % BUFFER_SIZE;  
}
```


Buffer

Initially



in, out

later



out in

■ Consumer

```
item next_consumed;  
while (true) {  
    while (in == out)  
        ; /* do nothing */  
    next_consumed = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
  
    /* consume the item in next_consumed */  
}
```


Interprocess Communication – Message Passing



- Mechanism for processes to communicate and to synchronize their actions
- Message system - processes communicate with each other without resorting to shared variables
- IPC facility provides two operations:
 - `send(message)` - message size fixed or variable
 - `receive(message)`
- If P and Q wish to communicate, they need to:
 - establish a ***communication link*** between them
 - exchange messages via send/receive
- Implementation of communication link
 - physical (e.g., shared memory, hardware bus)
 - logical (e.g., direct or indirect, synchronous or asynchronous, automatic or explicit buffering)
- Specifying a source/destination process
 - direct addressing, indirect addressing

Direct Communication

Case of msg passing



- Processes must name each other explicitly:
 - `send(P, message)` - send a message to process P
 - `receive(Q, message)` - receive a message from process Q

- Properties of communication link
 - Links are established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional

Indirect Communication

Case of msg passing



- Messages are directed to and received from mailboxes (also referred to as ports)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox

- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional

Indirect Communication (Cont.)



Case of msg passing

- Operations
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox

- Primitives are defined as:
 - `send(A, message)` - send a message to mailbox *A*
 - `receive(A, message)` - receive a message from mailbox *A*

Synchronization



Case of msg passing

- Message passing may be either **blocking** or **non-blocking**, also known as **synchronous** and **asynchronous**

- Types of `send()` and `receive()` primitives
 - **Blocking send** has the sender block until the message is received
 - **Non-blocking send** has the sender send the message and continue
 - **Blocking receive** has the receiver block until a message is available
 - **Non-blocking receive** has the receiver receive a valid message or null

Synchronization (Cont.)



- Different combinations possible
 - If both send and receive are blocking, we have a rendezvous (between the sender and the receiver)
- In the event of rendezvous, producer-consumer becomes trivial

```
message next_produced;
while (true) {
    /* produce an item in next_produced */
    send(next_produced);
}
```

```
message next_consumed;
while (true) {
    receive(next_consumed);

    /* consume the item in next_consumed */
}
```

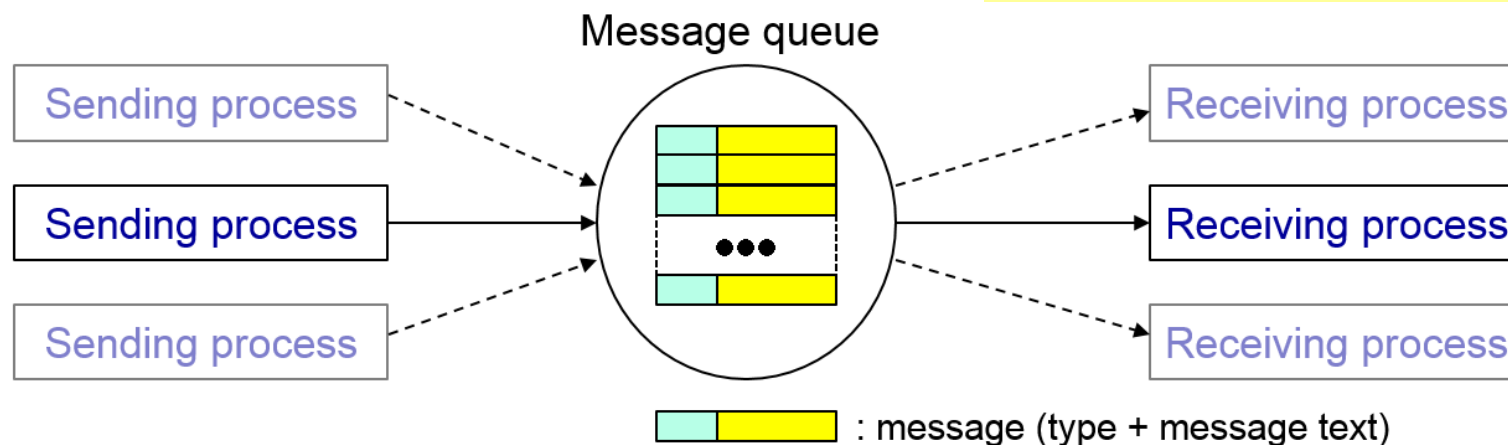
Buffering



Case of msg passing

- Queue of messages attached to the link; implemented in one of three ways
 1. Zero capacity - 0 messages
Sender must wait for receiver (rendezvous)
 2. Bounded capacity - finite length of n messages
Sender must wait if link full
 3. Unbounded capacity - infinite length
Sender never waits

There may be multiple senders and/or multiple receivers.





Examples of IPC Systems

- Example of IPC system - POSIX
- POSIX Shared Memory
 - Process first creates shared memory segment
`shm_fd = shm_open(name, O_CREAT |
O_RDWR, 0666);`
 - Also used to open an existing segment to share it
 - Set the size of the object
`ftruncate(shm_fd, 4096);`
 - Now the process could write to the shared memory
`sprintf(shared_memory, "Writing to
shared memory");`

IPC POSIX Producer



```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
```

```
int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* strings written to shared memory */
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* create the shared memory object */
    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);

    /* configure the size of the shared memory object */
    ftruncate(shm_fd, SIZE);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);

    /* write to the shared memory object */
    sprintf(ptr,"%s",message_0);
    ptr += strlen(message_0);
    sprintf(ptr,"%s",message_1);
    ptr += strlen(message_1);

    return 0;
}
```

Shared memory

IPC POSIX Consumer



```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>

int main()
{
    /* the size (in bytes) of shared memory object */
    const int SIZE 4096;
    /* name of the shared memory object */
    const char *name = "OS";
    /* shared memory file descriptor */
    int shm_fd;
    /* pointer to shared memory object */
    void *ptr;

    /* open the shared memory object */
    shm_fd = shm_open(name, O_RDONLY, 0666);

    /* memory map the shared memory object */
    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0);

    /* read from the shared memory object */
    printf("%s", (char *)ptr);

    /* remove the shared memory object */
    shm_unlink(name);

    return 0;
}
```

Shared memory



Examples of IPC Systems - Mach

- Mach communication is message-based
 - Even system calls are messages
 - Each task gets two mailboxes at creation - Kernel for kernel ↔ task comm. and Notify for event notification
 - Only three system calls needed for message transfer
`msg_send()`, `msg_receive()`, `msg_rpc()`
 - Mailboxes (called "port" in Mach) needed for communication, created via
`port_allocate()`
 - Send and receive are flexible, for example four options if mailbox full:
 - Wait indefinitely
 - Wait at most n milliseconds
 - Return immediately
 - Temporarily cache a message

Mach OS

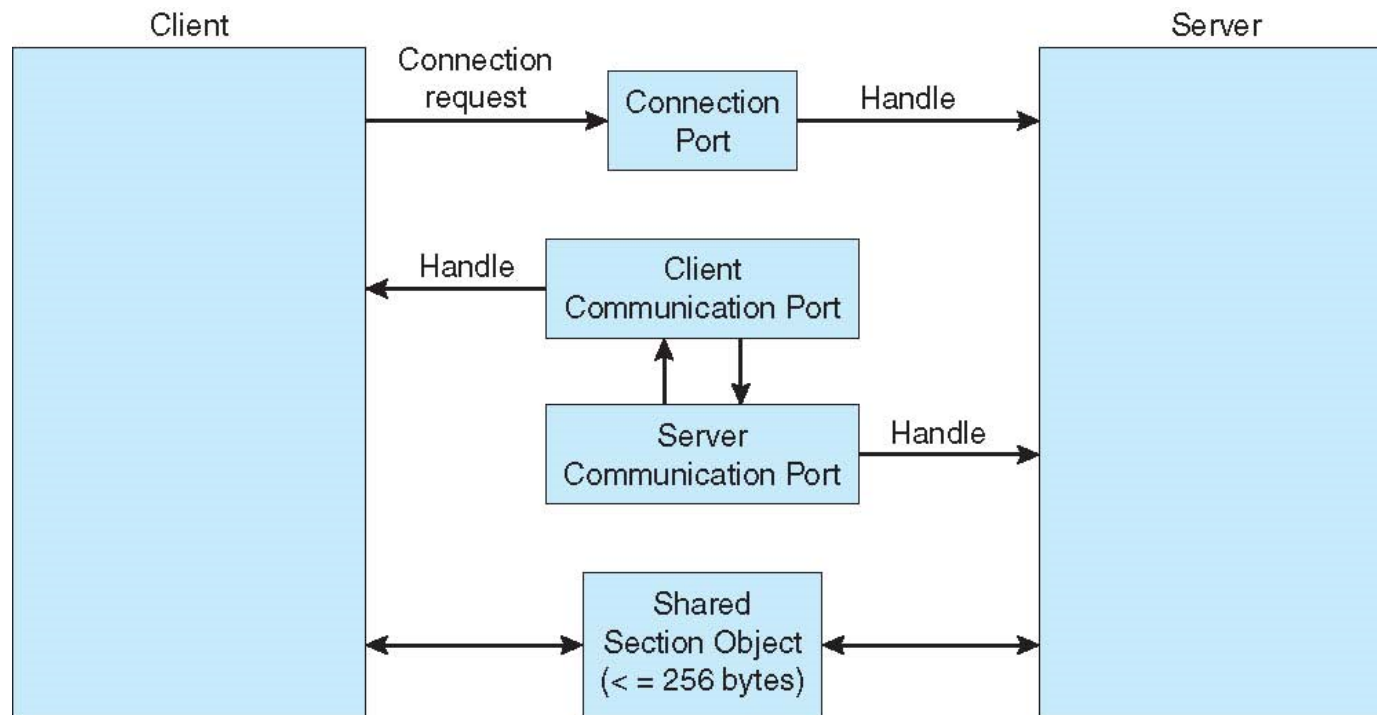
- A microkernel developed at CMU as a replacement for the kernel in the BSD version of UNIX
- Provides the basis of the OS kernel in Mac OS X (*not* microkernel)

Examples of IPC Systems – Windows



- Message-passing centric via **advanced local procedure call (LPC)** facility
 - Only works between processes on the same system
 - Uses ports (like mailboxes) to establish and maintain communication channels
 - Communication works as follows:
 - The client opens a handle to the subsystem's **connection port** object.
 - The client sends a connection request.
 - The server creates two private **communication ports** and returns the handle to one of them to the client.
 - The client and server use the corresponding port handle to send messages or callbacks and to listen for replies.

Local Procedure Calls in Windows XP



Summary



- 프로세스에 대하여
 - 프로세스란 프로그램이 실행되는 동안의 모습을 말한다
 - 프로세스는 컴퓨터가 계산하는데 있어 중심역할을 수행하며, 온갖 상태 정보를 다 PCB에 유지한다
- 프로세스의 상태
 - CPU에서 실행될 때와 CPU 밖에서 실행을 기다리고 있을 때로 나뉜다
 - running, ready, wait states
- 실행의 궤적(쓰레드)?
 - 프로그램에서 명령이 실행되는 일련의 흐름 혹은 궤적을 쓰레드라 부름
 - 혹은 이렇게 실행을 주도하는 주체를 쓰레드라고도 함
- 스케줄러
 - 프로세스가 CPU 상에서 수행될 순서를 정하는 주체는 스케줄러이다.
 - 스케줄러는 각종 큐를 유지한다
- On the process
 - "Process" refers to the appearance of a program while in execution.
 - A process plays the central role in computing activity of a computer. It maintains state information in PCB.
- Process states
 - A process is in two different states: while executing on CPU and while waiting off CPU
 - running, ready, wait states
- Threads of execution?
 - Thread: a flow or locus of instructions executed in a program
 - Or, the entity that generates such a flow
- Scheduler
 - Ordering the processes to execute on CPU is performed by a scheduler.
 - Schedulers maintains various queues.



Summary (Cont.)

- 단기 스케줄러 혹은 디스패처는 프로세스를 CPU에서 실제 실행되게 함
- 문맥교환(context switching)은 멀티프로그래밍 시스템에서 생기는 오버헤드
- 프로세스에 대한 연산
 - 프로세스의 생성(fork), 종료(exit), 실행(혹은 적재)(exec), 대기(wait), 등
- 프로세스간 통신(IPC)
 - 왜 필요한가? – 데이터/정보 공유 및 교환, 서비스의 요청/제공, 명령의 전달, 이벤트의 통지, 등
 - 공유기억(shared memory) 이용
 - 메시지 이용: 직접(link), 간접(mailbox)
 - 동기화: blocking(동기적), nonblocking(비동기적)
 - 버퍼링
- IPC 예
 - POSIX, Mach, Windows
- A short-term scheduler (dispatcher) lets a process run on CPU.
- Context switching is the overhead incurred in multiprogramming systems.
- Operation on processes
 - process creation (fork), termination (exit), execution or loading (exec), waiting (wait), etc.
- Interprocess communication (IPC)
 - why necessary? – sharing and exchange of data/information, service request/provision, delivery of commands, event notification, etc.
 - shared-memory approach
 - message-passing approach: direct(link), indirect(mailbox)
 - synchronization: blocking(synchronous), nonblocking(asynch)
 - buffering
- IPC examples
 - POSIX, Mach, Windows