

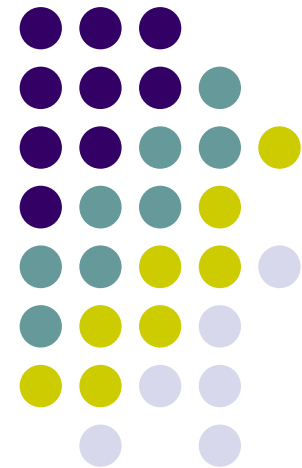
Chapter 13: I/O Systems

WHAT'S AHEAD:

- Overview
- I/O Hardware
- Application I/O Interface
 - Kernel I/O Subsystem
- Transforming I/O Requests to Hardware Operations
 - Performance

WE AIM:

- Explore the structure of an OS's I/O subsystem
- Discuss the principles of I/O hardware and its complexity
- Provide details of the performance aspects of I/O hardware and software



Note: These lecture materials are based on the lecture notes prepared by the authors of the book titled *Operating System Concepts*, 9e (Wiley)

핵심·요점

입출력 기법과 입출력 시스템




전형적인 입출력 기법 유형

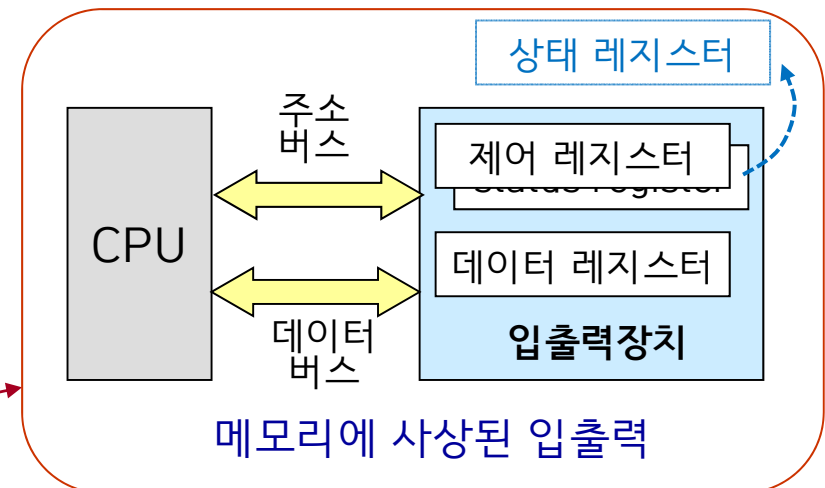

- ✓ 전적으로 프로그램에 의존 → **프로그램에 의한 입출력** (1)
- ✓ 인터럽트 하드웨어가 입출력 촉발 → **인터럽트에 의한 입출력** (2)
- ✓ CPU 도움 없이 → **직접 메모리 접근 (DMA)** (3)

입출력을 위한 주요 기능

- ✓ 목표 장치의 **초기화**
- ✓ 데이터 입출력 수행: **읽기/쓰기, 수신/송신**
- ✓ **상태** 정보 읽기, **제어** 정보 쓰기
- ✓ 입출력 **오류**의 검출 및 처리



입출력 장치 vs. 주변 장치(기기)




Core Ideas I/O Techniques & Systems



Typical types of I/O techniques

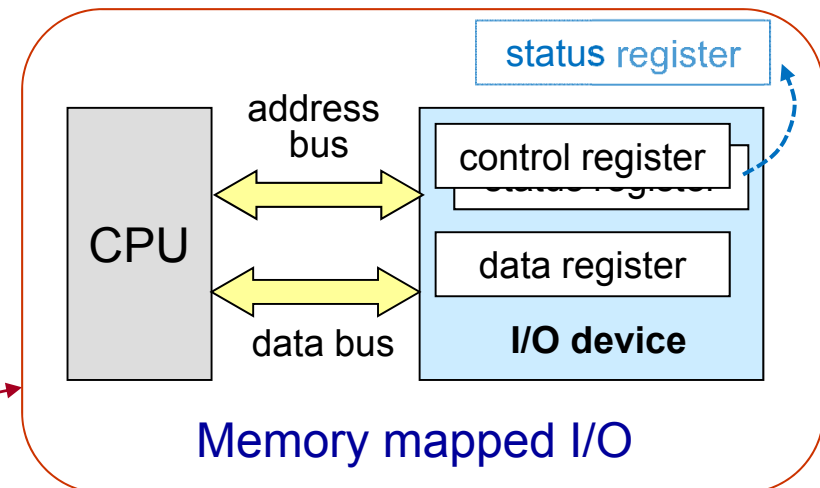

- ✓ The program does it all → **Programmed I/O**
- ✓ Interrupt hardware triggers I/O → **Interrupt I/O**
- ✓ CPU is left out → **Direct memory access (DMA)**



Main functions for I/O

- ✓ **Initialize** the target device
- ✓ Perform data I/O such as **read/write** or **send/receive**
- ✓ Read **status** information and write **control** data
- ✓ Detect and handle I/O **errors** properly

I/O device vs. peripheral device



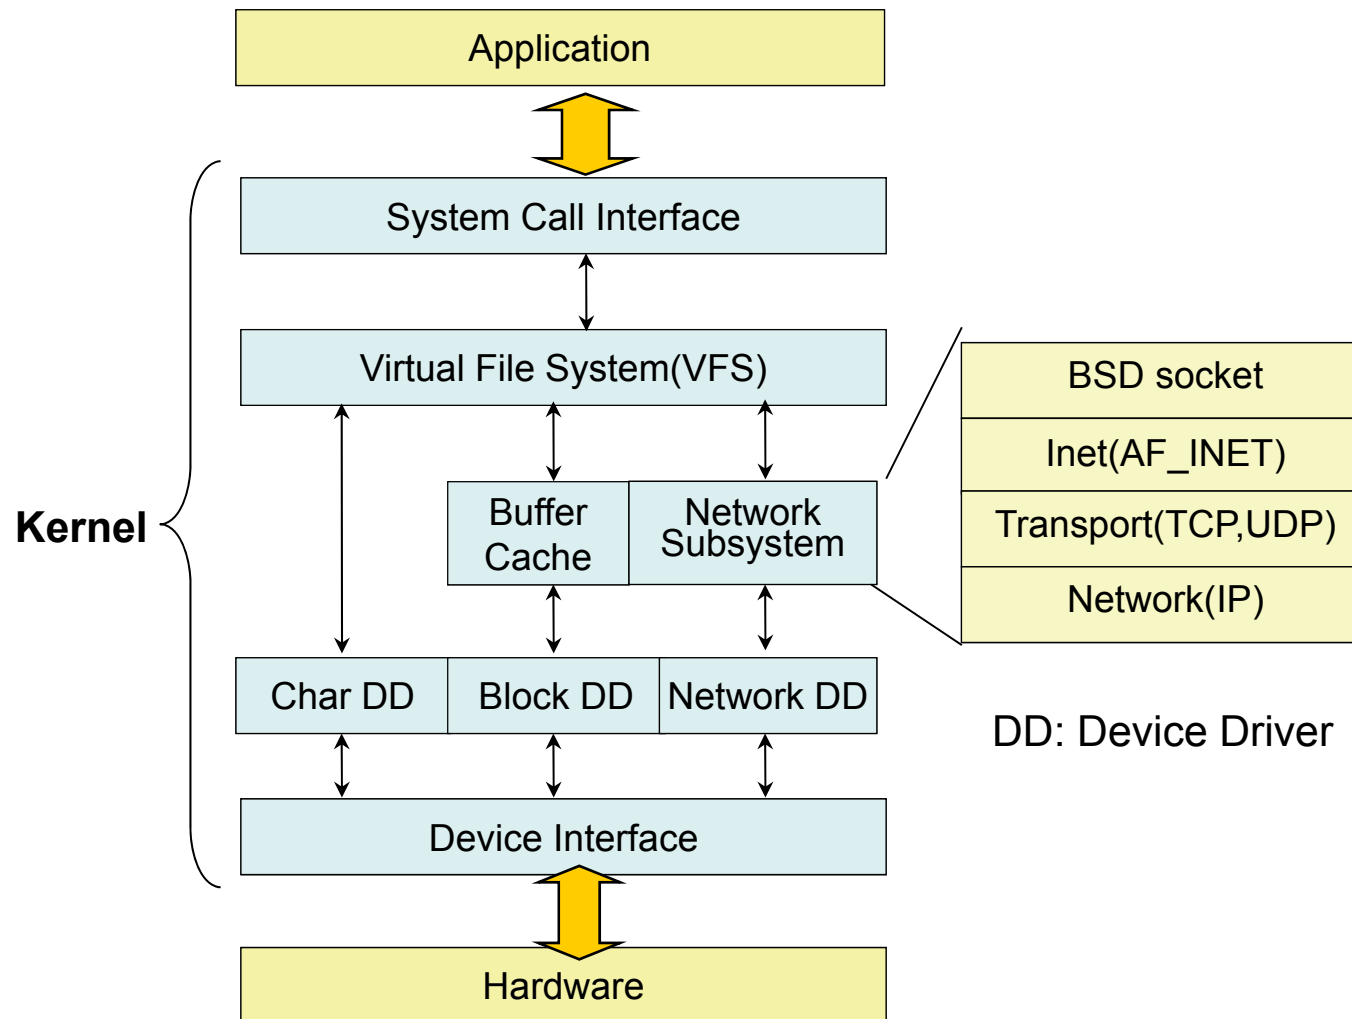
Overview



- I/O management is a major component of operating system design and operation
 - Important aspect of computer operation
 - I/O devices vary greatly
 - Various methods to control them
 - Performance management
 - New types of devices frequent
- Ports, busses, device controllers connect to various devices
- **Device drivers** encapsulate device details
 - Present uniform device-access interface to I/O subsystem

Caveat
I/O device
≠ peripheral device

Linux I/O Structure

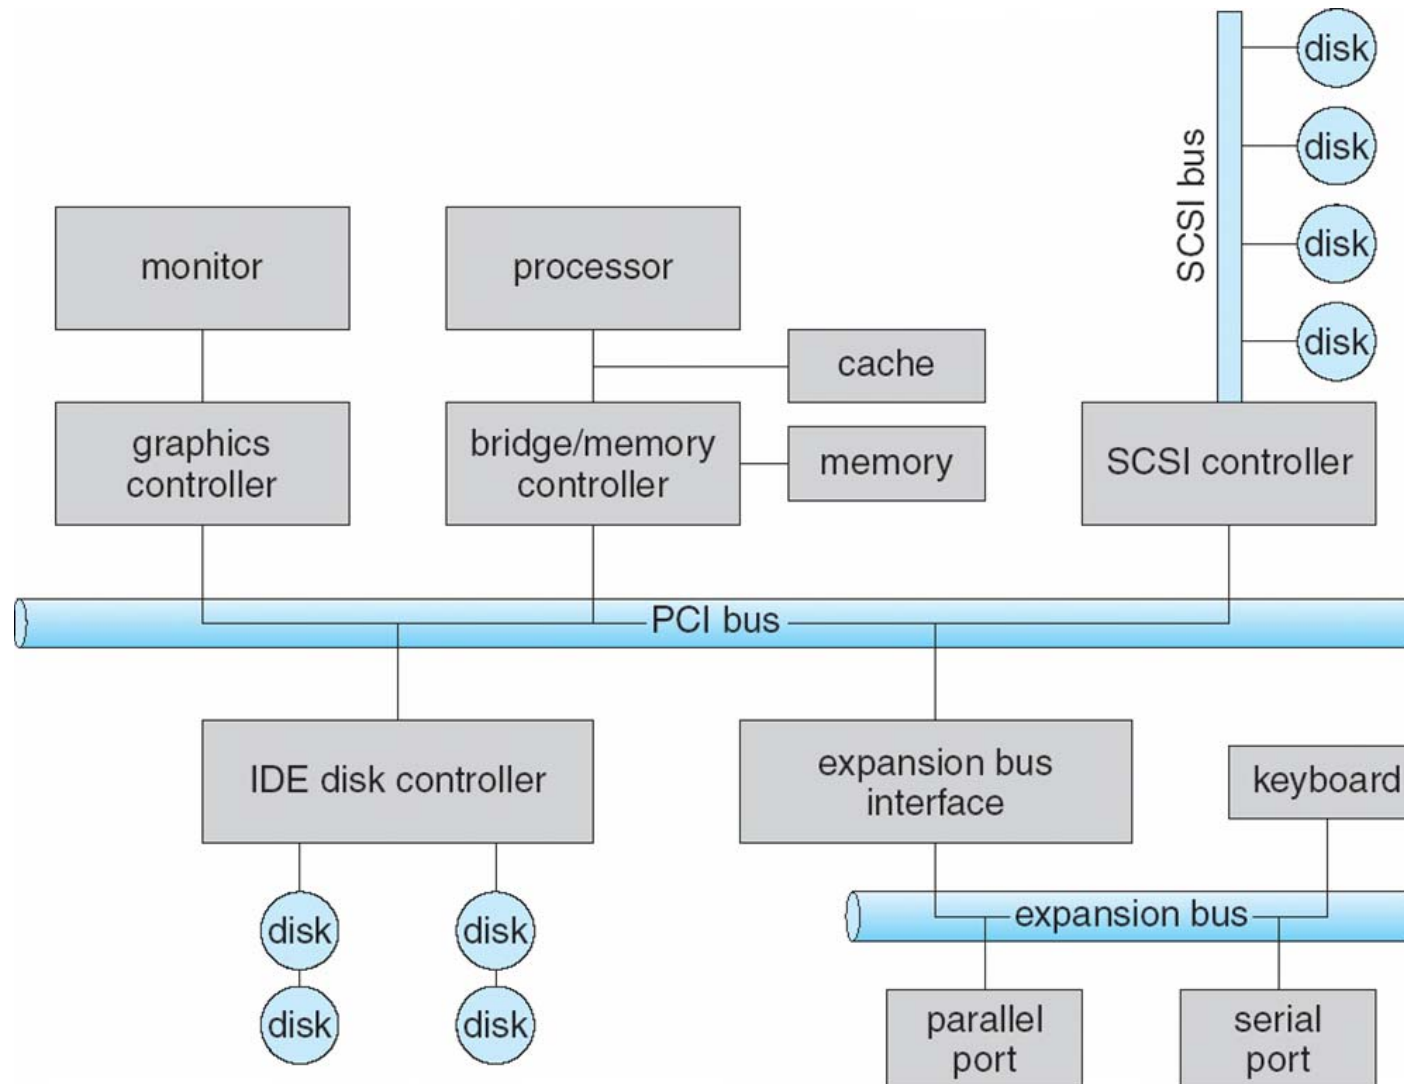


I/O Hardware



- Incredible variety of I/O devices
 - Storage
 - Communication, networking
 - Human-interface
 - Sensors
- Common concepts - signals from I/O devices interface with computer
 - **Port** - connection point for device
 - **Bus** - shared direct access
 - **Controller (host adapter)** - electronics that operate port, bus, device
 - Sometimes integrated
 - Sometimes separate circuit board (host adapter)
 - Contains processor, microcode, private memory, bus controller, etc
 - Some talk to per-device controller with bus controller, microcode, memory, etc.

A Typical PC Bus Structure





Polling for Programmed I/O

- Programmed I/O
 - Direct connection between CPU and I/O devices
 - All I/O operations performed on CPU by a program
- Request for I/O
 - Read or input: check to see if data is available in I/O device by polling
 - Write or output: check to see if output device is ready by polling
- Example: Polling loop

```
repeat
    for 1 millisecond    /* input from keyboard */
        check i/o control register to see if a key has been pressed
        then
            input a character into a CPU data register
            process the character
        end for
    if we need to output to screen then
        check i/o control register to see if screen is ready
        then output a character
    if we are printing then
        check i/o control register to see if the printer is ready
        then output a character
until forever
```




I/O Hardware (Cont.)

- I/O instructions control devices
- Devices usually have registers where device driver places commands, addresses, and data to write, or read data from registers after command execution
 - Data-in register, data-out register, status register, control register
 - Typically 1-4 bytes, or FIFO buffer
- Devices have addresses, used by
 - Isolated I/O: Use direct I/O instructions
 - Separate I/O address space
 - **Memory-mapped I/O**
 - No separate address space for I/O
 - Device data and command registers mapped to processor address space
 - Especially for large address spaces (graphics)



Device I/O Port and Memory Locations

Device I/O Port and Memory Locations on Linux (partial)

Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000-00F	DMA controller
020-021	interrupt controller
040-043	timer
200-20F	game controller
2F8-2FF	serial port (secondary)
320-32F	hard-disk controller
378-37F	parallel port
3D0-3DF	graphics controller
3F0-3F7	diskette-drive controller
3F8-3FF	serial port (primary)

```
[root]# cat /proc/ioports
```

```
0000-001f : dma1
0020-003f : pic1
0040-005f : timer
0060-006f : keyboard
0070-007f : rtc
0080-008f : dma page reg
00a0-00bf : pic2
00c0-00df : dma2
00f0-00ff : fpu
0170-0177 : ide1
01f0-01f7 : ide0
```

```
[root]# cat /proc/iomem
```

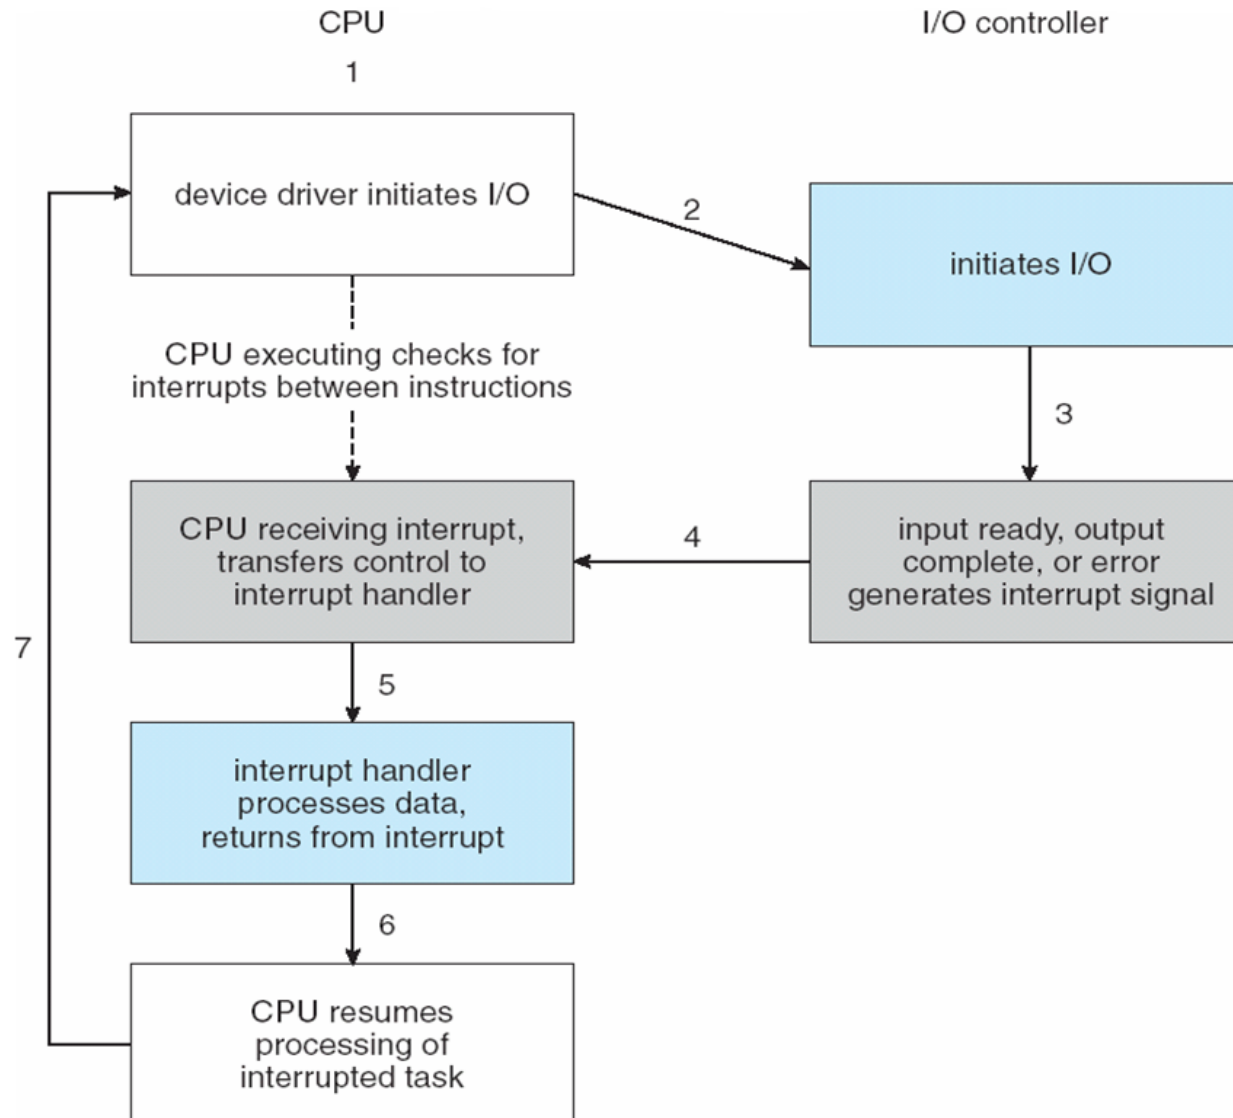
```
00000000-0009b7ff : System RAM
0009b800-0009ffff : reserved
000a0000-000bffff : Video RAM area
...
000f0000-000fffff : System ROM
00100000-3fedffff : System RAM
00100000-002d1daf : Kernel code
002d1db0-0038e2ff : Kernel data
...
d0100000-d02fffff : PCI Bus #02
```



Interrupts

- Polling can happen in 3 instruction cycles
 - Read status, logical-and to extract status bit, branch if not zero
 - How to be more efficient if non-zero infrequently? → Use hardware mechanism (i.e., interrupt)
- CPU **Interrupt-request line** triggered by I/O device
 - Checked by processor after each instruction
- **Interrupt handler** receives interrupts
 - **Maskable** to ignore or delay some interrupts
- Interrupt vector to dispatch interrupt to correct handler
 - Context switch at start and end
 - Based on priority
 - Some **nonmaskable**
 - Interrupt chaining if more than one device at same interrupt number

Interrupt-Driven I/O Cycle



Intel Pentium Processor Event-Vector Table



vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts



Interrupts (Cont.)

- Interrupt mechanism also used for exceptions
 - Terminate process, crash system due to hardware error
- Page fault executes when memory access error
- System call executes via trap to trigger kernel to execute request
- Multi-CPU systems can process interrupts concurrently
 - If operating system designed to handle it
- Used for time-sensitive processing, frequent, must be fast

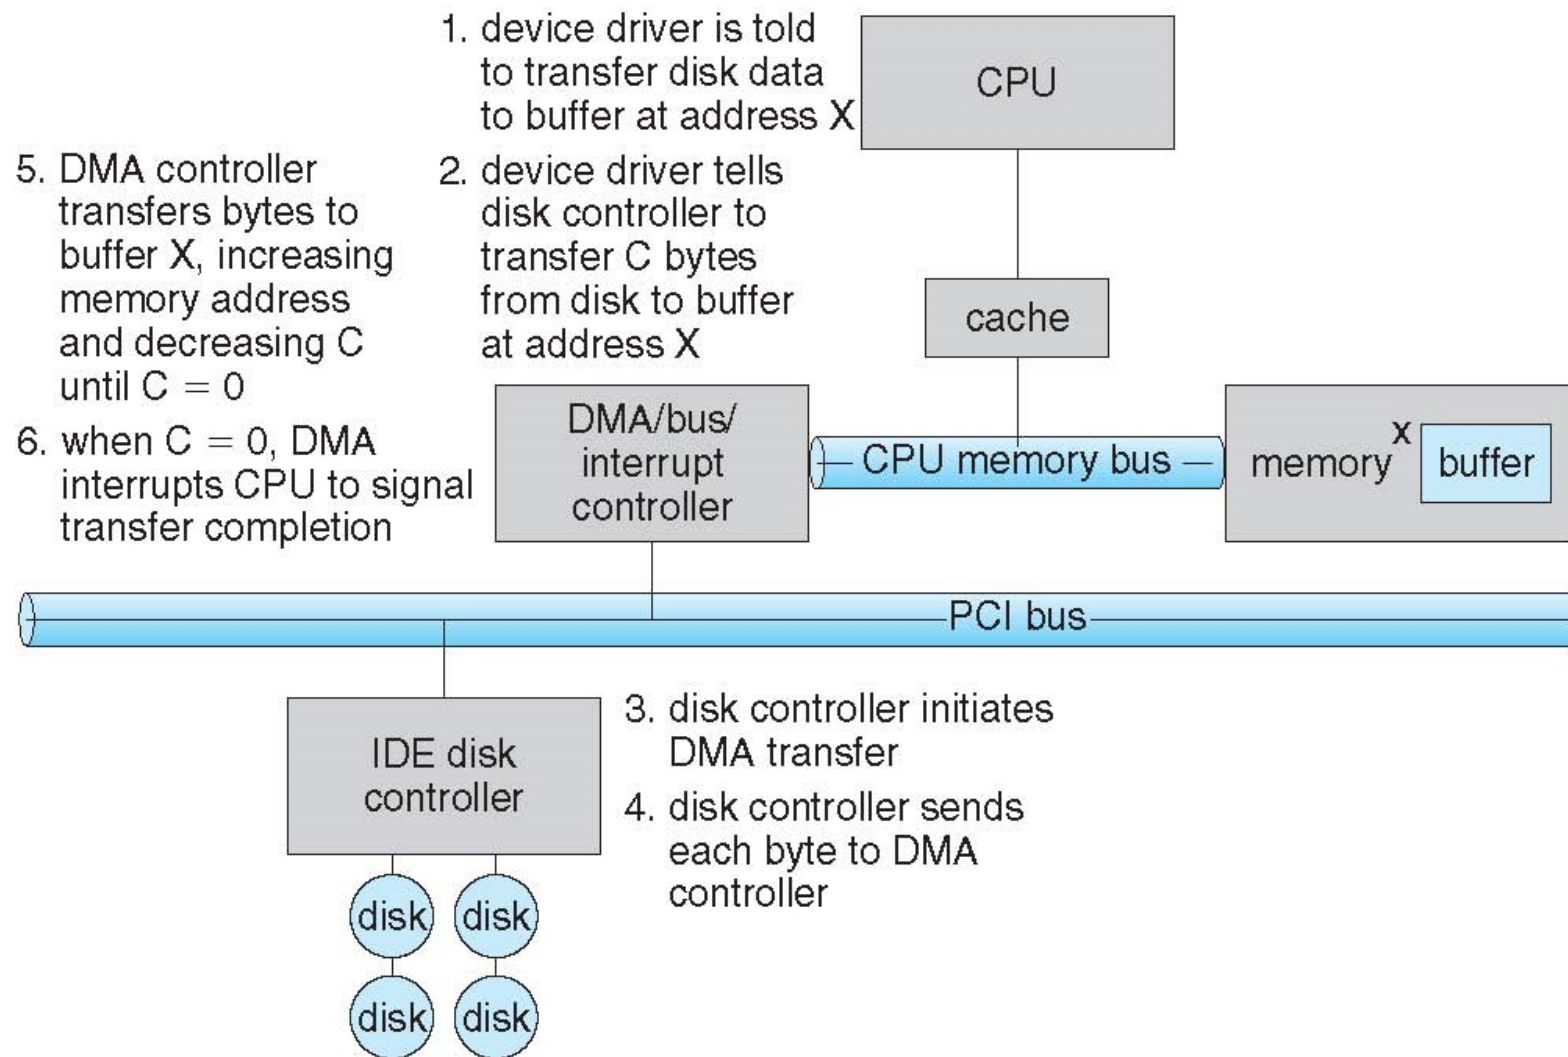


Direct Memory Access

- Used to avoid **programmed I/O** (one byte at a time) for large data movement
- Requires **DMA** controller
- Bypasses CPU to transfer data directly between I/O device and memory
- OS writes DMA command block into memory
 - Source and destination addresses
 - Read or write mode
 - Count of bytes
 - Writes location of command block to DMA controller
 - Bus mastering of DMA controller - grabs bus from CPU
 - When done, interrupts to signal completion



Steps in a DMA Transfer

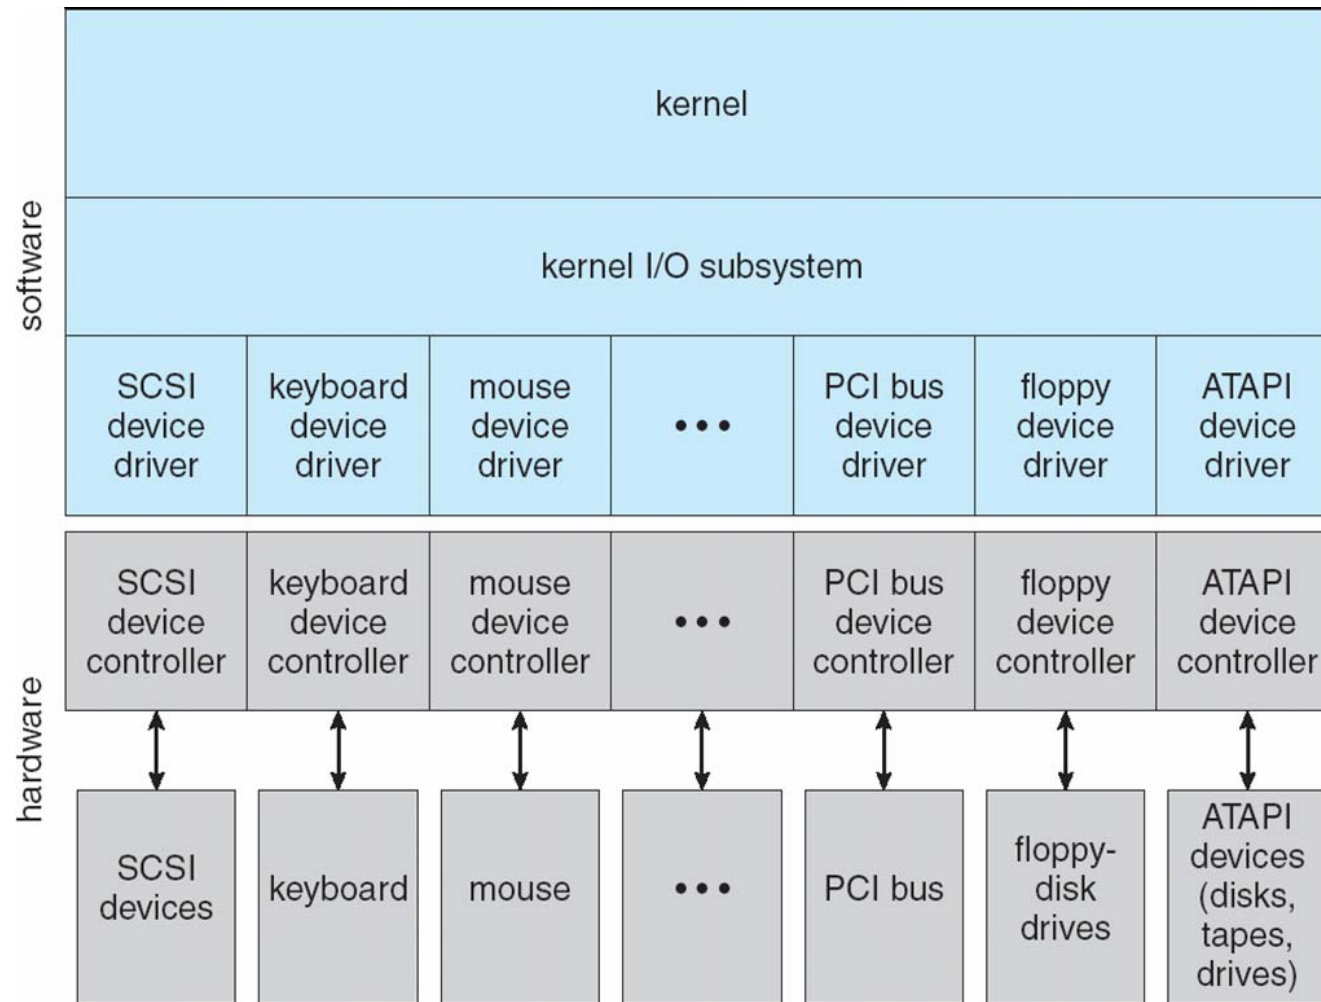


Application I/O Interface



- I/O system calls encapsulate device behaviors in generic classes
- Device-driver layer hides differences among I/O controllers from kernel
- New devices talking already-implemented protocols need no extra work
- Each OS has its own I/O subsystem structures and device driver frameworks
- Devices vary in many dimensions
 - **Character-stream** or **block**
 - **Sequential** or **random-access**
 - **Synchronous** or **asynchronous** (or both)
 - **Sharable** or **dedicated**
 - **Speed of operation**
 - **read-write**, **read only**, or **write only**

A Kernel I/O Structure





Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read–write	CD-ROM graphics controller disk

Characteristics of I/O Devices (Cont.)



- Subtleties of devices handled by device drivers
- Broadly I/O devices can be grouped by the OS into
 - Block I/O
 - Character I/O (Stream)
 - Memory-mapped file access
 - Network sockets
- For direct manipulation of I/O device specific characteristics, usually an escape / back door
 - Unix `ioctl()` call to send arbitrary bits to a device control register and data to device data register



Block and Character Devices

- Block devices include disk drives
 - Commands include read, write, seek
 - **Raw I/O, direct I/O**, or file-system access
 - Memory-mapped file access possible
 - File mapped to virtual memory and clusters brought via demand paging
 - DMA

- Character devices include keyboards, mice, serial ports
 - Commands include `get()`, `put()`
 - Libraries layered on top allow line editing



Network Devices

- Varying enough from block and character to have own interface
- Unix and Windows NT/9x/2000 include **socket** interface
 - Separates network protocol from network operation
 - Includes `select()` functionality
- Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)



Clocks and Timers

- Provide current time, elapsed time, timer
- Normal resolution about 1/60 second
- Some systems provide higher-resolution timers
- **Programmable interval timer** used for timings, periodic interrupts
- `ioctl()` (on UNIX) covers odd aspects of I/O such as clocks and timers



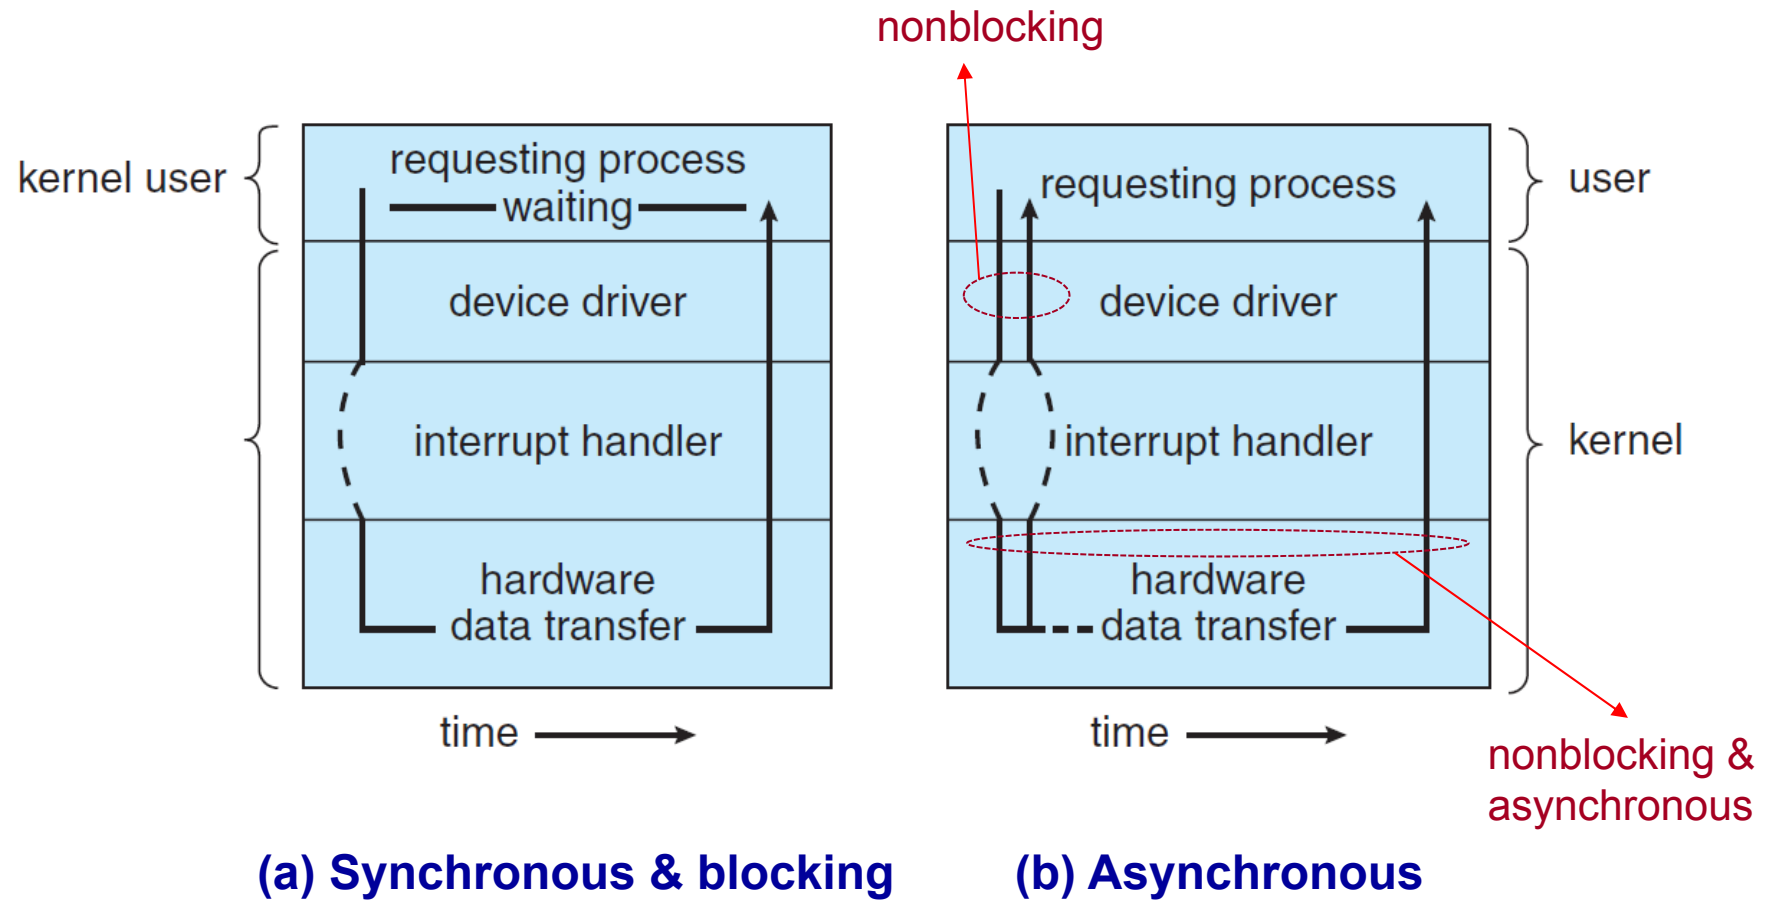
Blocking and Nonblocking I/O

- **Blocking** - process suspended until I/O completed
 - Easy to use and understand
 - Insufficient for some needs

- **Nonblocking** - I/O call returns with as much as available
 - User interface, data copy (buffered I/O)
 - Implemented via multi-threading
 - Returns quickly with count of bytes read or written
 - `select()` to find if data ready then `read()` or `write()` to transfer

- **Asynchronous** - process runs while I/O executes
 - Difficult to use
 - I/O subsystem signals process when I/O completed

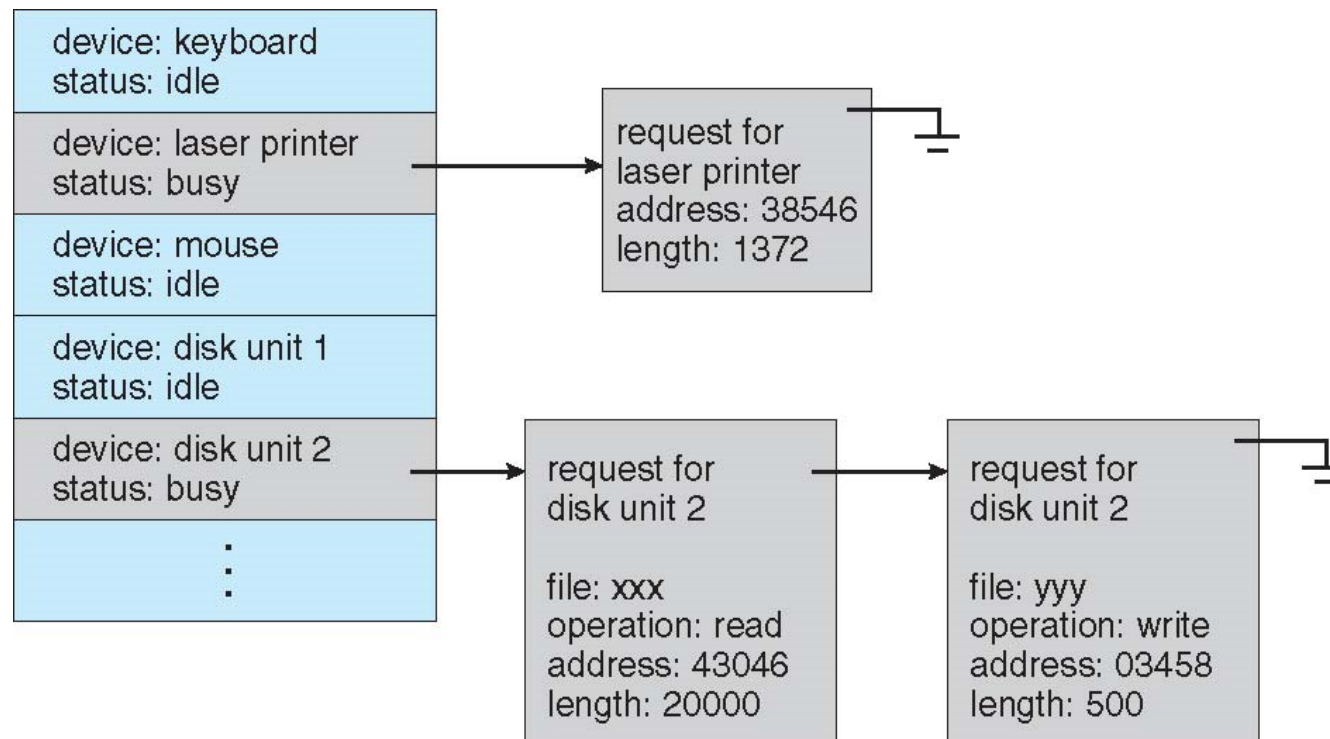
Two I/O Methods



Kernel I/O Subsystem



- Scheduling
 - Some I/O request ordering via per-device queue
 - Some OSs try fairness
 - Some implement Quality of Service (i.e. IP QoS)
- Device status table





Kernel I/O Subsystem (Cont.)

- Buffering
 - Store data in memory while transferring between devices
 - To cope with device speed mismatch (say, $10^{-5} \sim 10^6$ bytes/s)
 - Mitigate the difference in transfer rates in various devices
 - Double buffering - two copies of the data
 - Kernel and user
 - Decouple the producer of data from the consumer
 - To cope with device transfer size mismatch
 - To maintain "copy semantics"
 - E.g., What if the application modifies the contents of a buffer while the kernel is transferring them to disk? → Copy semantics is not guaranteed → This problem can be resolved by maintaining a separate buffer in the kernel, so the buffer contents are copied into the kernel buffer in advance



Kernel I/O Subsystem (Cont.)

- **Caching** - faster device holding copy of data
 - Always just a copy
 - Key to performance
 - Sometimes combined with buffering

- **Spooling** - hold output for a device
 - If device can serve only one request at a time
 - i.e., Printing

- **Device reservation** - provides exclusive access to a device
 - System calls for allocation and de-allocation
 - Watch out for deadlock



Error Handling

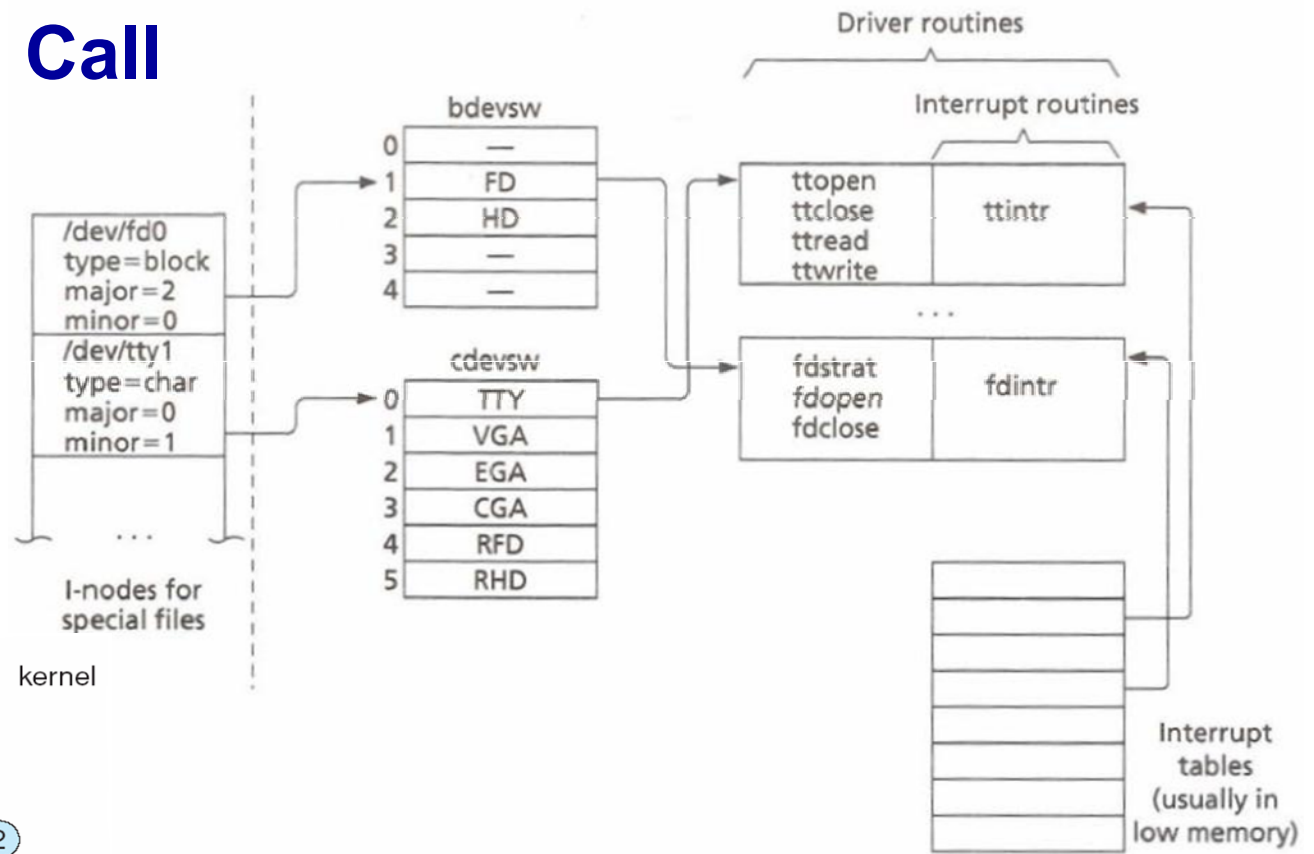
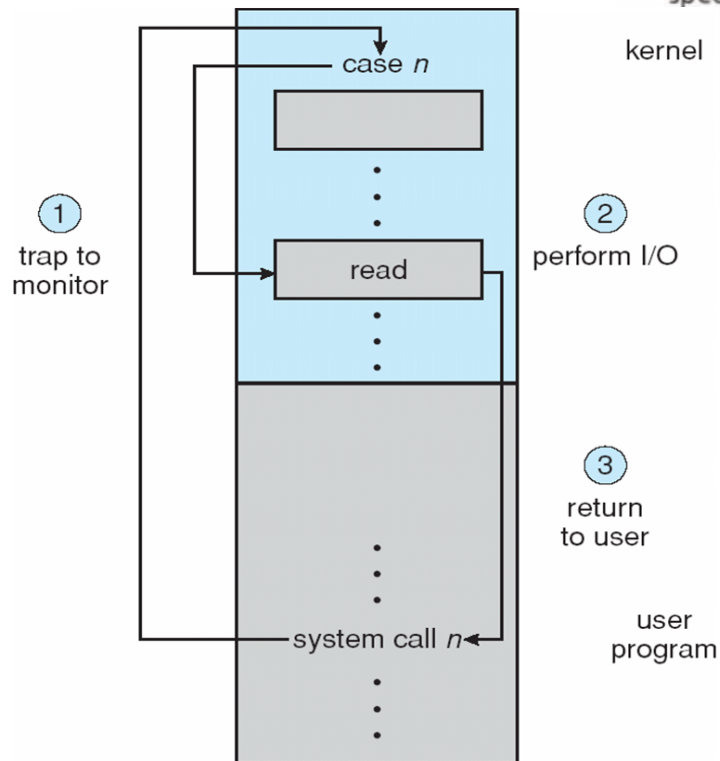
- OS can recover from disk read, device unavailable, transient write failures
 - Retry a read or write, for example
 - Some systems more advanced - Solaris FMA, AIX
 - Track error frequencies, stop using device with increasing frequency of retry-able errors
- Most return an error number or code when I/O request fails
- System error logs hold problem reports



I/O Protection

- User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions
 - All I/O instructions defined to be privileged
 - I/O must be performed via system calls
 - Memory-mapped and I/O port memory locations must be protected too

Use of a System Call and Kernel Data Structures to Perform I/O



↑ Kernel data structures for accessing peripheral devices
in Unix (K. Christian, S. Richter, *The UNIX Operating System*, 3e., Wiley, 1994)

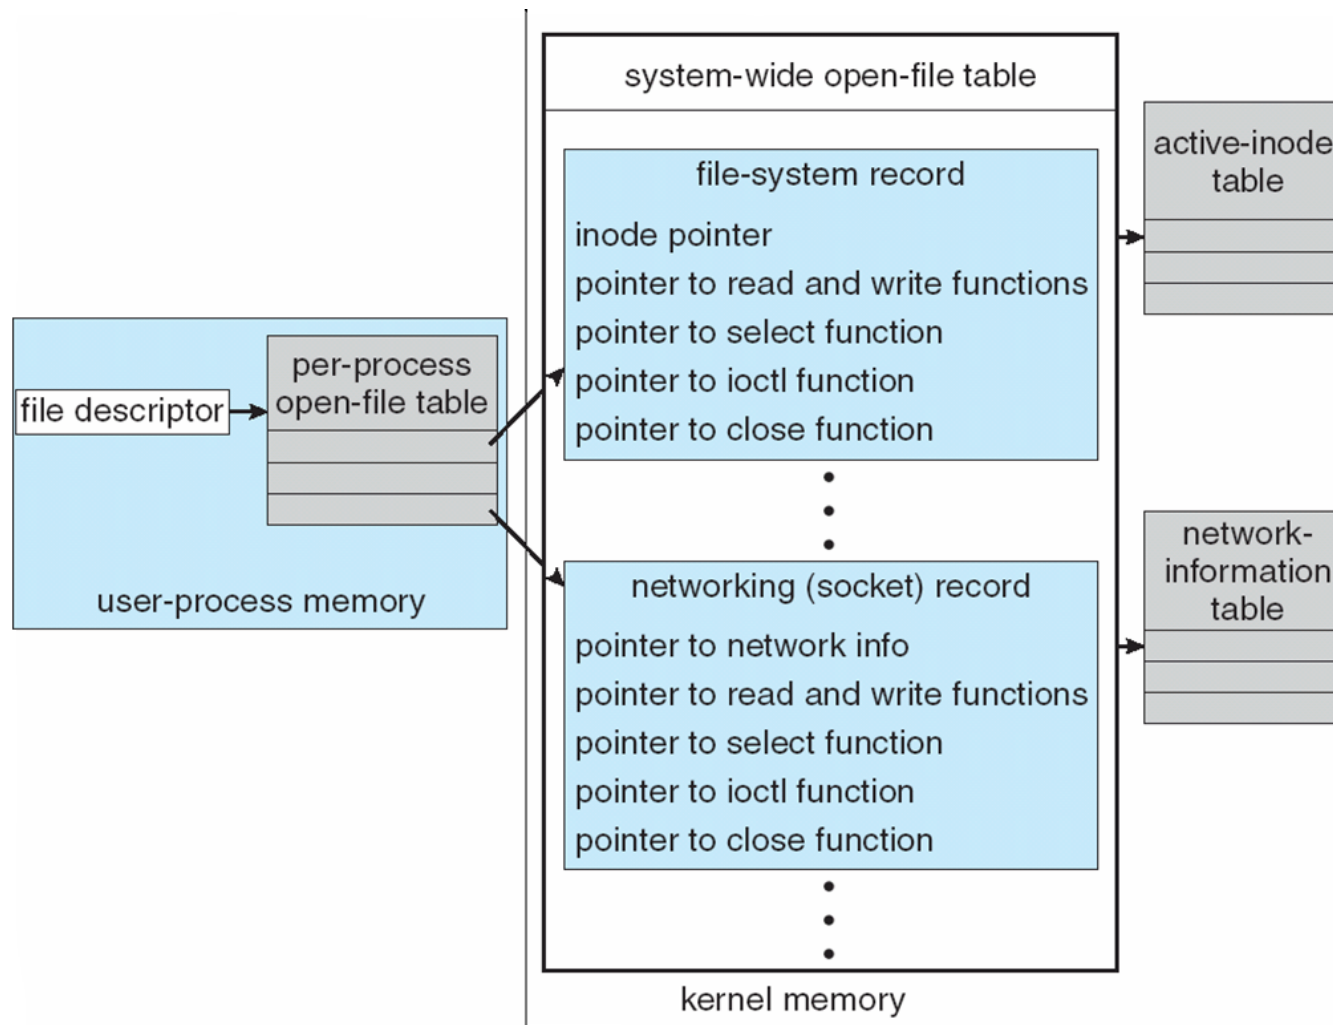
← Use of a system call to perform I/O



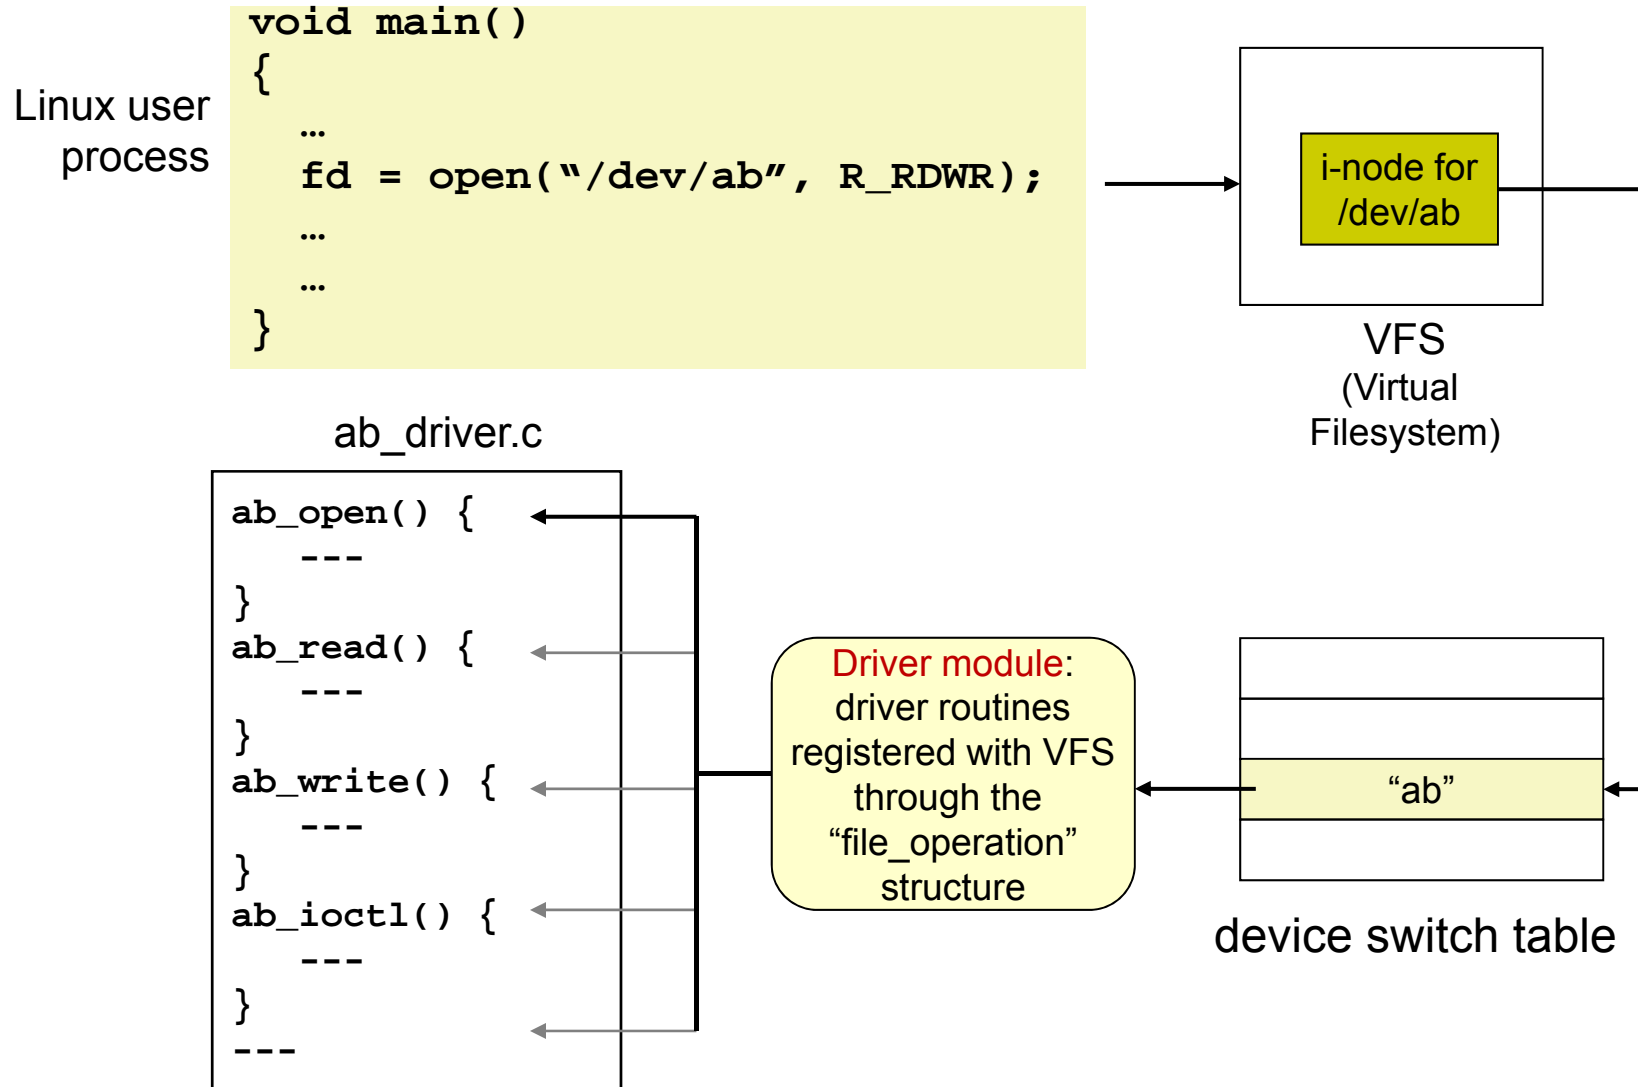
Kernel Data Structures

- Kernel keeps state info for I/O components, including open file tables, network connections, character device state
- Many, many complex data structures to track buffers, memory allocation, “dirty” blocks
- Some use object-oriented methods and message passing to implement I/O
 - Windows uses message passing
 - Message with I/O information passed from user mode into kernel
 - Message modified as it flows through to device driver and back to process
 - For input, the message contains the buffer to receive the data, and for output, the data to be written
 - Pros: Simplify the structure and the design of I/O system. Add flexibility
 - Cons: Overhead caused by shared data structures

UNIX I/O Kernel Structure



Conceptual Description of Linux I/O

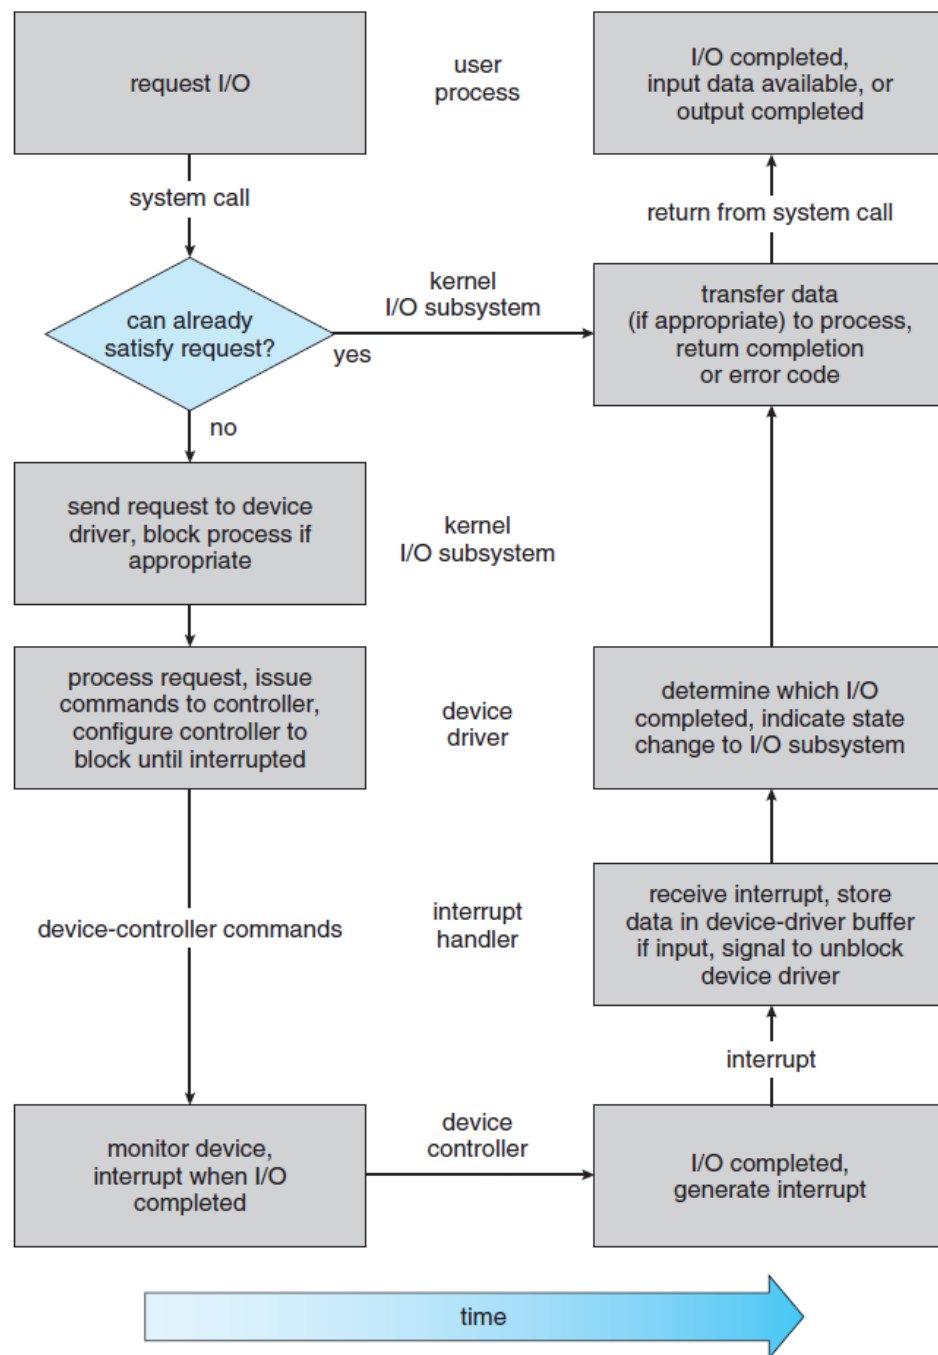


Transforming I/O Requests to Hardware Operations



- Consider reading a file from disk for a process:
 - Determine device holding file
 - Translate name to device representation
 - Physically read data from disk into buffer
 - Make data available to requesting process
 - Return control to process

Life Cycle of An I/O Request



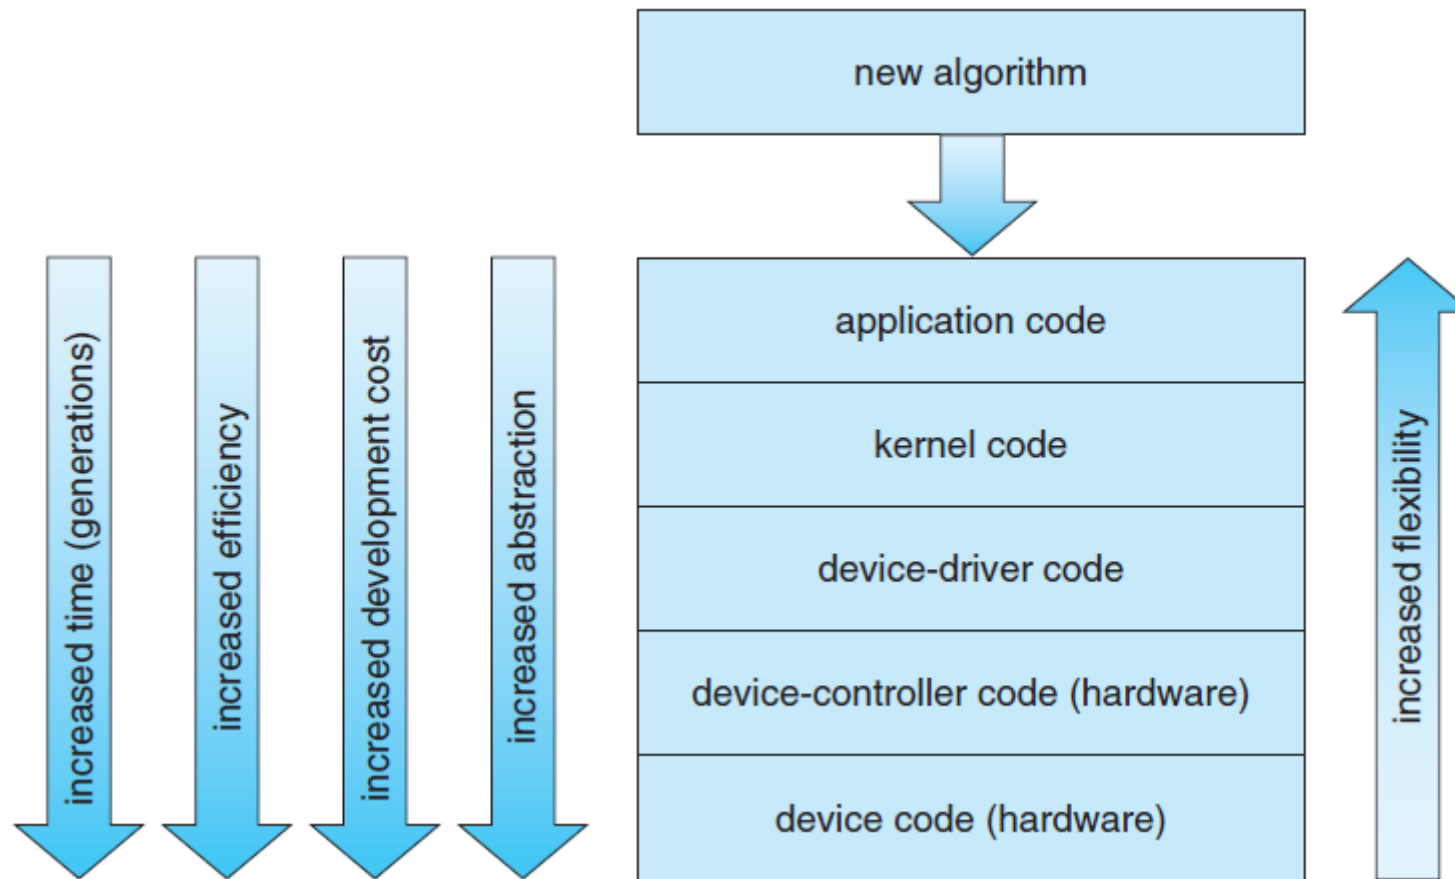
Performance



- I/O - a major factor in system performance:
 - Demands CPU to execute device driver, kernel I/O code
 - Context switches due to interrupts
 - Data copying
 - Network traffic especially stressful

- Improving performance
 - Reduce number of context switches
 - Reduce data copying
 - Reduce interrupts by using large transfers, smart controllers, polling
 - Use DMA
 - Use smarter hardware devices
 - Balance CPU, memory, bus, and I/O performance for highest throughput
 - Move user-mode processes / daemons to kernel threads

Device-Functionality Progression



Summary



- 입출력시스템 개관
 - 입출력장치란?
 - 디바이스 드라이버 역할
- 입출력 하드웨어
 - 입출력시스템의 구성
 - 입출력장치의 구성, 역할
 - 입출력기법: 폴링, 인터럽트, DMA
- 응용 레벨의 입출력 인터페이스
 - 입출력장치의 유형 및 특징
 - block vs. character vs. network devices
 - clocks, timers
 - blocking/nonblocking, synchronous/asynchronous
- Overview of I/O systems
 - what is an I/O device?
 - role of device drivers
- I/O hardware
 - organization of I/O system
 - configuration and role of I/O device
 - I/O techniques: polling, interrupt, DMA
- Application I/O interface
 - types, characteristics of I/O devices
 - block vs. character vs. network devices
 - clocks, timers
 - blocking/nonblocking, synchronous/asynchronous



Summary (Cont.)

- 커널 입출력시스템
 - 스케줄링, 버퍼링, 캐싱, 스푼링
 - 에러 처리, 입출력 보호
 - 커널 데이터 구조
- 입출력 요청부터 하드웨어 동작까지
 - 입출력의 계층적인 구조상에서 명령과 데이터의 흐름
- 성능 개선
 - 문맥교체 수 줄임
 - 데이터 복사 줄임
 - 인터럽트 수 줄임
 - DMA 사용
 - 커널 쓰레드 사용
 - CPU 부터 입출력장치 사이의 데이터 흐름의 최적화
- Kernel I/O subsystem
 - scheduling, buffering, caching, spooling
 - error processing, protection of I/O
- Transforming I/O requests to hardware operations
 - flows of data and commands in the I/O hierarchical structure
- Performance enhancement
 - reduce the number of context switches
 - reduce data copying
 - reduce the number of interrupts
 - use DMA
 - use kernel threads
 - optimize the data flow between CPU through I/O devices