

Running Programs on a System

Dynamic Memory Allocation II



Recap: Dynamic Memory Allocation Basics

■ Dynamic Memory Allocation

- by user-space allocator
- for data structures whose size is unknown at compile time


```
p1 = malloc(4)
```

```
p2 = malloc(5)
```

```
p3 = malloc(6)
```

```
free(p2)
```

```
p4 = malloc(2)
```




■ Fragmentation

- internal
- external


■ Allocation Methods:

- Implicit free lists
 - ▶ allocation: $O(n)$
 - ▶ free: $O(1)$, even with coalescing




Recap: Allocation and brk

- 1000 random allocations, followed by in-order 1000 deallocations



Recap: Allocation and brk

- 1000 random allocations, followed by reverse order 1000 deallocations




Dynamic Memory Allocation - Advanced Concepts

- **Dynamic memory allocation**
 - Implicit free lists
 - **Explicit free lists**
 - Segregated free lists
- Garbage collection
- Memory-related perils and pitfalls


Acknowledgement: slides based on the cs:app2e material

Keeping Track of Free Blocks

- Method 1: Implicit free list using length—links all blocks




- Method 2: Explicit free list among the free blocks using pointers




- Method 3: Segregated free list
 - Different free lists for different size classes
- Method 4: Blocks sorted by size
 - Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

Explicit Free Lists

Allocated (as before)



Free




- Maintain list(s) of **free** blocks, not **all** blocks
 - The “next” free block could be anywhere
 - ▶ So we need to store forward/back pointers, not just sizes
 - Still need boundary tags for coalescing
 - Luckily we track only free blocks, so we can use payload area

Explicit Free Lists

- Logically:



- Physically: blocks can be in any order



Allocating From Explicit Free Lists


conceptual graphic

Before



After


(with splitting)




Freeing With Explicit Free Lists

- **Insertion policy:** Where in the free list do you put a newly freed block?
 - LIFO (last-in-first-out) policy
 - ▶ Insert freed block at the beginning of the free list
 - ▶ Pro: simple and constant time
 - ▶ Con: studies suggest fragmentation is worse than address ordered
 - Address-ordered policy
 - ▶ Insert freed blocks so that free list blocks are always in address order:
 $\text{addr}(\text{prev}) < \text{addr}(\text{curr}) < \text{addr}(\text{next})$
 - ▶ Con: requires search
 - ▶ Pro: studies suggest fragmentation is lower than LIFO


Freeing With a LIFO Policy (Case 1)




- Insert the freed block at the root of the list




Freeing With a LIFO Policy (Case 2)




- Splice out predecessor block, coalesce both memory blocks, and insert the new block at the root of the list




Freeing With a LIFO Policy (Case 3)




- Splice out successor block, coalesce both memory blocks and insert the new block at the root of the list



Freeing With a LIFO Policy (Case 4)



- Splice out predecessor and successor blocks, coalesce all 3 memory blocks and insert the new block at the root of the list




Explicit List Summary


- Comparison to implicit list:
 - Allocate is linear time in number of free blocks instead of all blocks
 - ▶ Much faster when most of the memory is full
 - Slightly more complicated allocate and free since needs to splice blocks in and out of the list
 - Some extra space for the links (2 extra words needed for each block)
 - ▶ Does this increase internal fragmentation?
- Most common use of linked lists is in conjunction with segregated free lists
 - Keep multiple linked lists of different size classes, or possibly for different types of objects

Keeping Track of Free Blocks

- Method 1: Implicit list using length—links all blocks



- Method 2: Explicit list among the free blocks using pointers



- Method 3: Segregated free list

- Different free lists for different size classes

- Method 4: Blocks sorted by size


- Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

Dynamic Memory Allocation - Advanced Concepts

- **Dynamic memory allocation**
 - Implicit free lists
 - Explicit free lists
 - **Segregated free lists**
- Garbage collection
- Memory-related perils and pitfalls

Segregated List (Seglist) Allocators

- Each size class of blocks has its own free list



- Often have separate classes for each small size
- For larger sizes: One class for each two-power size

Seglist Allocator

- Given an array of free lists, each one for some size class
- To allocate a block of size n :
 - Search appropriate free list for block of size $m > n$
 - If an appropriate block is found:
 - ▶ Split block and place fragment on appropriate list (optional)
 - If no block is found, try next larger class
 - Repeat until block is found
- If no block is found:
 - Request additional heap memory from OS (using `sbrk()`)
 - Allocate block of n bytes from this new memory
 - Place remainder as a single free block in largest size class.

Seglist Allocator (cont.)

- To free a block:
 - Coalesce and place on appropriate list (optional)
- Advantages of seglist allocators
 - Higher throughput
 - ▶ log time for power-of-two size classes
 - Better memory utilization
 - ▶ First-fit search of segregated free list approximates a best-fit search of entire heap.
 - ▶ Extreme case: Giving each block its own size class is equivalent to best-fit.

More Info on Allocators

- D. Knuth, “The Art of Computer Programming”, 2nd edition, Addison Wesley, 1973
 - The classic reference on dynamic storage allocation
- Wilson et al, “Dynamic Storage Allocation: A Survey and Critical Review”, Proc. 1995 Int’l Workshop on Memory Management, Kinross, Scotland, Sept, 1995.
 - Comprehensive survey
 - Available on eTL

Dynamic Memory Allocation - Advanced Concepts

- Dynamic memory allocation
 - Implicit free lists
 - Explicit free lists
 - Segregated free lists
- **Garbage collection**
- Memory-related perils and pitfalls

Implicit Memory Management: Garbage Collection

- **Garbage collection:** automatic reclamation of heap-allocated storage — application never has to free

```
void foo() {  
    int *p = malloc(128);  
    return; /* p block is now garbage */  
}
```

- Common in functional languages, scripting languages, and modern object oriented languages:
 - Lisp, ML, Java, Perl, Mathematica, Oberon
- Variants (“conservative” garbage collectors) exist for C and C++
 - However, cannot necessarily collect all garbage

Garbage Collection


- How does the memory manager know when memory can be freed?
 - In general we cannot know what is going to be used in the future since it depends on conditionals
 - But we can tell that certain blocks cannot be used if there are no pointers to them
- Must make certain assumptions about pointers
 - Memory manager can distinguish pointers from non-pointers
 - All pointers point to the start of a block
 - Cannot hide pointers
(e.g., by coercing them to an int, and then back again)

Classical GC Algorithms

- Mark-and-sweep collection (McCarthy, 1960)
 - Does not move blocks (unless you also “compact”)
- Reference counting (Collins, 1960)
 - Does not move blocks
- Copying collection (Minsky, 1963)
 - Moves blocks
- Generational Collectors (Lieberman and Hewitt, 1983)
 - Collection based on lifetimes
 - ▶ Most allocations become garbage very soon
 - ▶ So focus reclamation work on zones of memory recently allocated
- For more information:
Jones and Lin, “Garbage Collection: Algorithms for Automatic Dynamic Memory”,
John Wiley & Sons, 1996.

Memory as a Graph



- We view memory as a directed graph
 - Each block is a node in the graph
 - Each pointer is an edge in the graph
 - Locations not in the heap that contain pointers into the heap are called root nodes (e.g. registers, locations on the stack, global variables)




A block is *reachable* if there is a path from any root to that node. Non-reachable blocks are *garbage* (no (legal) access possible anymore by the application)

Mark and Sweep Collecting

- Can build on top of malloc/free package
 - Allocate using malloc until you “run out of space”
- When out of space:
 - Use extra mark bit in the head of each block
 - Mark: Start at roots and set mark bit on each reachable block
 - Sweep: Scan all blocks and free blocks that are not marked



Note: arrows here denote memory refs, not free list ptrs.



 Mark bit set

Assumptions For a Simple Implementation

■ Application

- `new(n)`: returns pointer to new block with all locations cleared
- `read(b, i)`: read location `i` of block `b` into register
- `write(b, i, v)`: write `v` into location `i` of block `b`

■ Each block will have a header word

- addressed as `b[-1]`, for a block `b`
- Used for different purposes in different collectors

■ Instructions used by the Garbage Collector

- `is_ptr(p)`: determines whether `p` is a pointer
- `length(b)`: returns the length of block `b`, not including the header
- `get_roots()`: returns all the roots

Mark and Sweep (cont.)

Mark using depth-first traversal of the memory graph

```
ptr mark(ptr p) {  
    if (!is_ptr(p)) return;           // do nothing if not pointer  
    if (markBitSet(p)) return;        // check if already marked  
    setMarkBit(p);                    // set the mark bit  
    for (i=0; i < length(p); i++)    // call mark on all words  
        mark(p[i]);                  // in the block  
    return;  
}
```

Sweep using lengths to find next block

```
ptr sweep(ptr p, ptr end) {  
    while (p < end) {  
        if markBitSet(p)  
            clearMarkBit();  
        else if (allocateBitSet(p))  
            free(p);  
        p = nextBlock(p);  
    }  
}
```


■ Instructions used by the Garbage Collector

- `is_ptr(p)`: determines whether `p` is a pointer
- `length(b)`: returns the length of block `b`, not including the header
- `get_roots()`: returns all the roots

Conservative Mark & Sweep in C


■ A “conservative garbage collector” for C programs

- `is_ptr()` determines if a word is a pointer by checking if it points to an allocated block of memory
- But, in C pointers can point to the middle of a block



■ So how to find the beginning of the block?

- Can use a balanced binary tree to keep track of all allocated blocks (key is start-of-block)
- Balanced-tree pointers can be stored in header (use two additional words)



Left: smaller addresses

Right: larger addresses


Mark & Sweep

■ Discussion

- no overhead on pointer manipulation
- no problems with circular data structures
- start/stop algorithm: non-interruptible, global memory traversal necessary during which no other process can run

Reference Counting

- Each memory block has an additional *reference count* field counting the number of references to this block:





- Invariant: reference count is always correct
 - allocation: simple
 - assignments: compiler

Reference Counting

■ Discussion

- garbage collection cost distributed, but relatively high
- immediate knowledge when a block is garbage
- cannot reclaim cyclic data structures



Copying Algorithm

- Heap divided into two semi-spaces
 - current data
 - obsolete data
- When running out of memory
 - roles of the spaces reversed
 - GC traverses reachable set and copies each reachable block from the (now) obsolete data space into the current data space
 - ▶ automatic compaction

Copying Algorithm

■ Discussion

- simplified & efficient heap management
 - ▶ no allocated/free block lists necessary, just an “end of allocated space pointer”.
 - ▶ no internal fragmentation
 - ▶ eliminates external fragmentation on each GC
- effectively usable memory space halved
 - ▶ with virtual memory, do we need to back the spaces with physical memory?
- GC needs to update pointer references

Exact Garbage Collection

- Treating every word in allocated blocks as a potential pointer is
 - slow
 - inexact
- Modern languages include information about data structures in the binary, so that the GC knows exactly
 - what registers contain pointers
 - for each data structure (on the stack, in the heap) where the pointers are

Dynamic Memory Allocation - Advanced Concepts

- Dynamic memory allocation
 - Implicit free lists
 - Explicit free lists
 - Segregated free lists
- Garbage collection
- **Memory-related perils and pitfalls**

Memory-Related Perils and Pitfalls

- Dereferencing bad pointers
- Reading uninitialized memory
- Overwriting memory
- Referencing nonexistent variables
- Freeing blocks multiple times
- Referencing freed blocks
- Failing to free blocks

Dereferencing Bad Pointers

- The classic scanf bug

```
int val;  
  
...  
  
scanf("%d", val);
```

Reading Uninitialized Memory

- Assuming that heap data is initialized to zero

```
/* return y = Ax */
int *matvec(int **A, int *x) {
    int *y = malloc(N*sizeof(int));
    int i, j;

    for (i=0; i<N; i++)
        for (j=0; j<N; j++)
            y[i] += A[i][j]*x[j];
    return y;
}
```


Overwriting Memory

- Allocating the (possibly) wrong sized object

```
int **p;  
  
p = malloc(N*sizeof(int));  
  
for (i=0; i<N; i++) {  
    p[i] = malloc(M*sizeof(int));  
}
```

Overwriting Memory

- Not checking the max string size

```
char s[8];  
int i;  
  
gets(s);  /* reads "123456789" from stdin */
```

- Basis for classic buffer overflow attacks
 - 1988 Internet worm
 - modern attacks on Web servers
 - AOL/Microsoft IM war

Overwriting Memory

- Misunderstanding pointer arithmetic

```
int *search(int *p, int val) {  
    while (*p && *p != val)  
        p += sizeof(int);  
  
    return p;  
}
```

Overwriting Memory

- Referencing a pointer instead of the object it points to

```
int *BinheapDelete(int **binheap, int *size) {  
    int *packet;  
    packet = binheap[0];  
    binheap[0] = binheap[*size - 1];  
    *size--;  
    Heapify(binheap, *size, 0);  
    return(packet);  
}
```

Referencing Nonexistent Variables

- Forgetting that local variables disappear when a function returns

```
int *foo () {  
    int val;  
  
    return &val;  
}
```

Freeing Blocks Multiple Times

■ Nasty!

```
x = malloc(N*sizeof(int));  
    <manipulate x>  
free(x);  
  
y = malloc(M*sizeof(int));  
    <manipulate y>  
free(x);
```

Referencing Freed Blocks

■ Evil!

```
x = malloc(N*sizeof(int));  
    <manipulate x>  
free(x);  
    ...  
y = malloc(M*sizeof(int));  
for (i=0; i<M; i++)  
    y[i] = x[i]++;
```

Failing to Free Blocks (Memory Leaks)

- Slow, long-term killer!

```
foo() {  
    int *x = malloc(N*sizeof(int));  
    ...  
    return;  
}
```


Failing to Free Blocks (Memory Leaks)

- Freeing only part of a data structure

```
struct list {
    int val;
    struct list *next;
};

foo() {
    struct list *head = malloc(sizeof(struct list));
    head->val = 0;
    head->next = NULL;
    <create and manipulate the rest of the list>
    ...
    free(head);
    return;
}
```

Dealing With Memory Bugs

- Conventional debugger (gdb)
 - Good for finding bad pointer dereferences
 - Hard to detect the other memory bugs
- Debugging malloc (UToronto CSRI malloc)
 - Wrapper around conventional malloc
 - Detects memory bugs at malloc and free boundaries
 - ▶ Memory overwrites that corrupt heap structures
 - ▶ Some instances of freeing blocks multiple times
 - ▶ Memory leaks
 - Cannot detect all memory bugs
 - ▶ Overwrites into the middle of allocated blocks
 - ▶ Freeing block twice that has been reallocated in the interim
 - ▶ Referencing freed blocks

Dealing With Memory Bugs (cont.)

- Some malloc implementations contain checking code
 - Linux glibc malloc: `setenv MALLOC_CHECK_ 2`
 - FreeBSD: `setenv MALLOC_OPTIONS AJR`
- Binary translator: valgrind (Linux), Purify
 - Powerful debugging and analysis technique
 - Rewrites text section of executable object file
 - Can detect all errors as debugging malloc
 - Can also check each individual reference at runtime
 - ▶ Bad pointers
 - ▶ Overwriting
 - ▶ Referencing outside of allocated block
- Garbage collection (Boehm-Weiser Conservative GC)
 - Let the system free blocks instead of the programmer.